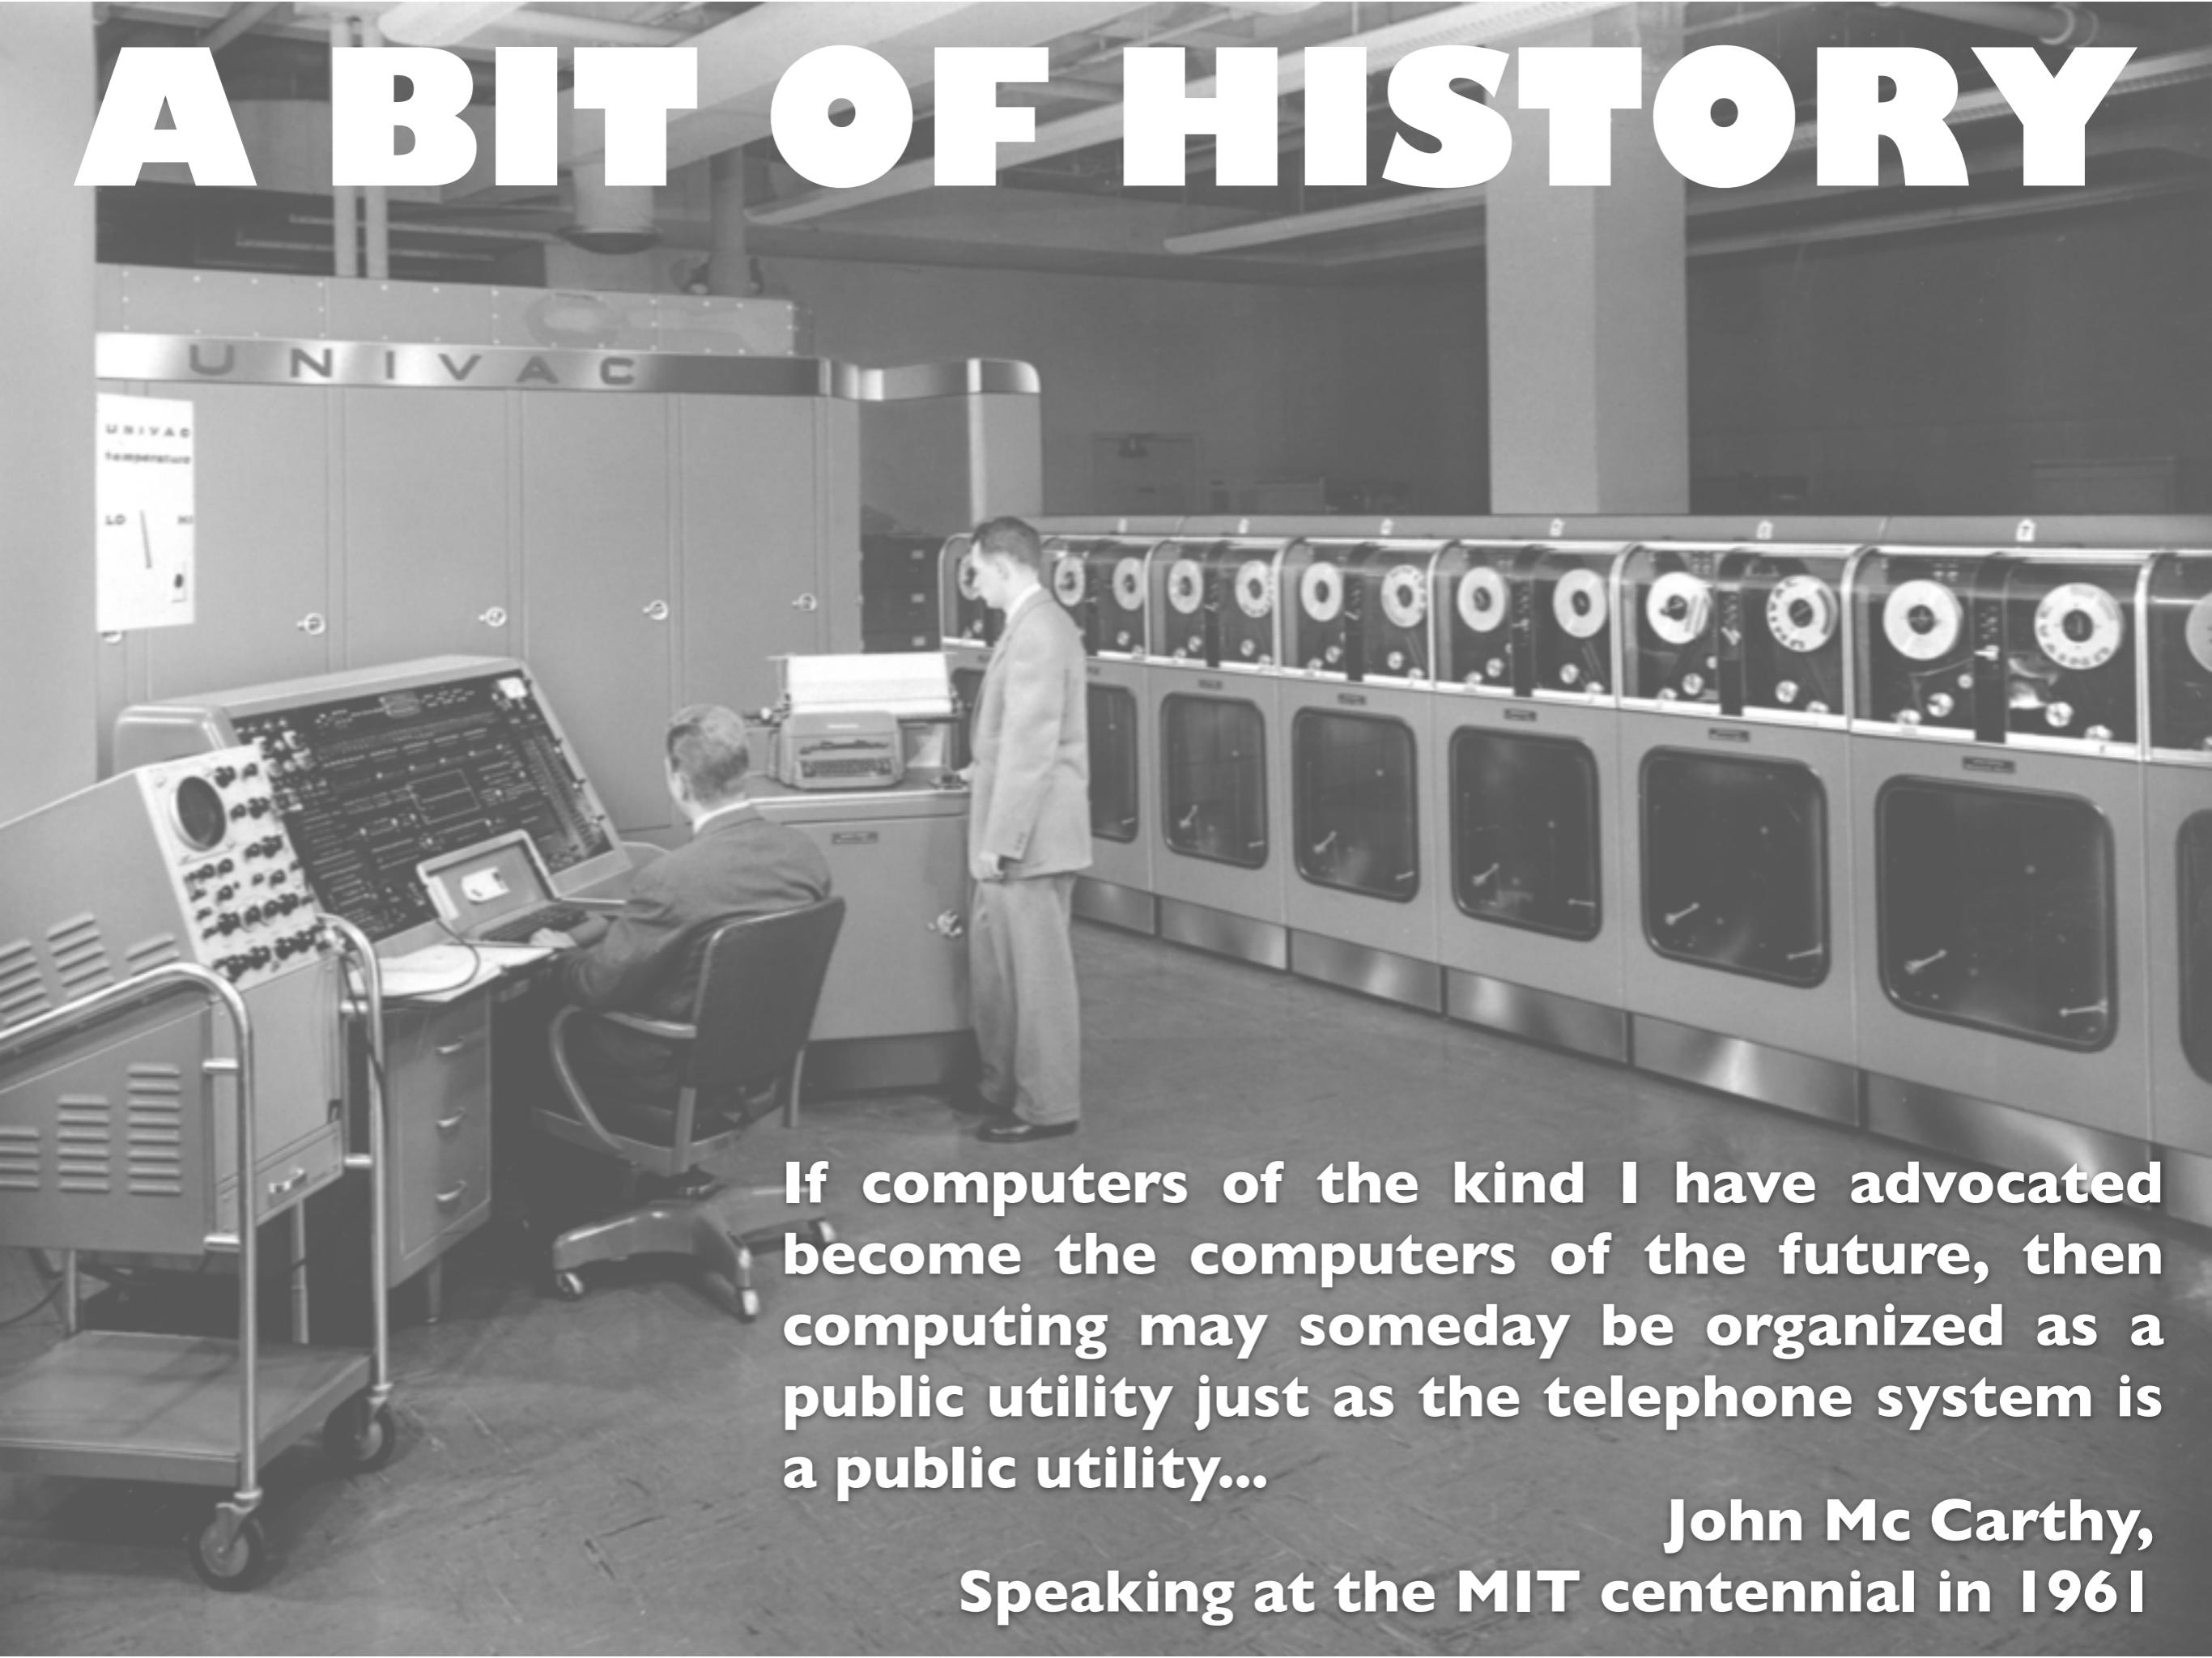


# Infrastructure-as-a-Service

# OpenStack

Adrien Lebre  
STACK Research Group

# A BIT OF HISTORY



If computers of the kind I have advocated become the computers of the future, then computing may someday be organized as a public utility just as the telephone system is a public utility...

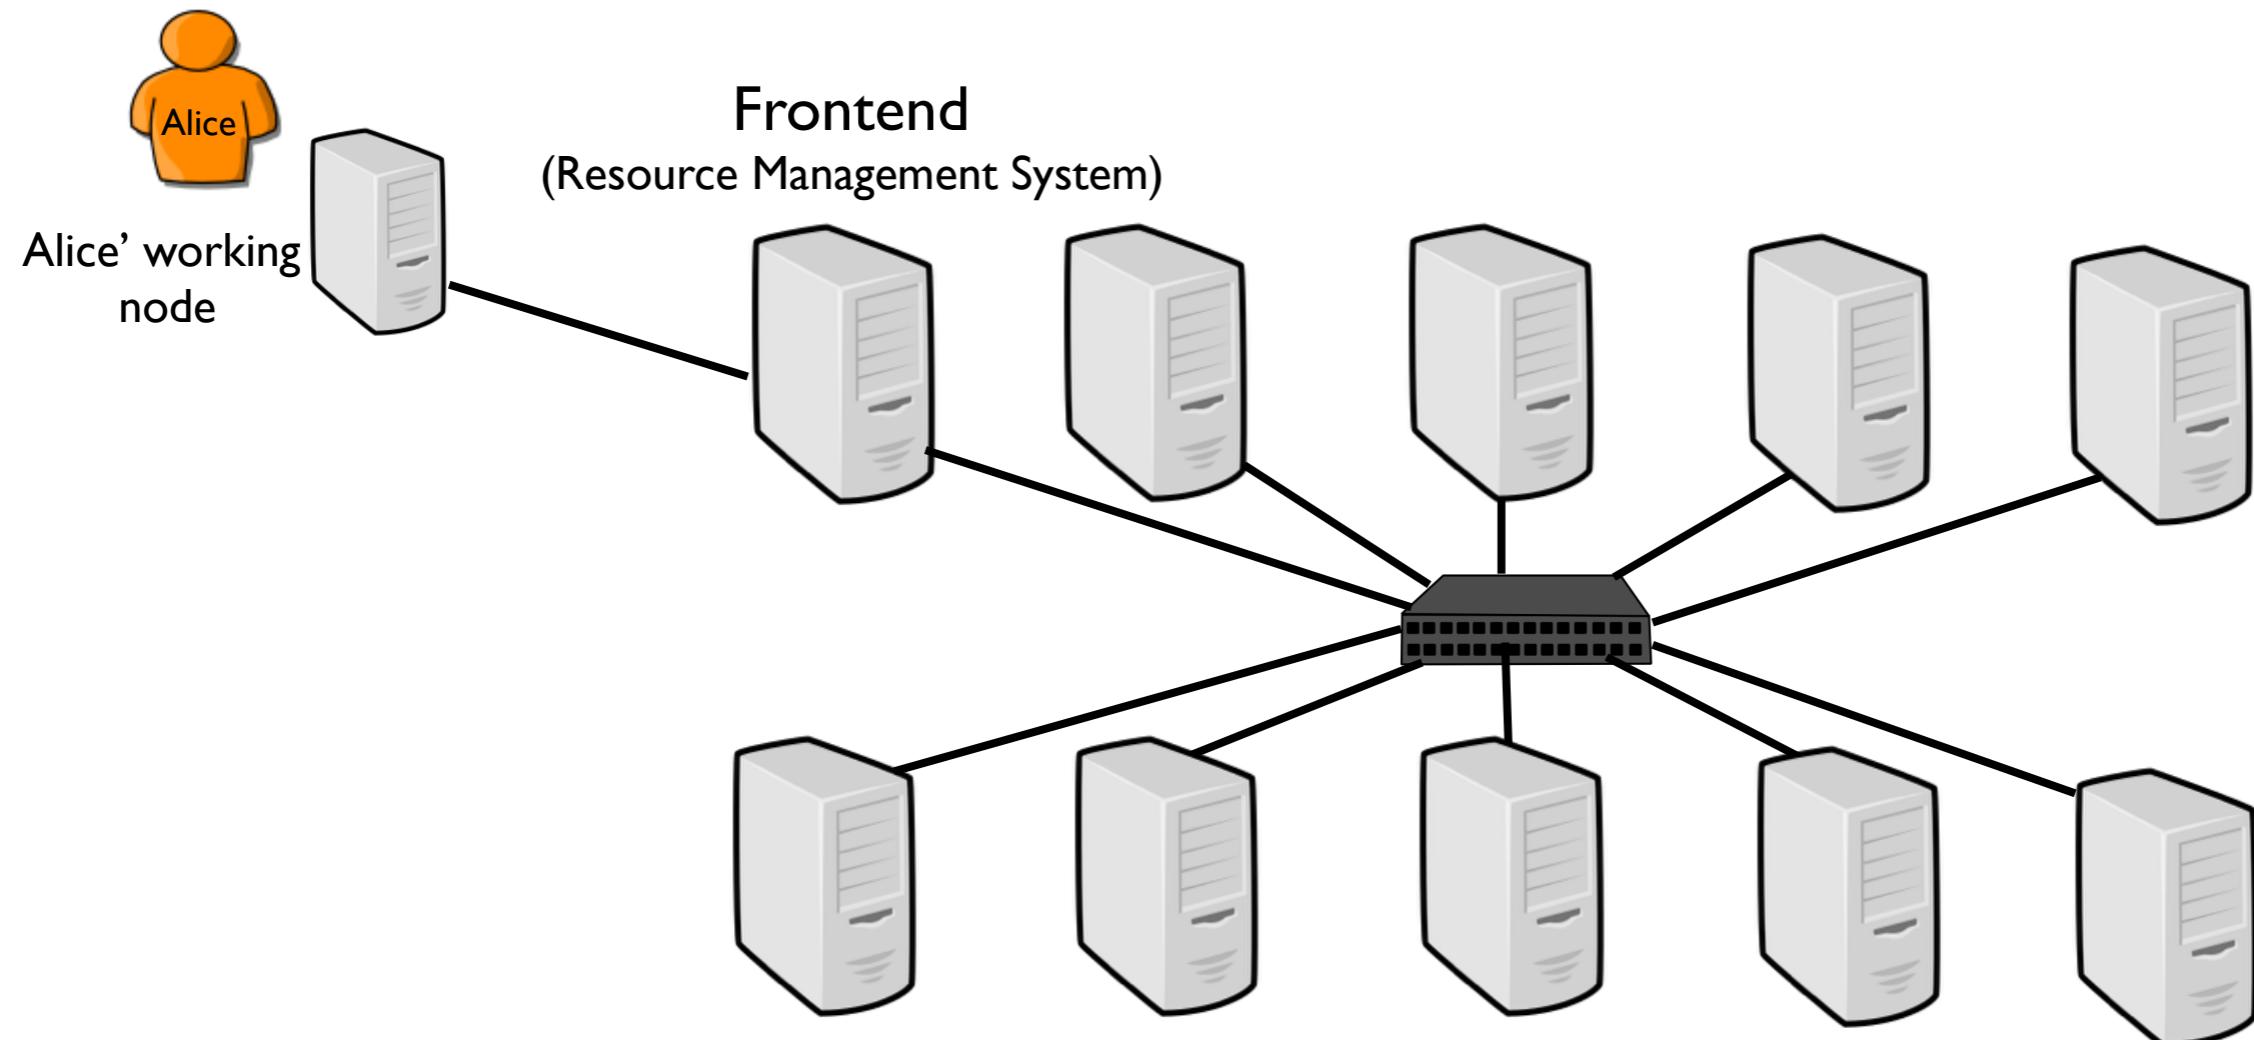
John Mc Carthy,  
Speaking at the MIT centennial in 1961

# Looking back...

- **xxx Computing**  
Meta / Cluster / Grid / Desktop / “Hive” / Cloud / Sky ...  
⇒ **xxx as Utility Computing**
- A common objective: provide computing resources  
(both hardware and software) in a flexible, transparent,  
secure, reliable, ... way
- **Challenges**
  - Software/Hardware heterogeneity
  - Security (Isolation between applications, ...)
  - Reliability / Resiliency
  - Data Sharing
  - Performance guarantees...

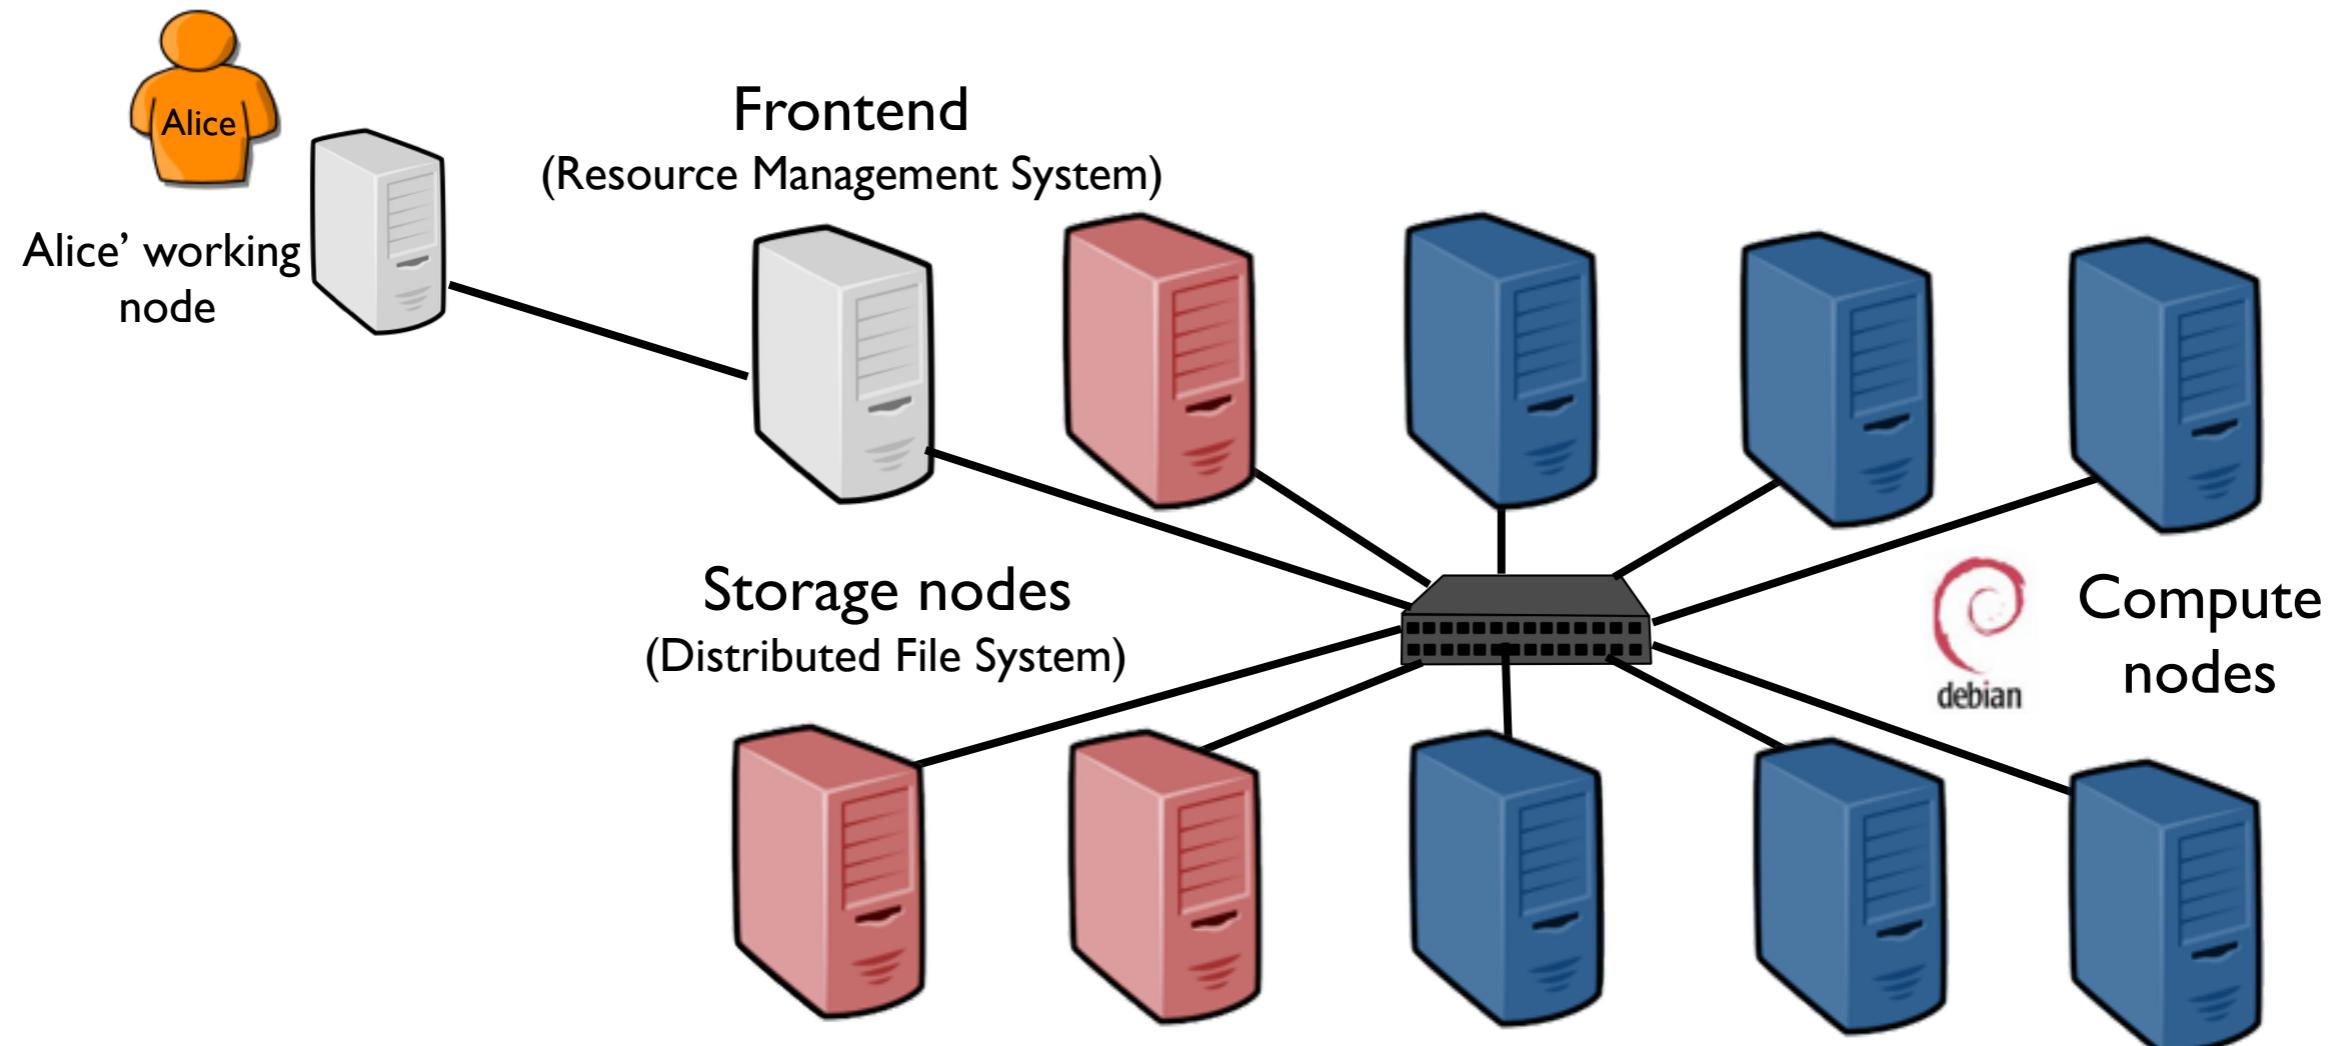
# Looking back...

- Network of Workstations 1990 / 20xx



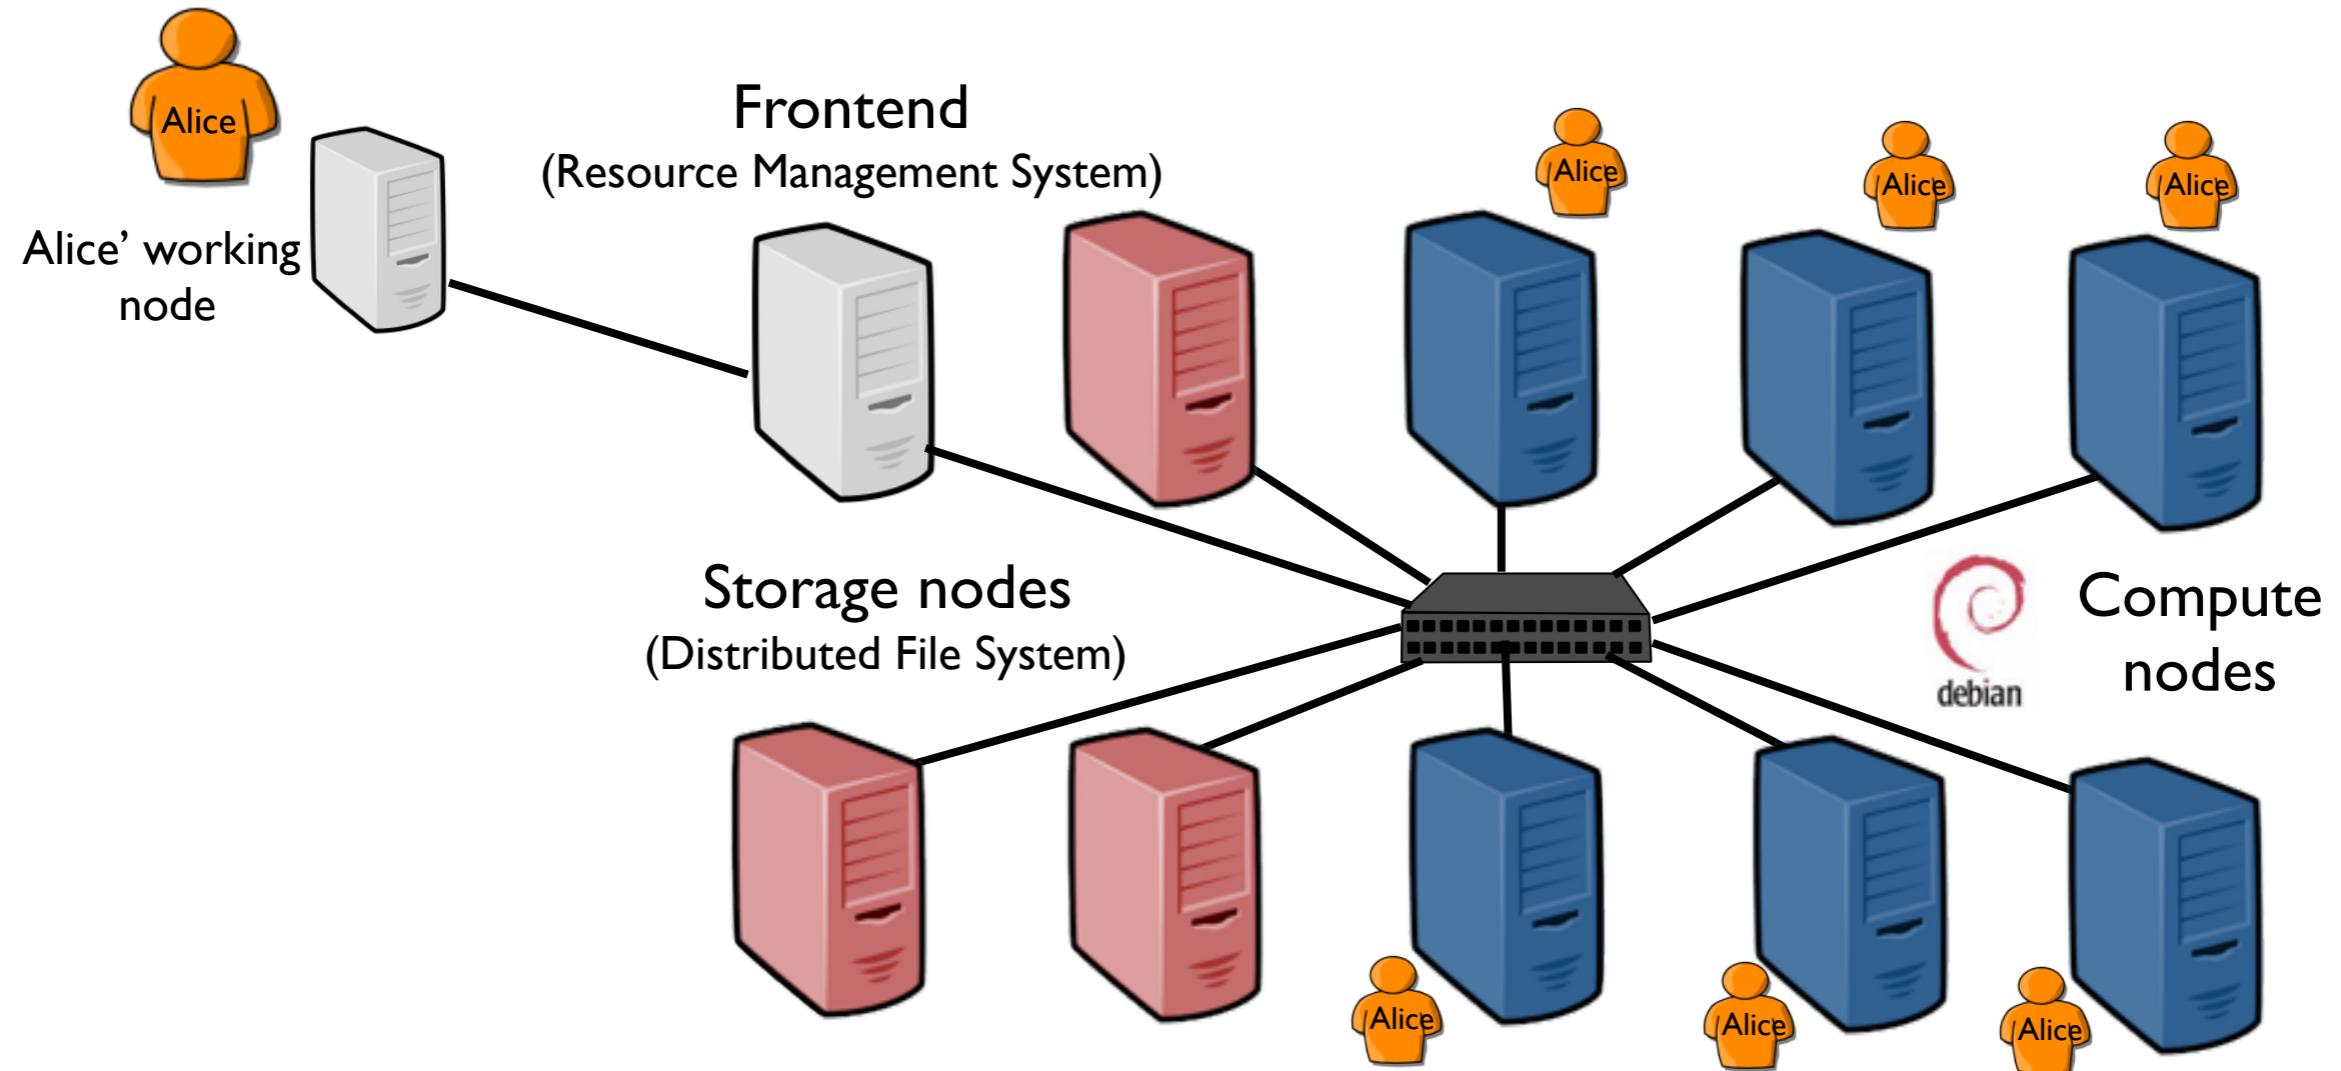
# Looking back...

- Network of Workstations 1990 / 20xx



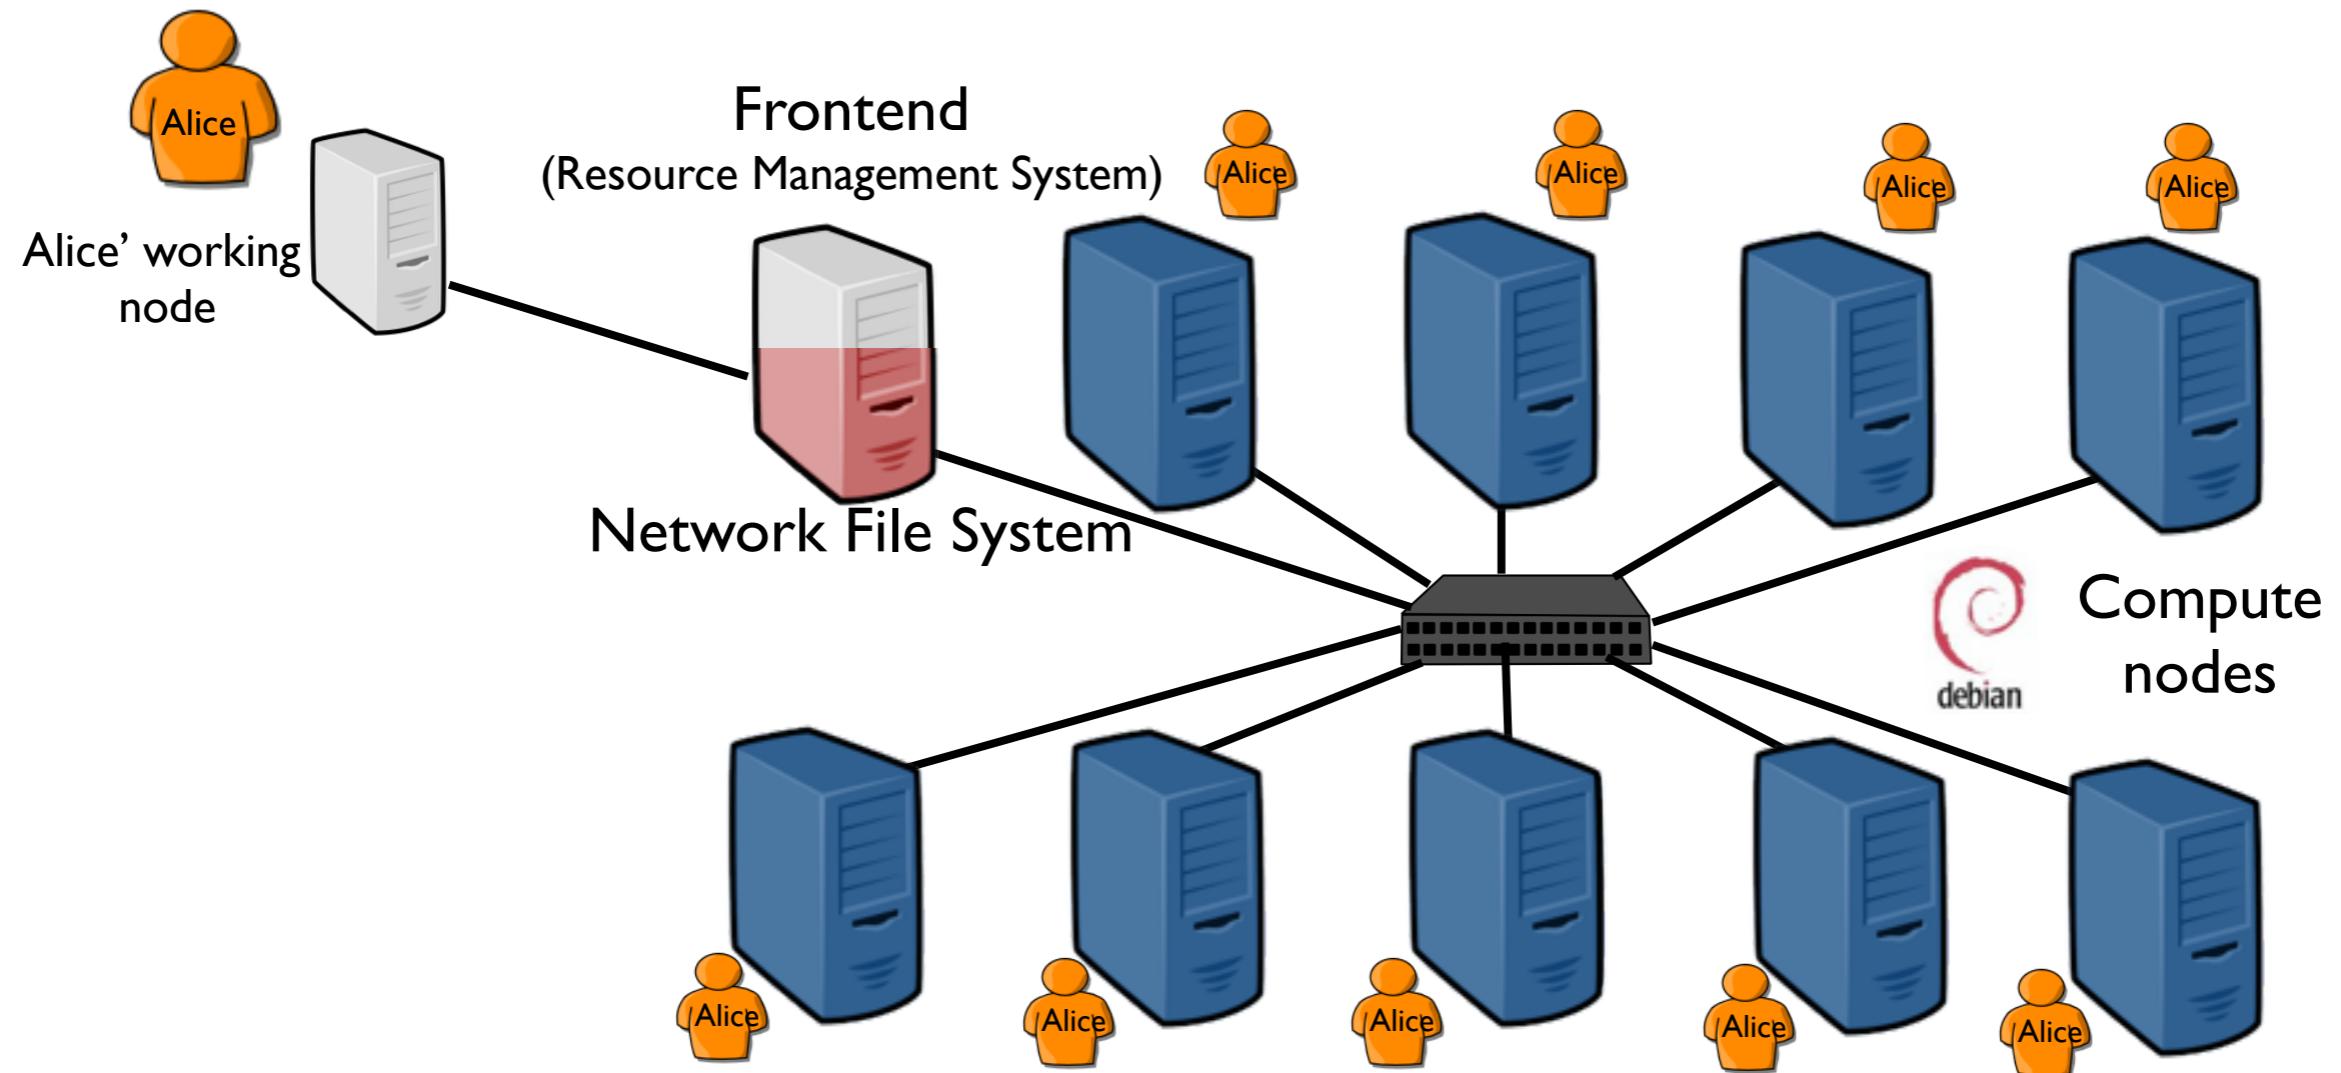
# Looking back...

- Network of Workstations 1990 / 20xx



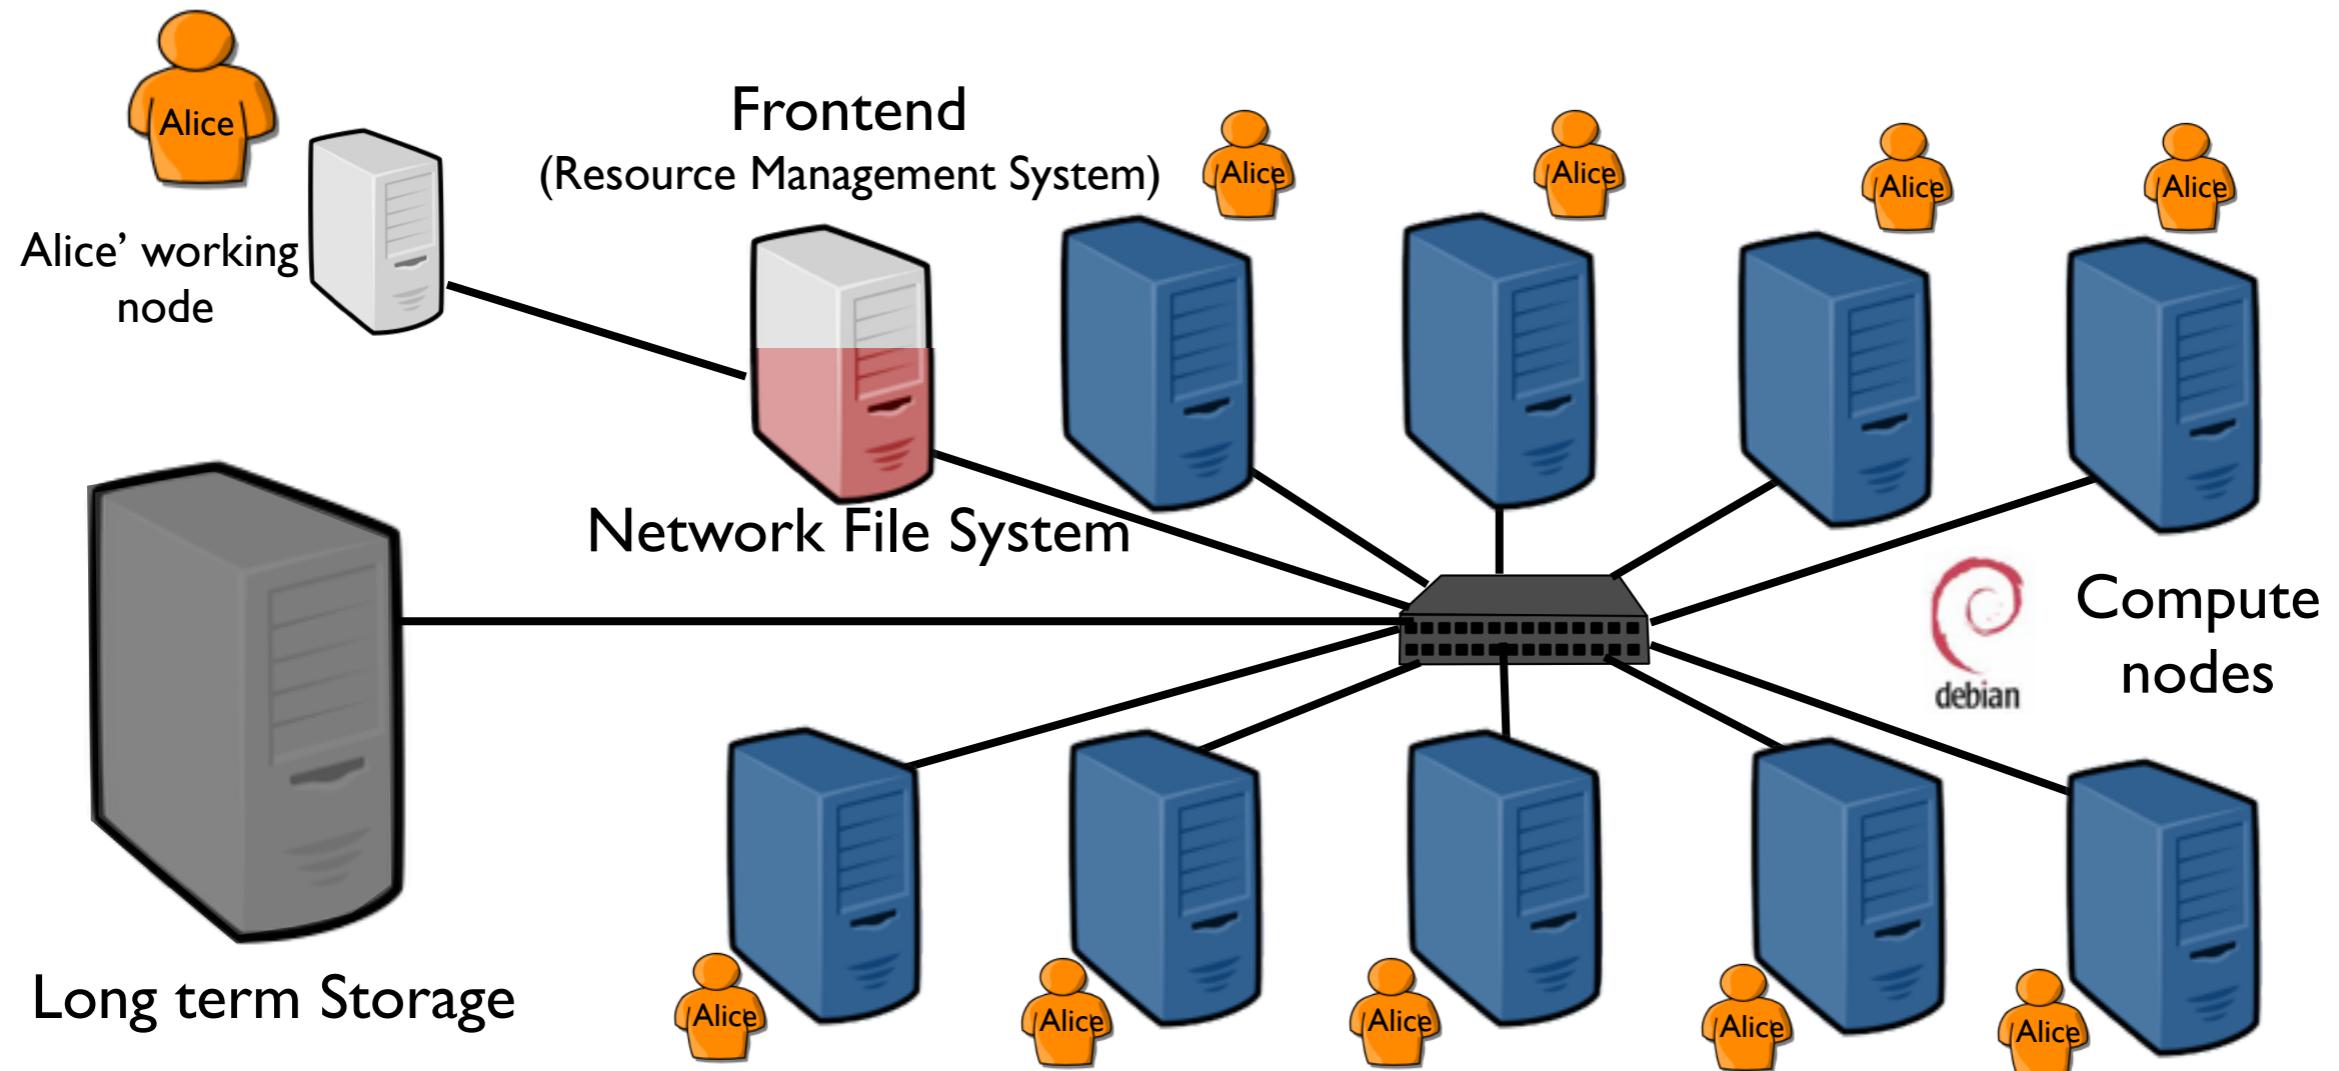
# Looking back...

- Network of Workstations 1990 / 20xx



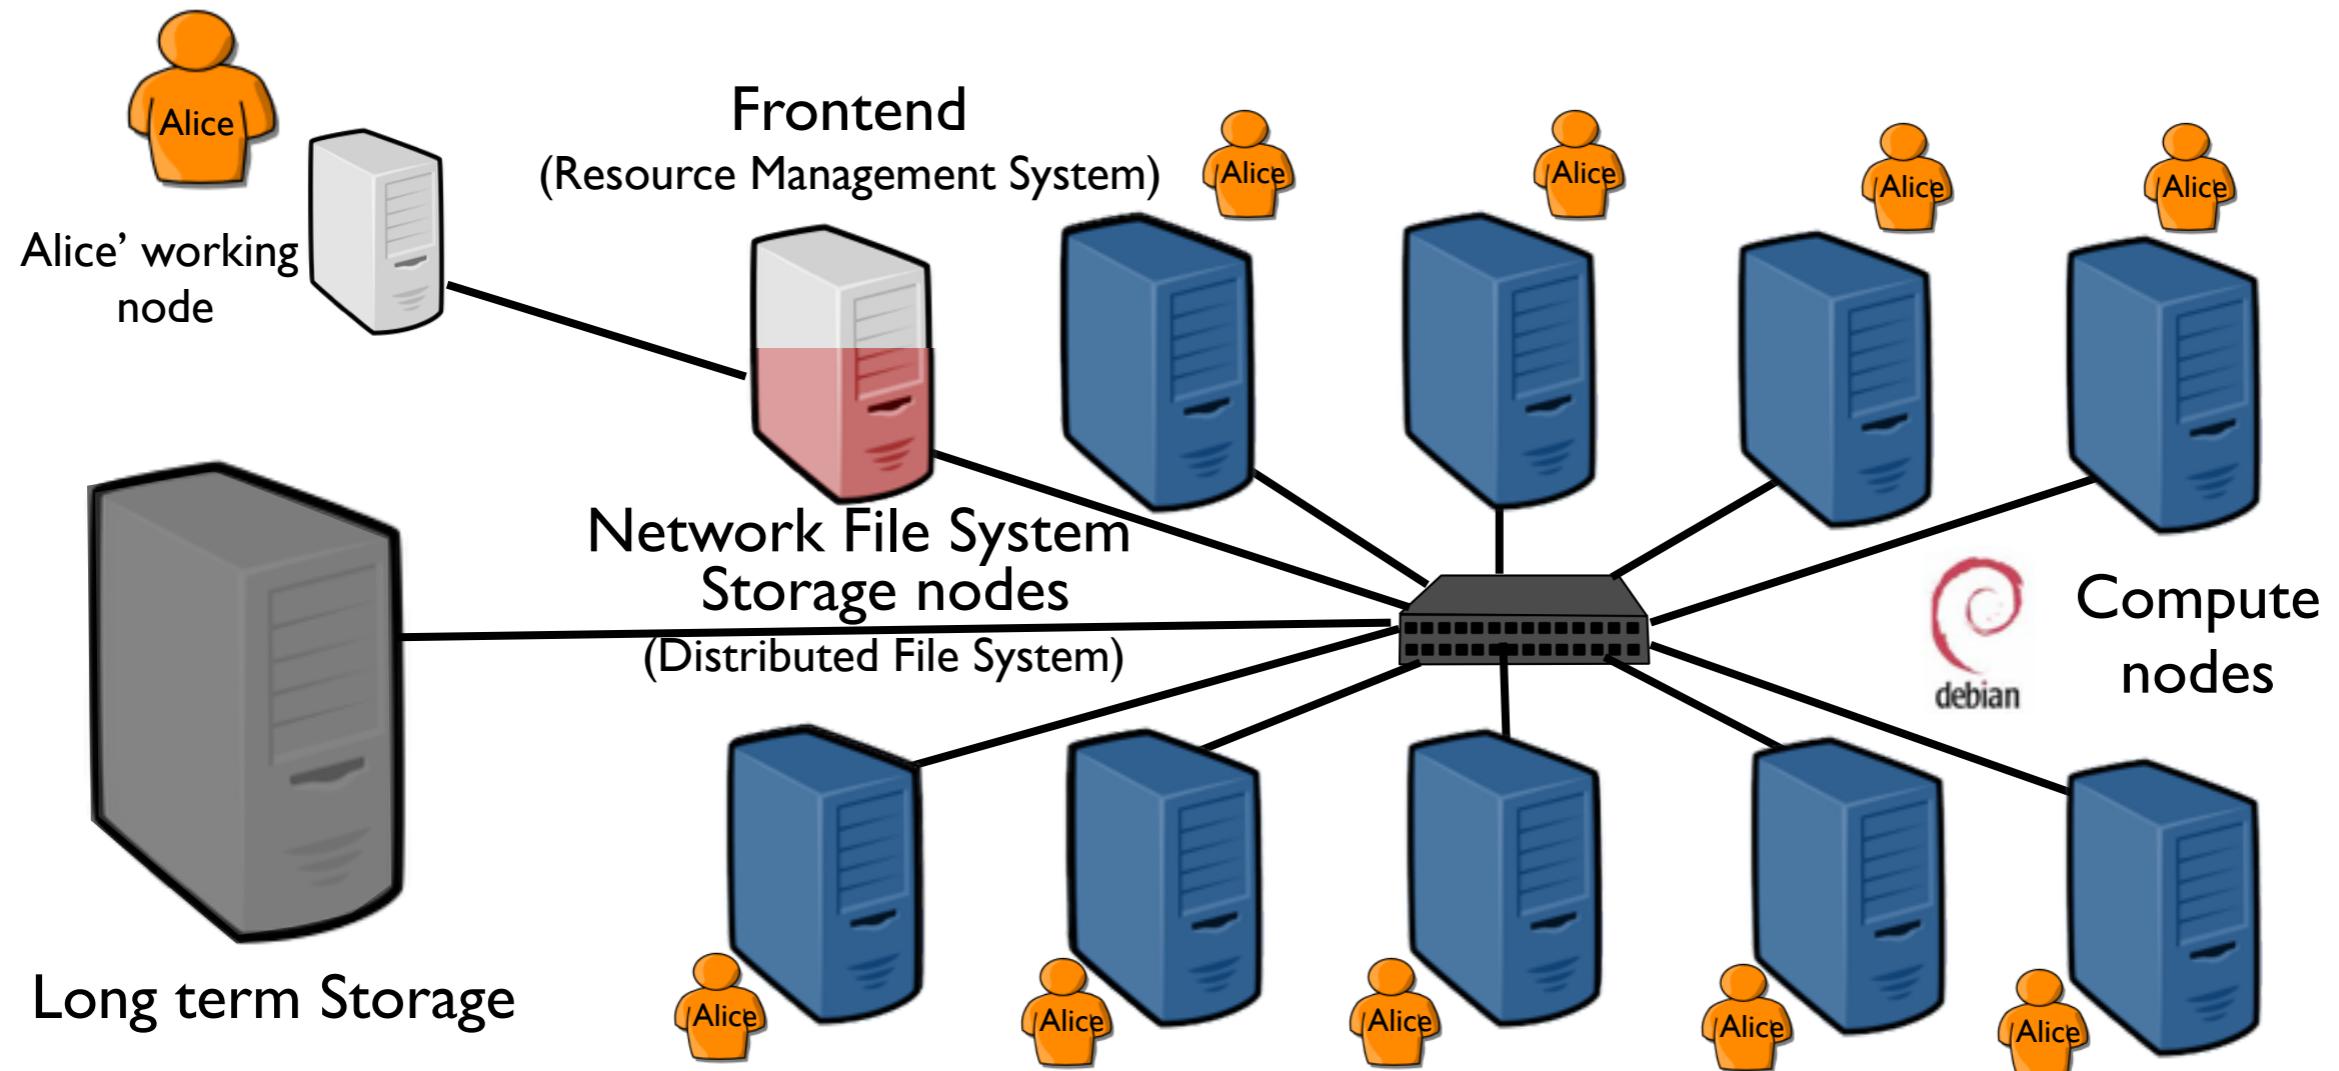
# Looking back...

- Network of Workstations 1990 / 20xx



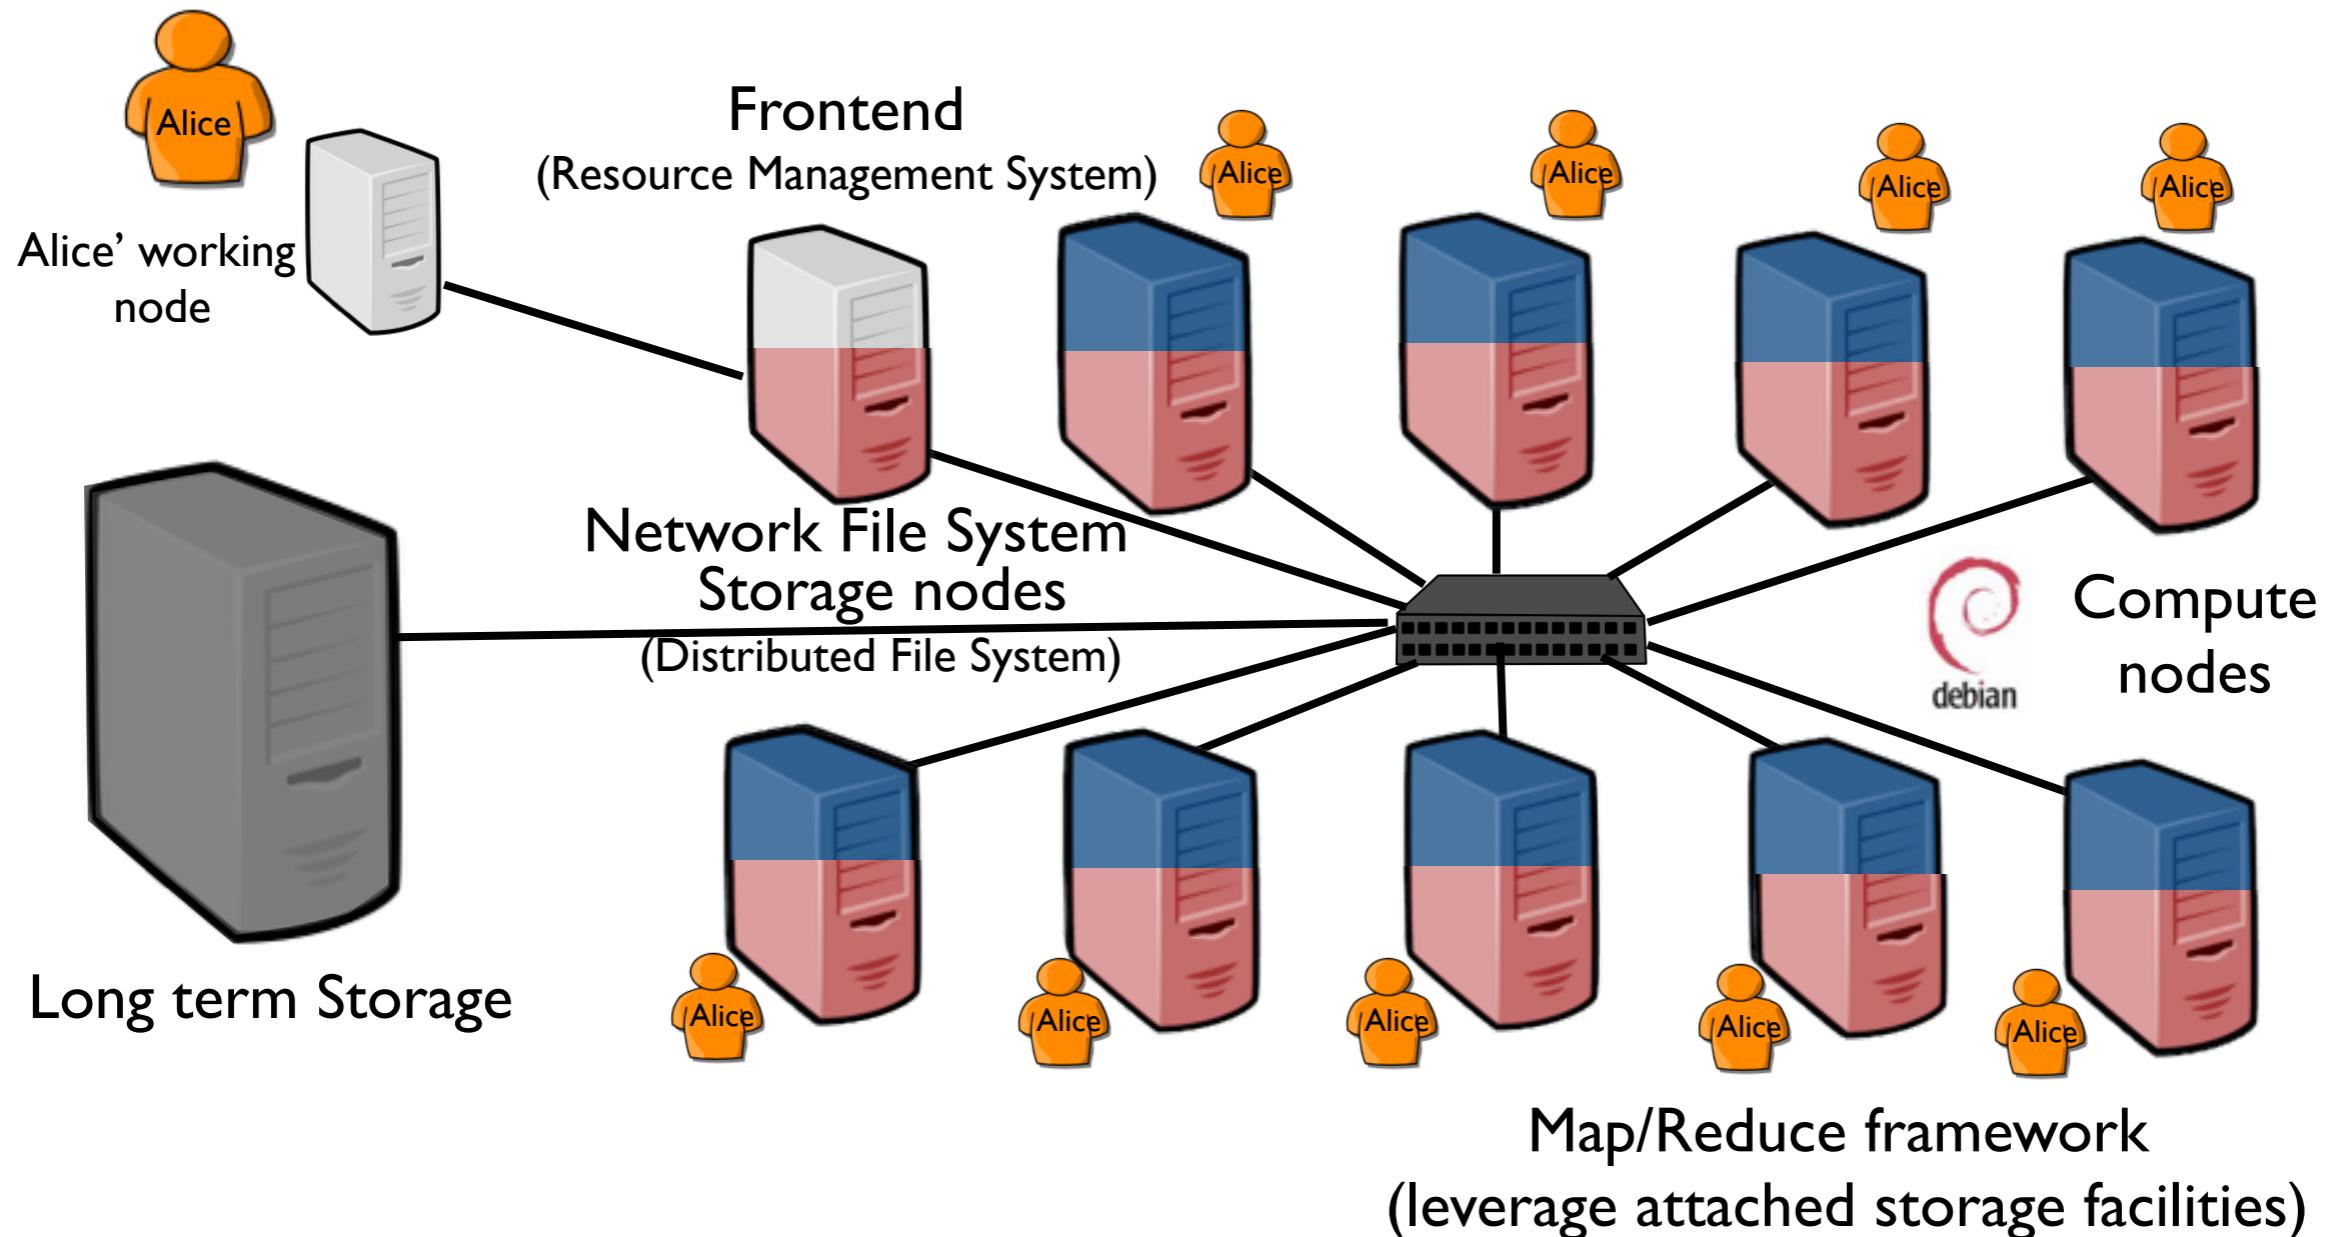
# Looking back...

- Network of Workstations 1990 / 20xx



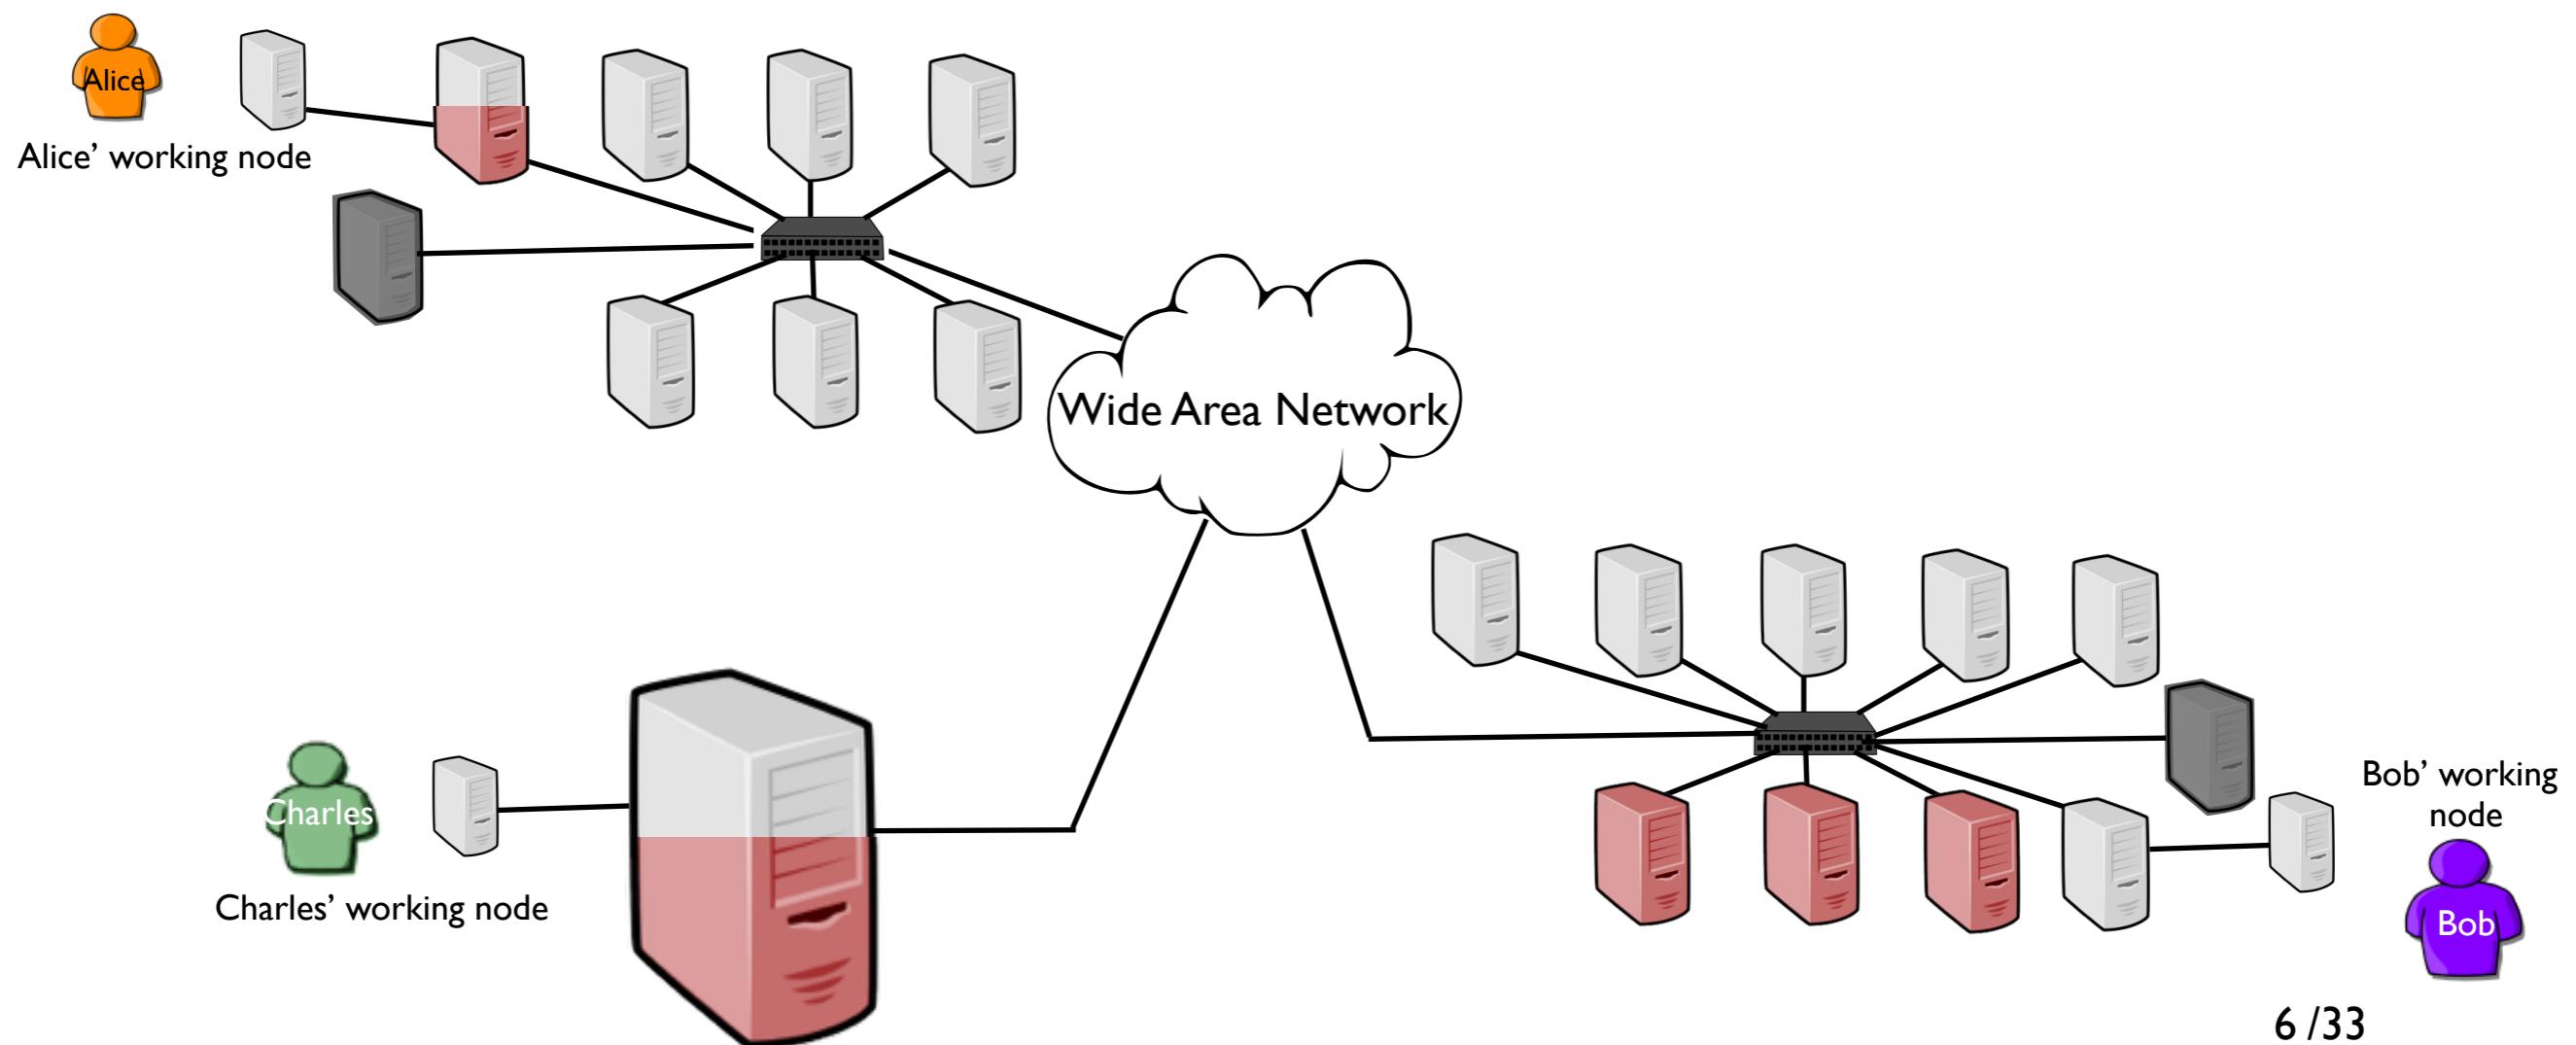
# Looking back...

- Network of Workstations 1990 / 20xx



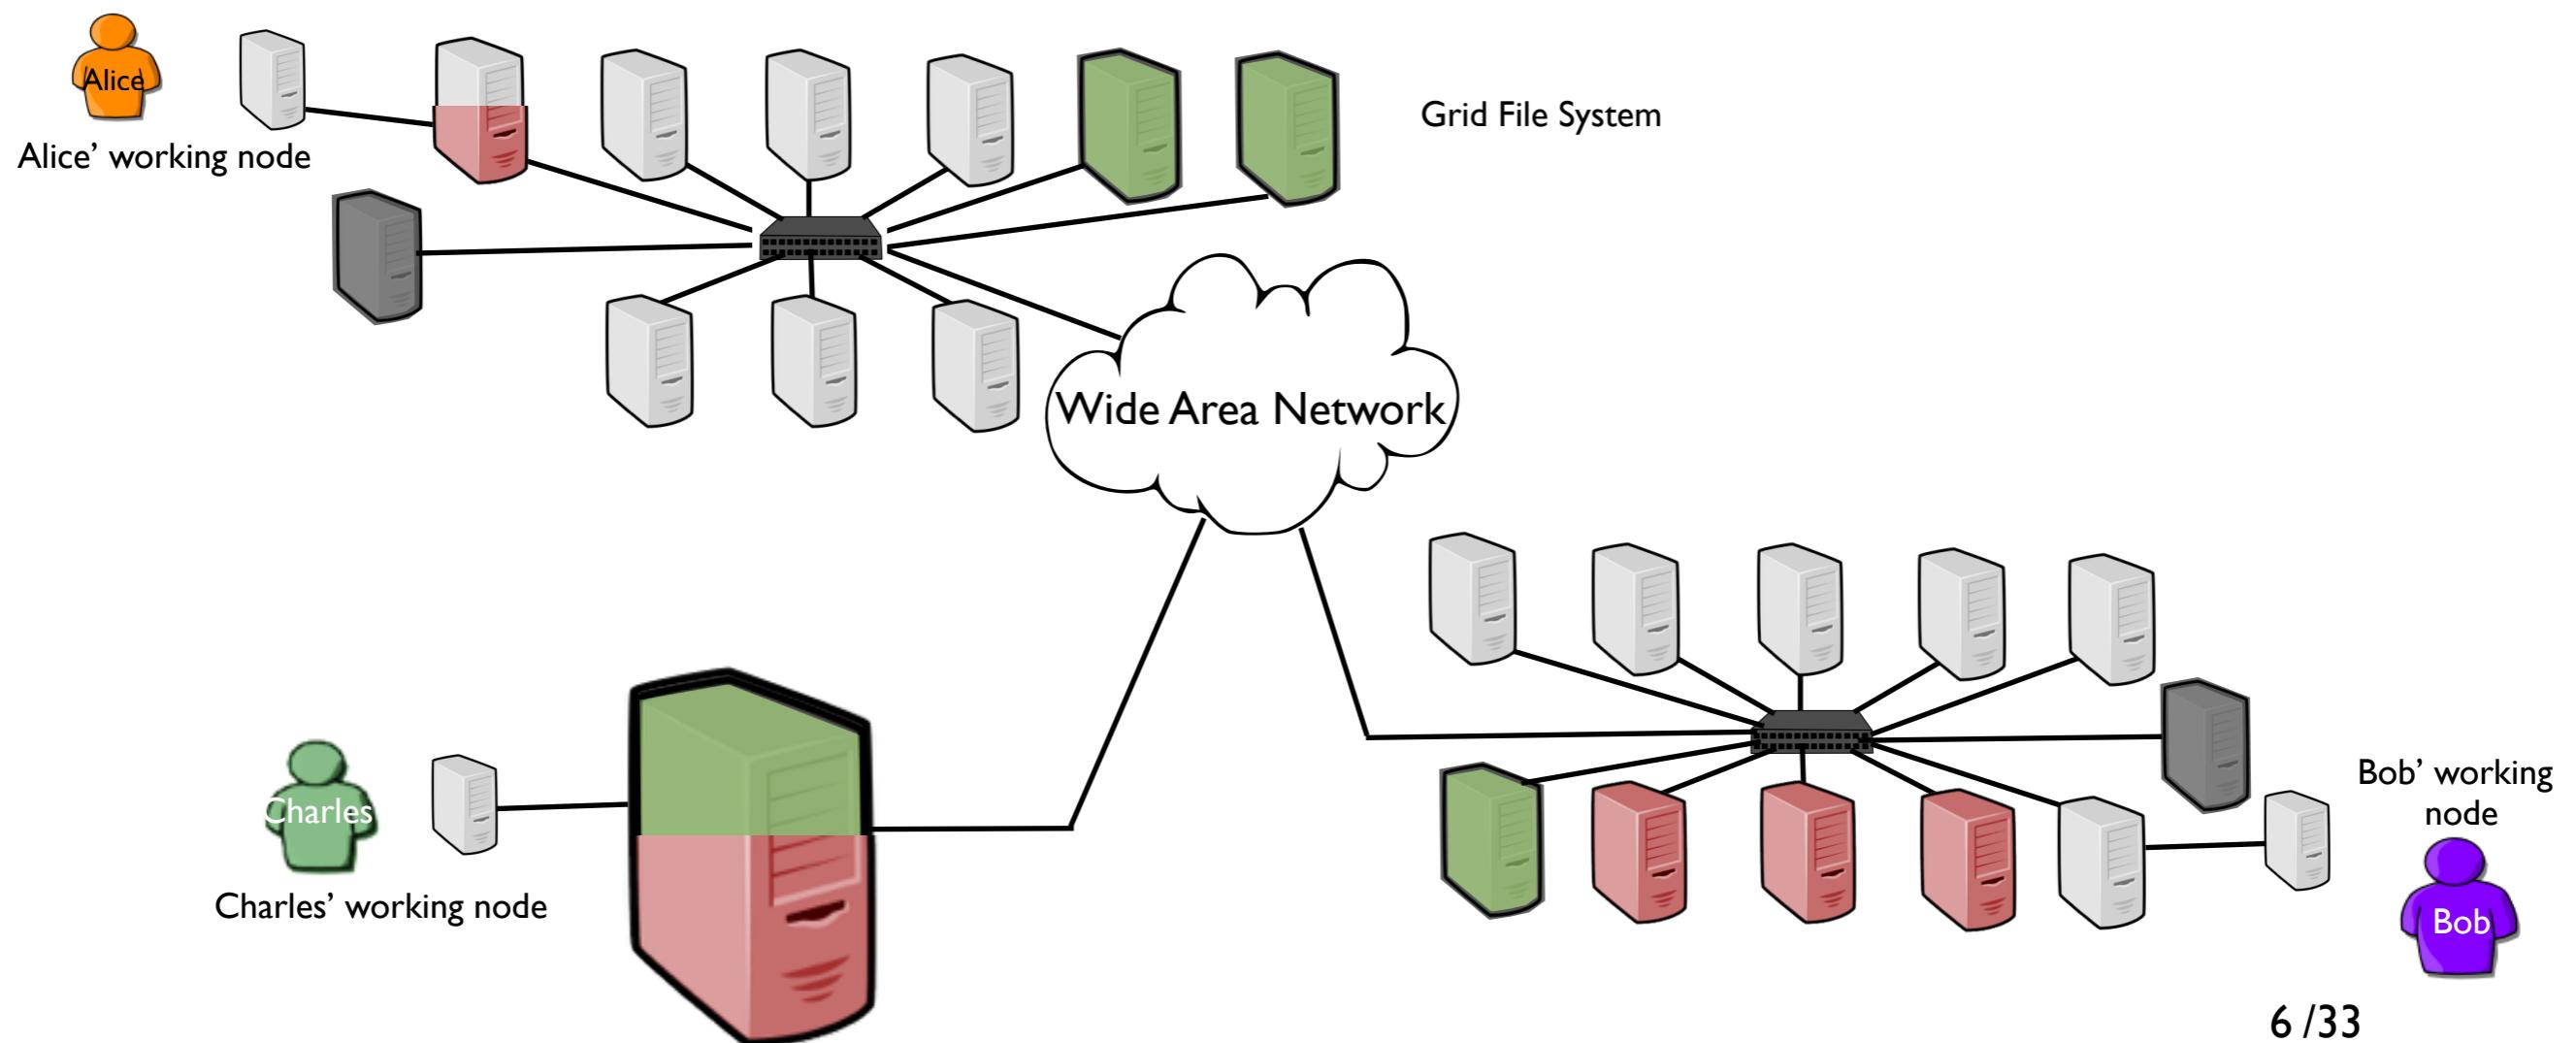
# Looking back...

- Network of Workstations 1990 / 20~~xx~~
- Desktop 1998 / 20~~xx~~
- Grid 1998 / 20~~xx~~



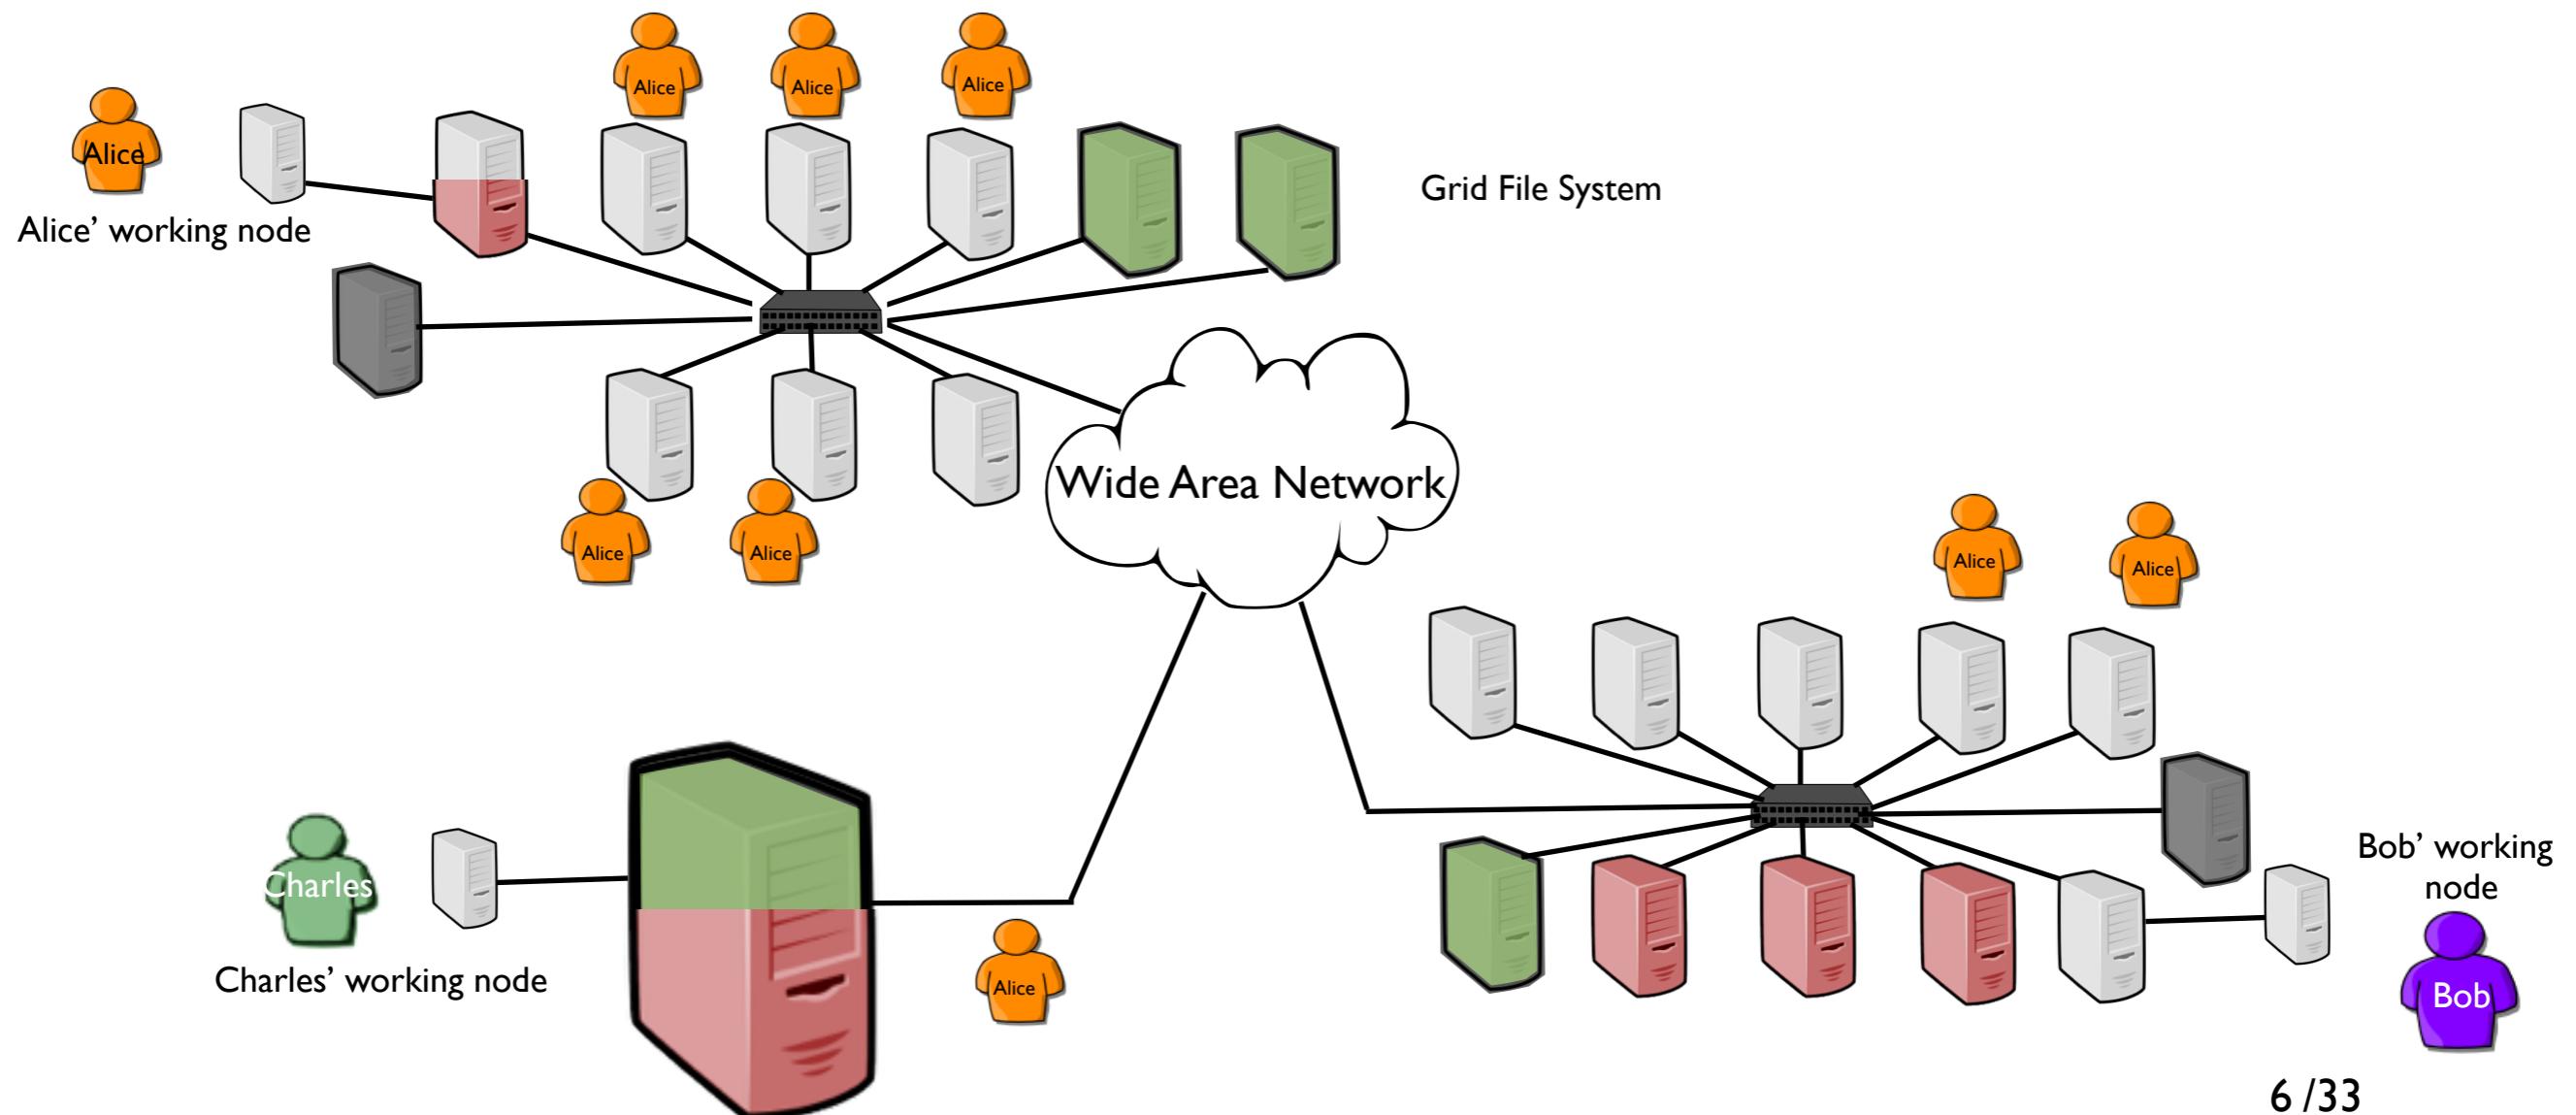
# Looking back...

- Network of Workstations 1990 / 20~~xx~~
- Desktop 1998 / 20~~xx~~
- Grid 1998 / 20~~xx~~



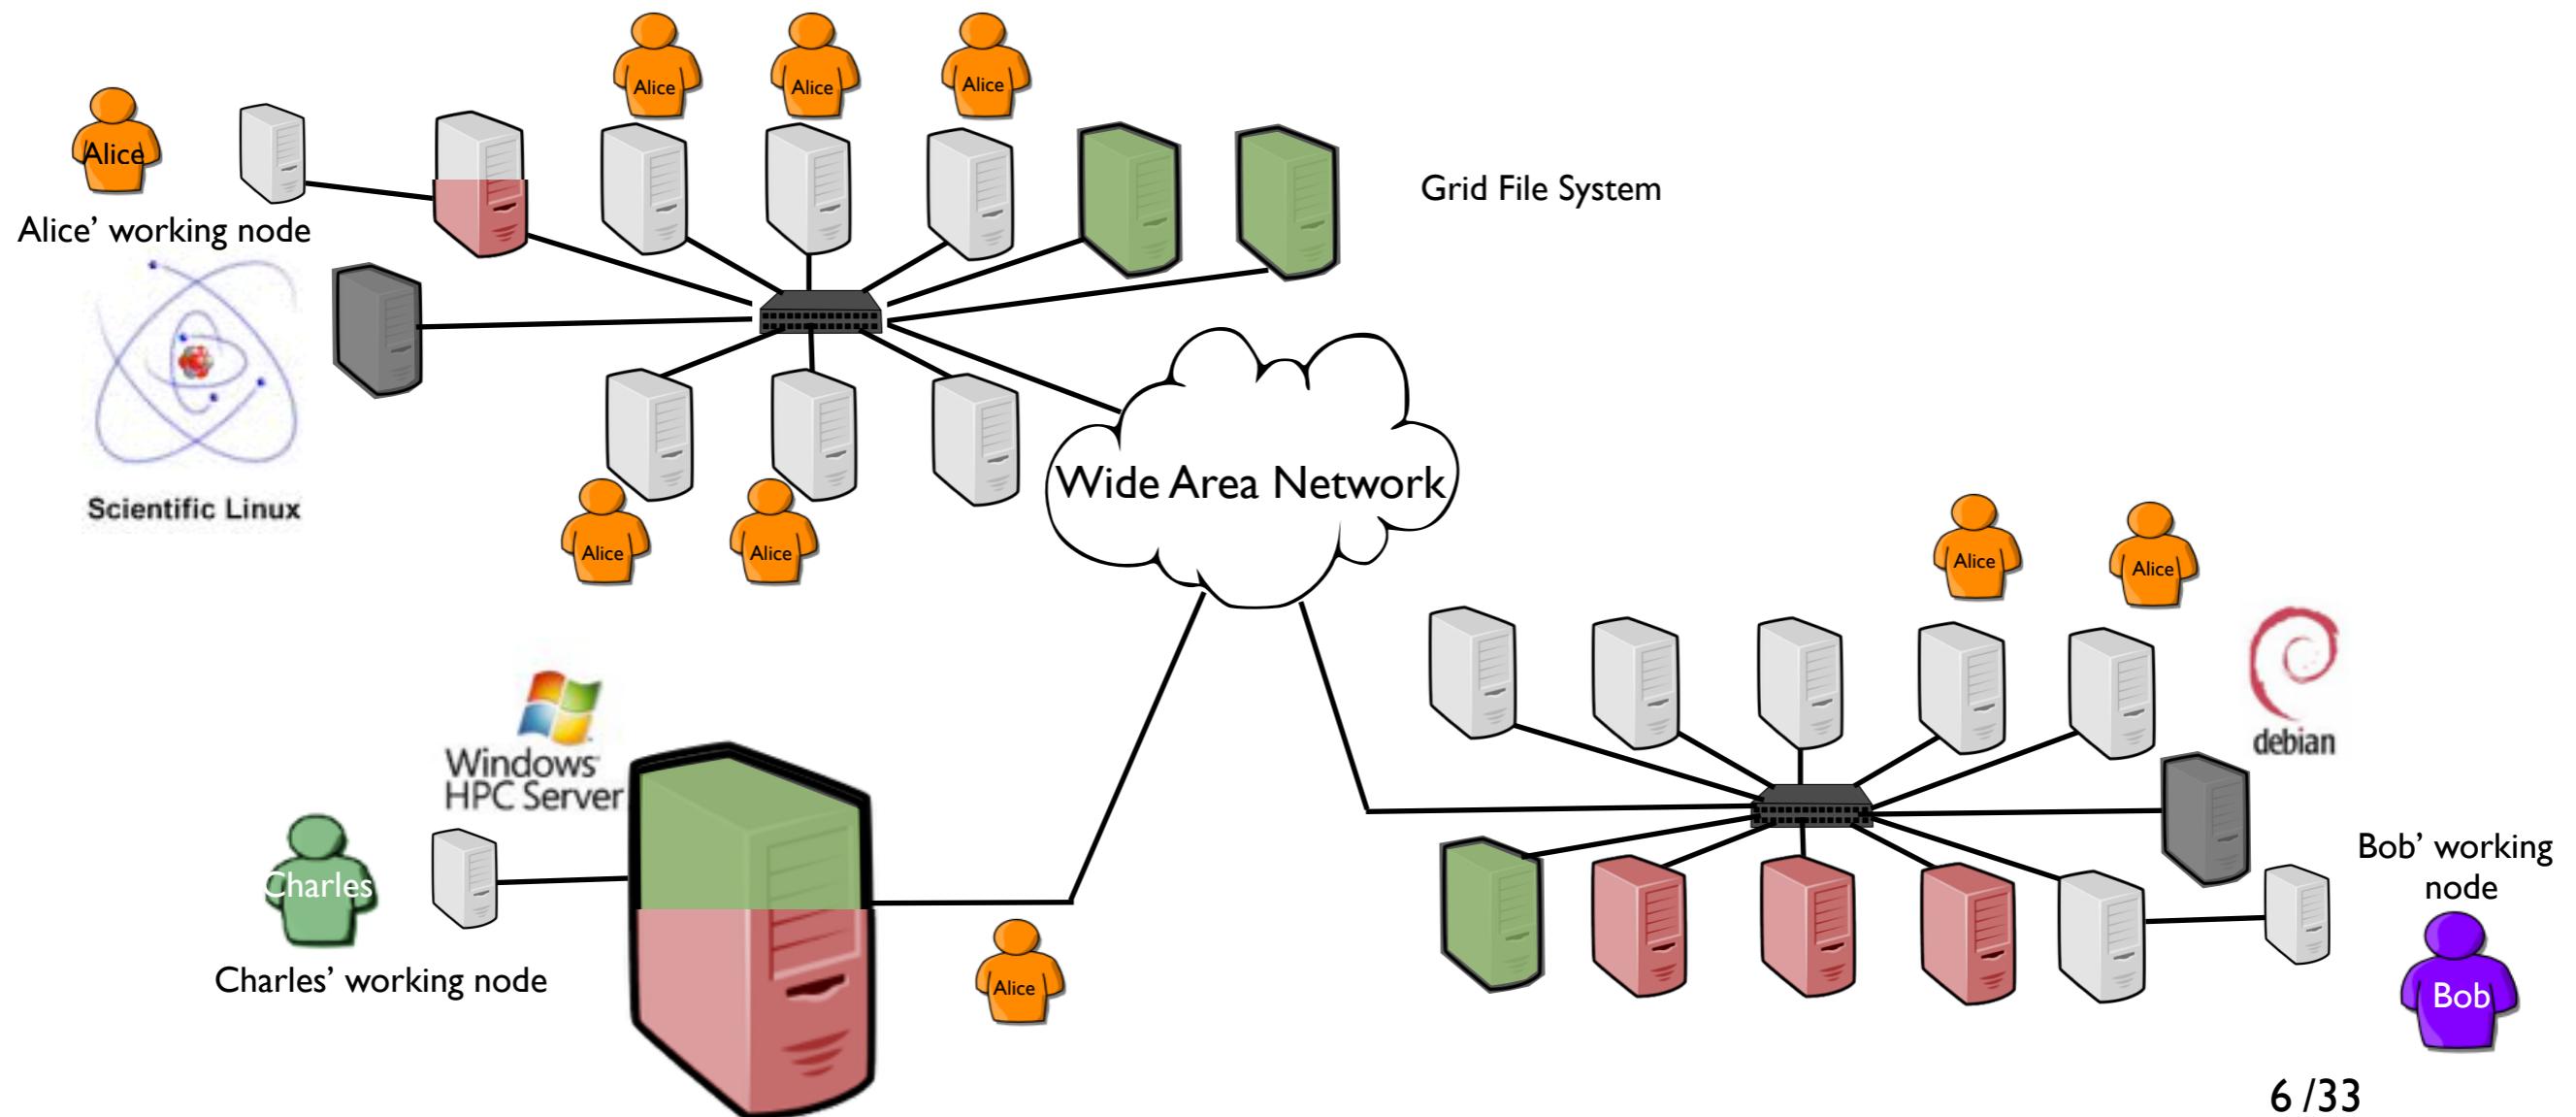
# Looking back...

- Network of Workstations 1990 / 20~~xx~~
- Desktop 1998 / 20~~xx~~
- Grid 1998 / 20~~xx~~



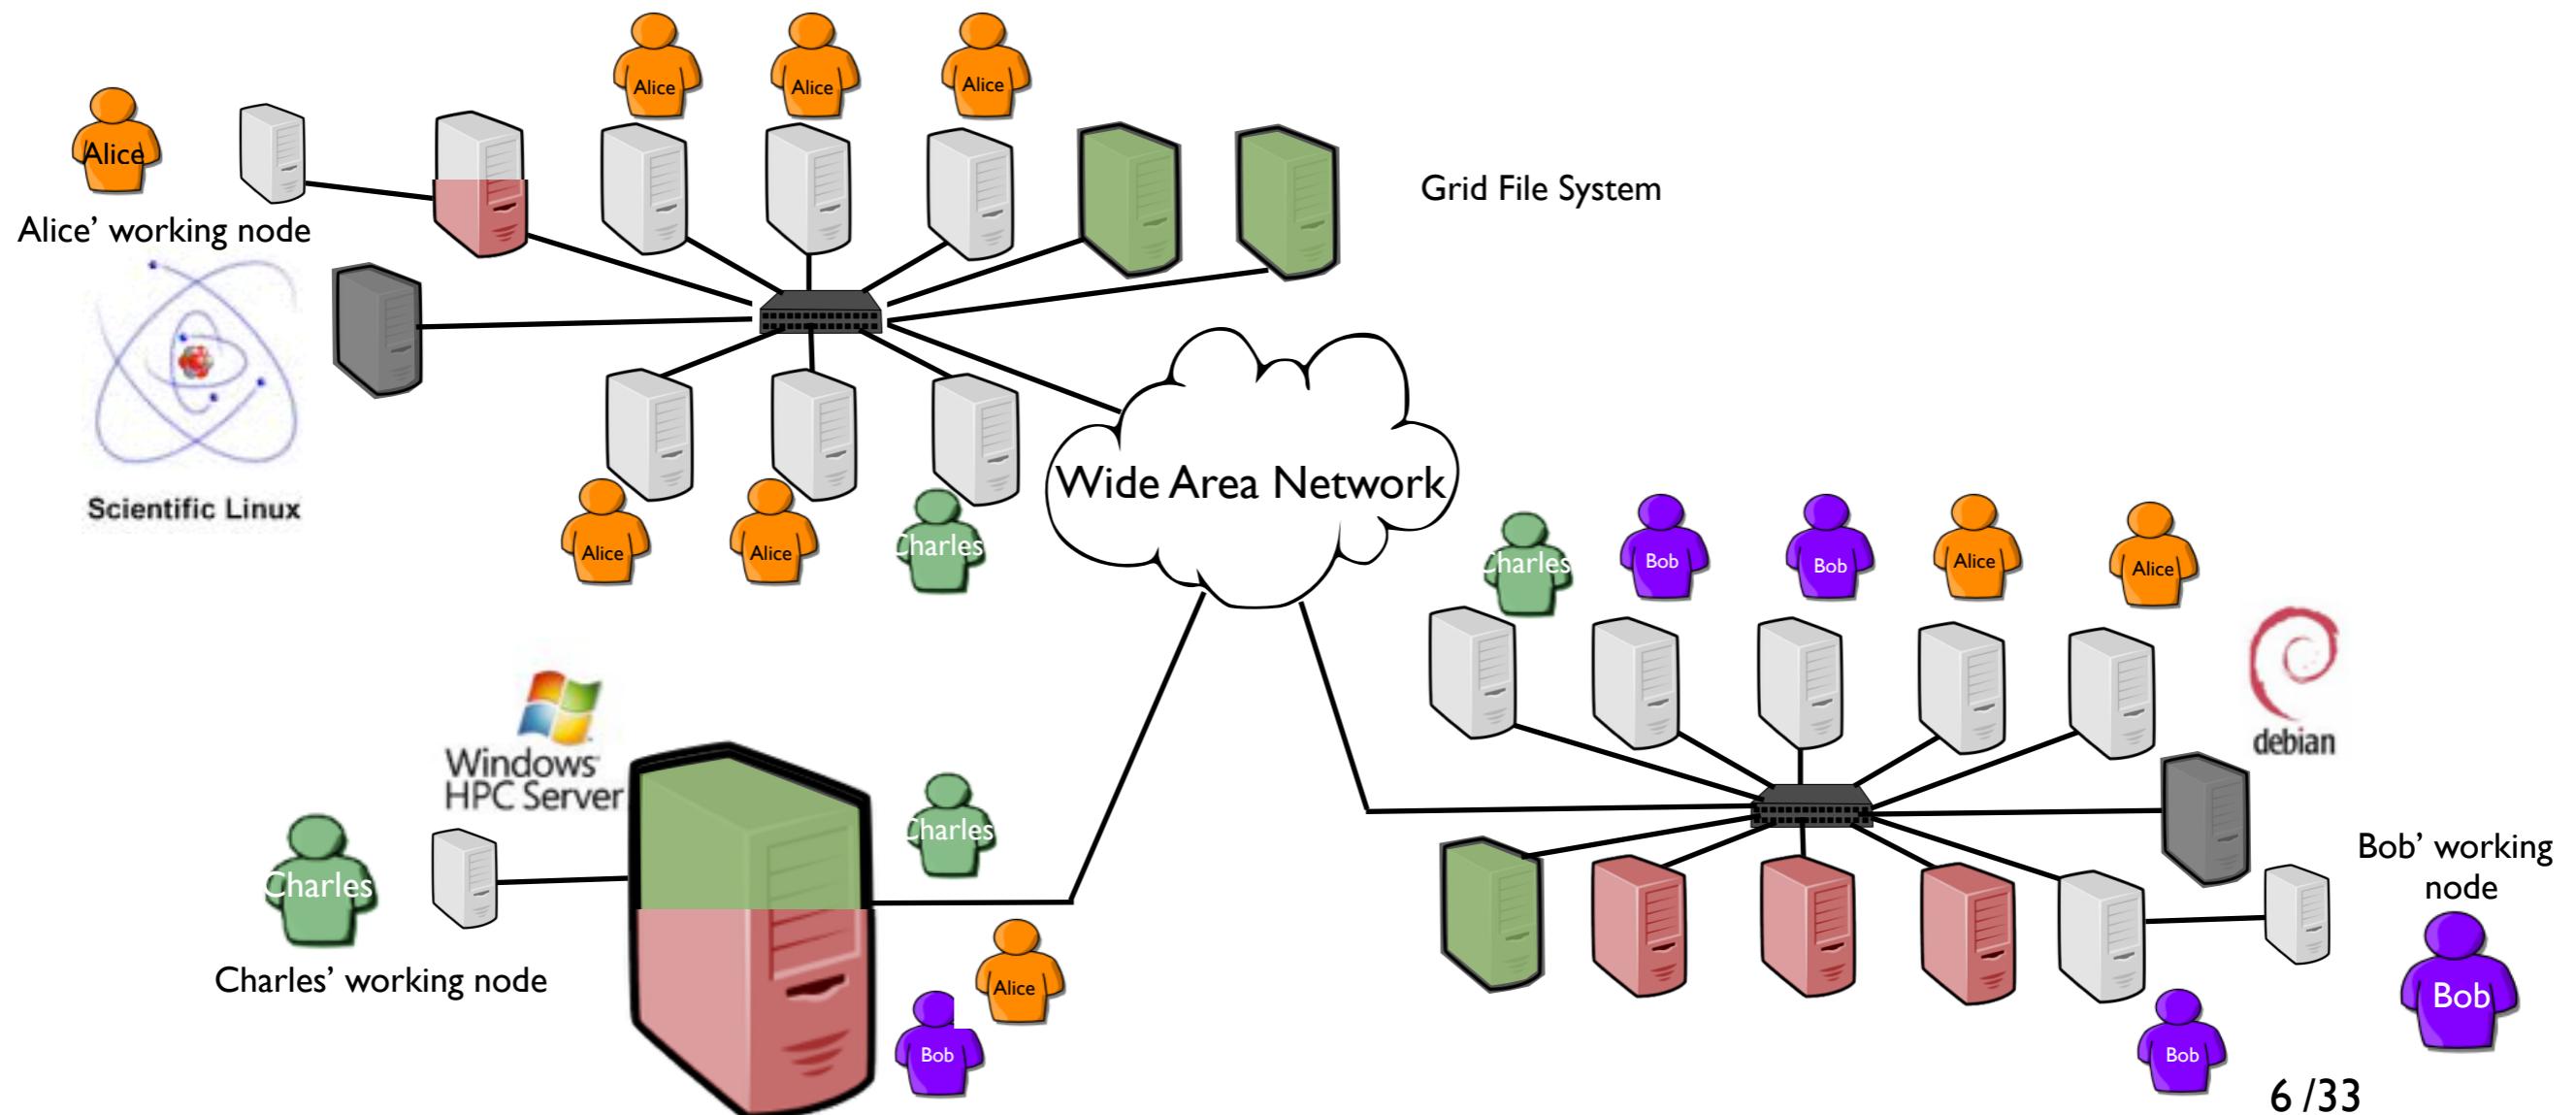
# Looking back...

- Network of Workstations 1990 / 20~~xx~~
- Desktop 1998 / 20~~xx~~
- Grid 1998 / 20~~xx~~



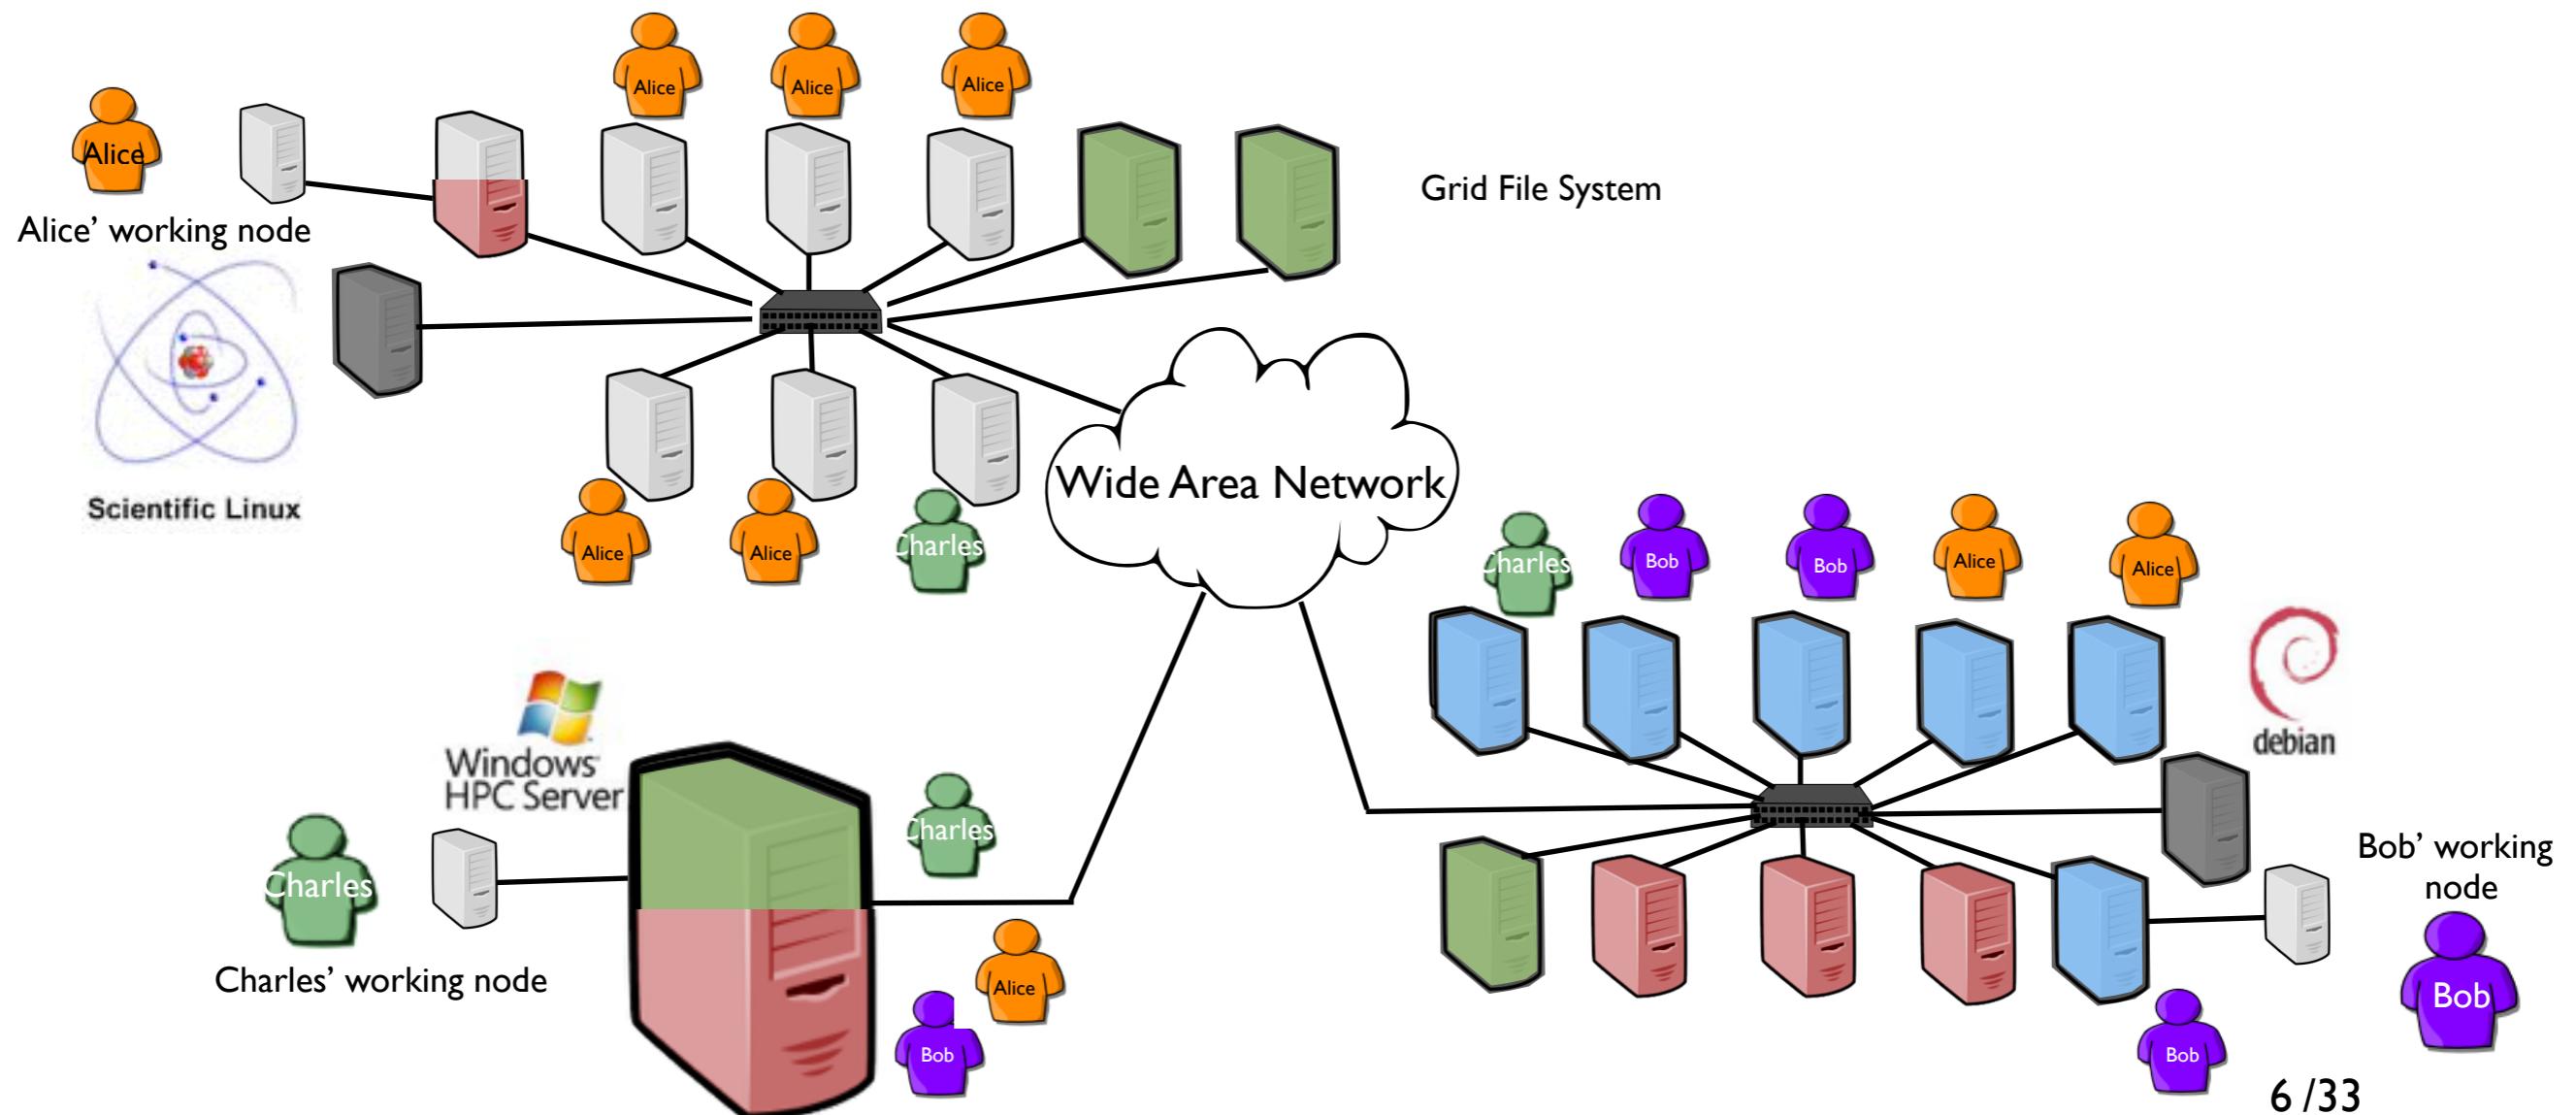
# Looking back...

- Network of Workstations 1990 / 20~~xx~~
- Desktop 1998 / 20~~xx~~
- Grid 1998 / 20~~xx~~



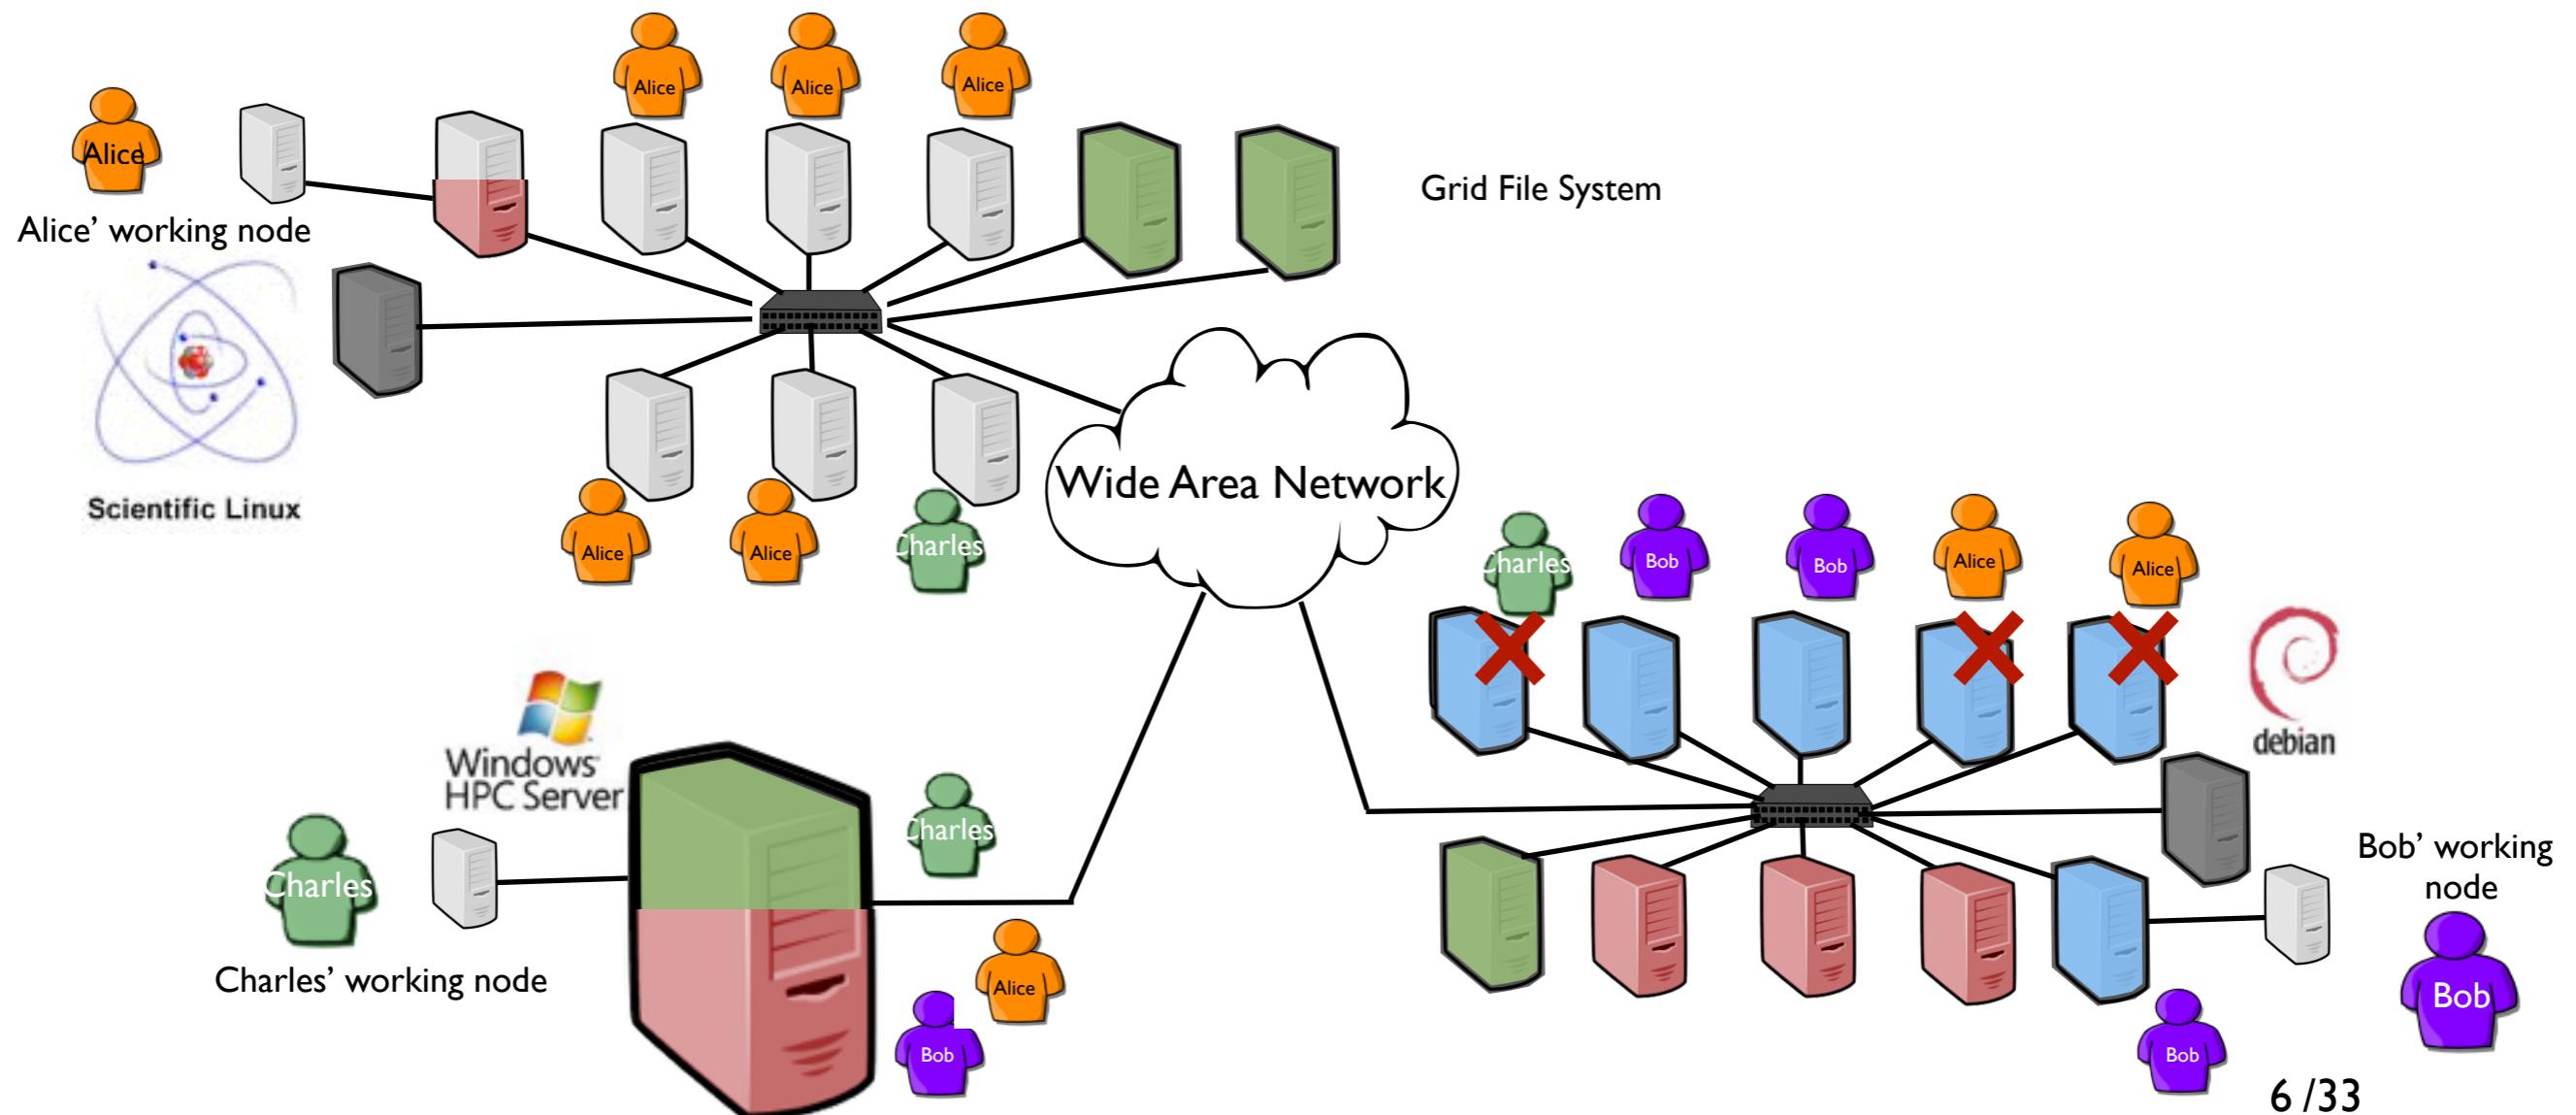
# Looking back...

- Network of Workstations 1990 / 20~~xx~~
- Desktop 1998 / 20~~xx~~
- Grid 1998 / 20~~xx~~



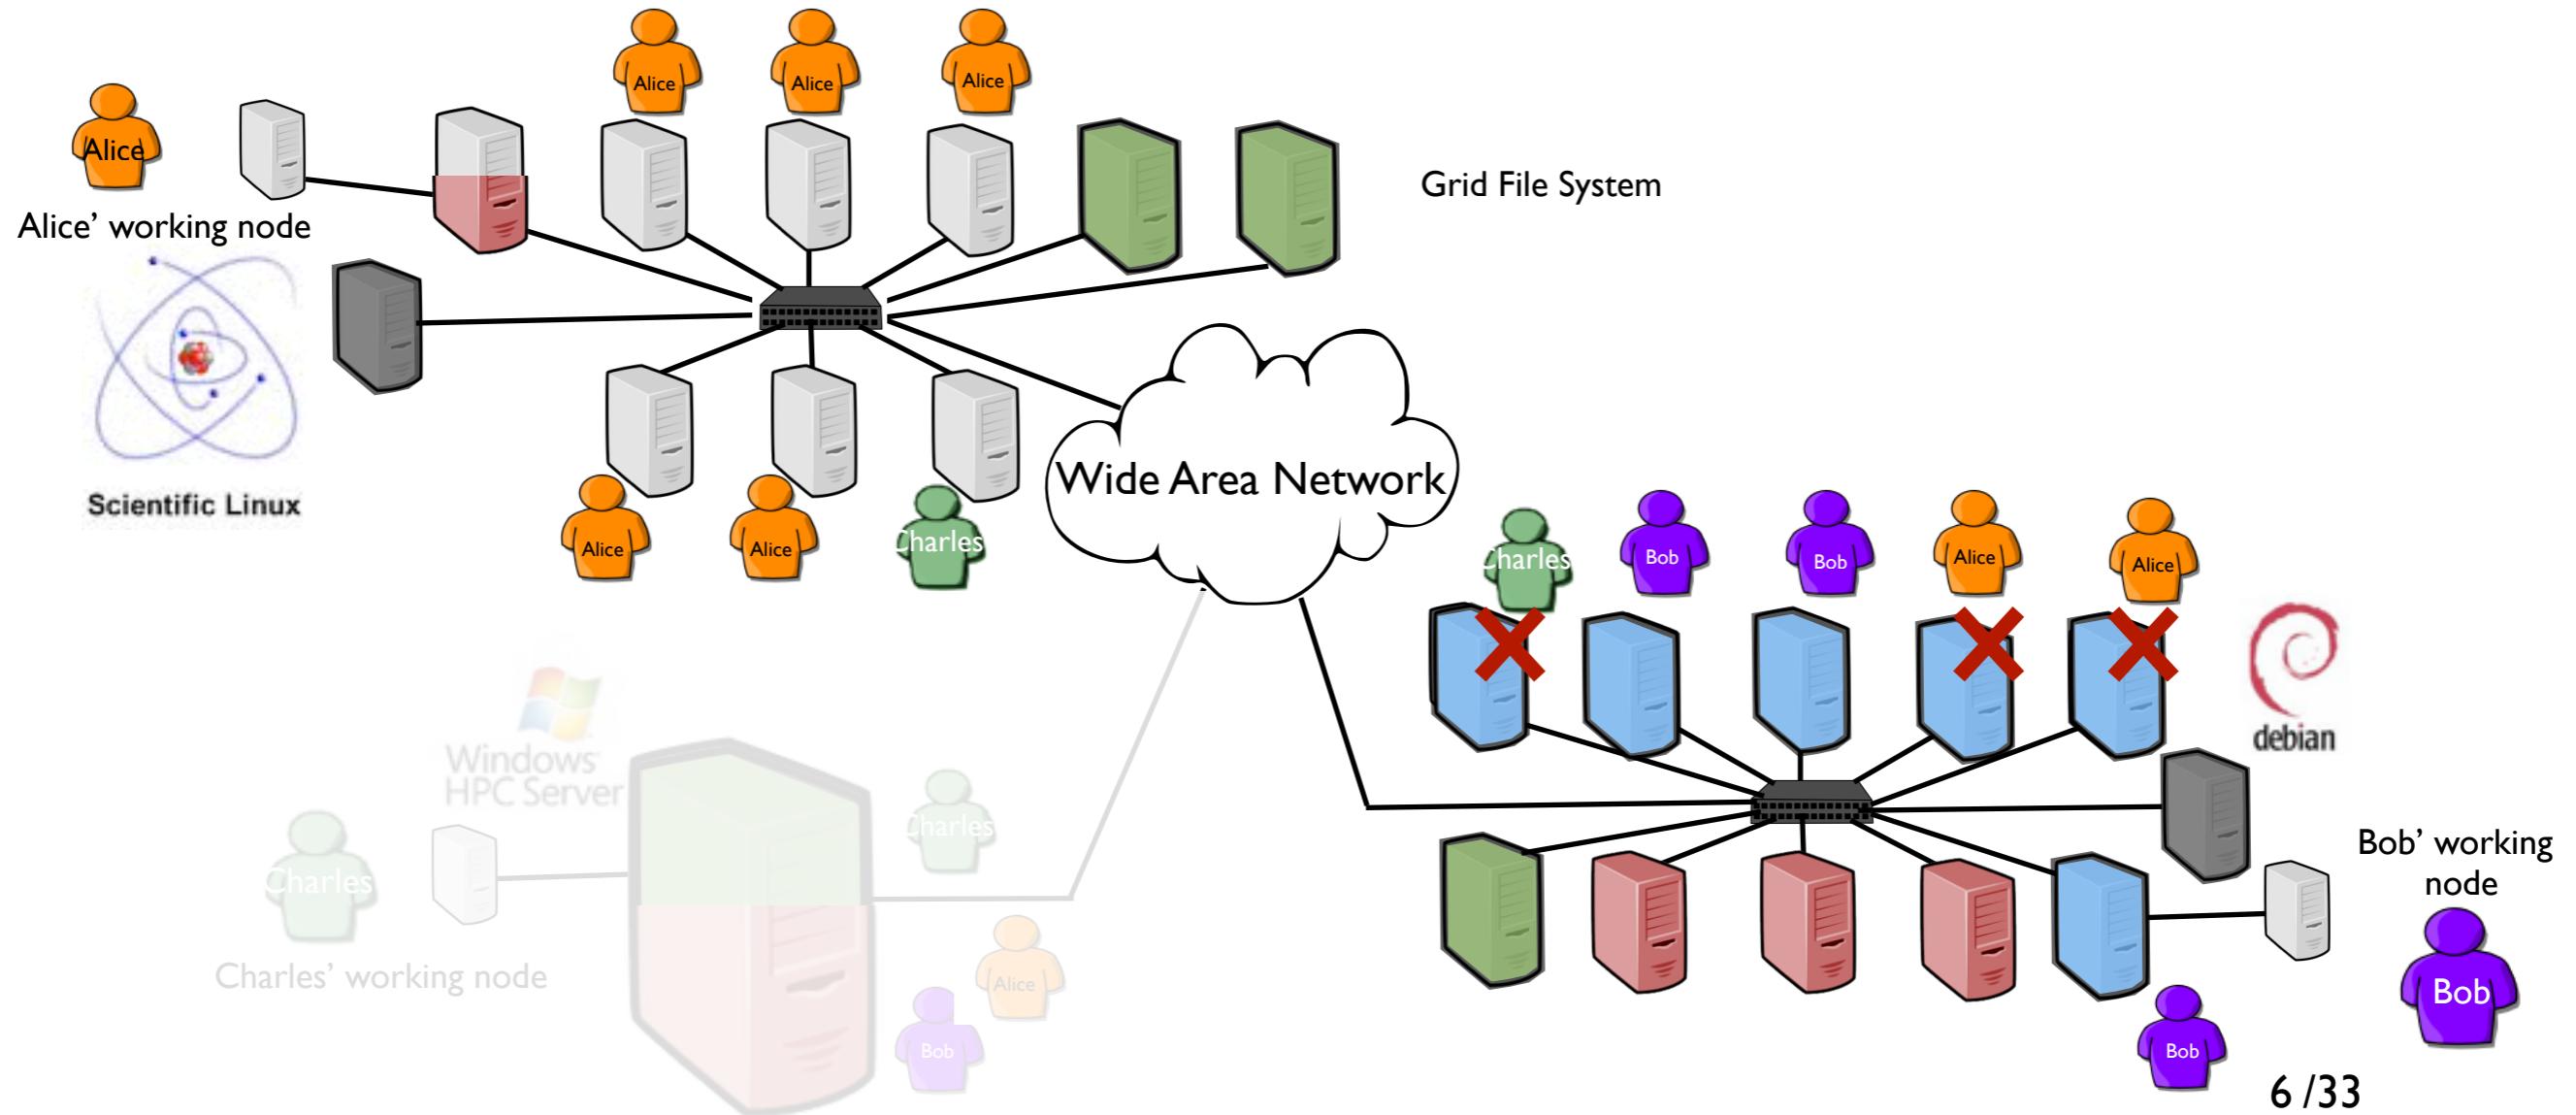
# Looking back...

- Network of Workstations 1990 / 20~~xx~~
- Desktop 1998 / 20~~xx~~
- Grid 1998 / 20~~xx~~



# Looking back...

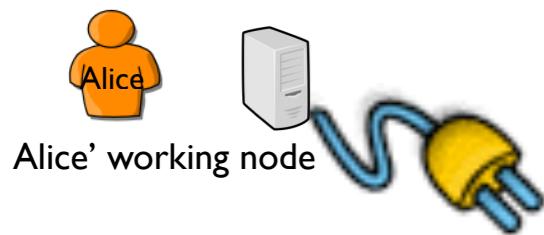
- Network of Workstations 1990 / 20~~xx~~
- Desktop 1998 / 20~~xx~~
- Grid 1998 / 20~~xx~~



# Looking back...

- Network of Workstations 1990 / 20~~xx~~
- Desktop 1998 / 20~~xx~~
- Grid 1998 / 20~~xx~~

## What a Grid ! ? !



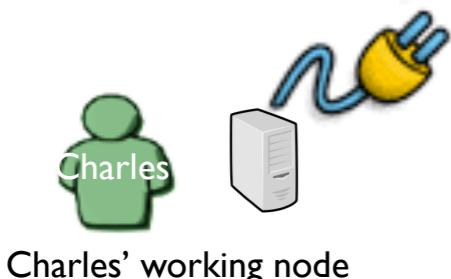
Resource booking (based on user's estimates)

Security concerns (job isolation)

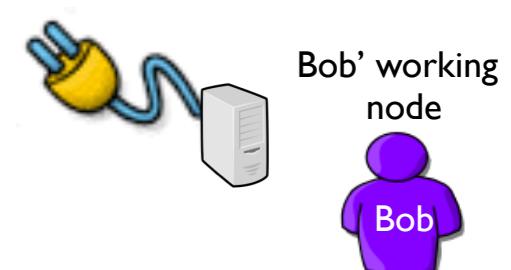
Heterogeneity concerns (hardware and software)

Scheduling limitations (a job cannot be easily relocated)

Fault tolerance issues



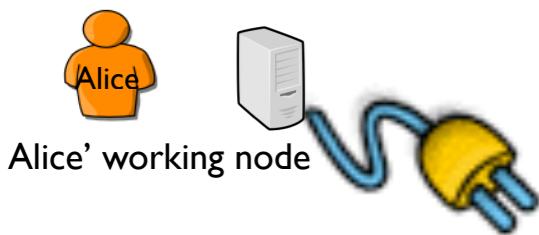
...



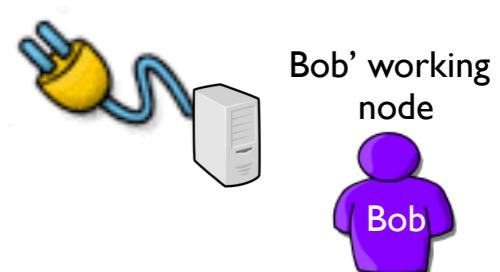
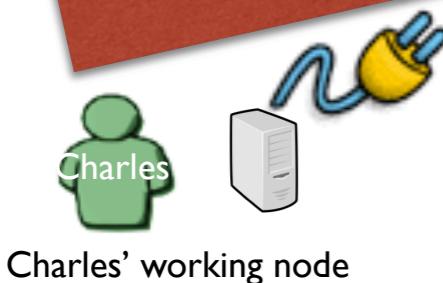
# Looking back...

- Network of Workstations 1990 / 20~~xx~~
- Desktop 1998 / 20~~xx~~
- Grid 1998 / 20~~xx~~

What a Grid ! ? !

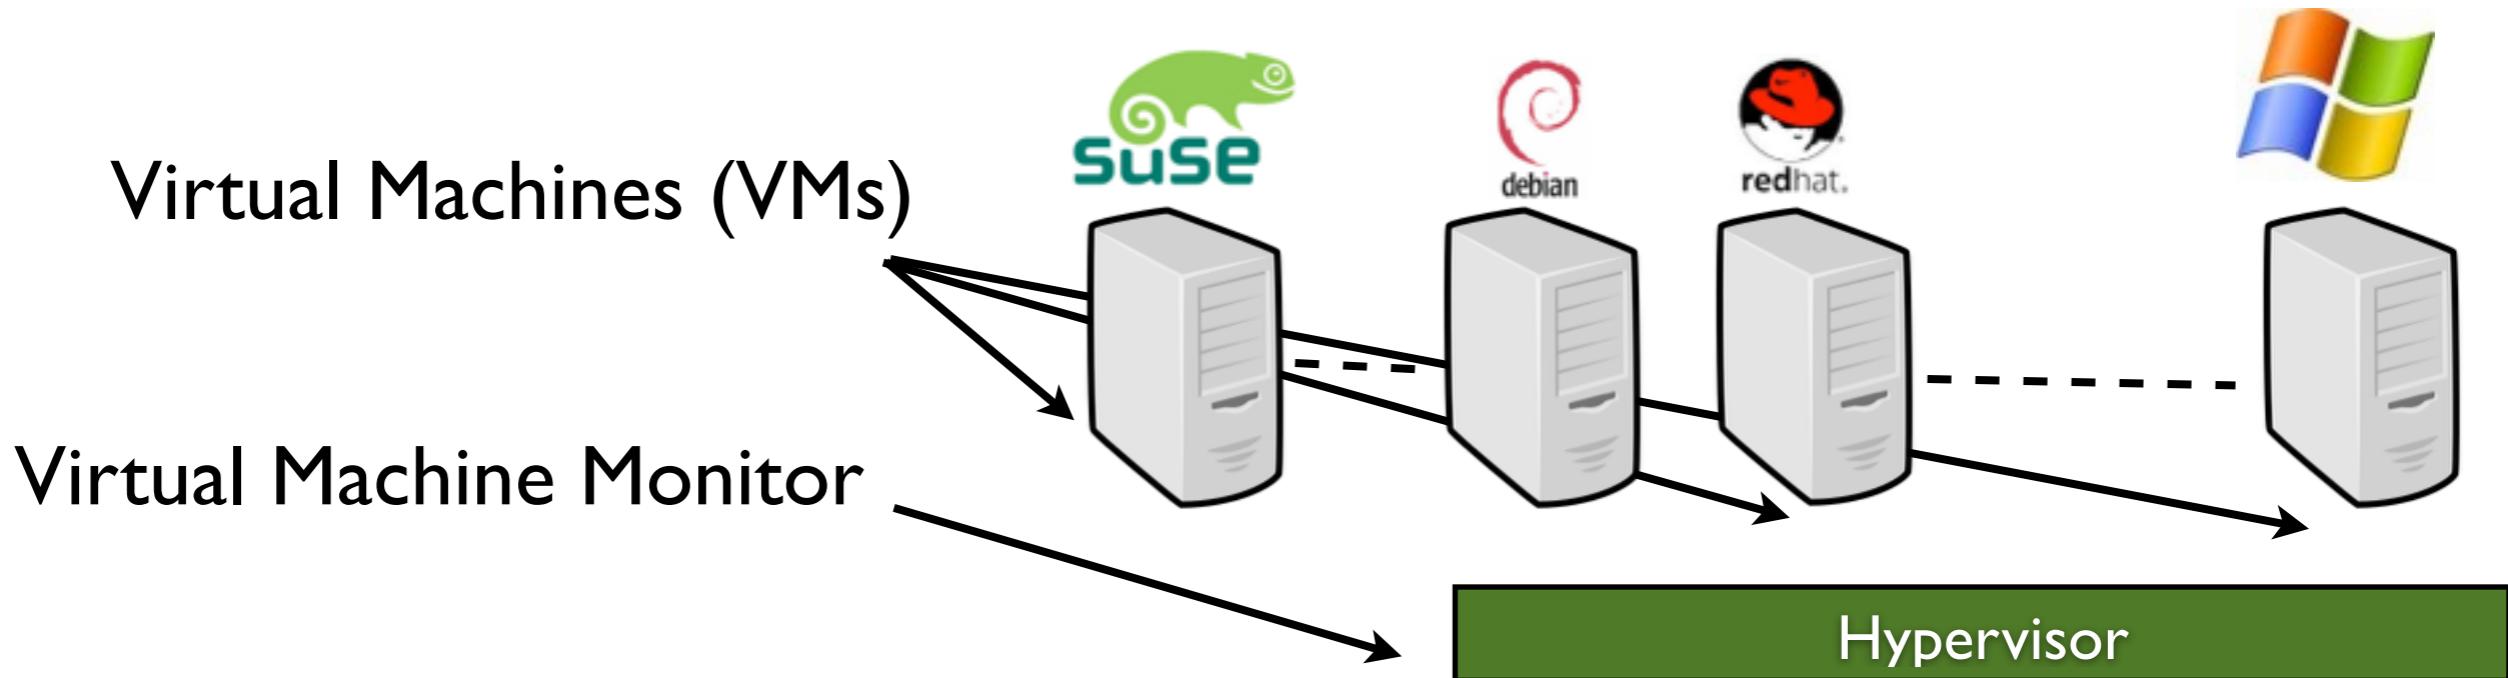


A lot of progress has been done since the 90's and several proposals partially addressed these concerns.



# Looking back...

- System virtualization: One to multiple OSes on a physical node thanks to a hypervisor (an operating system of OSes)



“A **virtual machine** (VM) provides a faithful implementation of a physical processor’s hardware running in a protected and isolated environment.

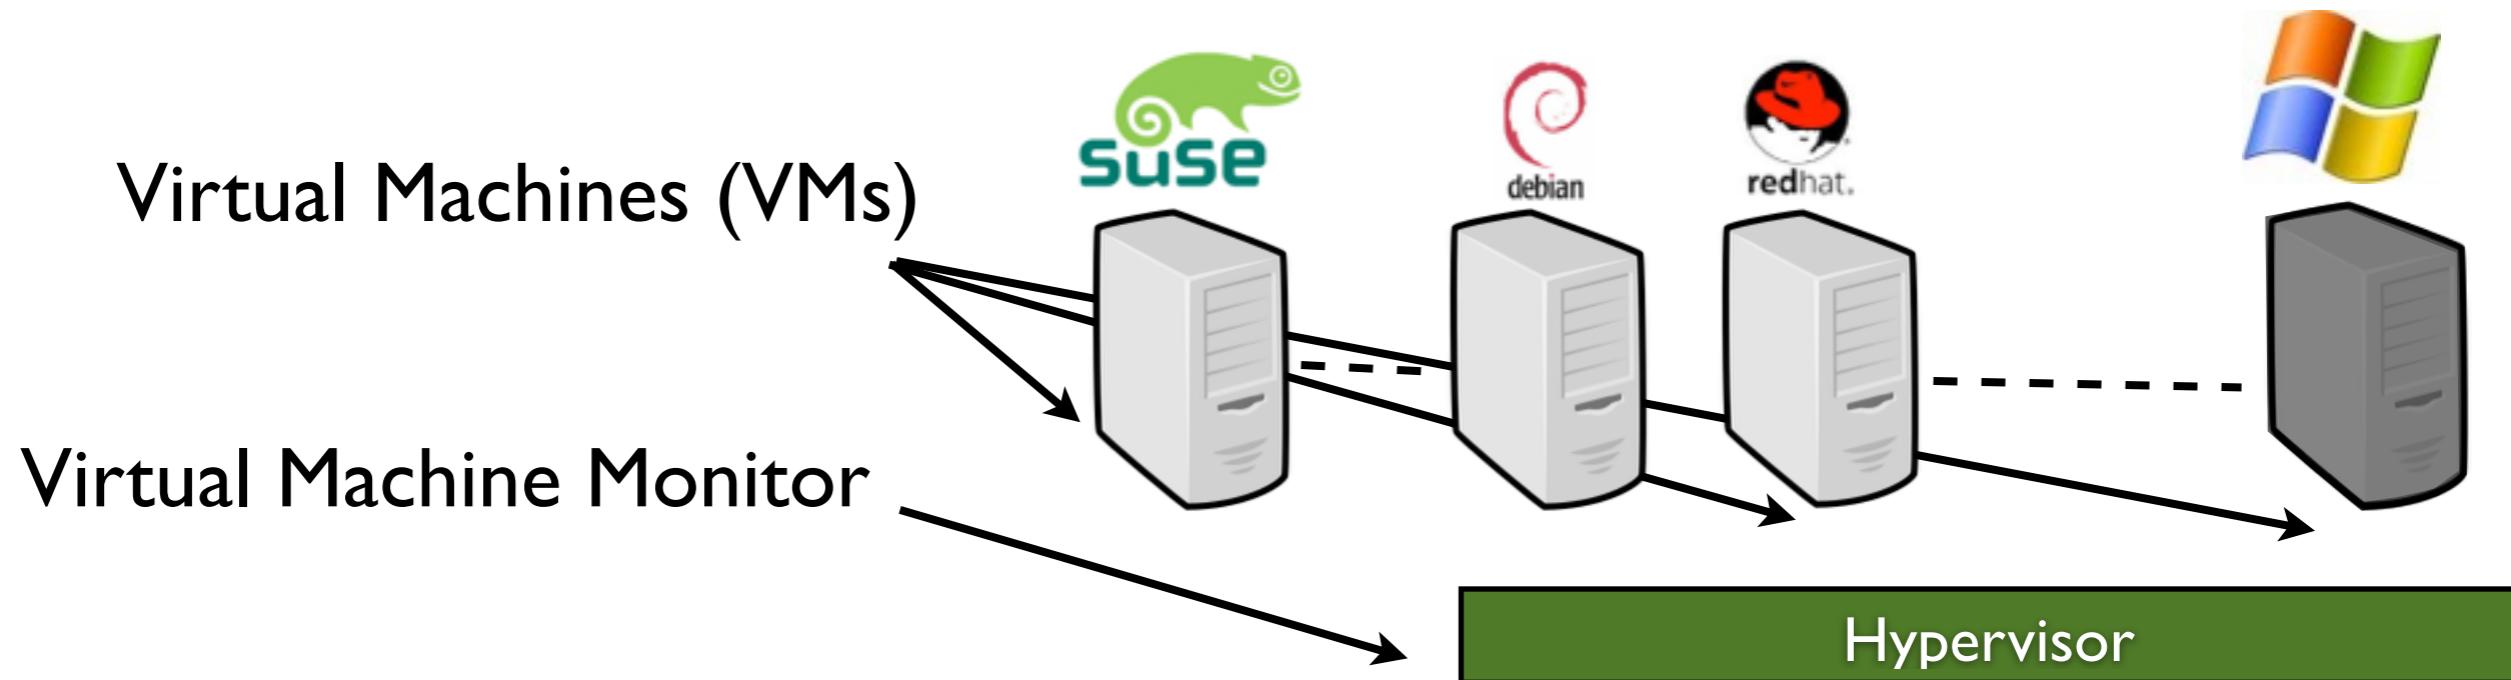
Virtual machines are created by a software layer called the **virtual machine monitor** (VMM) that runs as a privileged task on a physical processor.”



Physical Machine (PM)

# Looking back...

- System virtualization: One to multiple OSes on a physical node thanks to a hypervisor (an operating system of OSes)



“A **virtual machine** (VM) provides a faithful implementation of a physical processor’s hardware running in a protected and isolated environment.

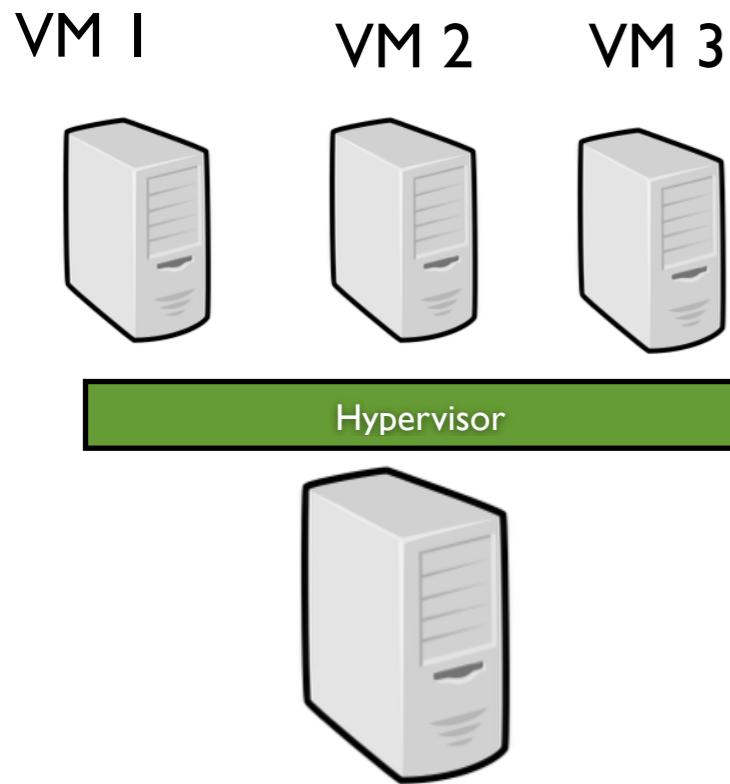
Virtual machines are created by a software layer called the **virtual machine monitor** (VMM) that runs as a privileged task on a physical processor.”



Physical Machine (PM)

# Looking back...

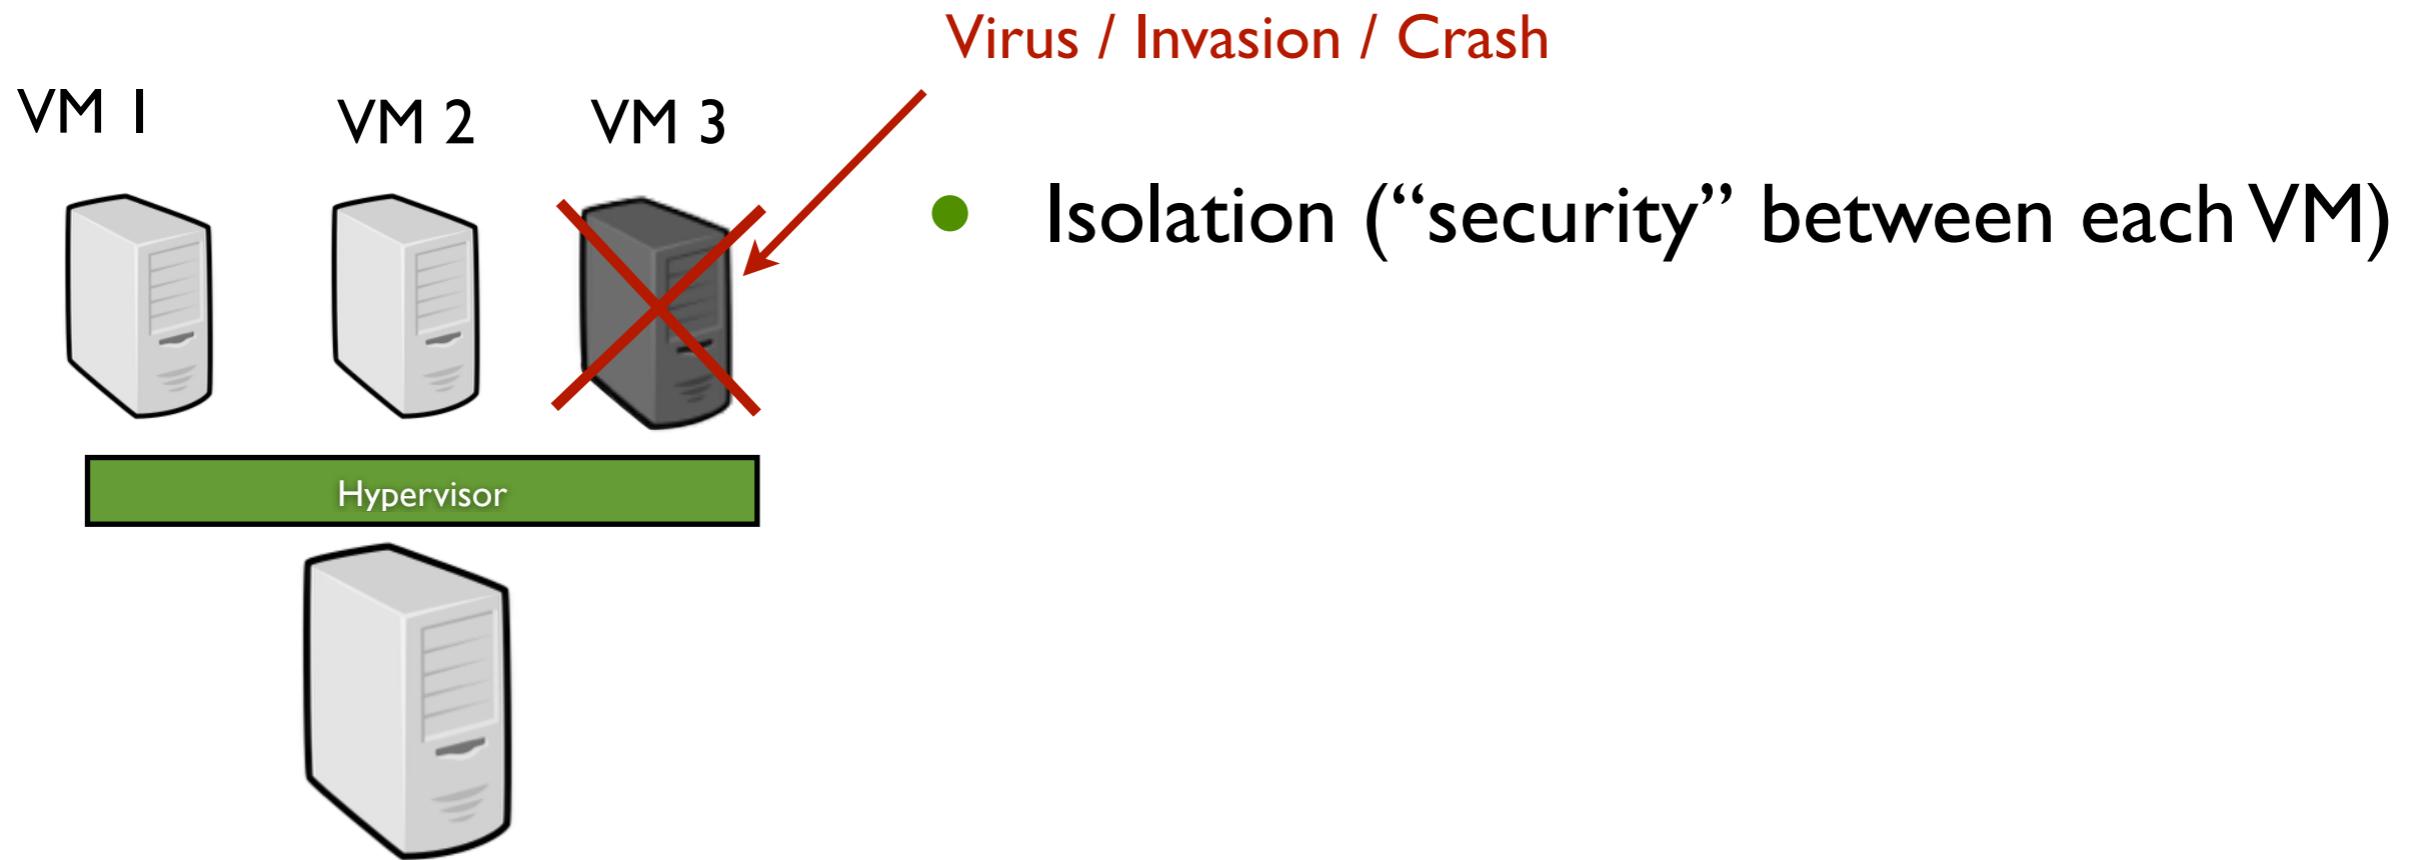
- System virtualization: a great sandbox



- Isolation (“security” between each VM)

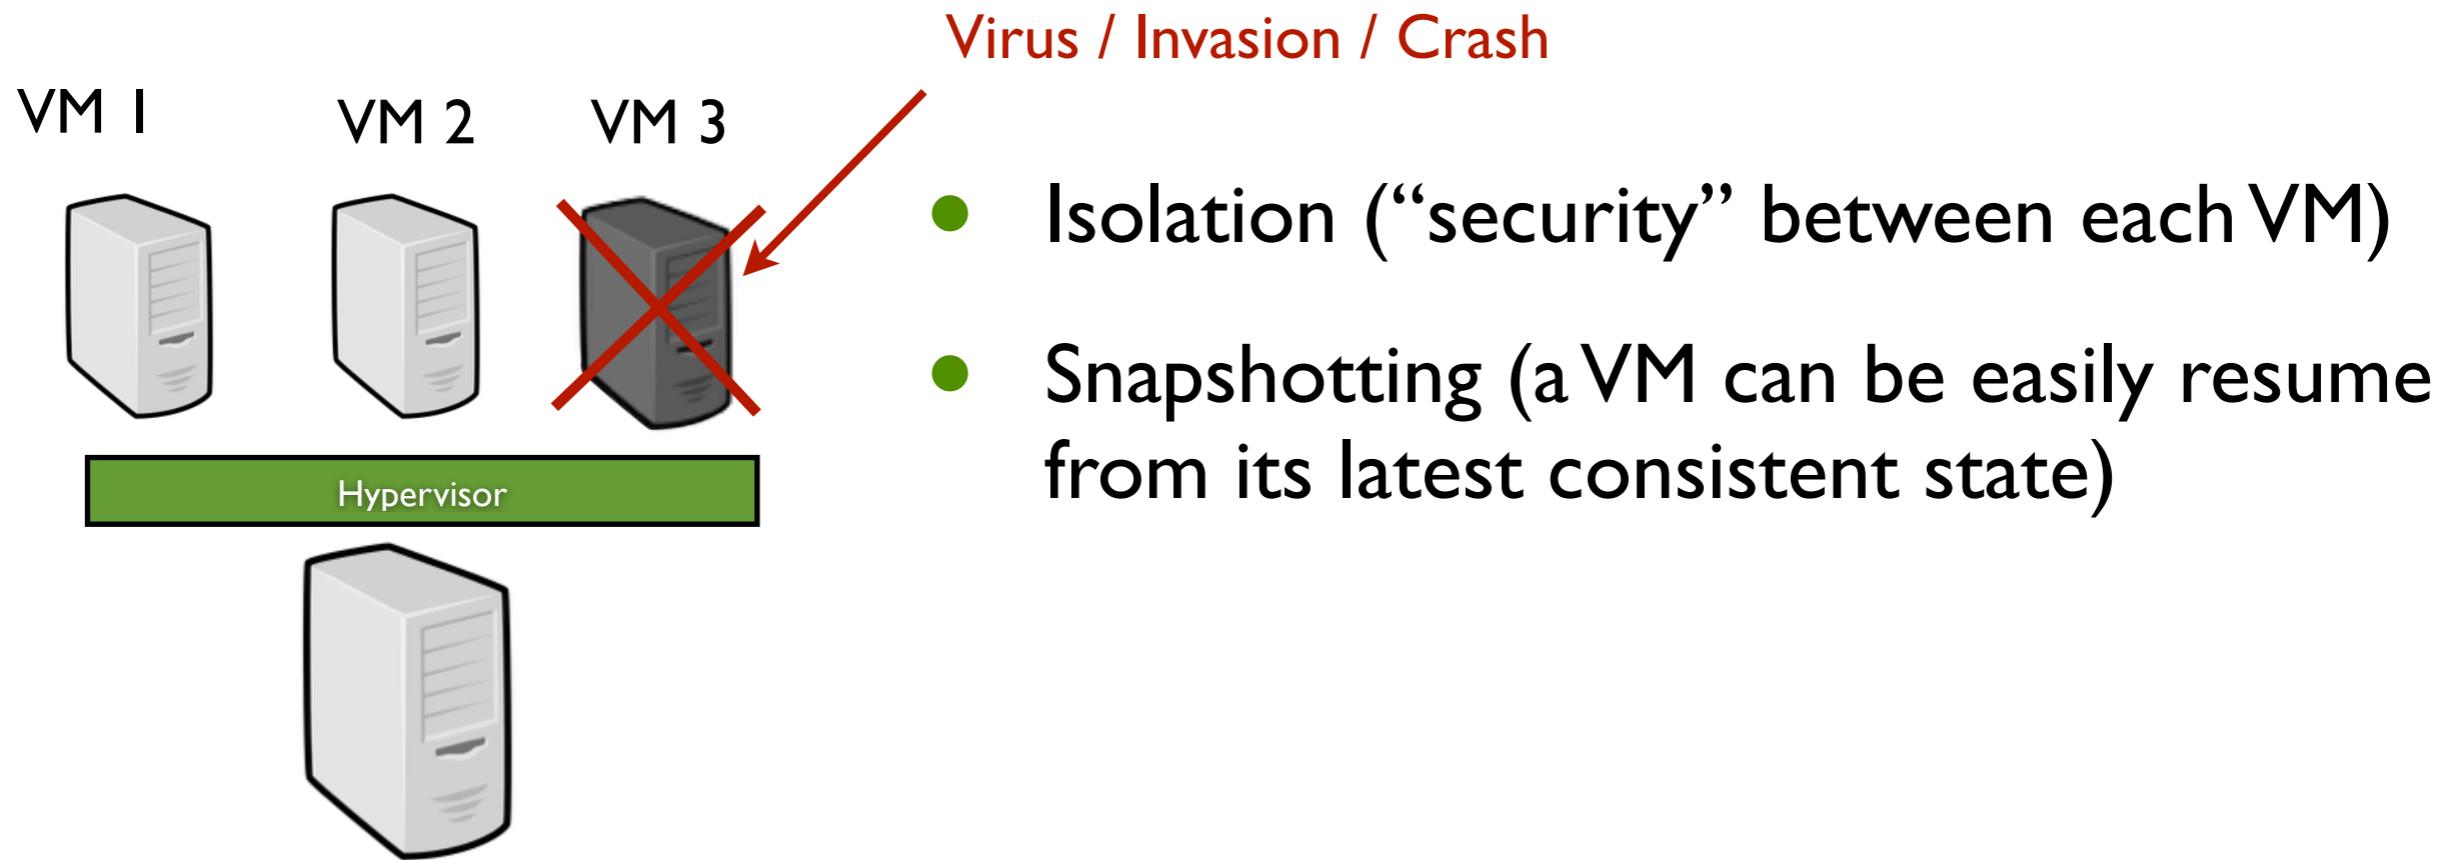
# Looking back...

- System virtualization: a great sandbox



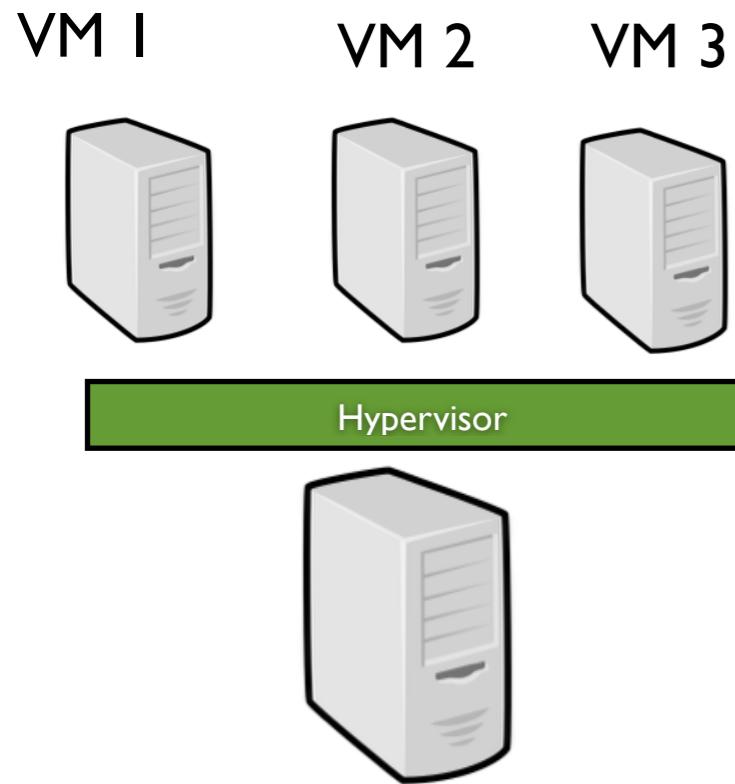
# Looking back...

- System virtualization: a great sandbox



# Looking back...

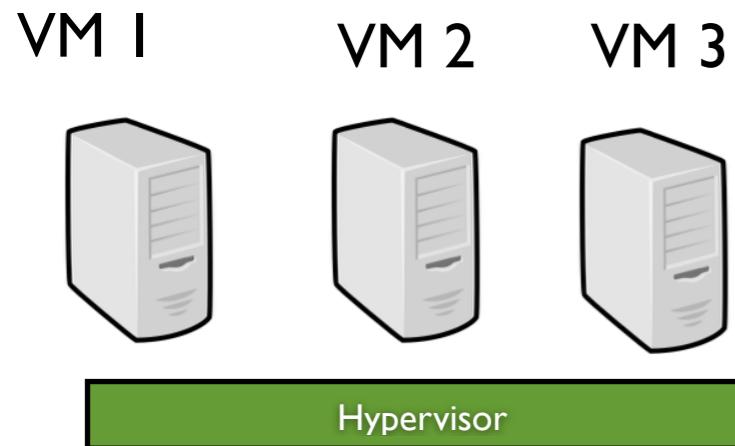
- System virtualization: a great sandbox



- Isolation (“security” between each VM)
- Snapshotting (a VM can be easily resume from its latest consistent state)

# Looking back...

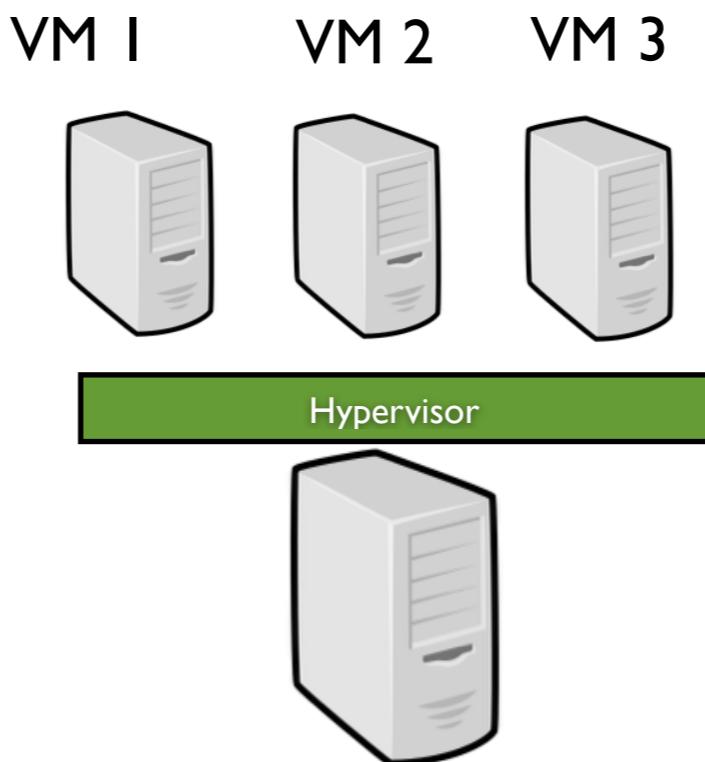
- System virtualization: a great sandbox



- Isolation (“security” between each VM)
- Snapshotting (a VM can be easily resume from its latest consistent state)

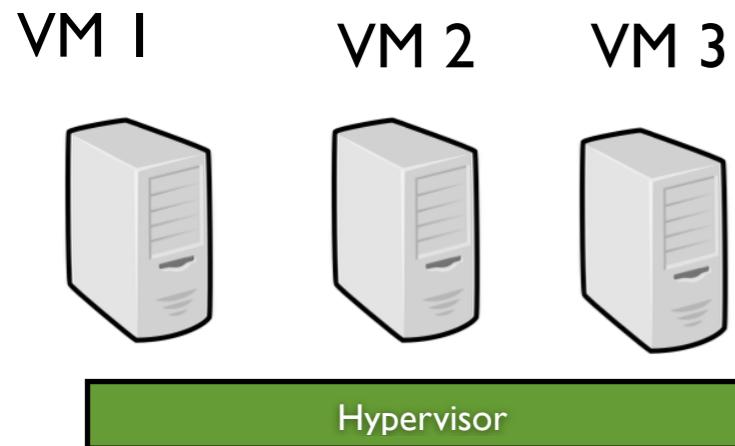


- Suspend/Resume



# Looking back...

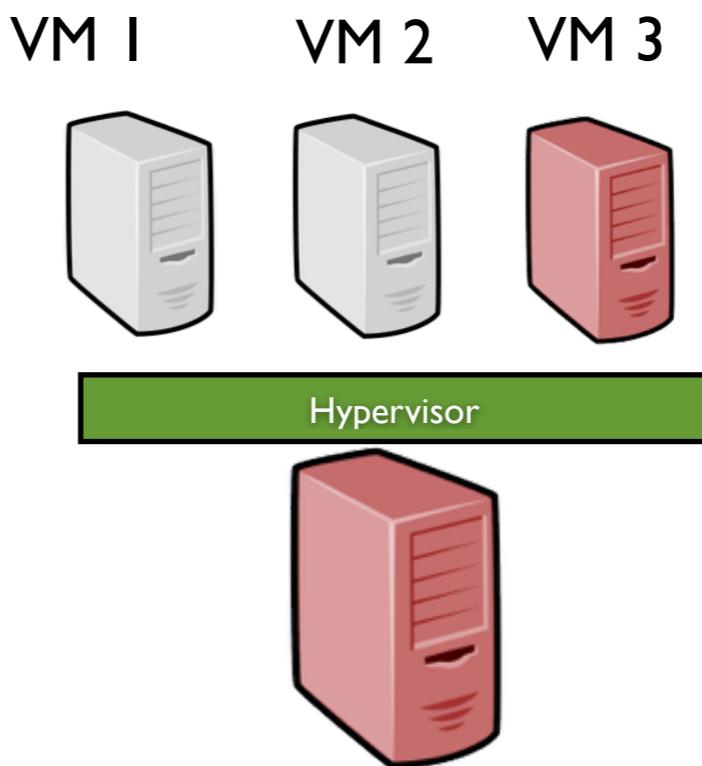
- System virtualization: a great sandbox



- Isolation (“security” between each VM)
- Snapshotting (a VM can be easily resume from its latest consistent state)

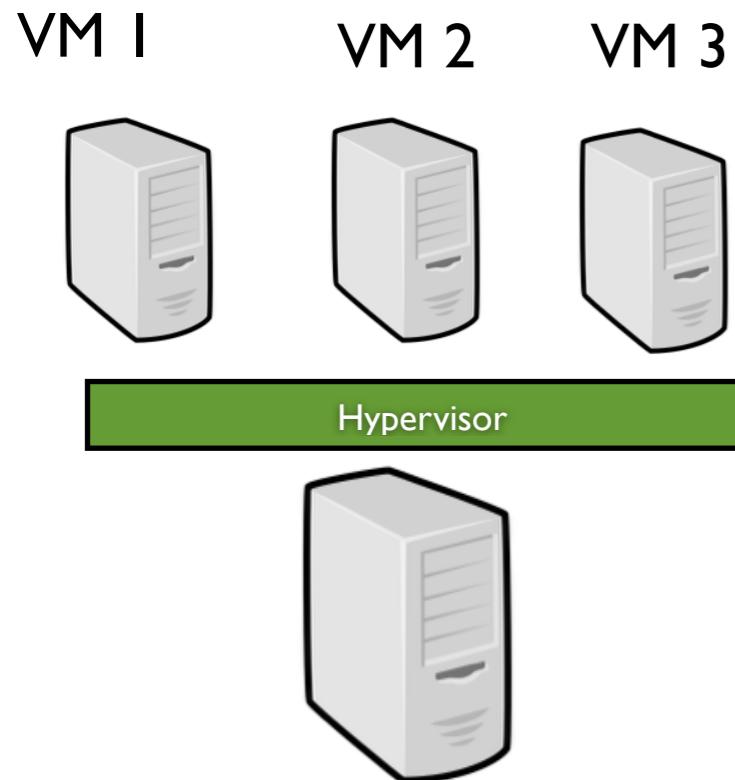


- Suspend/Resume



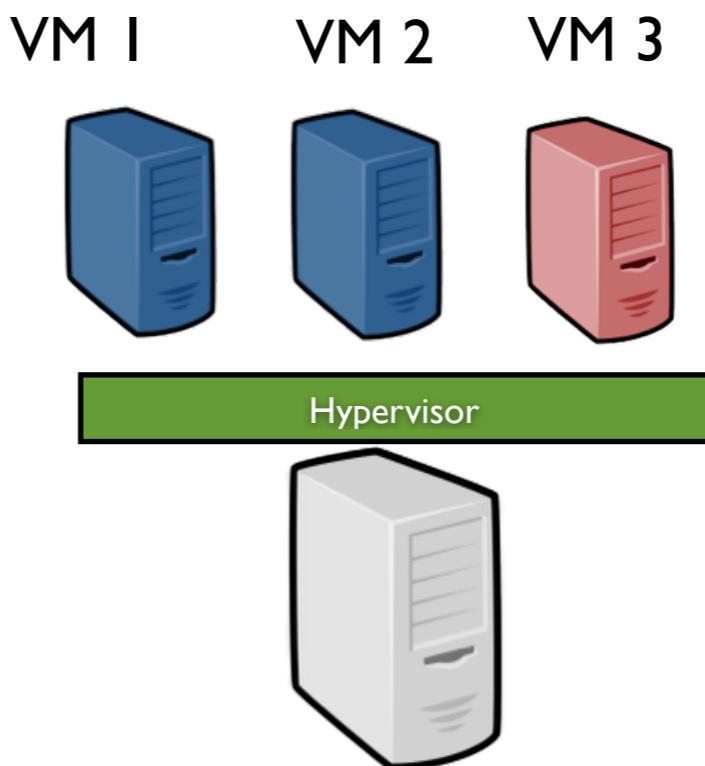
# Looking back...

- System virtualization: a great sandbox



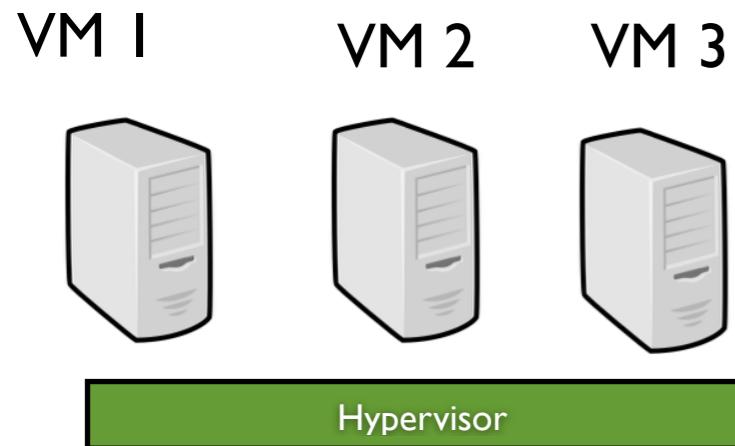
- Isolation (“security” between each VM)
- Snapshotting (a VM can be easily resume from its latest consistent state)

- Suspend/Resume

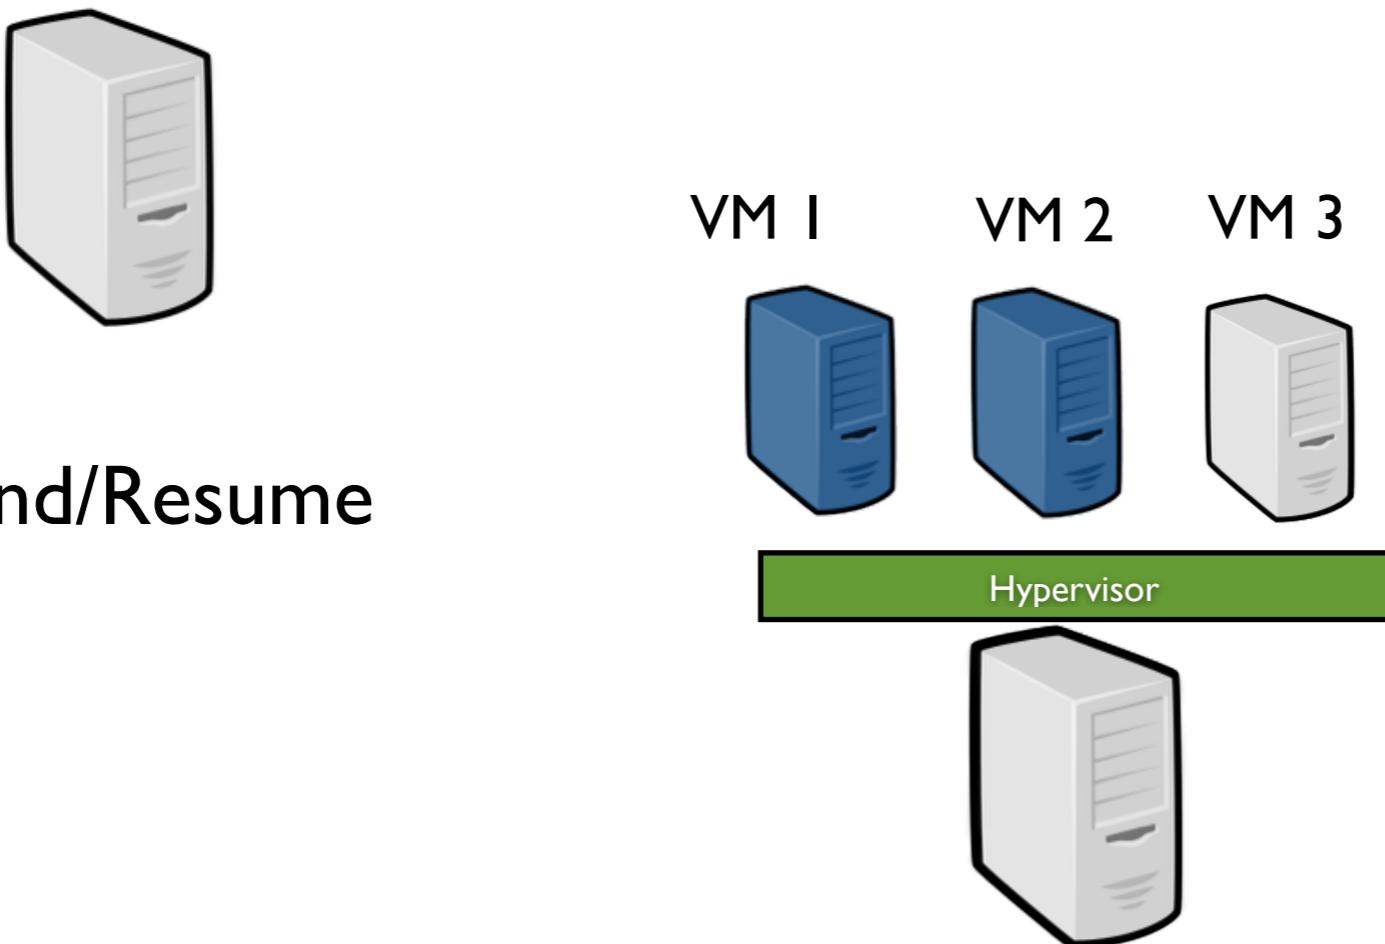


# Looking back...

- System virtualization: a great sandbox



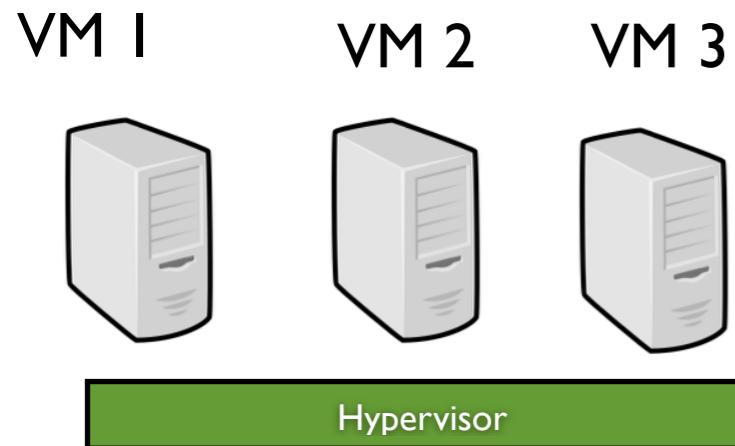
- Isolation (“security” between each VM)
- Snapshotting (a VM can be easily resume from its latest consistent state)



- Suspend/Resume

# Looking back...

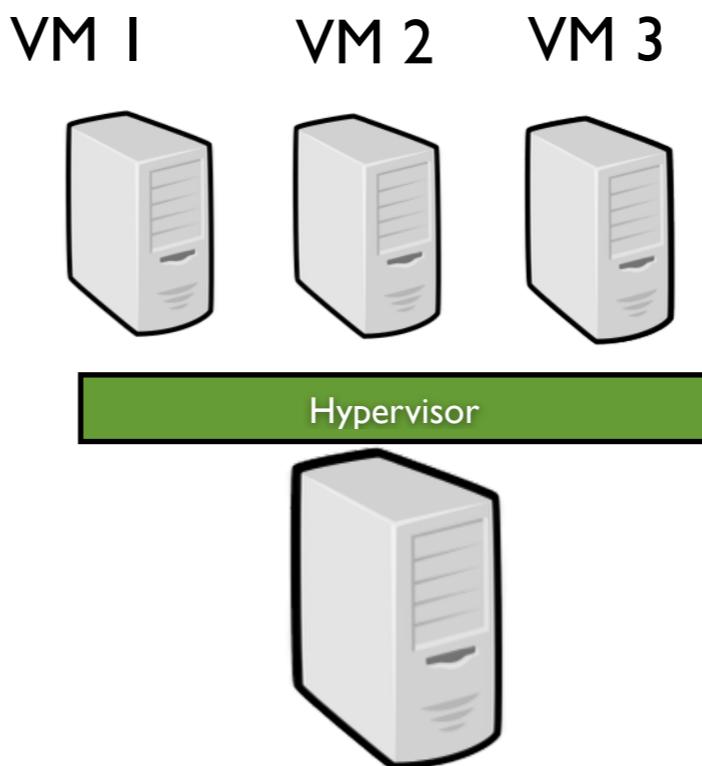
- System virtualization: a great sandbox



- Isolation (“security” between each VM)
- Snapshotting (a VM can be easily resume from its latest consistent state)

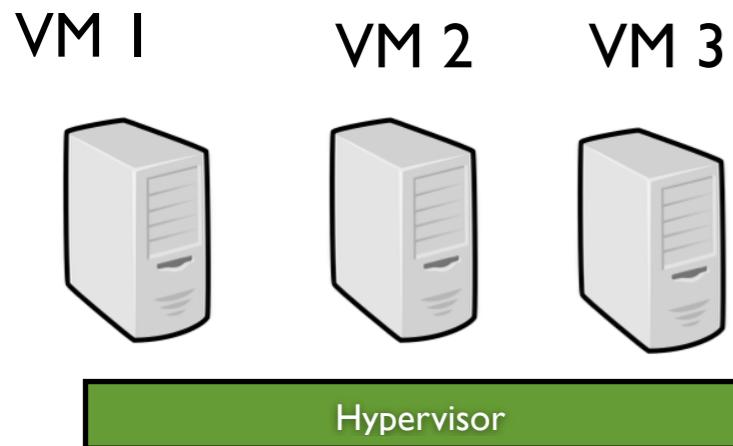


- Suspend/Resume



# Looking back...

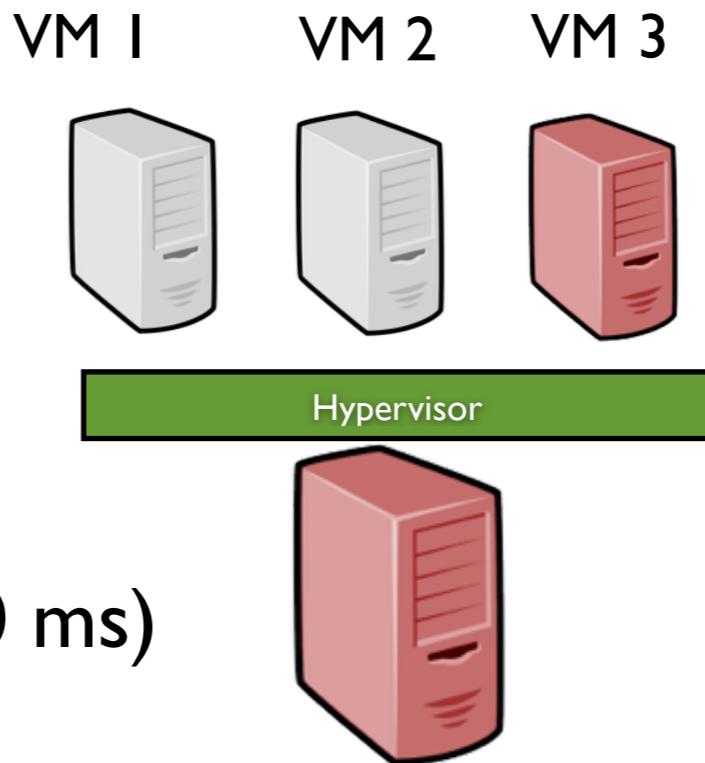
- System virtualization: a great sandbox



- Isolation (“security” between each VM)
- Snapshotting (a VM can be easily resume from its latest consistent state)

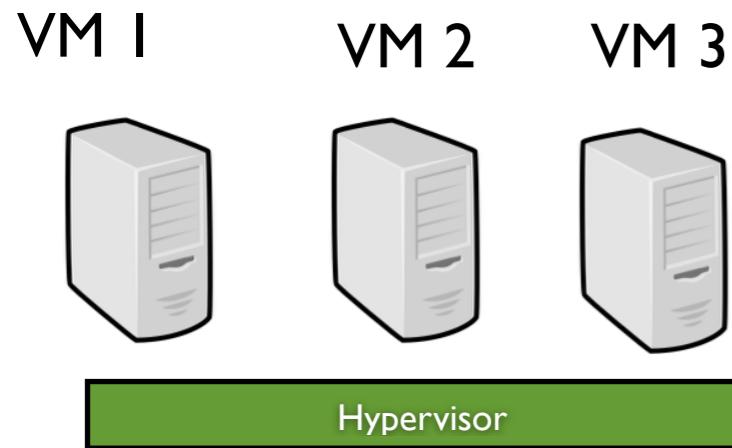


- Suspend/Resume
- Live migration  
(negligible downtime ~ 60 ms)  
Post/Pre Copy



# Looking back...

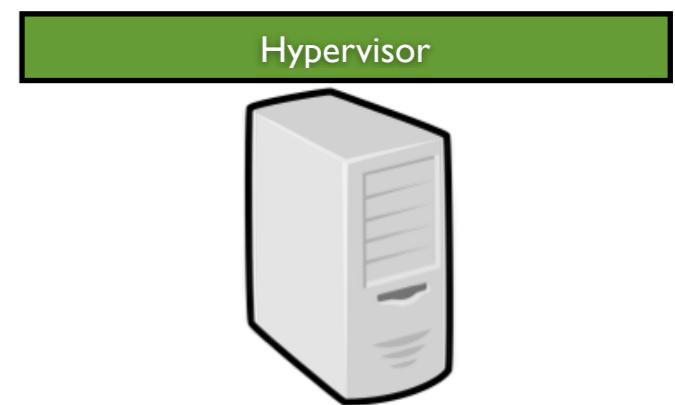
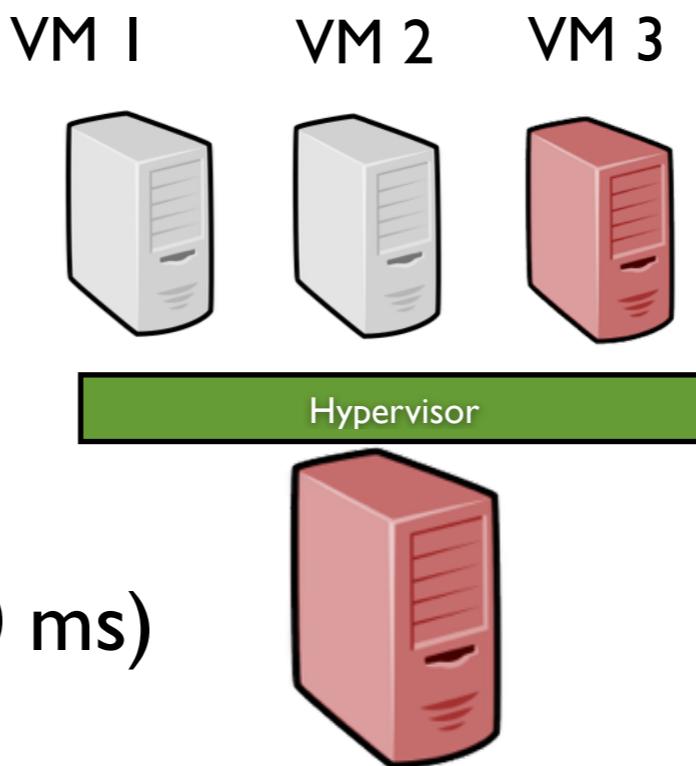
- System virtualization: a great sandbox



- Isolation (“security” between each VM)
- Snapshotting (a VM can be easily resume from its latest consistent state)

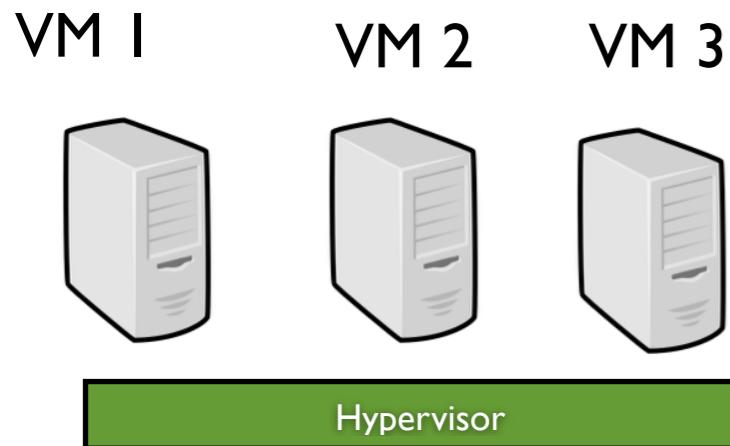


- Suspend/Resume
- Live migration  
(negligible downtime ~ 60 ms)  
Post/Pre Copy



# Looking back...

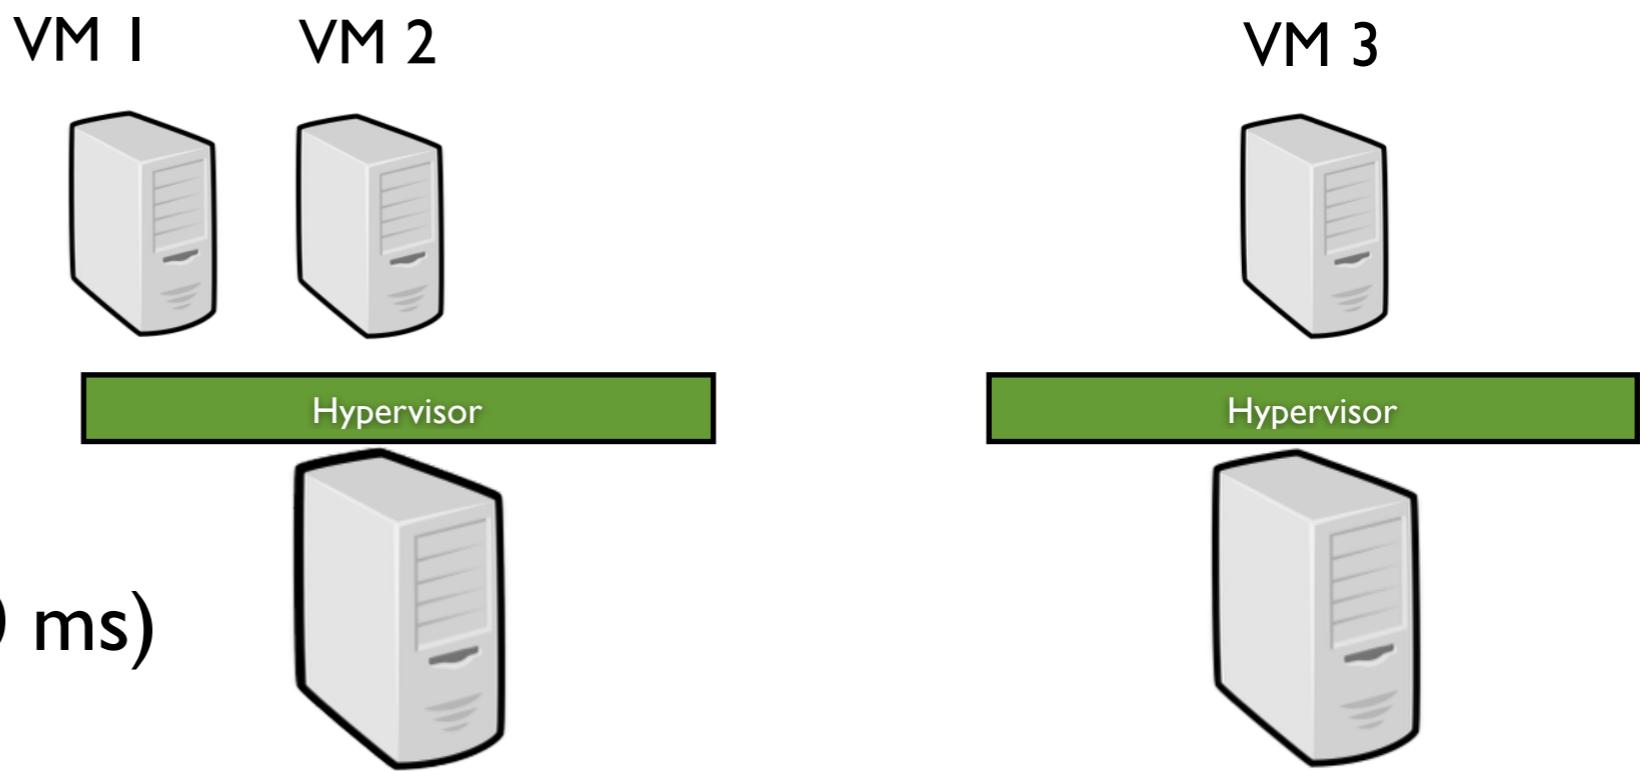
- System virtualization: a great sandbox



- Isolation (“security” between each VM)
- Snapshotting (a VM can be easily resume from its latest consistent state)

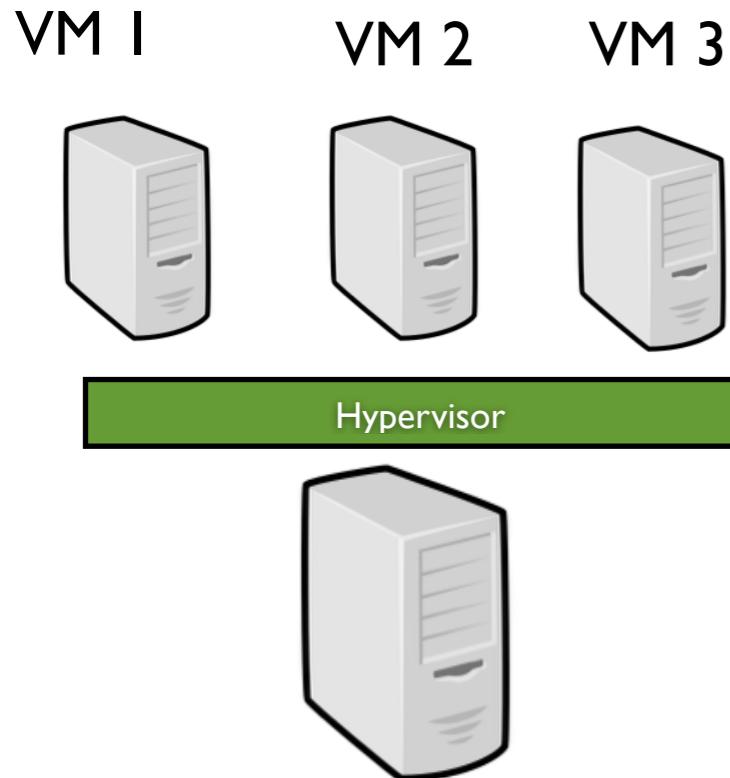


- Suspend/Resume
- Live migration  
(negligible downtime ~ 60 ms)  
Post/Pre Copy



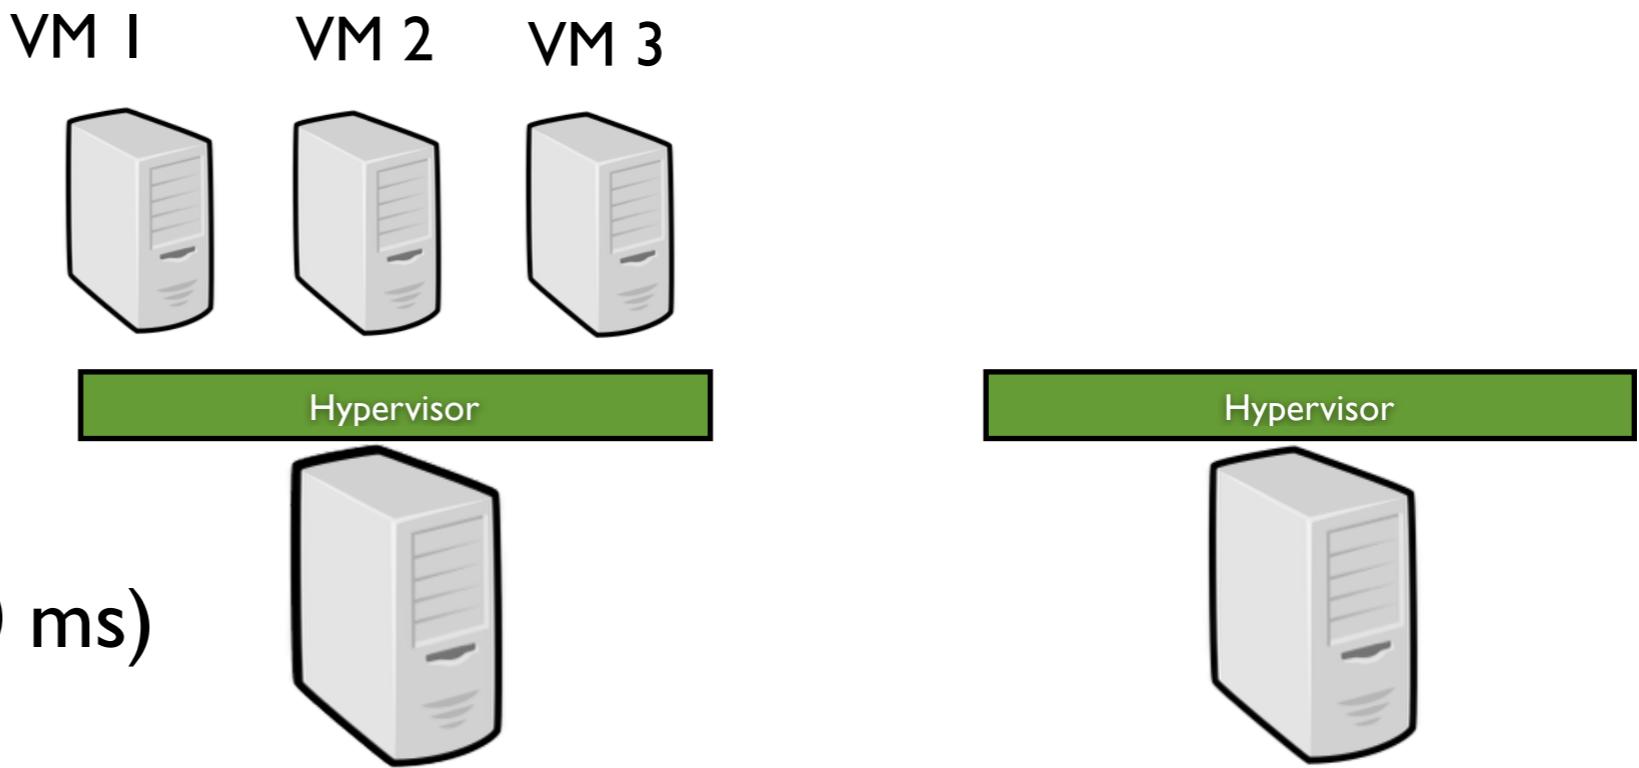
# Looking back...

- System virtualization: a great sandbox



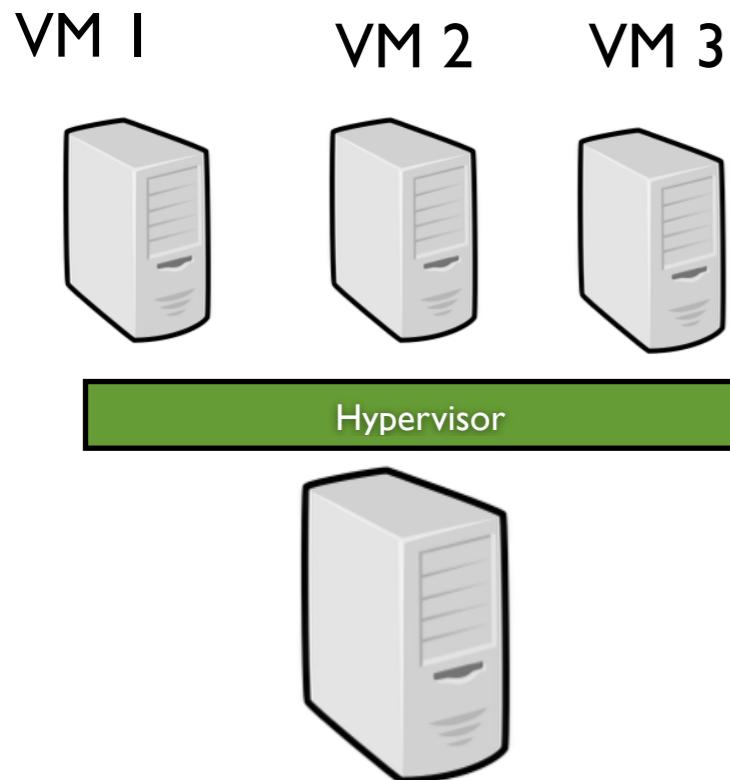
- Isolation (“security” between each VM)
- Snapshotting (a VM can be easily resume from its latest consistent state)

- Suspend/Resume
- Live migration  
(negligible downtime ~ 60 ms)  
Post/Pre Copy

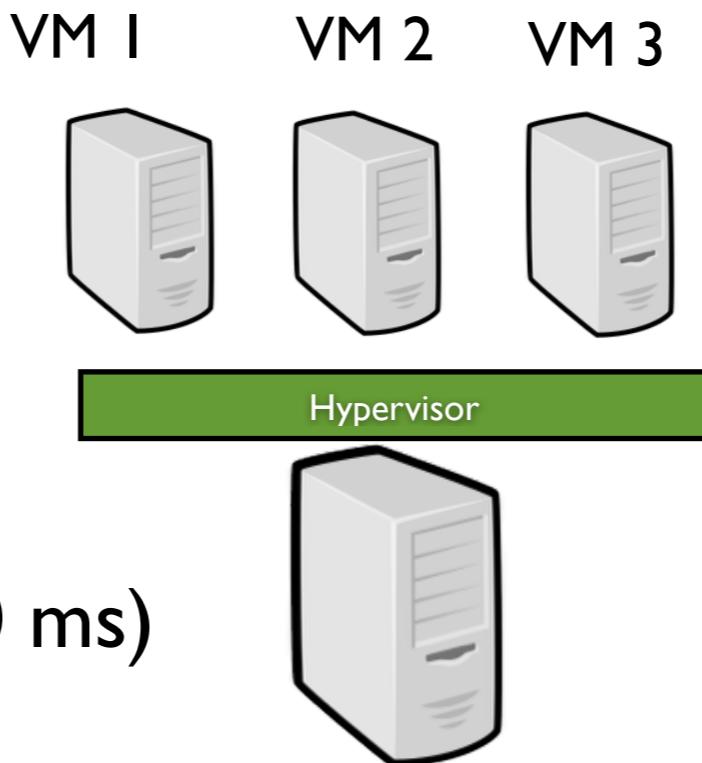


# Looking back...

- System virtualization: a great sandbox



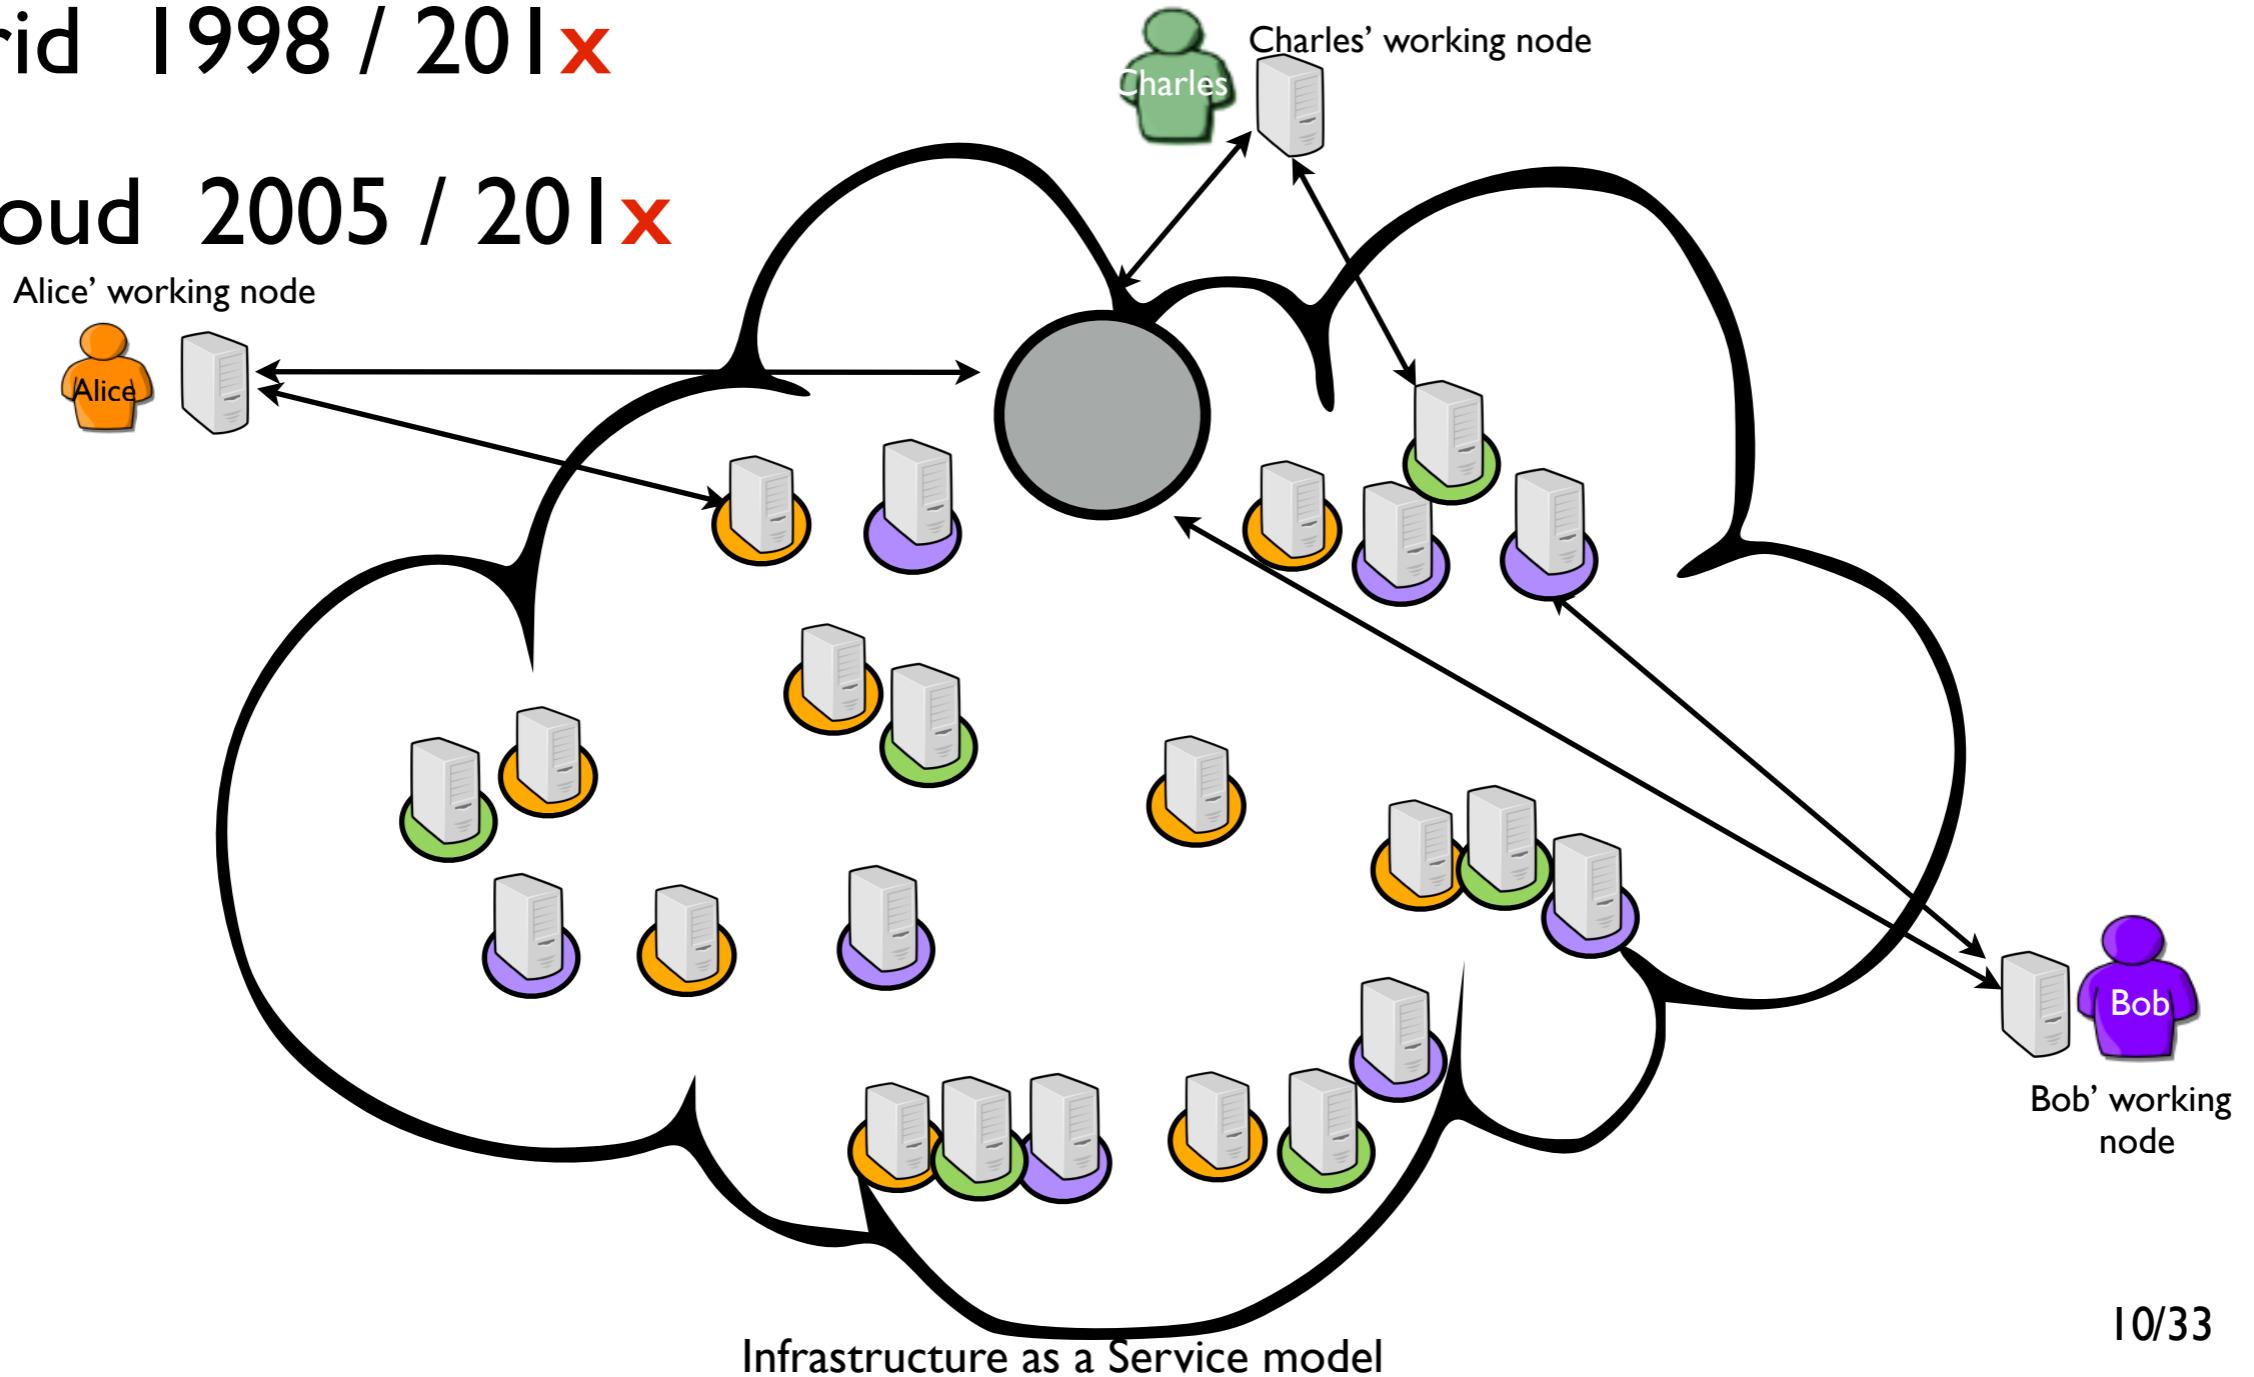
- Isolation (“security” between each VM)
- Snapshotting (a VM can be easily resume from its latest consistent state)



- Suspend/Resume
- Live migration  
(negligible downtime ~ 60 ms)  
Post/Pre Copy

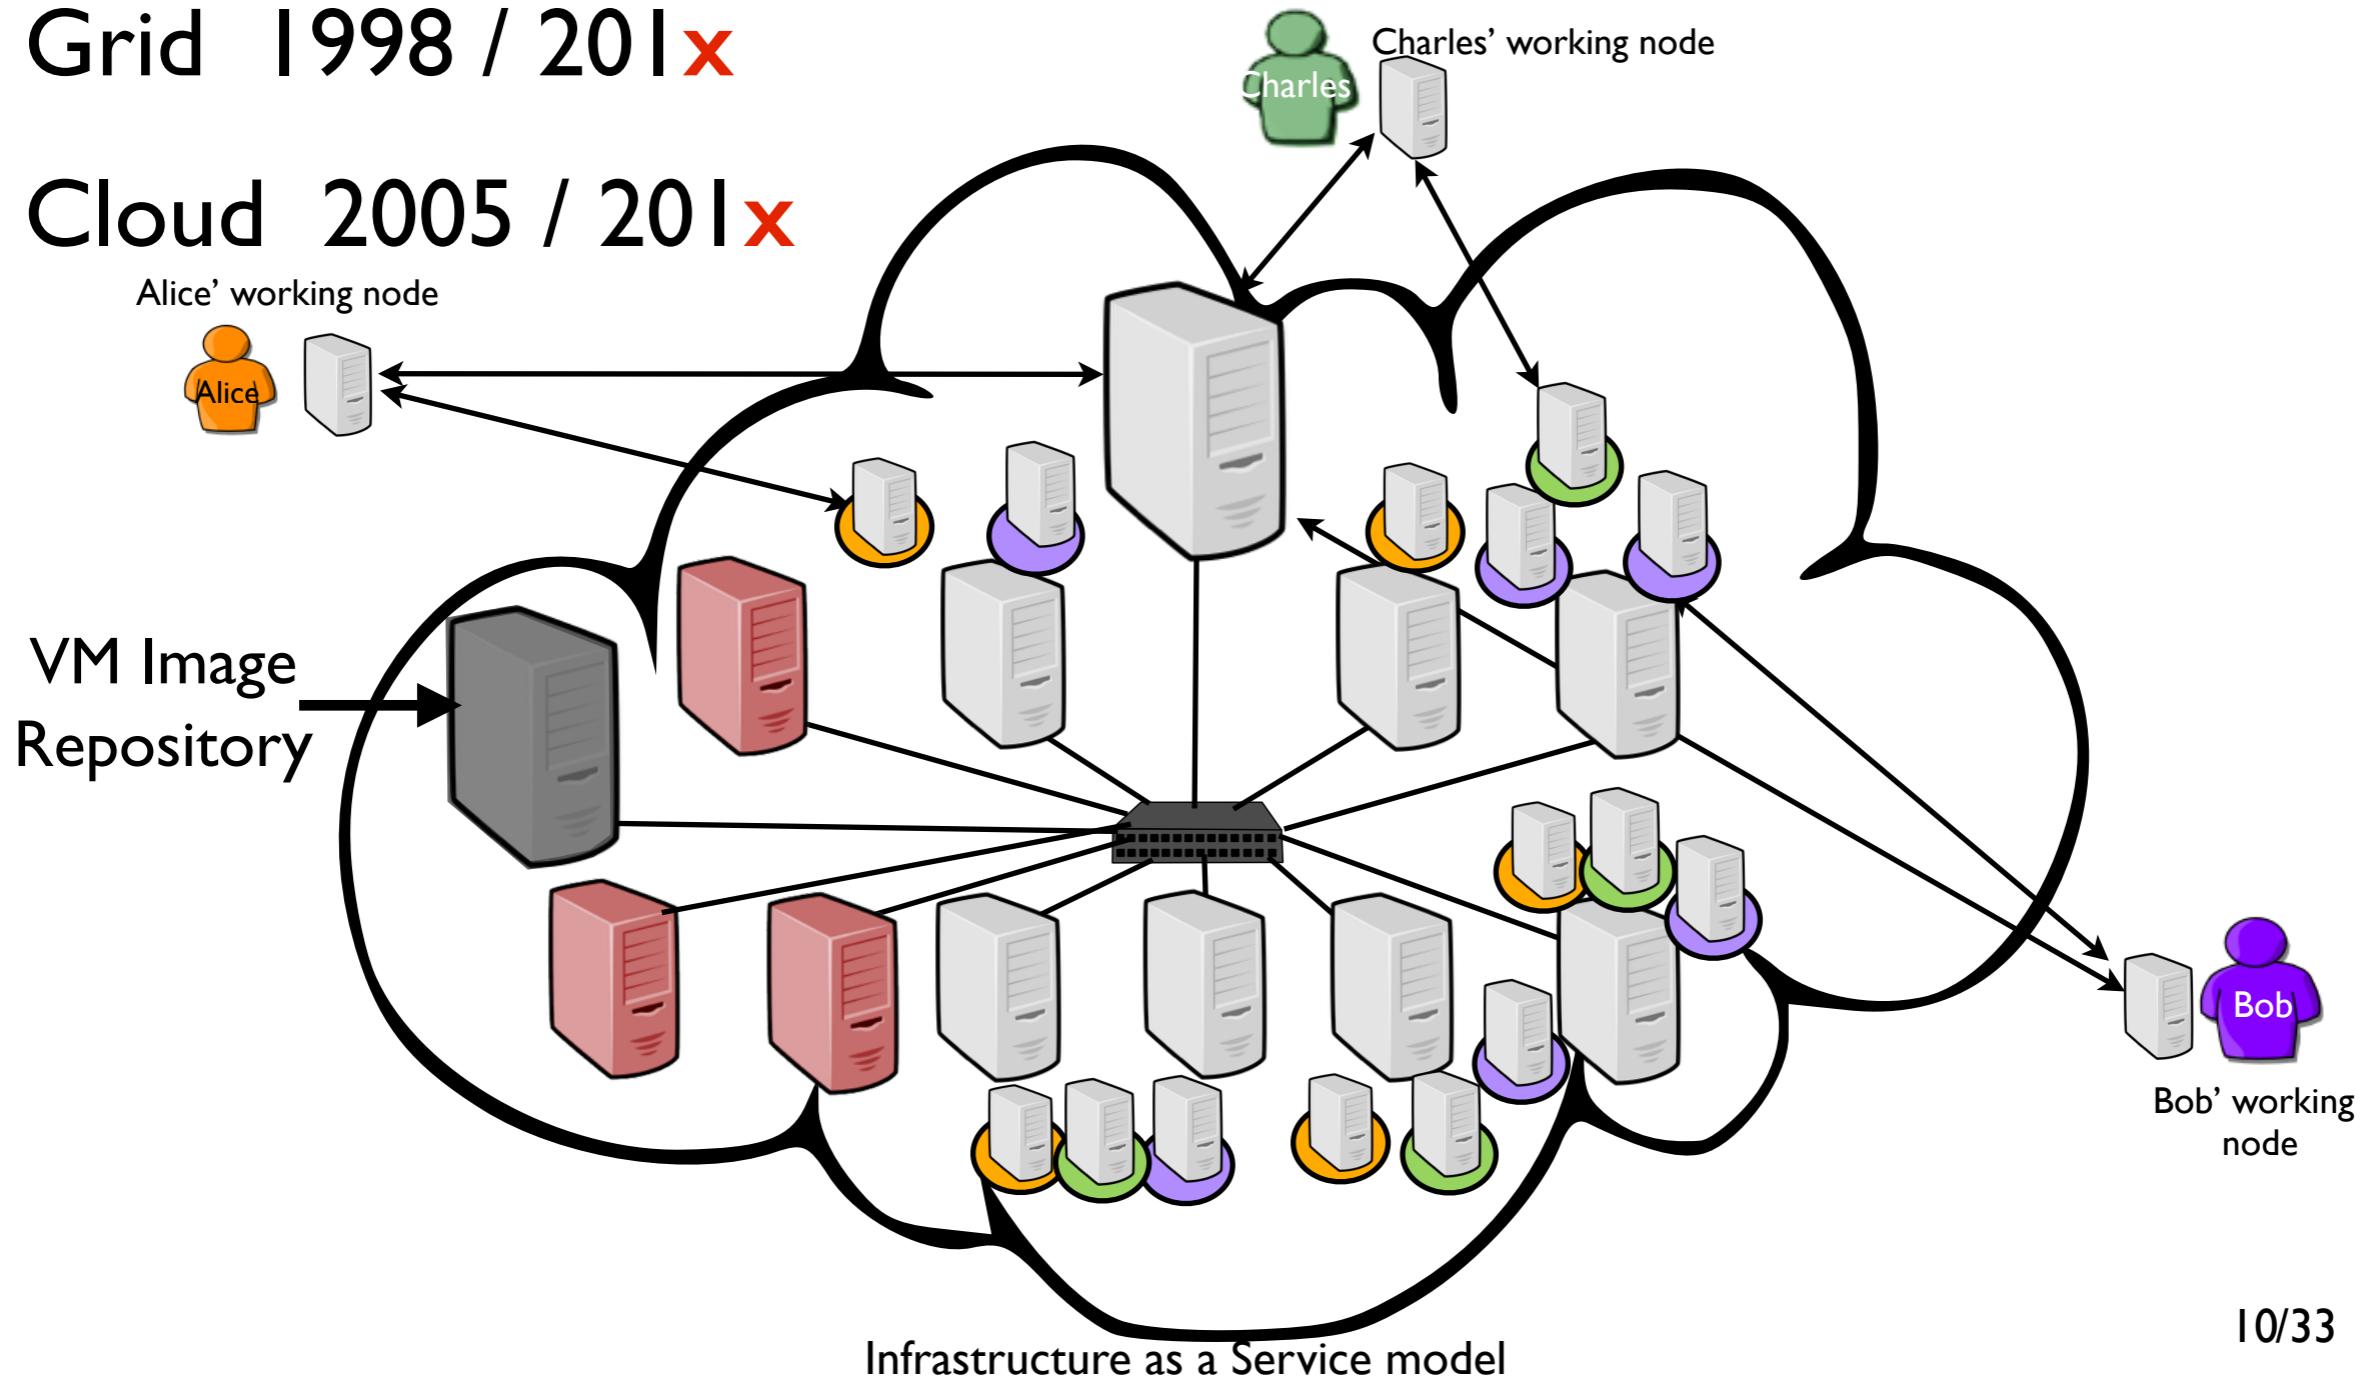
# Looking back ...

- Network of Workstations 1990 / 20~~xx~~
- Desktop Computing 1998/201~~x~~
- Grid 1998 / 201~~x~~
- Cloud 2005 / 201~~x~~



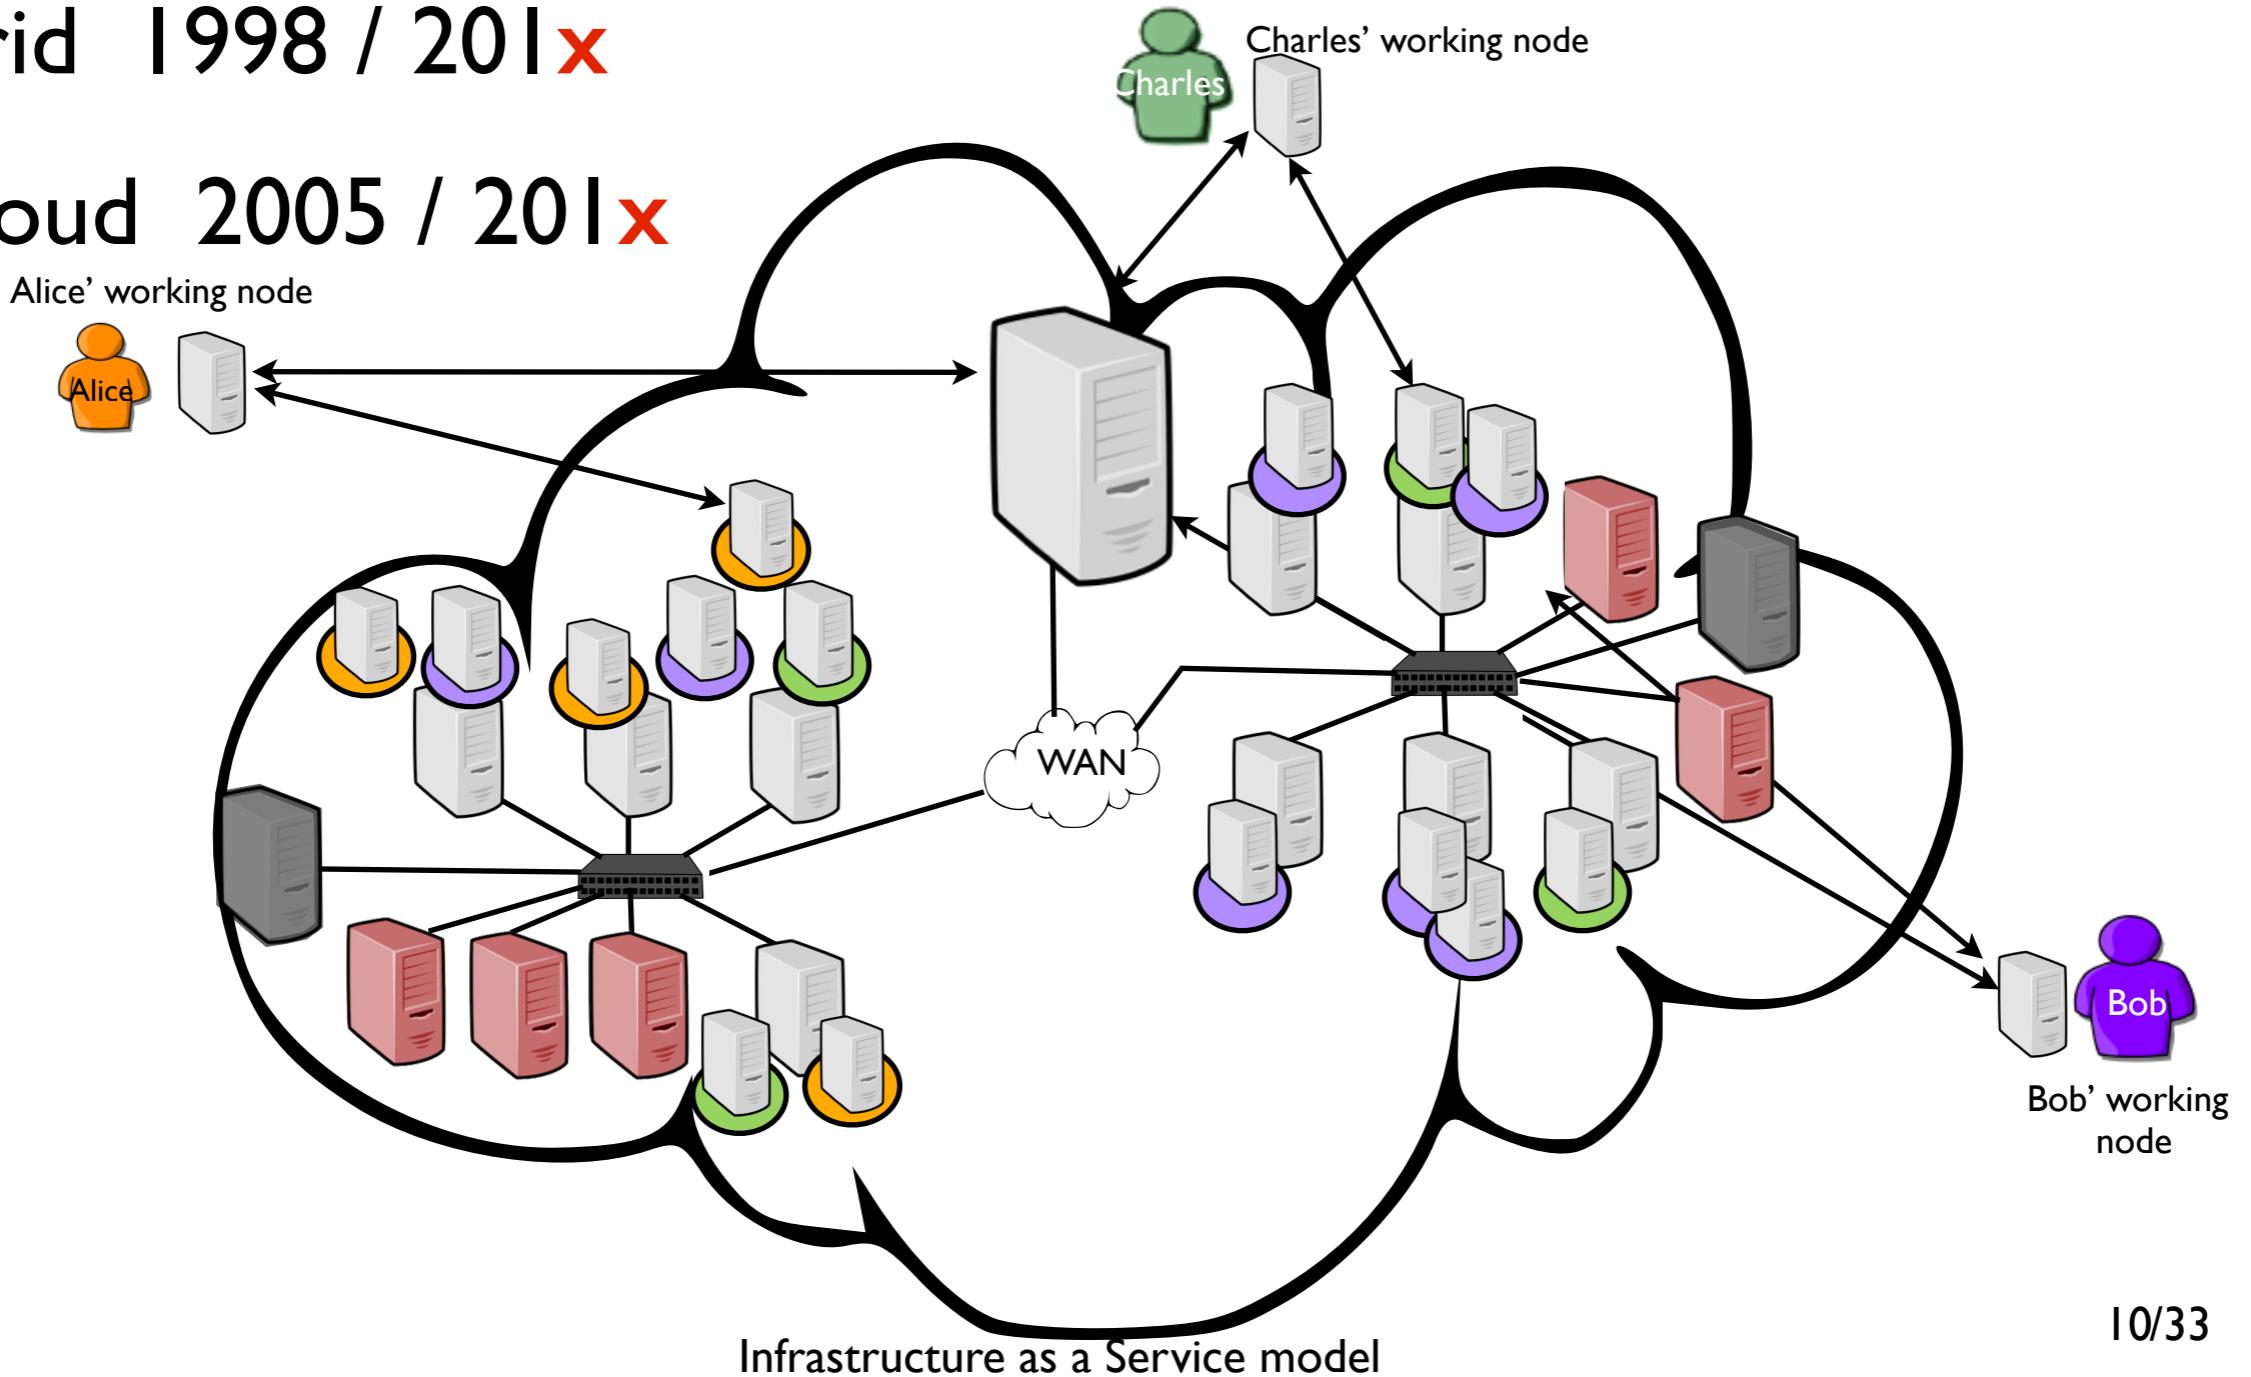
# Looking back ...

- Network of Workstations 1990 / 20~~xx~~
- Desktop Computing 1998/201~~x~~
- Grid 1998 / 201~~x~~
- Cloud 2005 / 201~~x~~



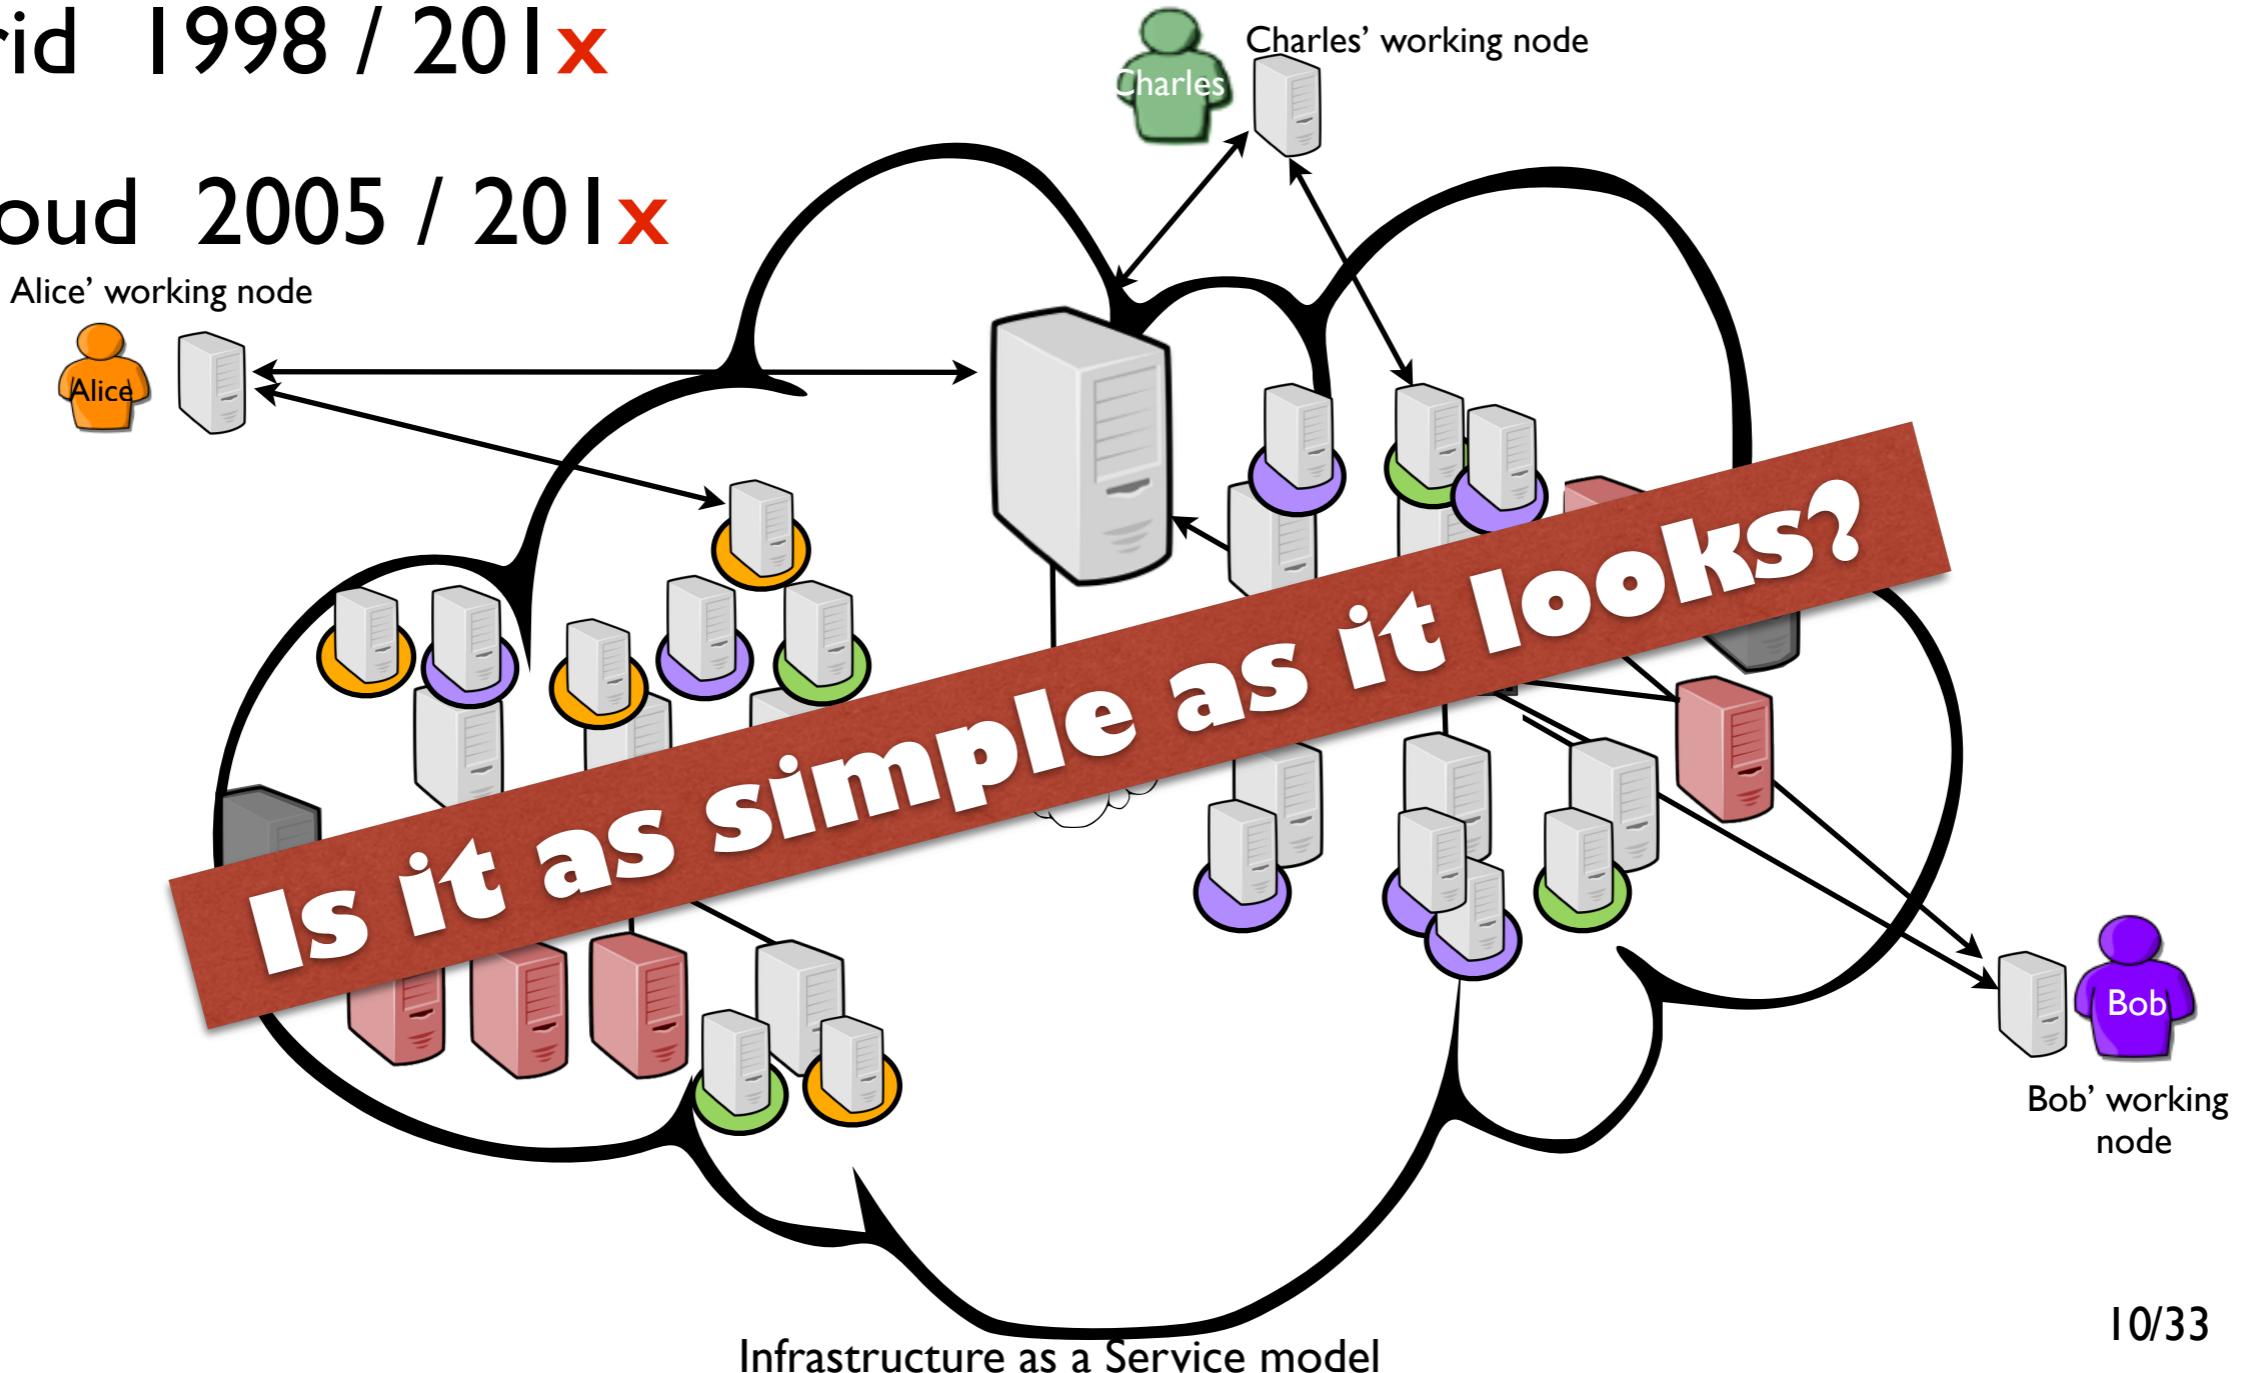
# Looking back ...

- Network of Workstations 1990 / 20~~xx~~
- Desktop Computing 1998/201~~x~~
- Grid 1998 / 201~~x~~
- Cloud 2005 / 201~~x~~



# Looking back ...

- Network of Workstations 1990 / 20~~xx~~
- Desktop Computing 1998/201~~x~~
- Grid 1998 / 201~~x~~
- Cloud 2005 / 201~~x~~



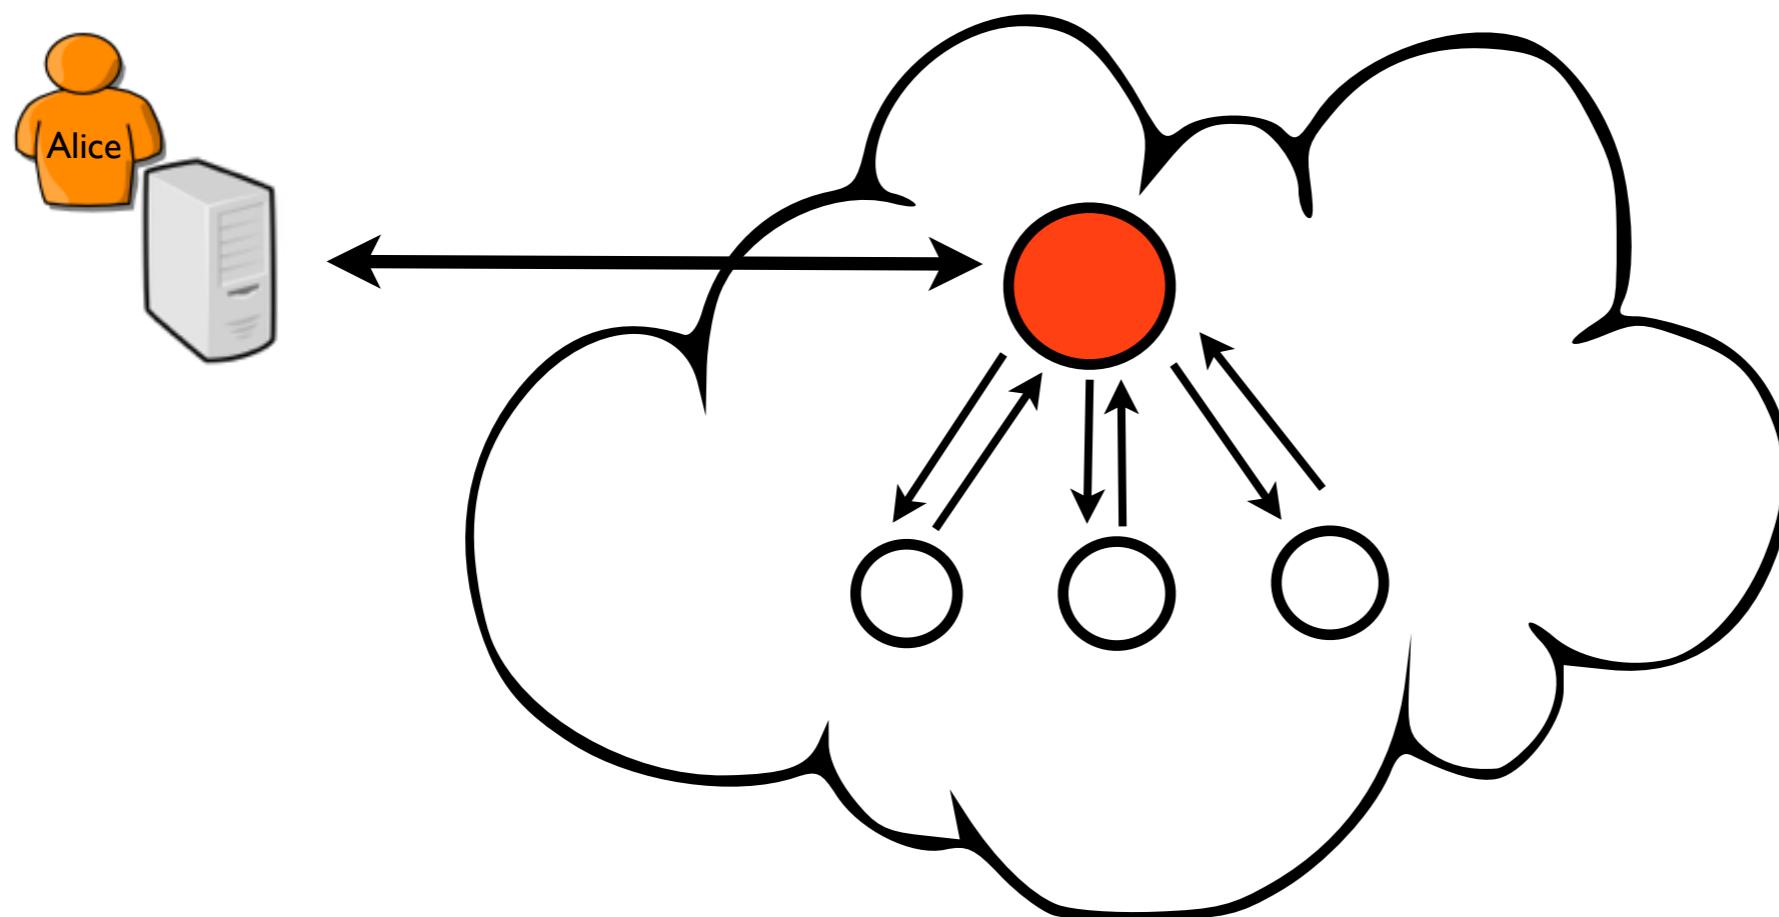
# IaaS Resource Management Systems

- An Operating System for Cloud infrastructures (aka Cloudkits)

Configuration of Virtual Environments (VEs)  
(contextualization, network...)

Images management/deployment

“Secure” accesses to the VEs



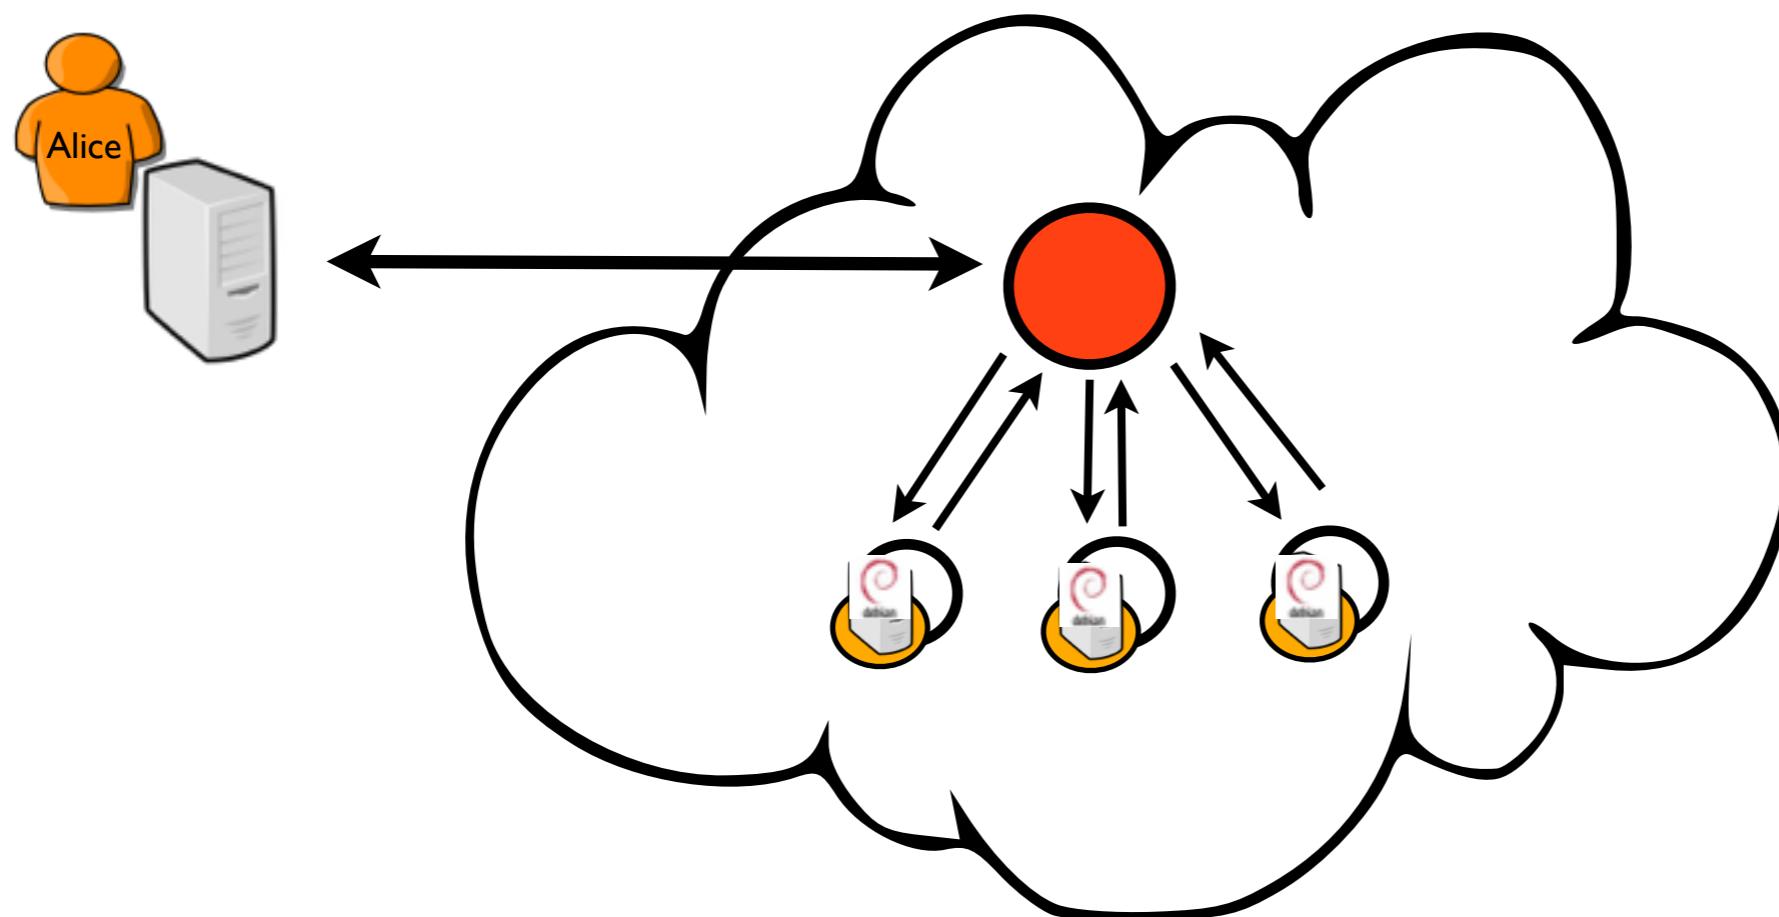
# IaaS Resource Management Systems

- An Operating System for Cloud infrastructures (aka Cloudkits)

Configuration of Virtual Environments (VEs)  
(contextualization, network...)

Images management/deployment

“Secure” accesses to the VEs



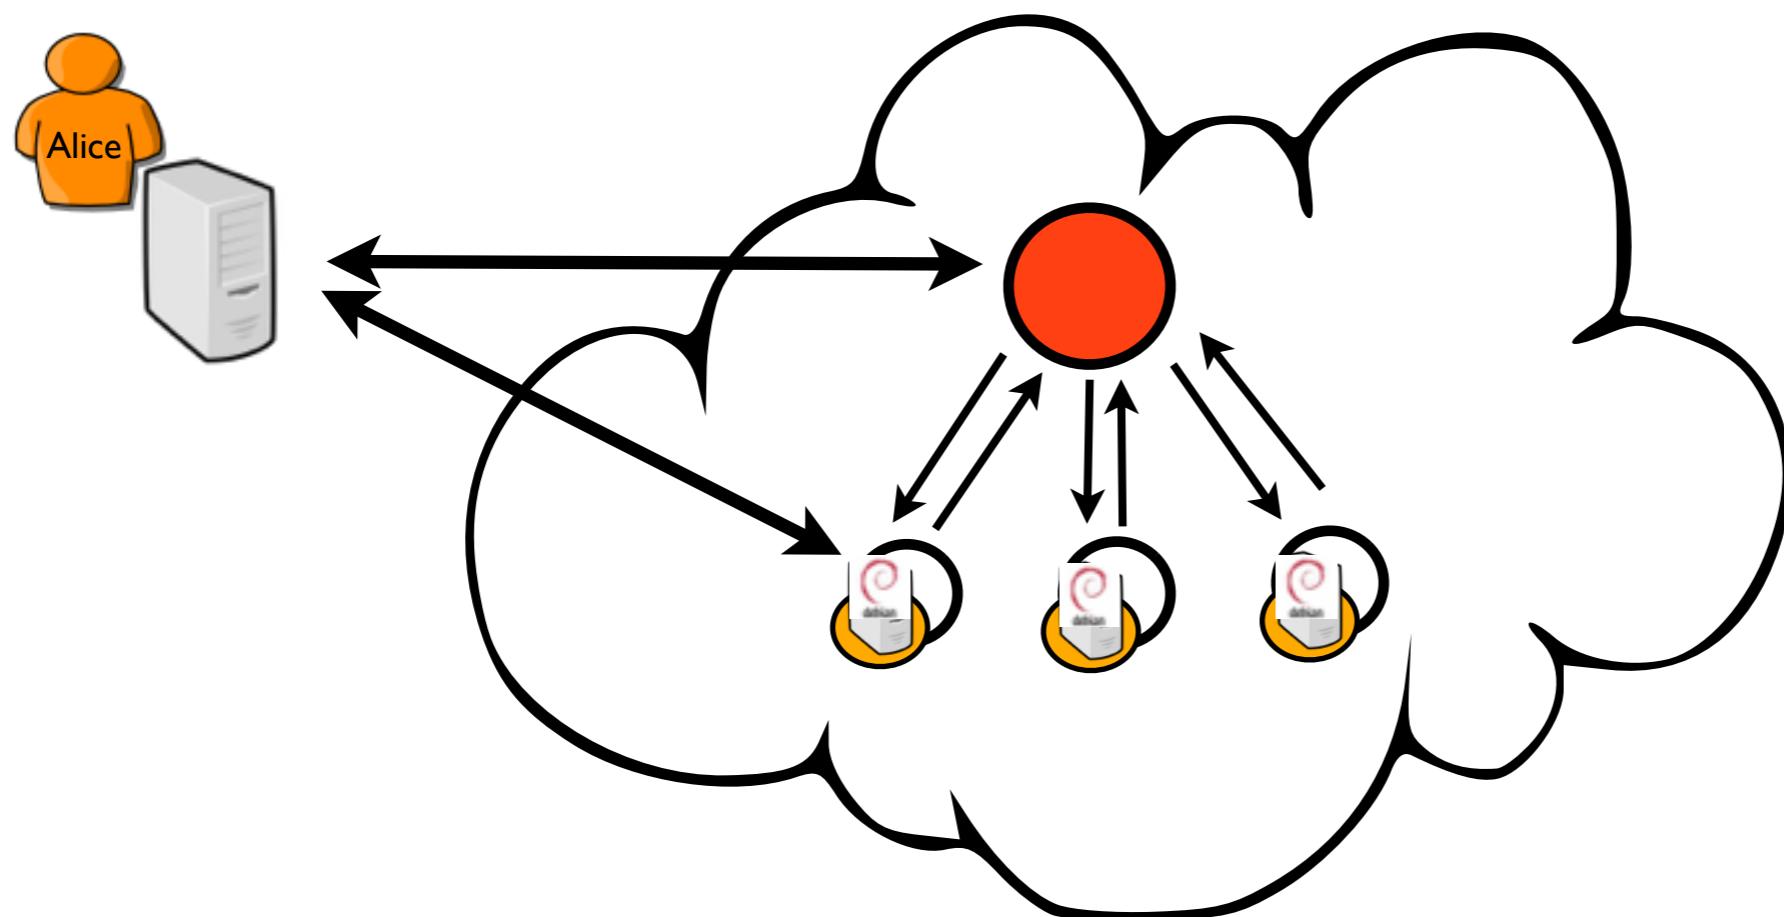
# IaaS Resource Management Systems

- An Operating System for Cloud infrastructures (aka Cloudkits)

Configuration of Virtual Environments (VEs)  
(contextualization, network...)

Images management/deployment

“Secure” accesses to the VEs



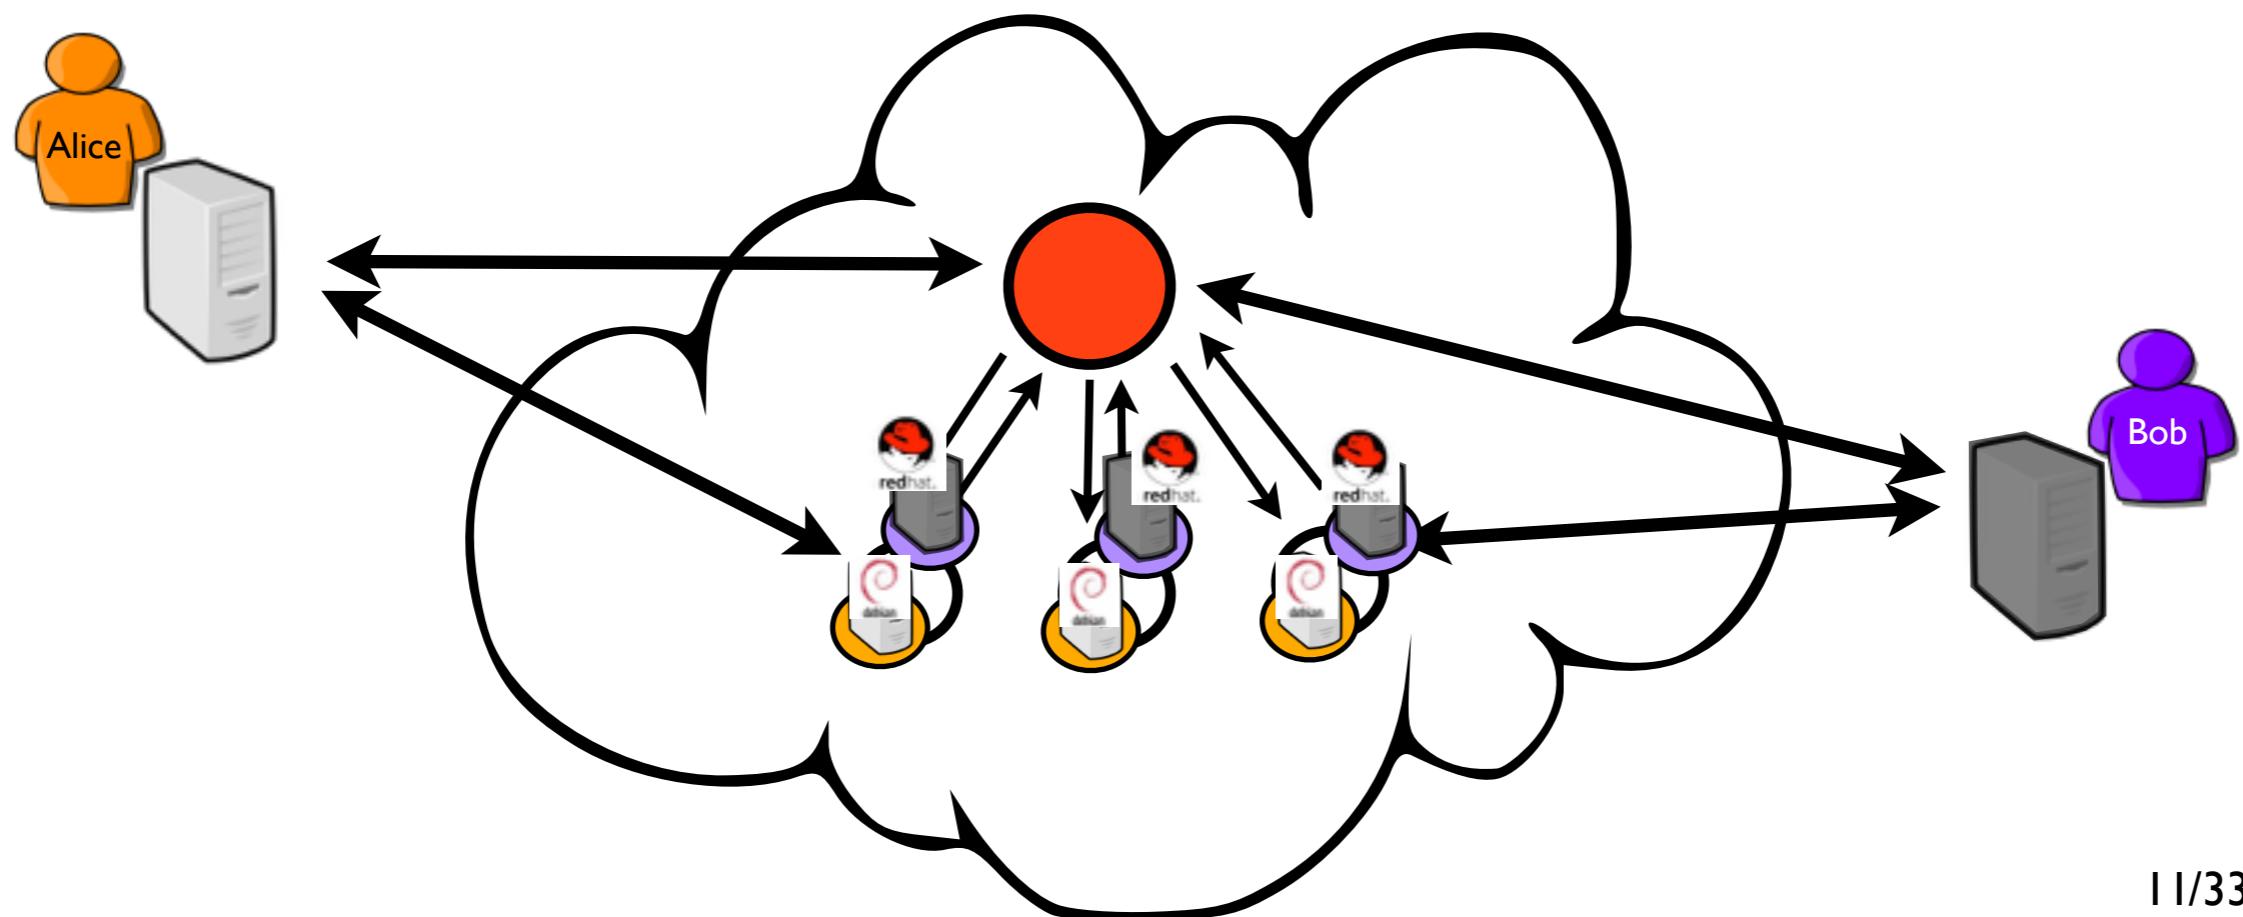
# IaaS Resource Management Systems

- An Operating System for Cloud infrastructures (aka Cloudkits)

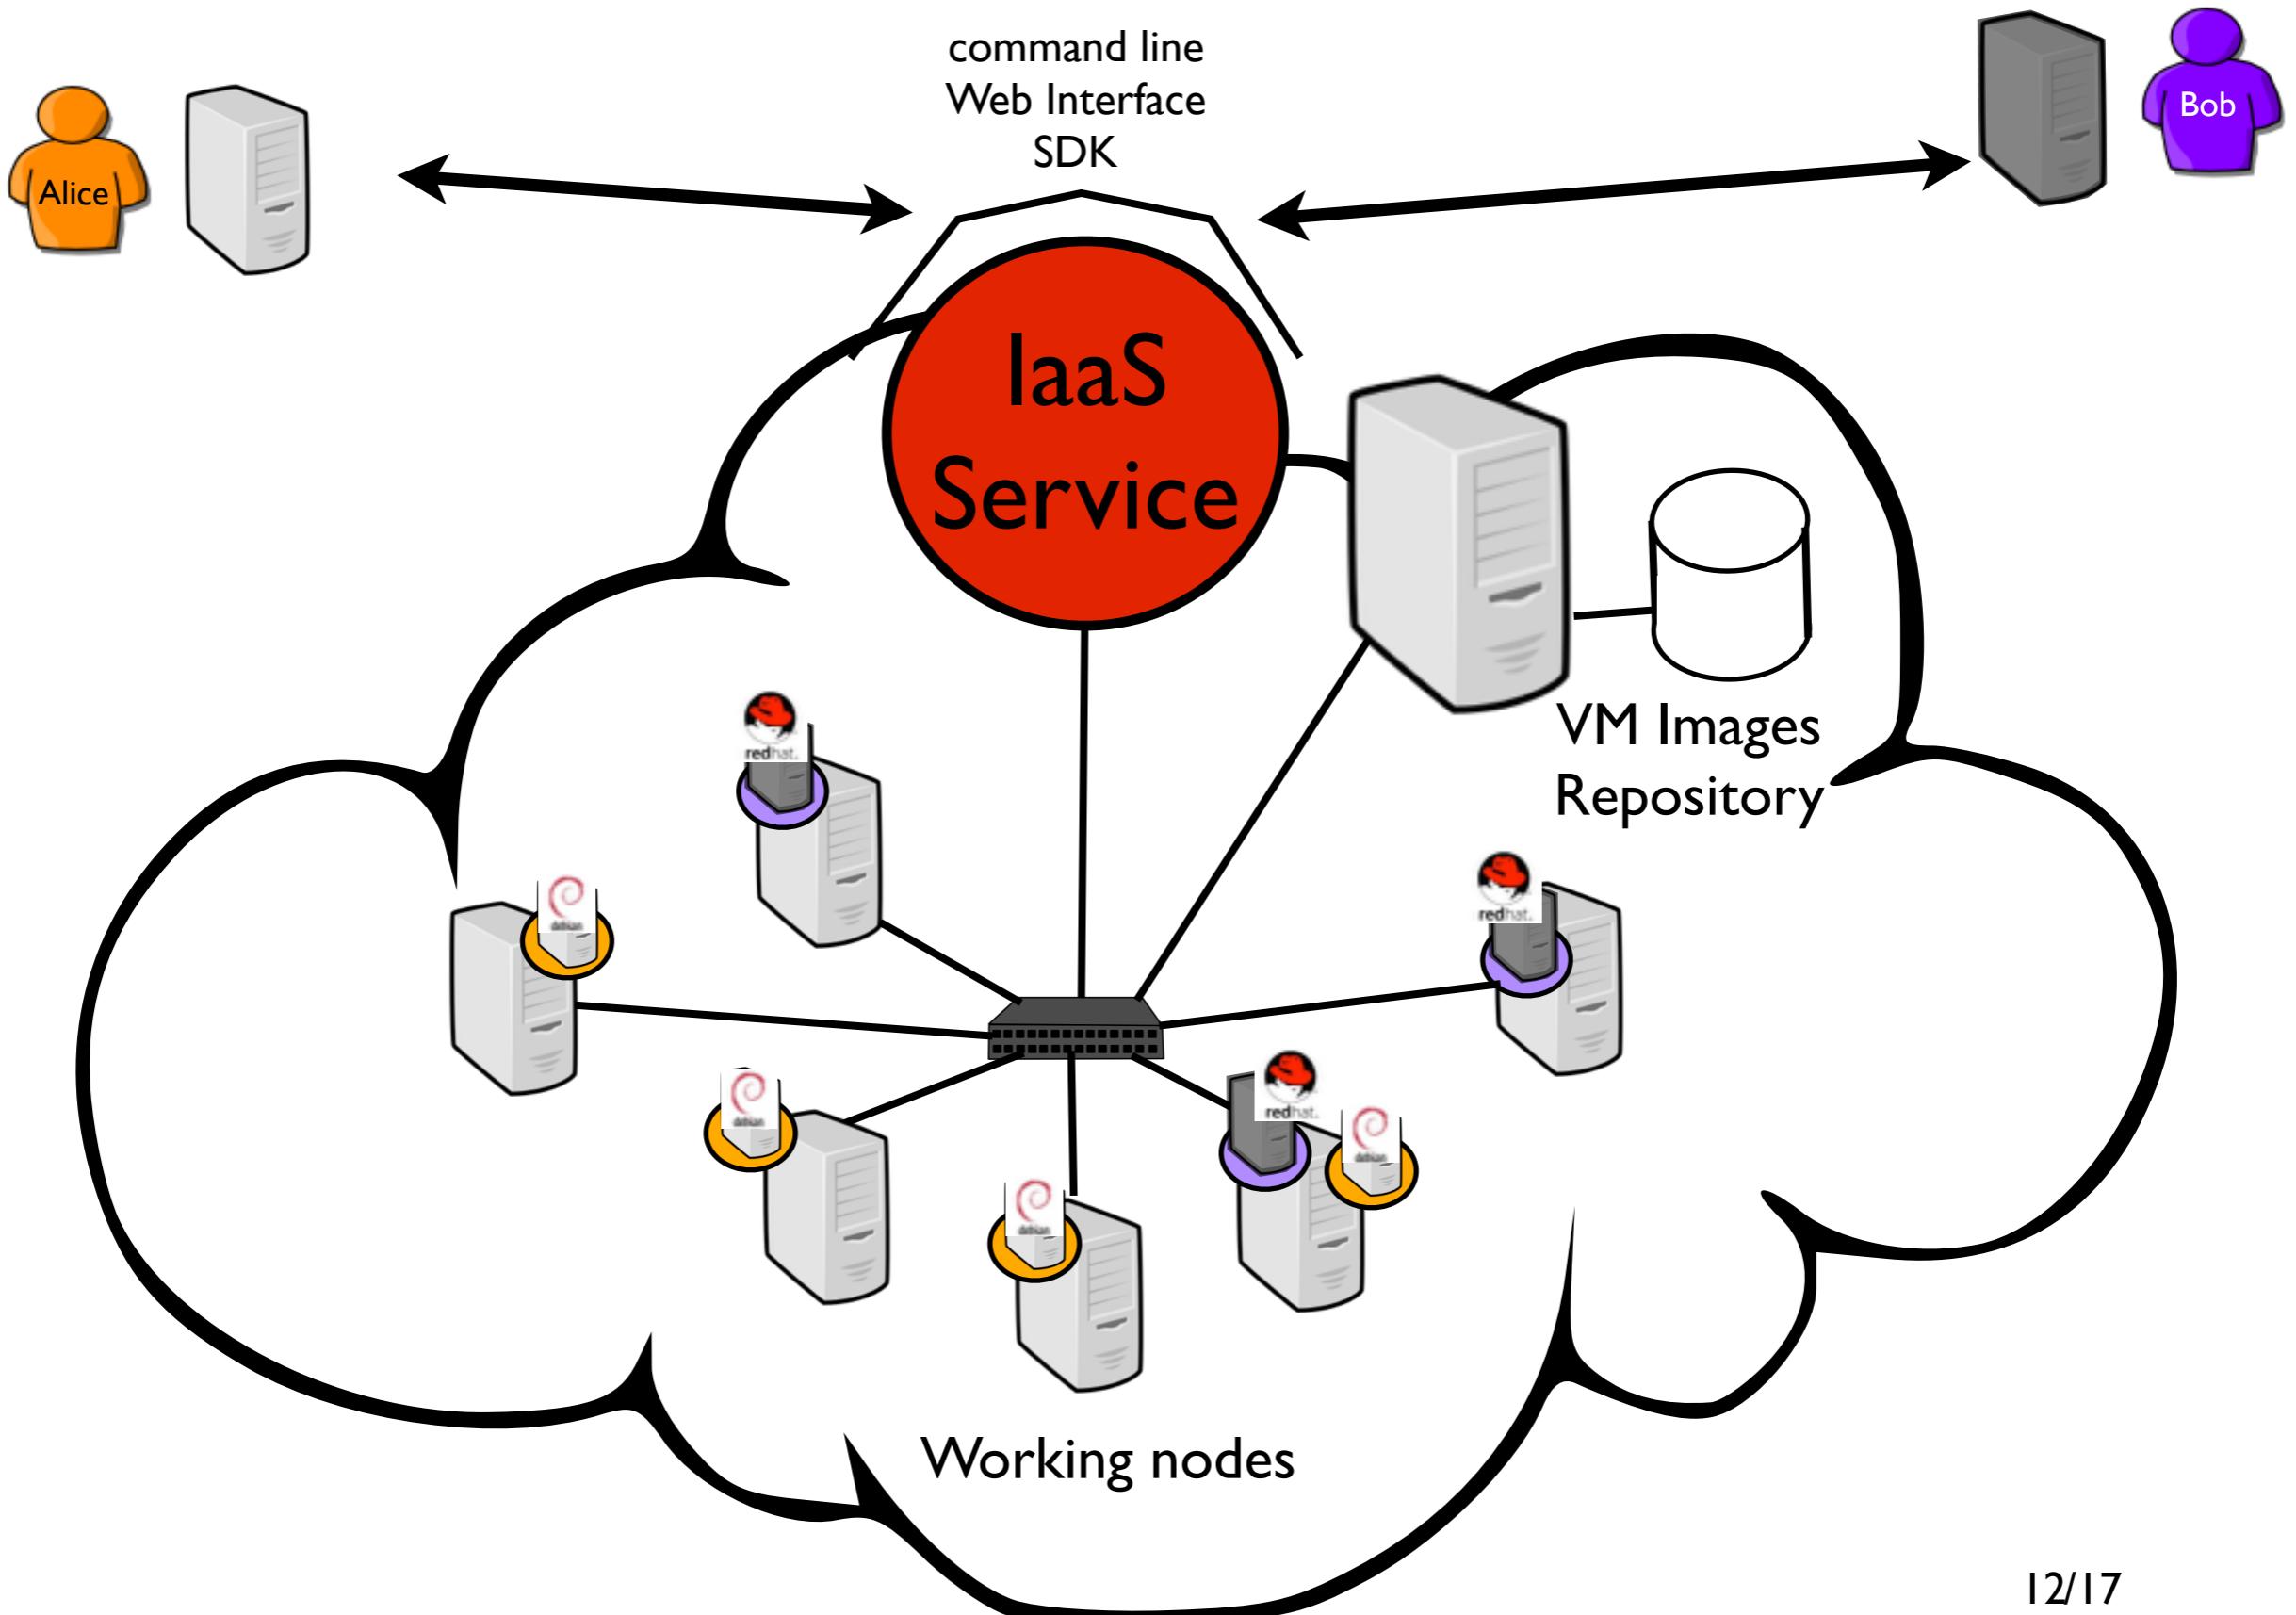
Configuration of Virtual Environments (VEs)  
(contextualization, network...)

Images management/deployment

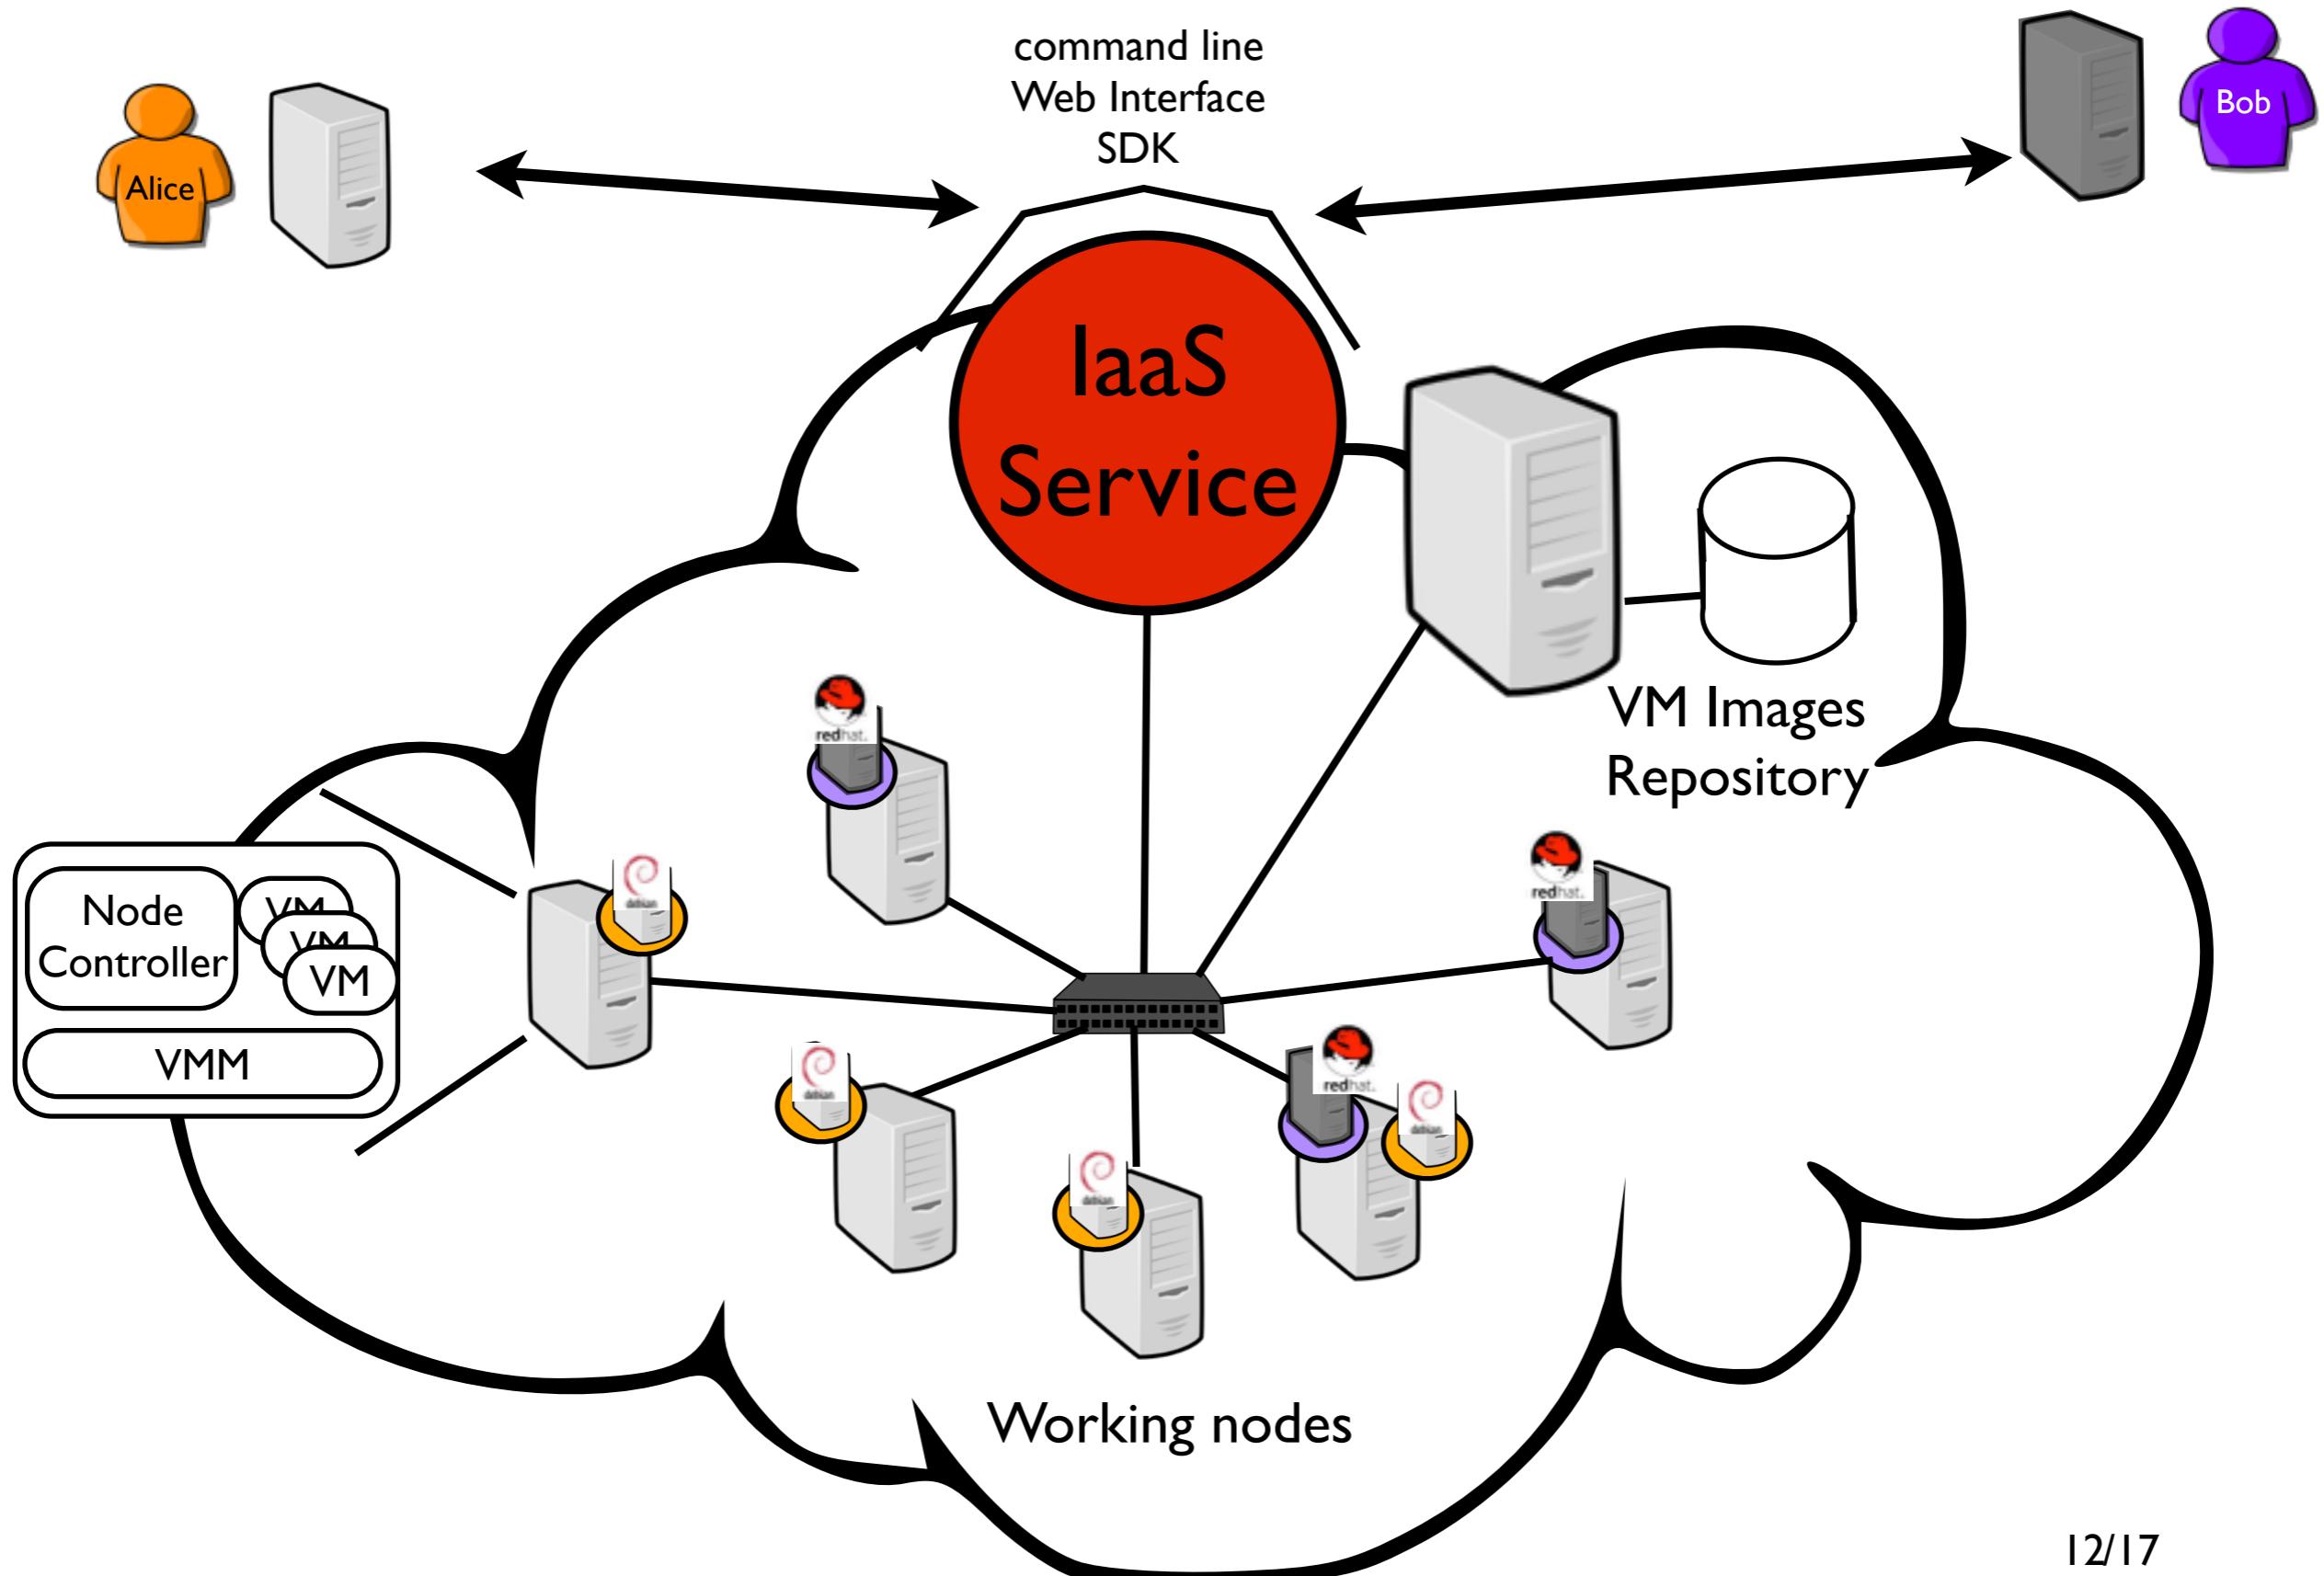
“Secure” accesses to the VEs



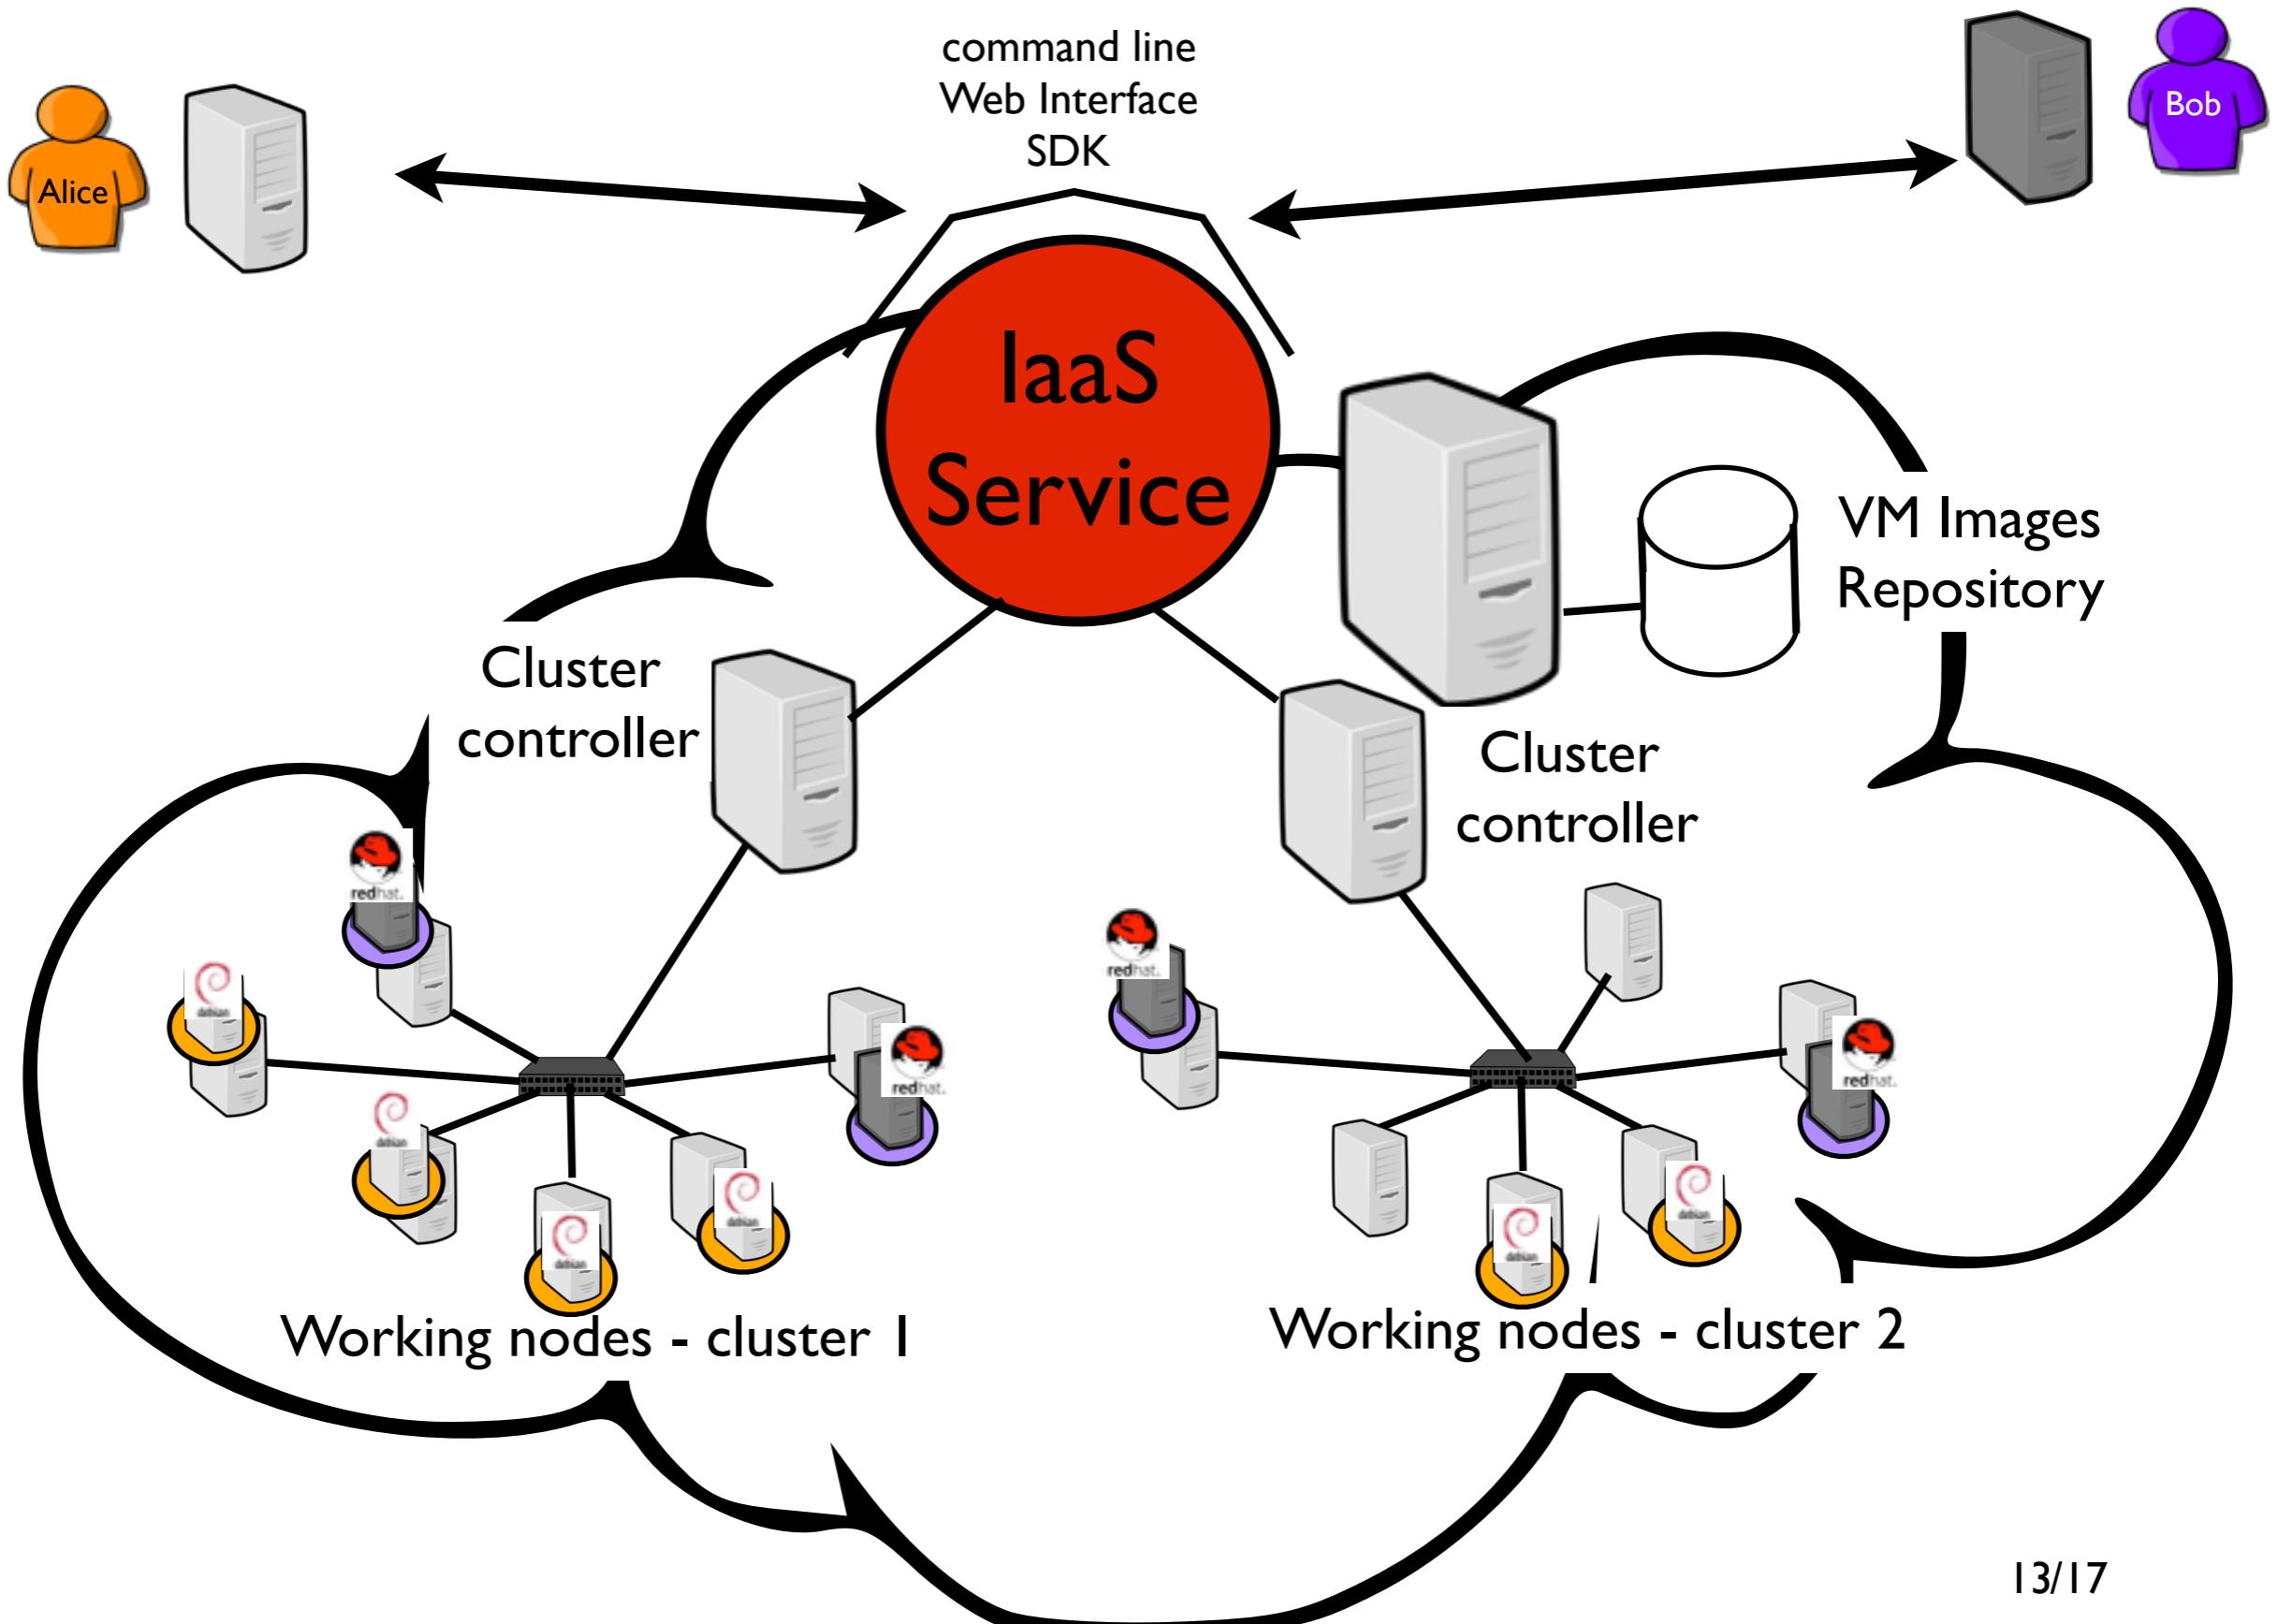
# An Overview of the IaaS Internals



# An Overview of the IaaS Internals



# An Overview of the IaaS Internals



# Managing IaaS - OpenSource solutions

- Open Nebula

[OpenNebula.org](http://OpenNebula.org)

2008-20XX

Results of the RESERVOIR project (mainly used in EU)

Montero & Llorente, DSA-Research at UCM

C++ / set of scripts

- CloudStack

2010-20XX

Apache project (in 2011)

Java Based



- Open Stack

2010-20XX

Supported by several industrials

The defacto open-source solution

Python

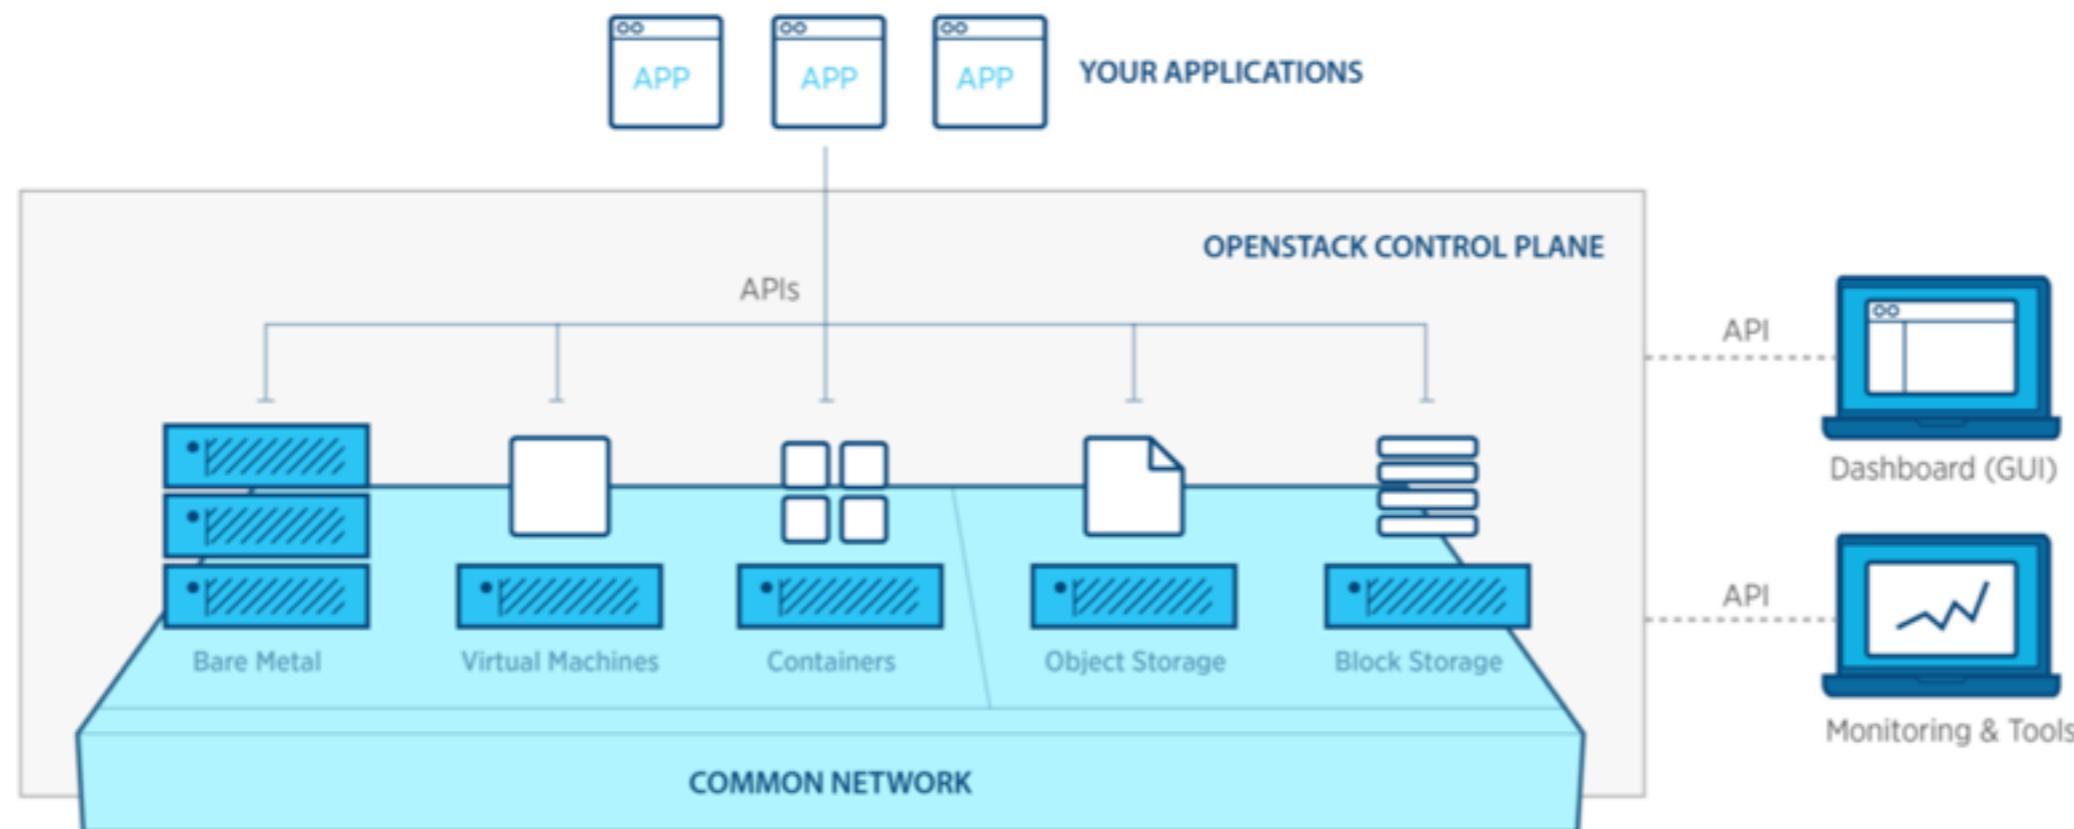
 openstack.

*You said OpenStack...  
You said OpenStack...*

# OpenStack



## One platform for bare metal, VMs and containers



OpenStack provides one platform to orchestrate bare metal, containers, and virtual machines on a single network, allowing private users to optimize for their application without creating more silos in their datacenters, and giving service providers more delivery options.

# OpenStack



# OpenStack



More than 70,000 registered community members

- 649 supporting organizations
- 181 countries represented
- 116 global user groups

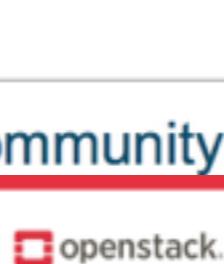


# OpenStack

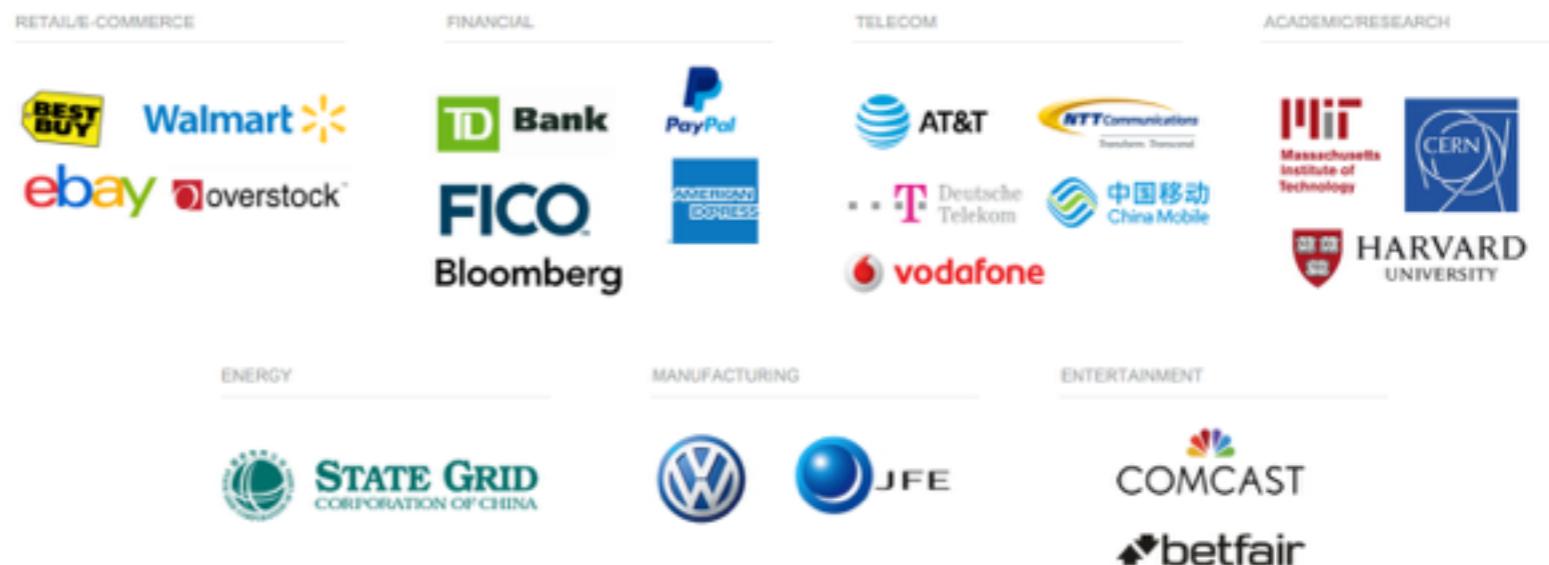


More than 70,000 registered community members

- 649 supporting organizations
  - 181 countries represented
  - 116 global user groups



# OpenStack users span industries



# OpenStack

 openstack.

More than 70,000 registered community members

- 649 supporting organizations
- 181 countries represented

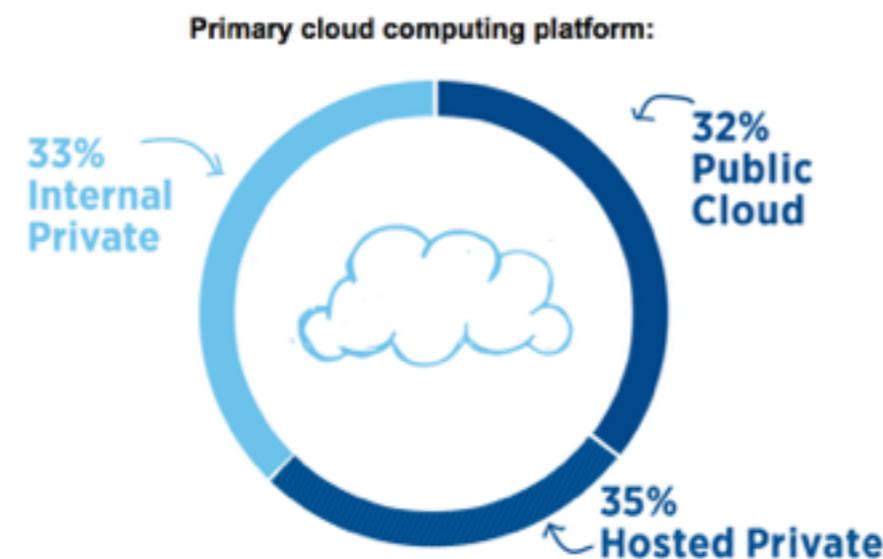


 openstack.

OpenStack users span industries

 openstack.

Private clouds make up 2/3 of cloud platforms



ICOM



Deutsche Telekom



vodafone

ACADEMIC/RESEARCH



ENTERTAINMENT



FE

# OpenStack



More than 70,000 registered community members

- 649 supporting organizations
- 181 countries represented



OpenStack users span industries



ICOM

ACADEMIC/RESEARCH



Private clouds make up 2/3 of cloud platforms

Primary cloud cc



33%  
Internal  
Private



- The first OpenStack PTG was held in February in Atlanta, Georgia. Over 500 developers representing nearly 50 teams attended to collaborate and work on the OpenStack Pike release.
- The second PTG will be held in Denver, September 11-14. There are a few sponsorship opportunities available.

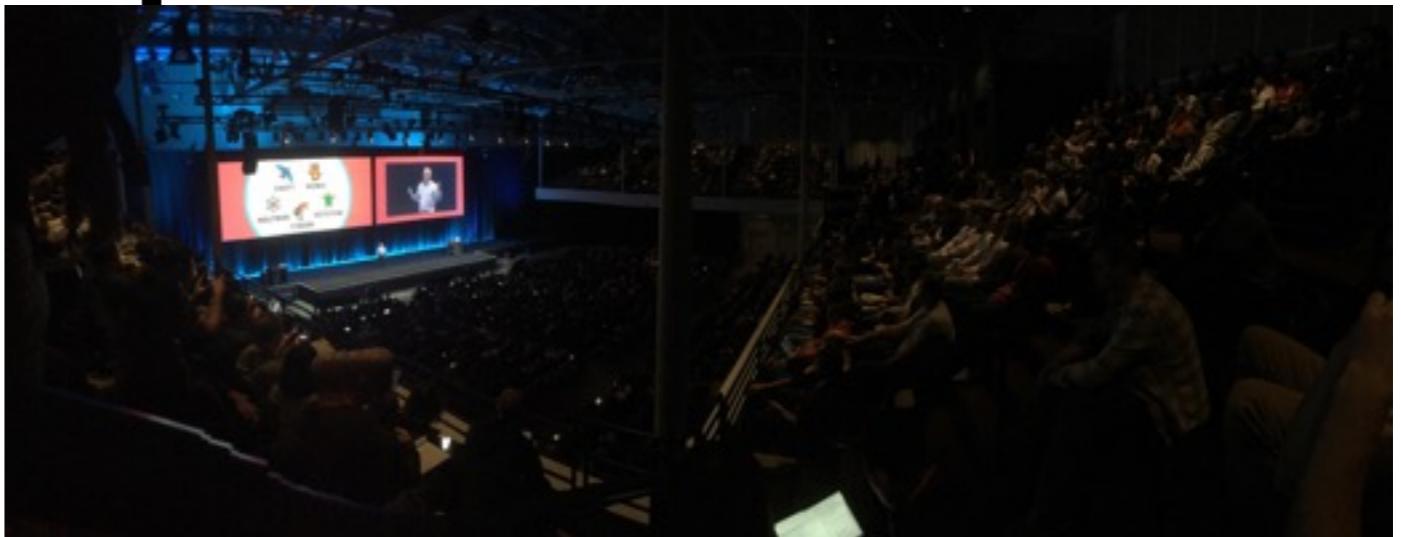
Source: Forrester Research, 2017

# The OOO principles

- Open source
- Open development: access to each contribution/ logs of meetings,...
- Open design: the community is listened to set the direction of OpenStack
- Open Community: anyone can raise to leadership position

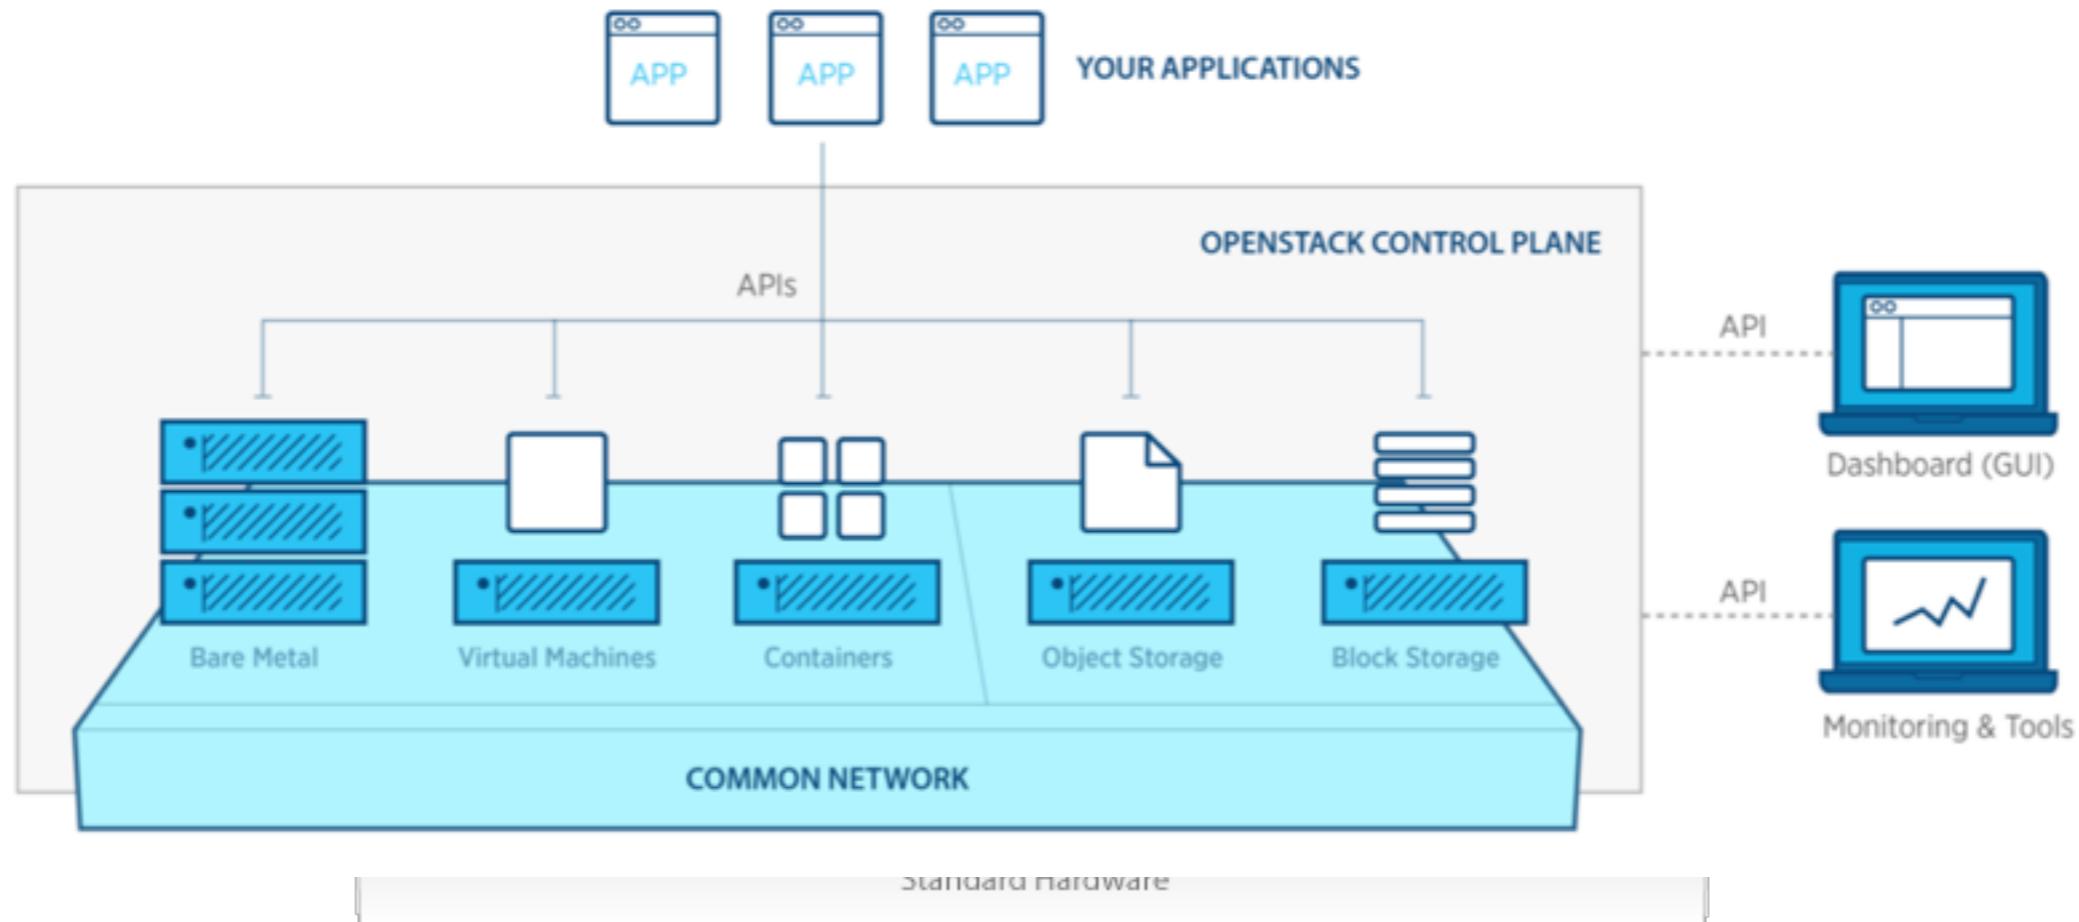
# What is OpenStack?

- an open-source project
- a common goal
- a coalition of organizations
- a foundation
- a trademark
- an interoperability standard
- a set of events (OpenStack Summit/PTG/OpenStack Days...)
- a governance model
- a job market
- a single project / a set of projects
- a set of principles
- a development community
- a big tent
- a bunch of python code
- a way to produce software
- a very active open-source project
- a success story
- ok so what it is....



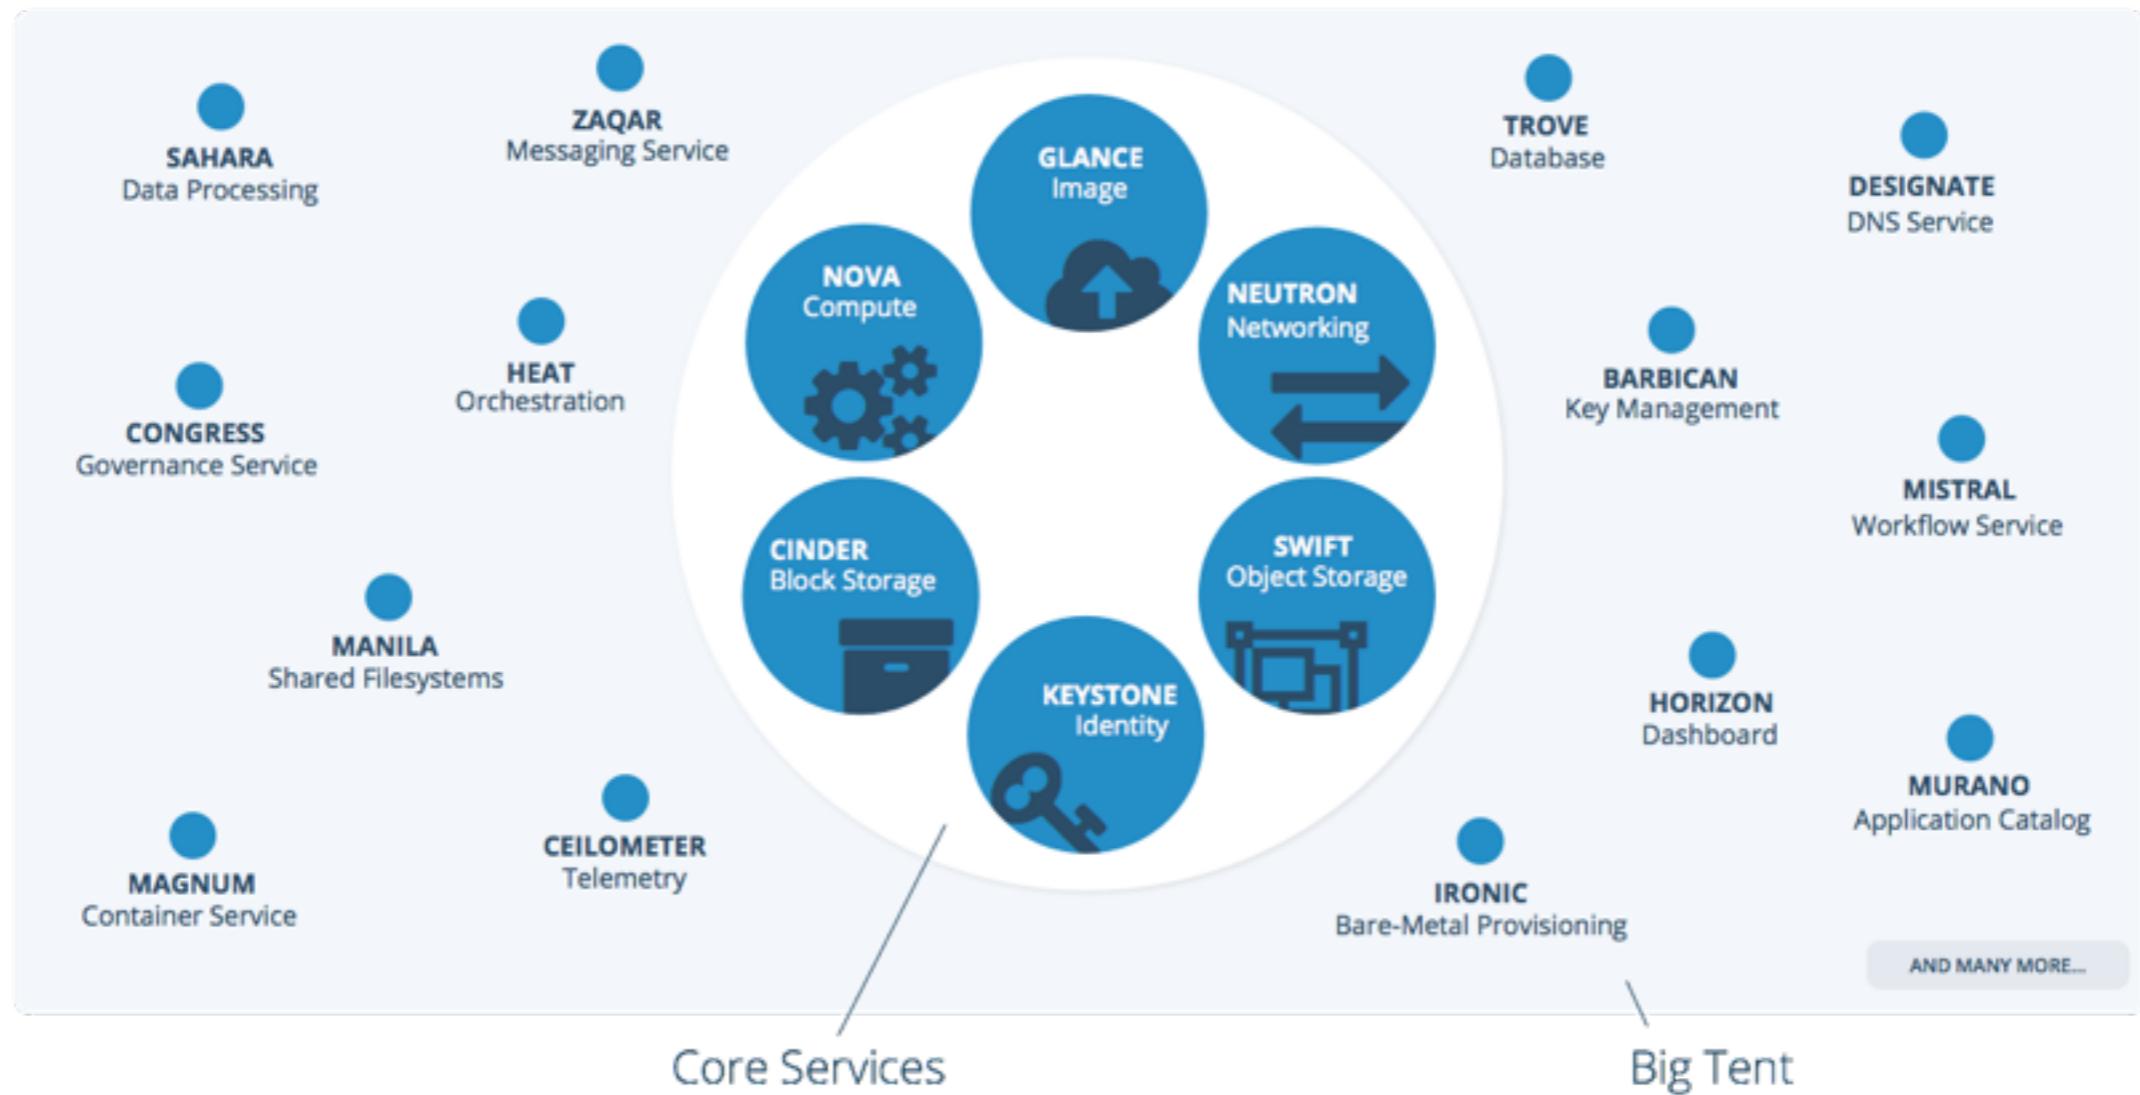
**A lot of Fun !**

# A Rich (and Complex) Ecosystem



- **20 Millions of LoC, 164 services**, some services are composed of sub-services (e.g. nova-scheduler, nova-conductor, ...)

# A Rich (and Complex) Ecosystem



- **20 Millions of LoC, 164 services, some services are composed of sub-services (e.g. nova-scheduler, nova-conductor, ...)**

# The User/Admin Viewpoints

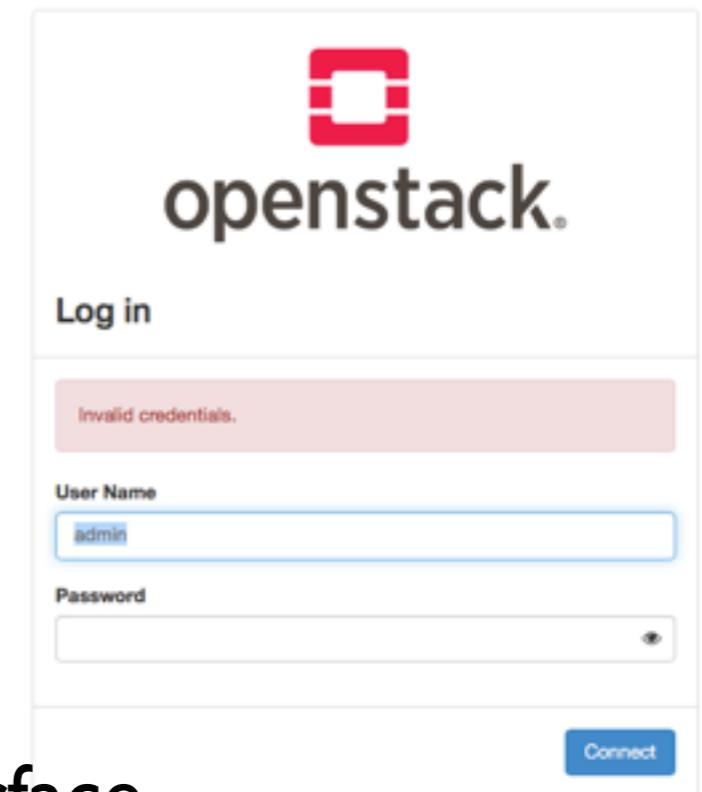
- Everything goes through the API (including the **HORIZON** dashboard)
- APIs: REST / one per service

Through HTTP (curl)

Through SDKs and broker libraries

Through **Horizon** or the command line interface

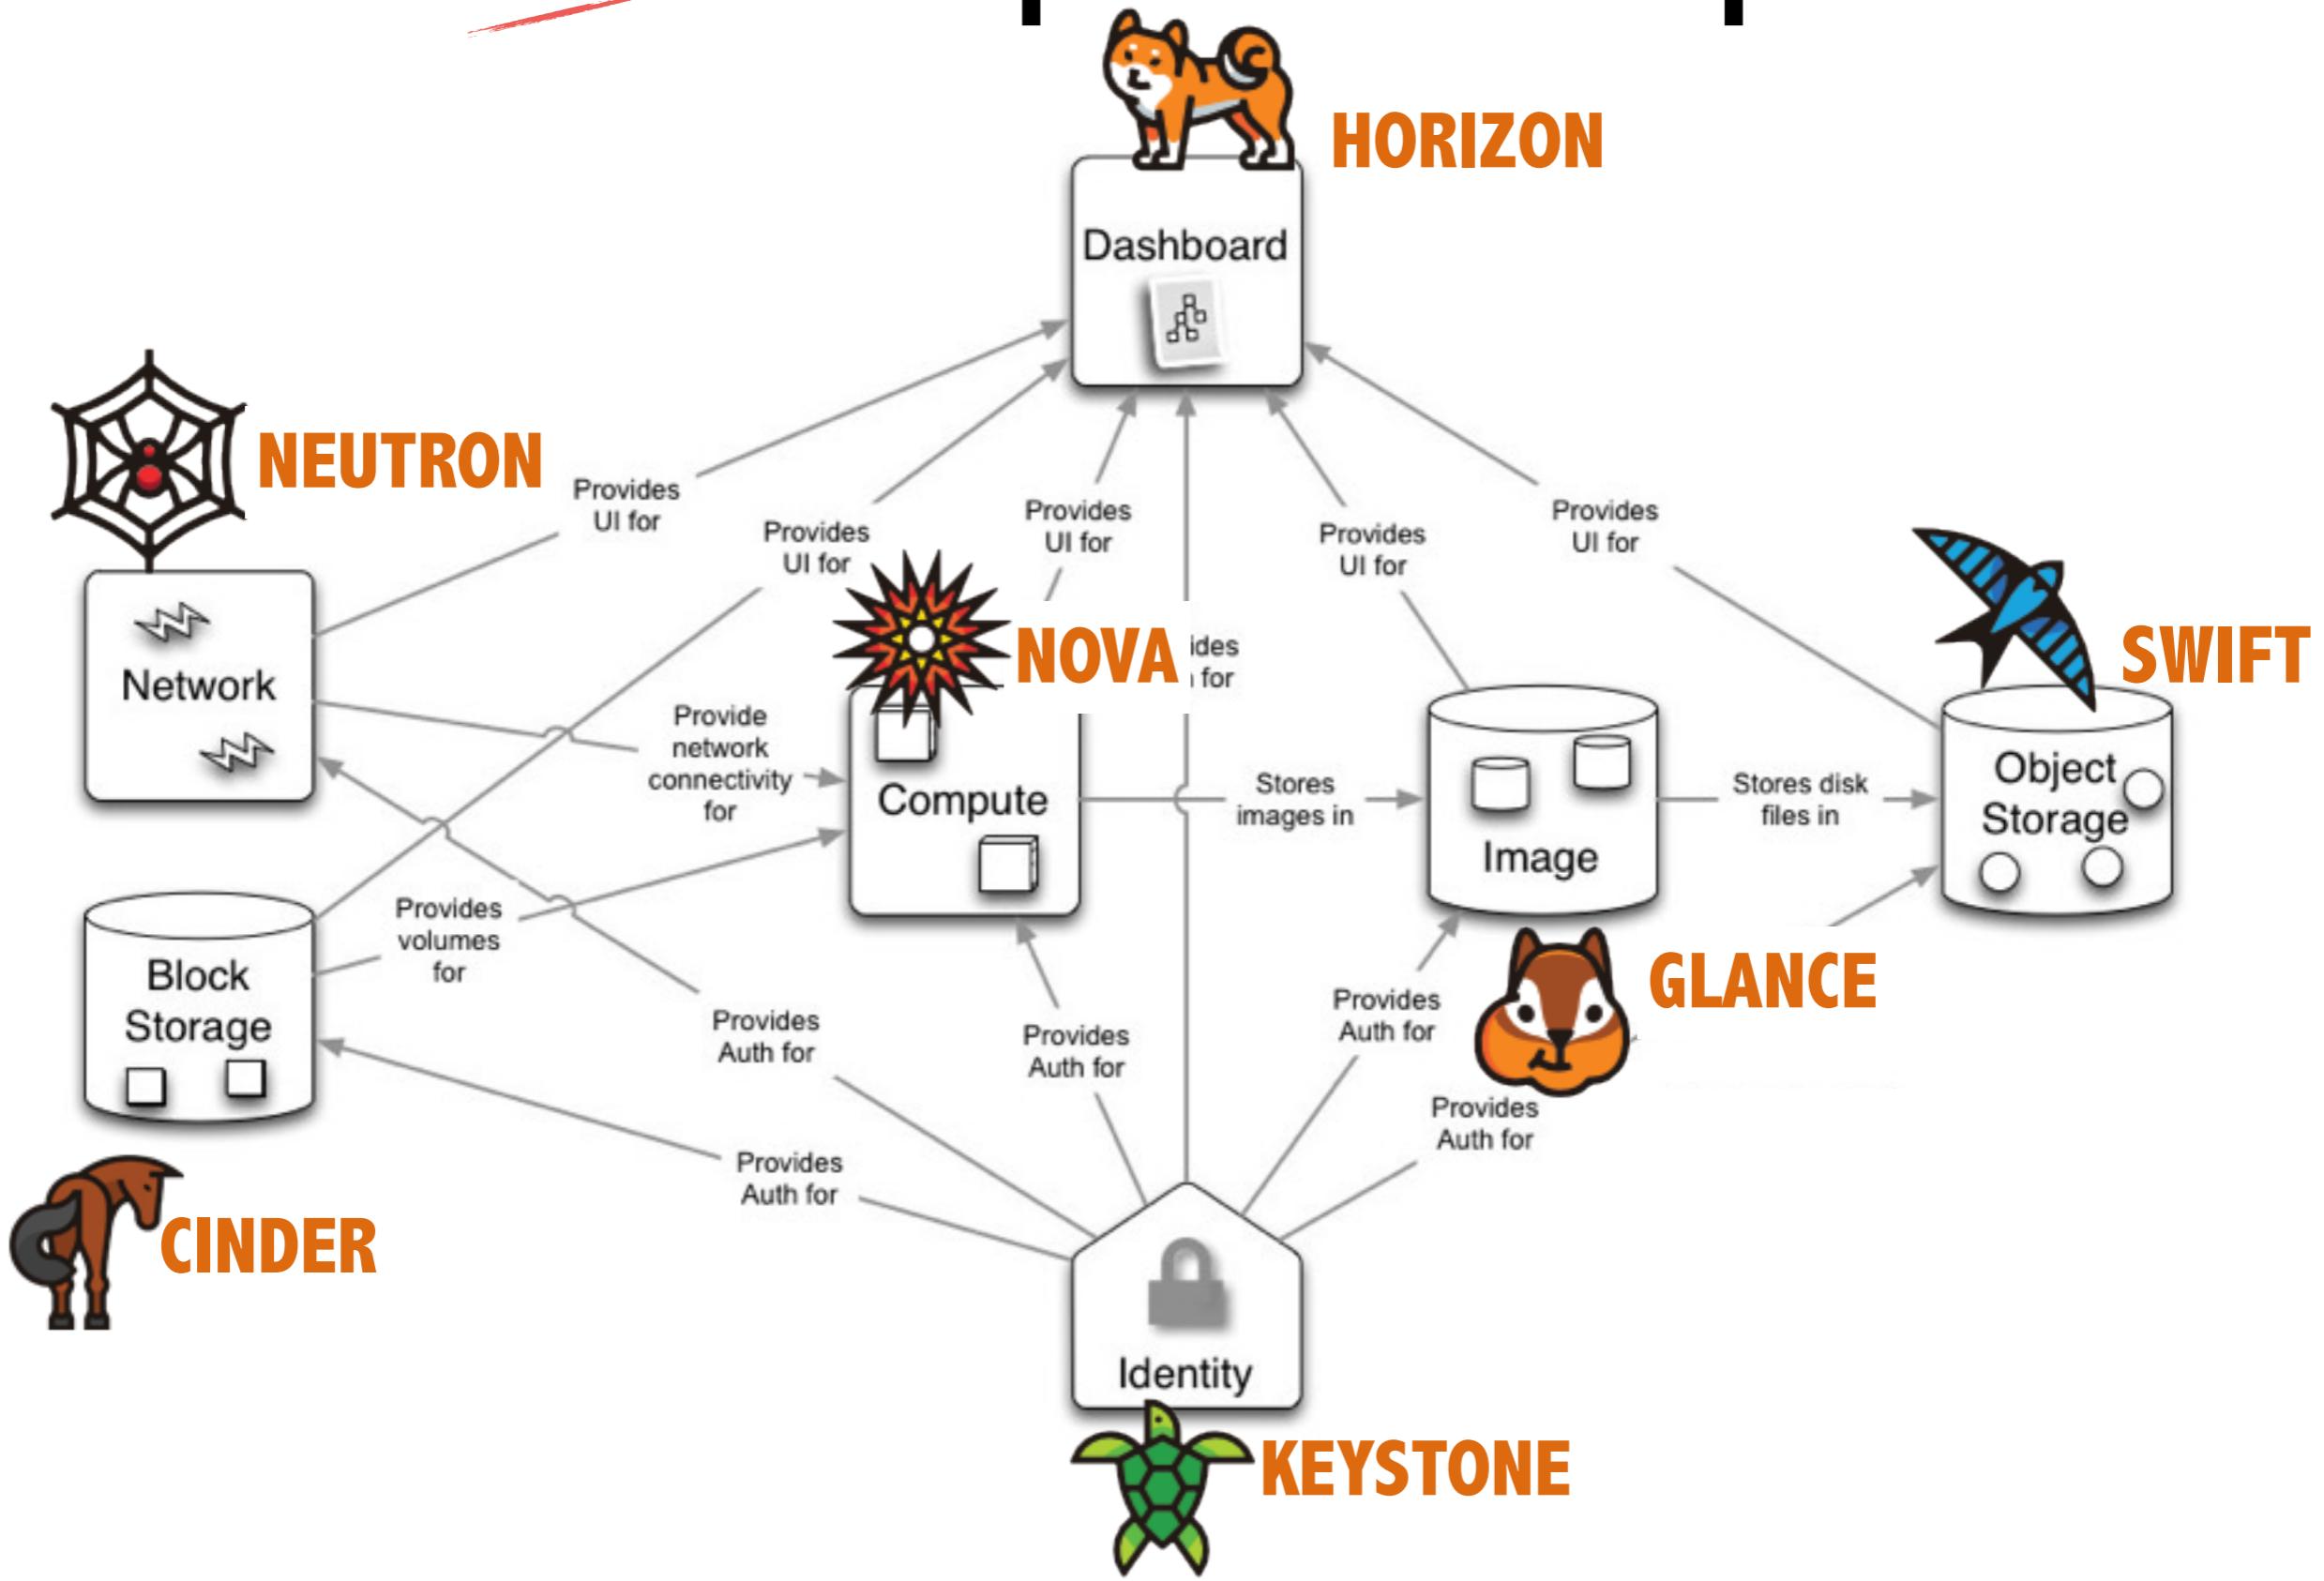
Through **HEAT**



- You need specific credentials (delivered by **KEYSTONE**)

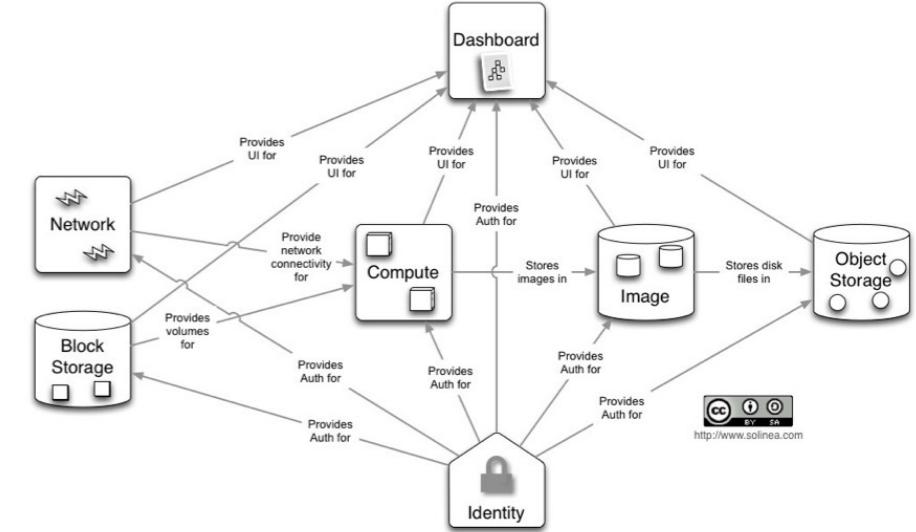
**R-A Cherrueau / D. Pertin**  
**Next Week !**

# The Developer Viewpoint



# OpenStack core-services

- Each core-service is divided into several sub-services
- Services communicate through a communication bus (AMQP)
- System states are stored in a SQL DB (MySQL/MariaDB)
- Python for all projects
- APIs: OpenStack and AWS-like

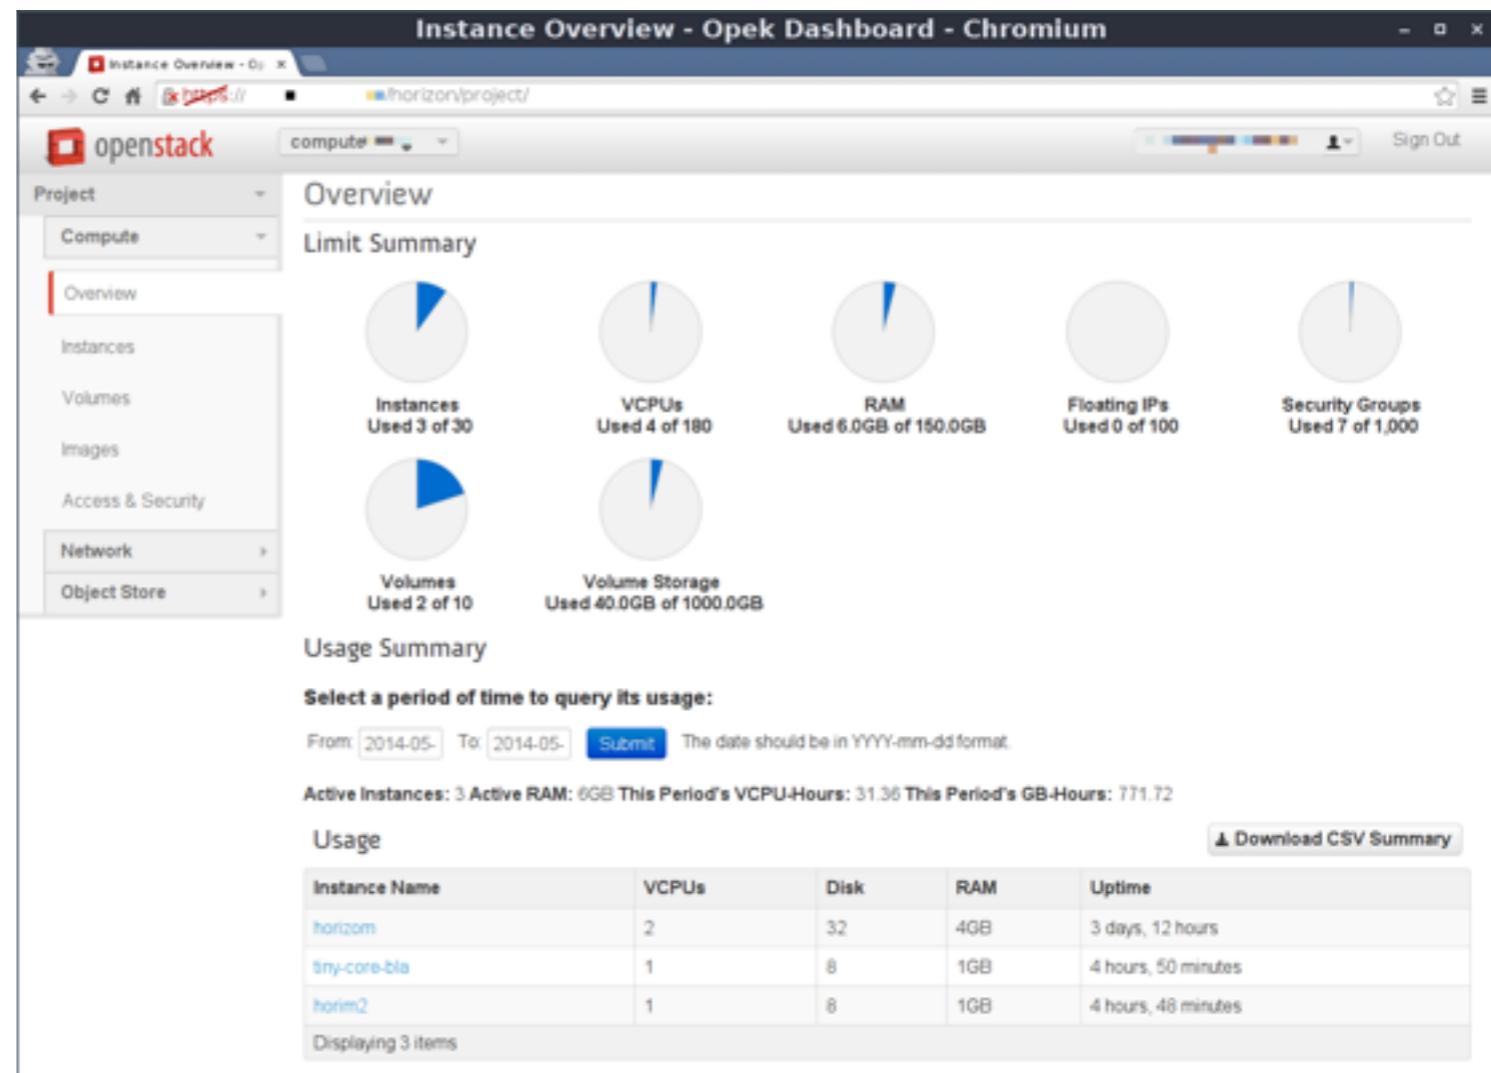


# HORIZON



# Dashboard

- Provides a web based user interface to OpenStack services (a Django web application)
- Three central dashboards, a “User Dashboard”, a “System Dashboard”, and a “Settings” dashboard.



# HORIZON



# Dashboard

- Provides a web based user interface to OpenStack services (a Django web application)
- Three central dashboards, a “User Dashboard”, a “System Dashboard”, and a “Settings” dashboard.

The screenshot shows the OpenStack User Dashboard. At the top, there's a navigation bar with the OpenStack logo, "USER DASHBOARD", "SYSTEM PANEL", and a user dropdown showing "admin as admin". Below the navigation is a sidebar titled "Manage Compute" with links for "Overview", "Instances", "Images", "Snapshots", "Keypairs", "Floating IPs", and "Security Groups". The main content area is titled "Overview" and contains three cards: "CPU" (5 CORES Active, 870.0 CPU-HR Used), "RAM" (4.0 GB Active), and "Disk" (20.0 GB Active, 2683.4 GB-HR Used). Below these cards is a section titled "Server Usage Summary" with a "Show Terminated" link and a "Download CSV" button. A table lists five servers with columns: ID, Name, User, VCPUs, Ram Size, Disk Size, Flavor, Uptime, and Status. The servers are:

ID	Name	User	VCPUs	Ram Size	Disk Size	Flavor	Uptime	Status
39	x	admin	1	512MB	0GB	m1.tiny	1 day	Active
40	somkinda	admin	1	512MB	0GB	m1.tiny	7 hours, 16 minutes	Active
41	test	admin	1	512MB	0GB	m1.tiny	4 hours, 47 minutes	Active
42	jakedevstack	admin	1	2GB	20GB	m1.small	15 hours, 56 minutes	Active
49	dt-vdi	admin	1	512MB	0GB	m1.tiny	20 hours, 40 minutes	Active

# KEYSTONE Authentification

- Keystone provides API client authentication, service discovery, and distributed multi-tenant authorization by implementing OpenStack's Identity API.
- It supports LDAP, OAuth, OpenID Connect, SAML and SQL.PIs: REST / one per service

# Nova Compute Service

- Implement services and associated libraries to provide massively scalable, on demand, self service access to compute instances
- Nova supports creating virtual machines and baremetal servers, through the use of Ironic.

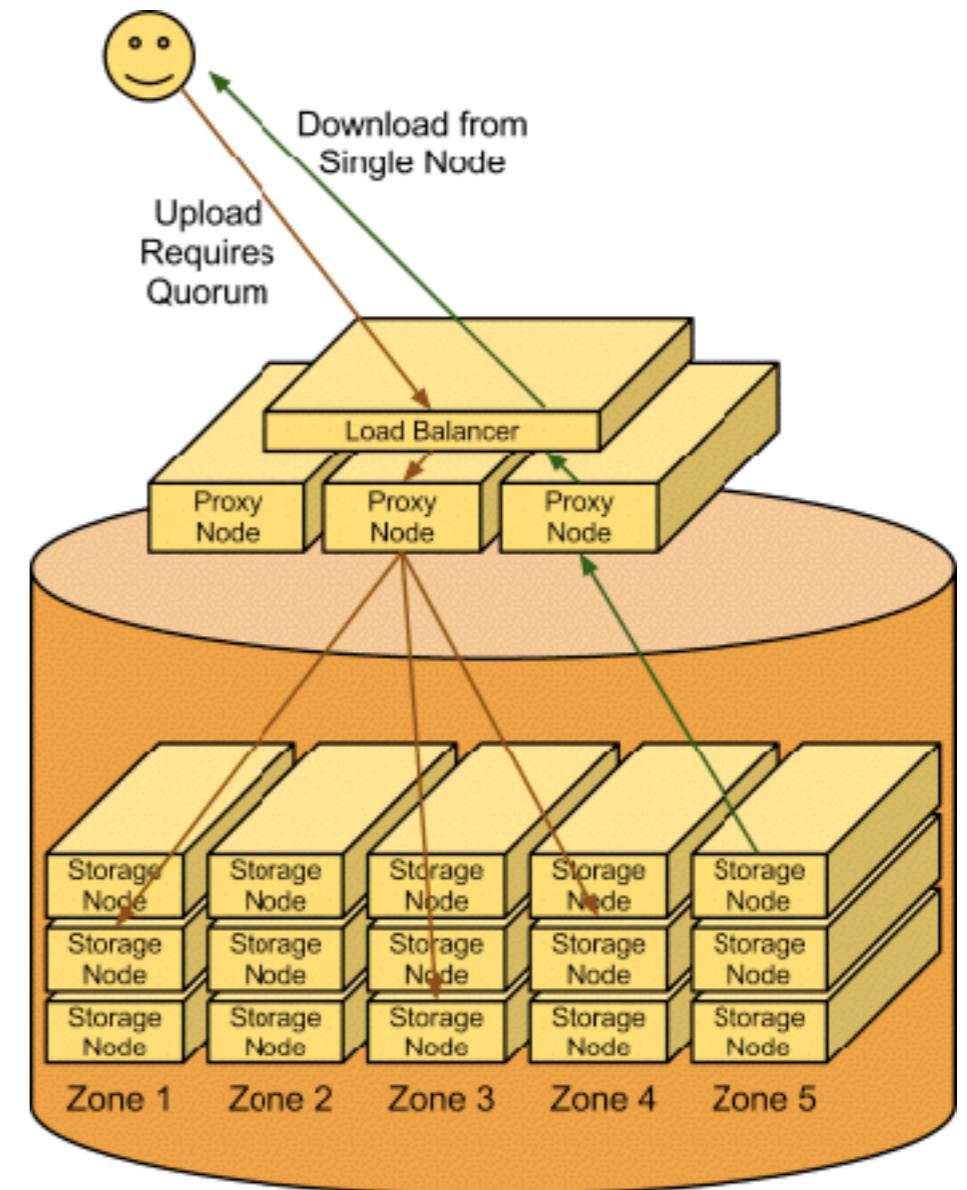
More details later...

# Glance Image Service

- VM images contain a virtual disk that holds a bootable operating system on it.
- Each launched instance runs from a copy of the base image. Any changes made to the instance do not affect the base image. Snapshots capture the state of an instances running disk.
- Users can create a snapshot, and build a new image based on these snapshots
- Glance has a RESTful API that allows querying of VM image metadata as well as retrieval of the actual image.
- VM images made available through Glance can be stored in a variety of locations from simple filesystems to object-storage systems like the OpenStack Swift project.

# Swift Object Store

- Swift is a highly available, distributed, eventually consistent object/blob store (a S3-like system)
- Organizations can use Swift to store lots of data efficiently, safely, and cheaply. It's built for scale and optimized for durability, availability, and concurrency across the entire data set.
- Swift is ideal for storing unstructured data that can grow without bound.
- The only service that does not leverage a central DB



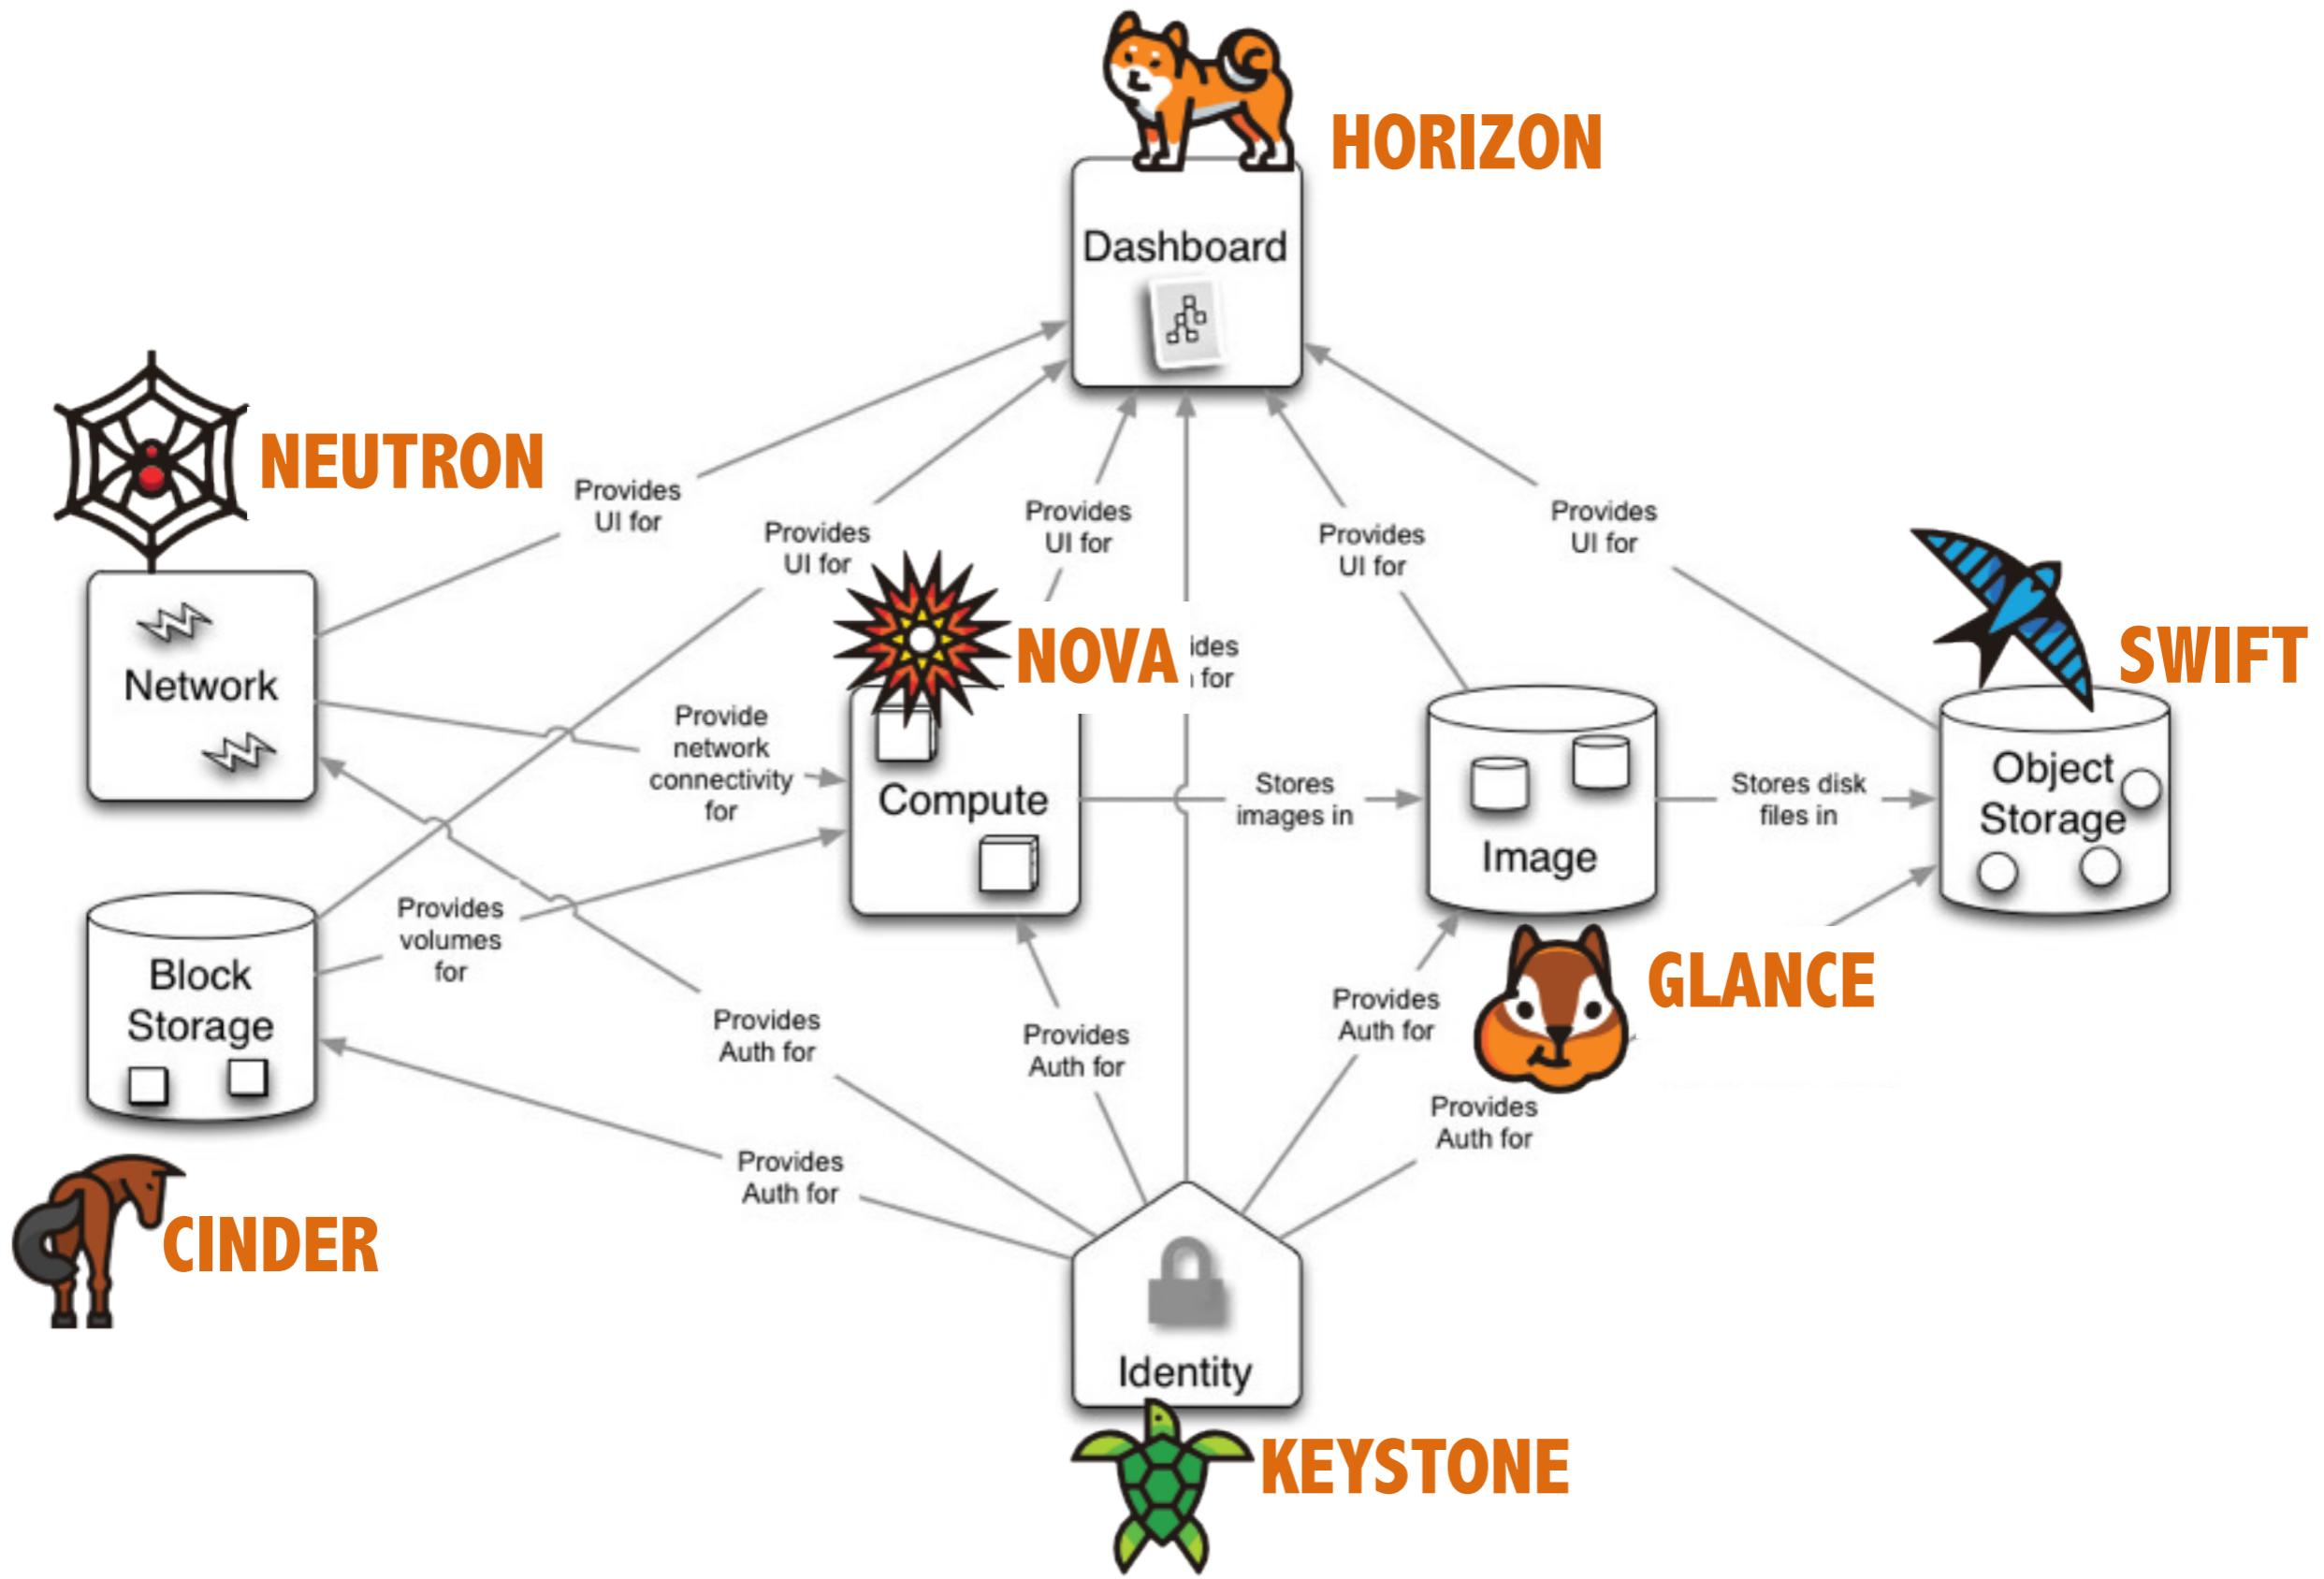
# Cinder Block Storage

- You can add and remove additional resources from running instances, such as persistent volume storage
- The Cinder-volume service provides persistent block storage, instead of the ephemeral storage provided by the base image (i.e. the Glance one)
- Cinder virtualizes the management of block storage devices and provides end users with a self service API to request and consume those resources without requiring any knowledge of where their storage is actually deployed or on what type of device.
- This is done through the use of either a reference implementation (LVM) or plugin drivers for other storage.

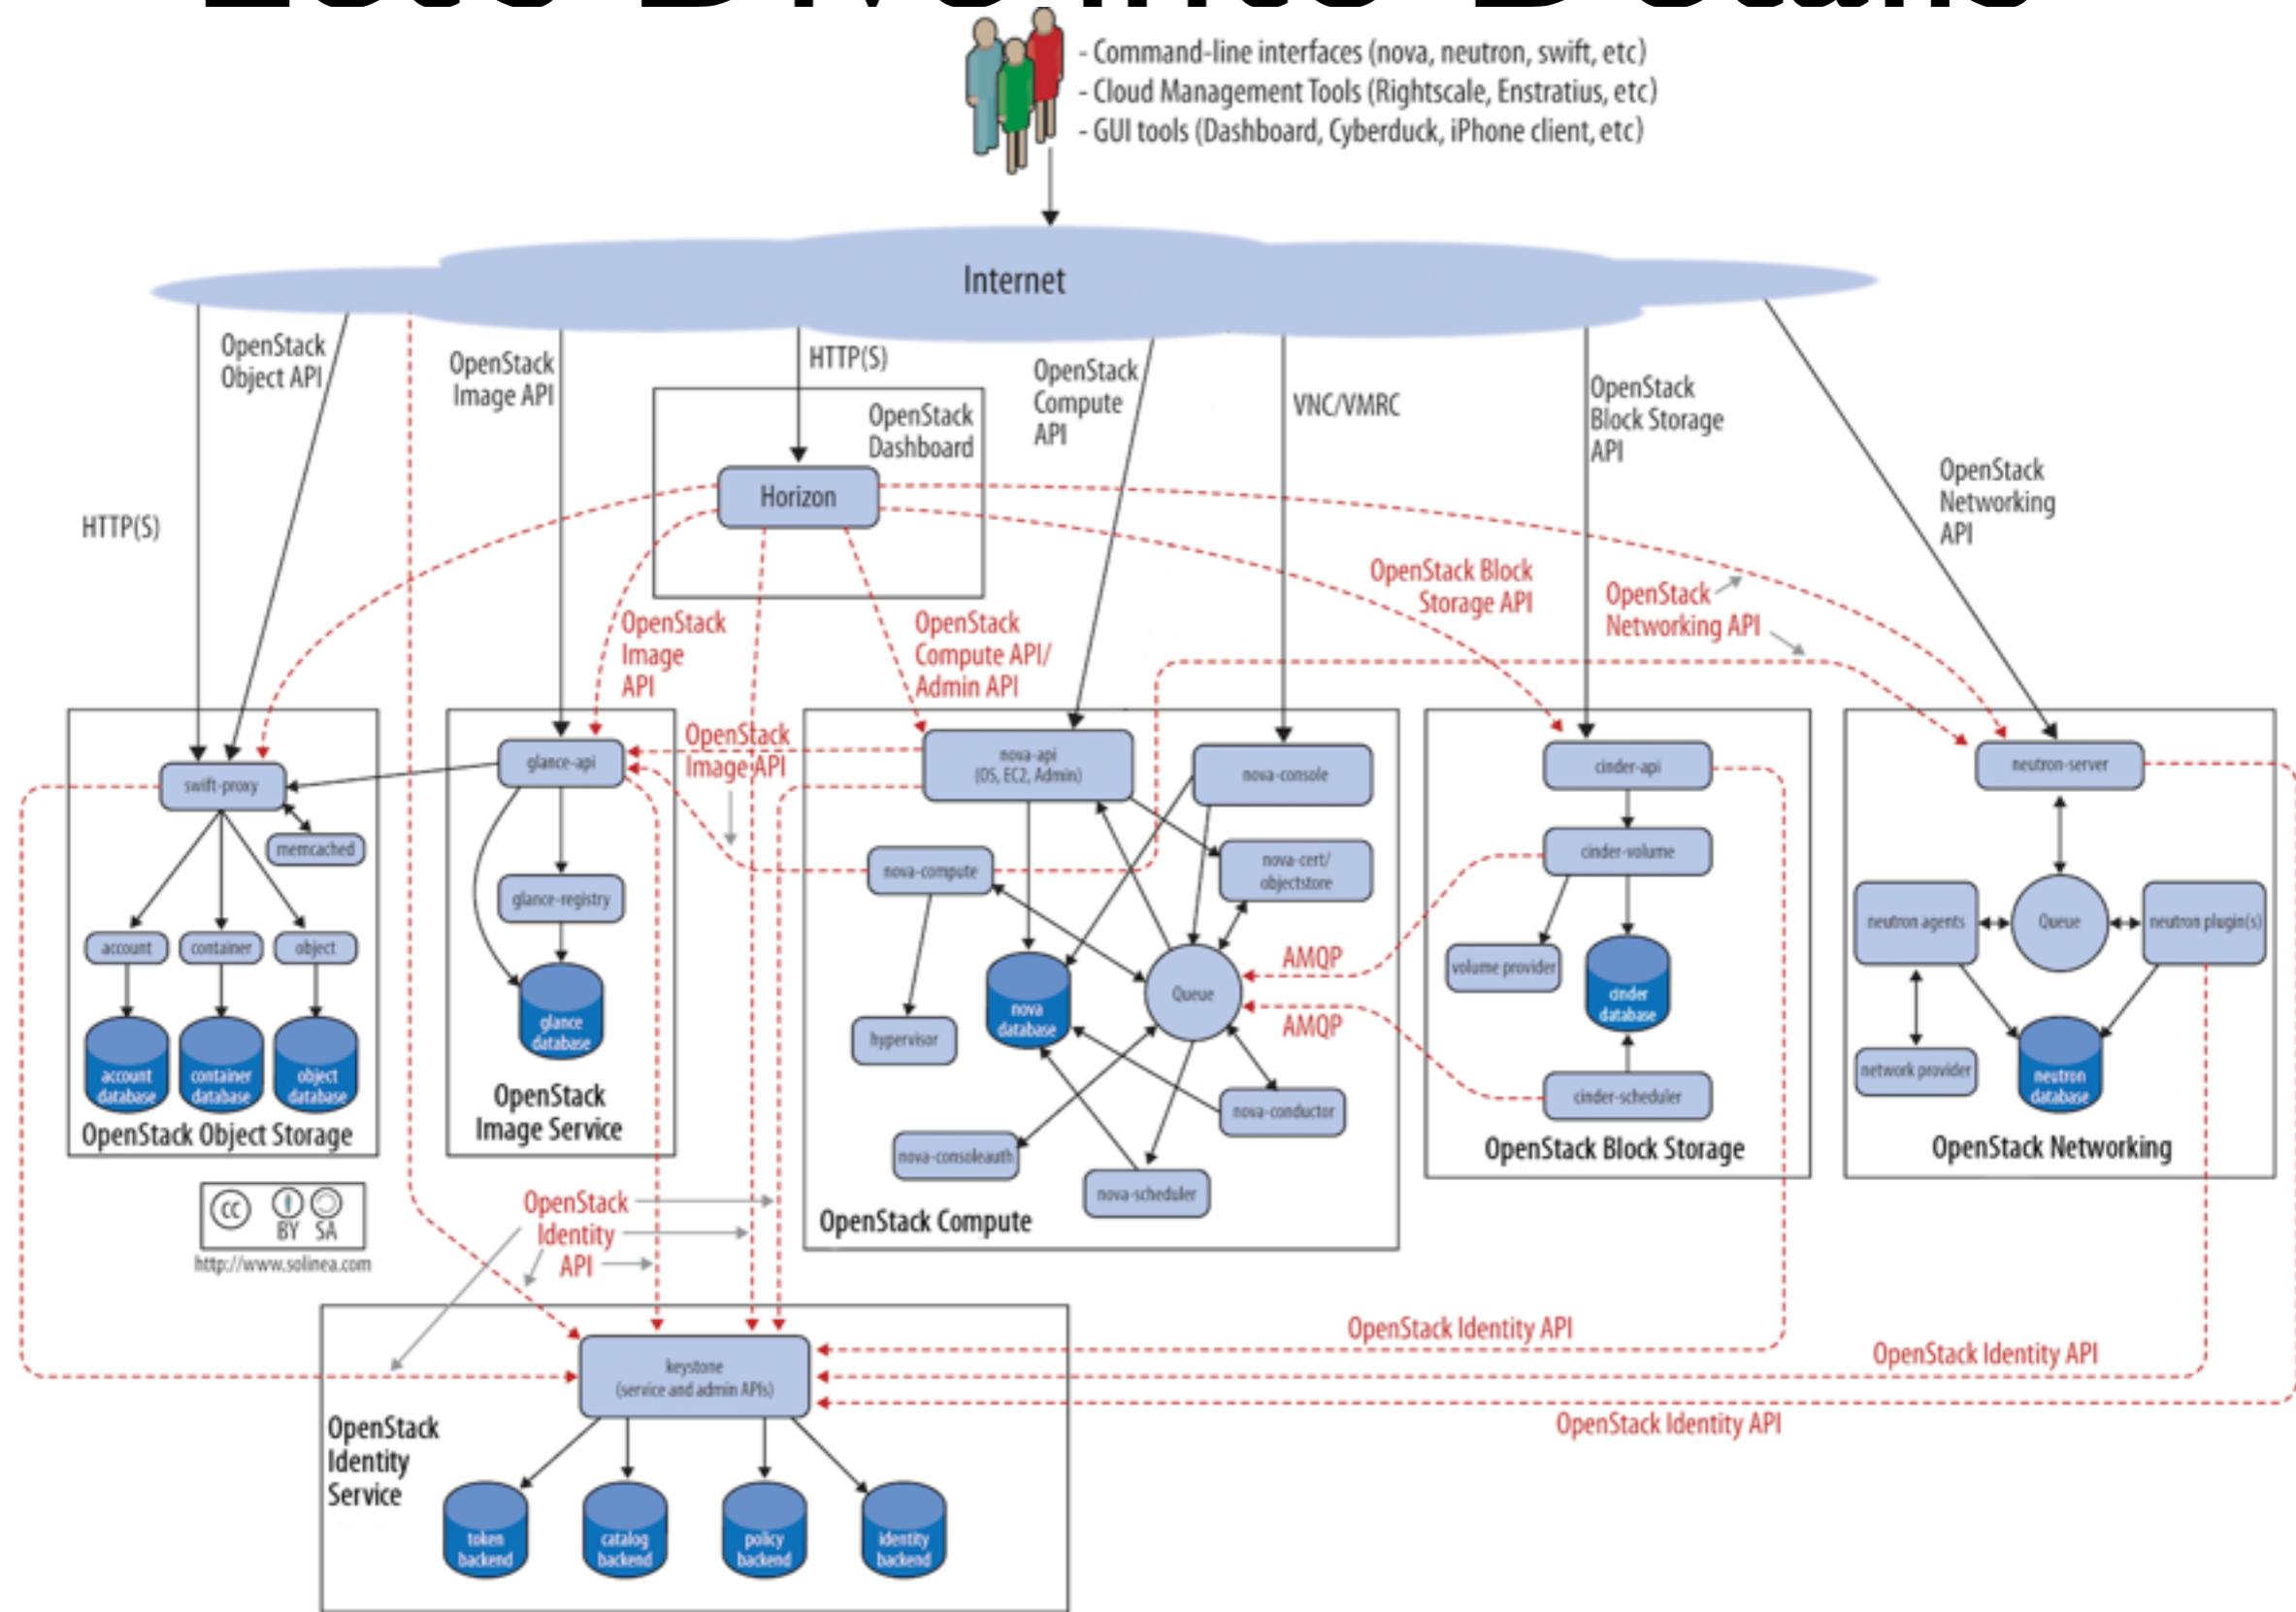
# Neutron Networking

- OpenStack Neutron is an SDN networking project focused on delivering networking-as-a-service (NaaS) in virtual compute environments.
- Composed of several sub-services
  - Neutron-server: API service
  - Agent DHCP: DHCP service for instances
  - Agent L3 : Routing service
- In addition to elementary services (L2/L3), Neutron provides additional mechanisms (load-balancing, firewalls, VPN...)
- A lot of plugins (LinuxBridge default one)

# Let's Dive into Details



# Let's Dive into Details

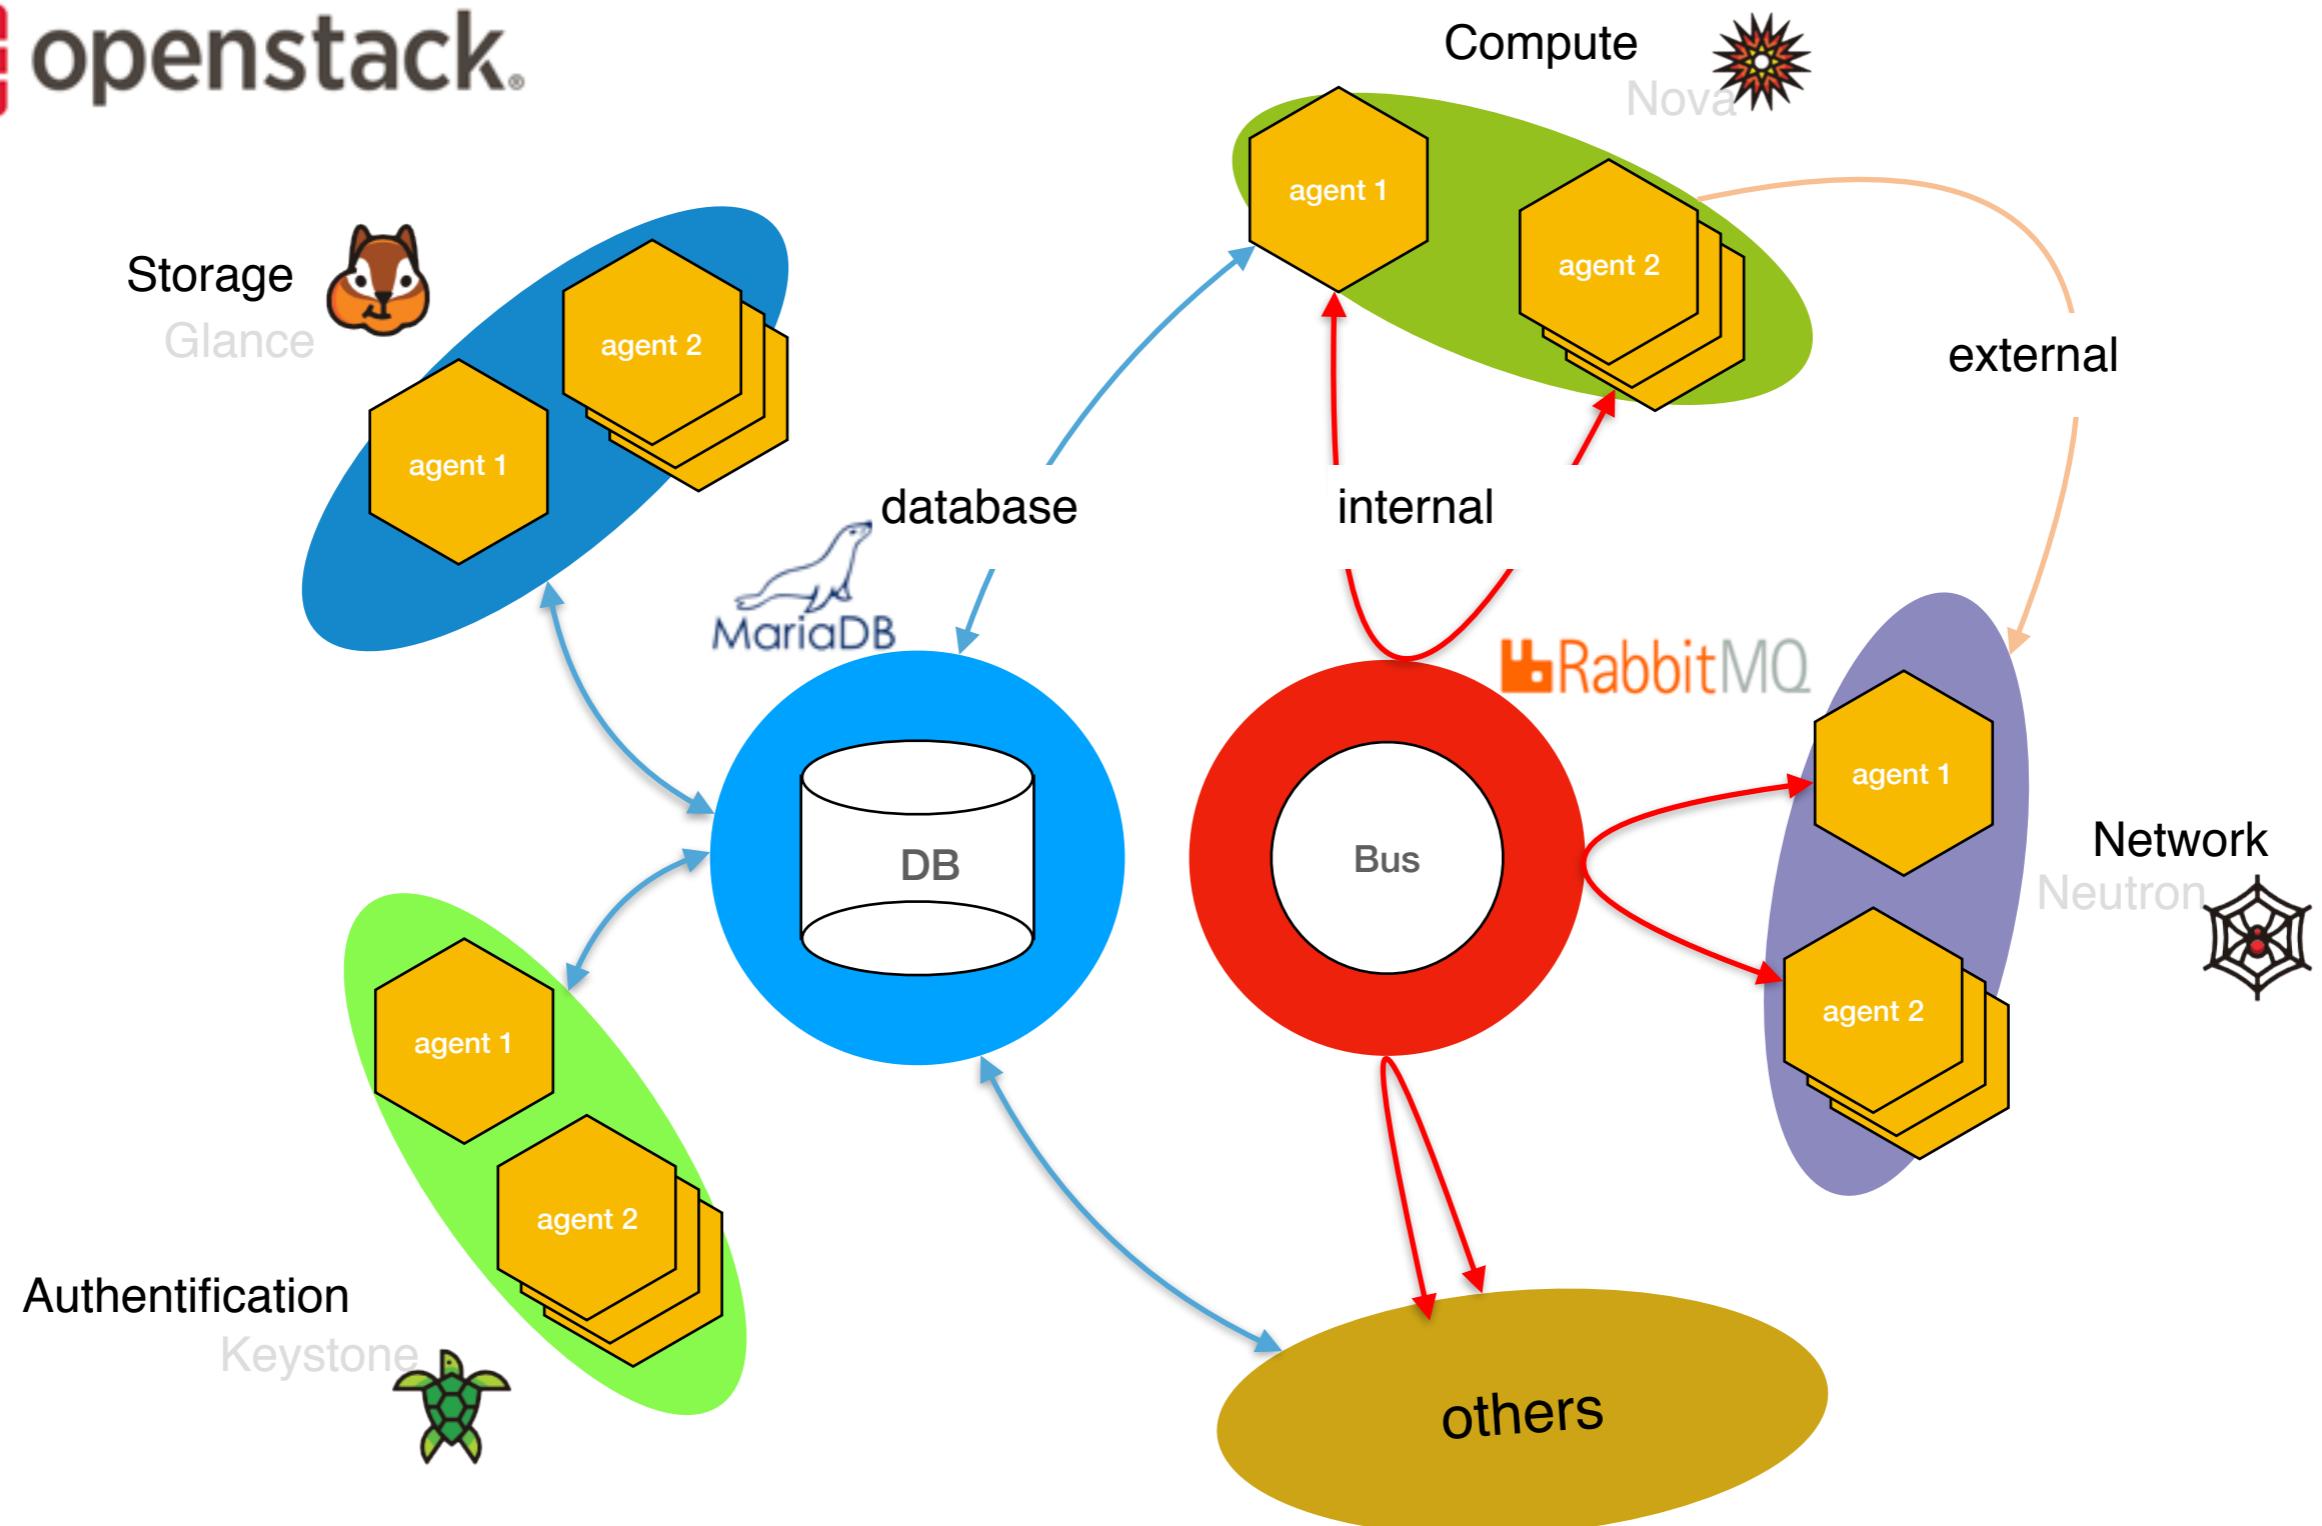




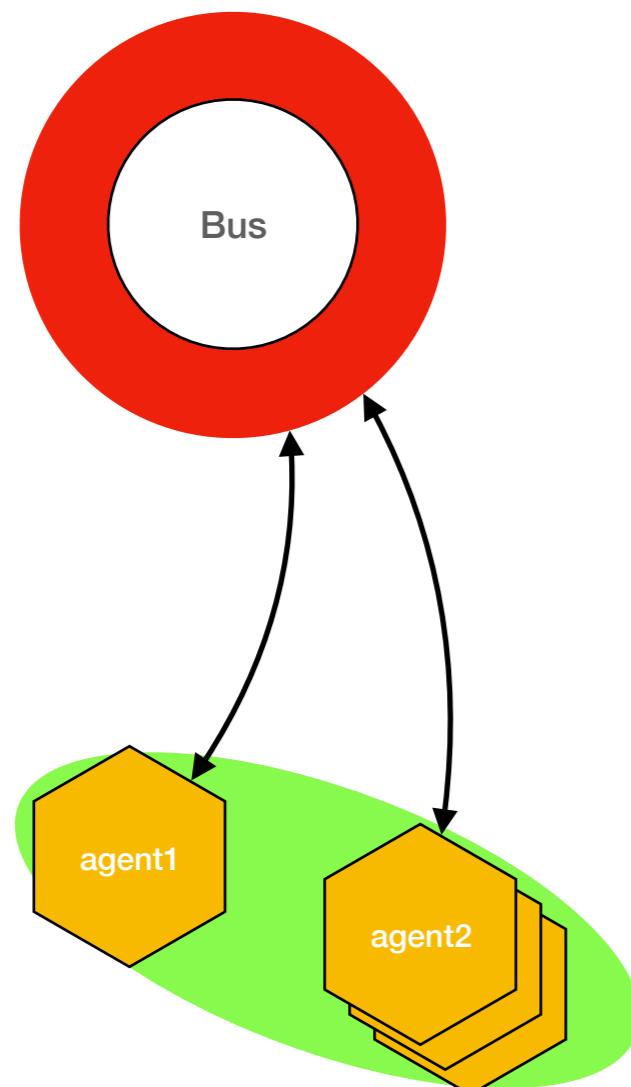
# DBs/Communication Bus

RabbitMQ

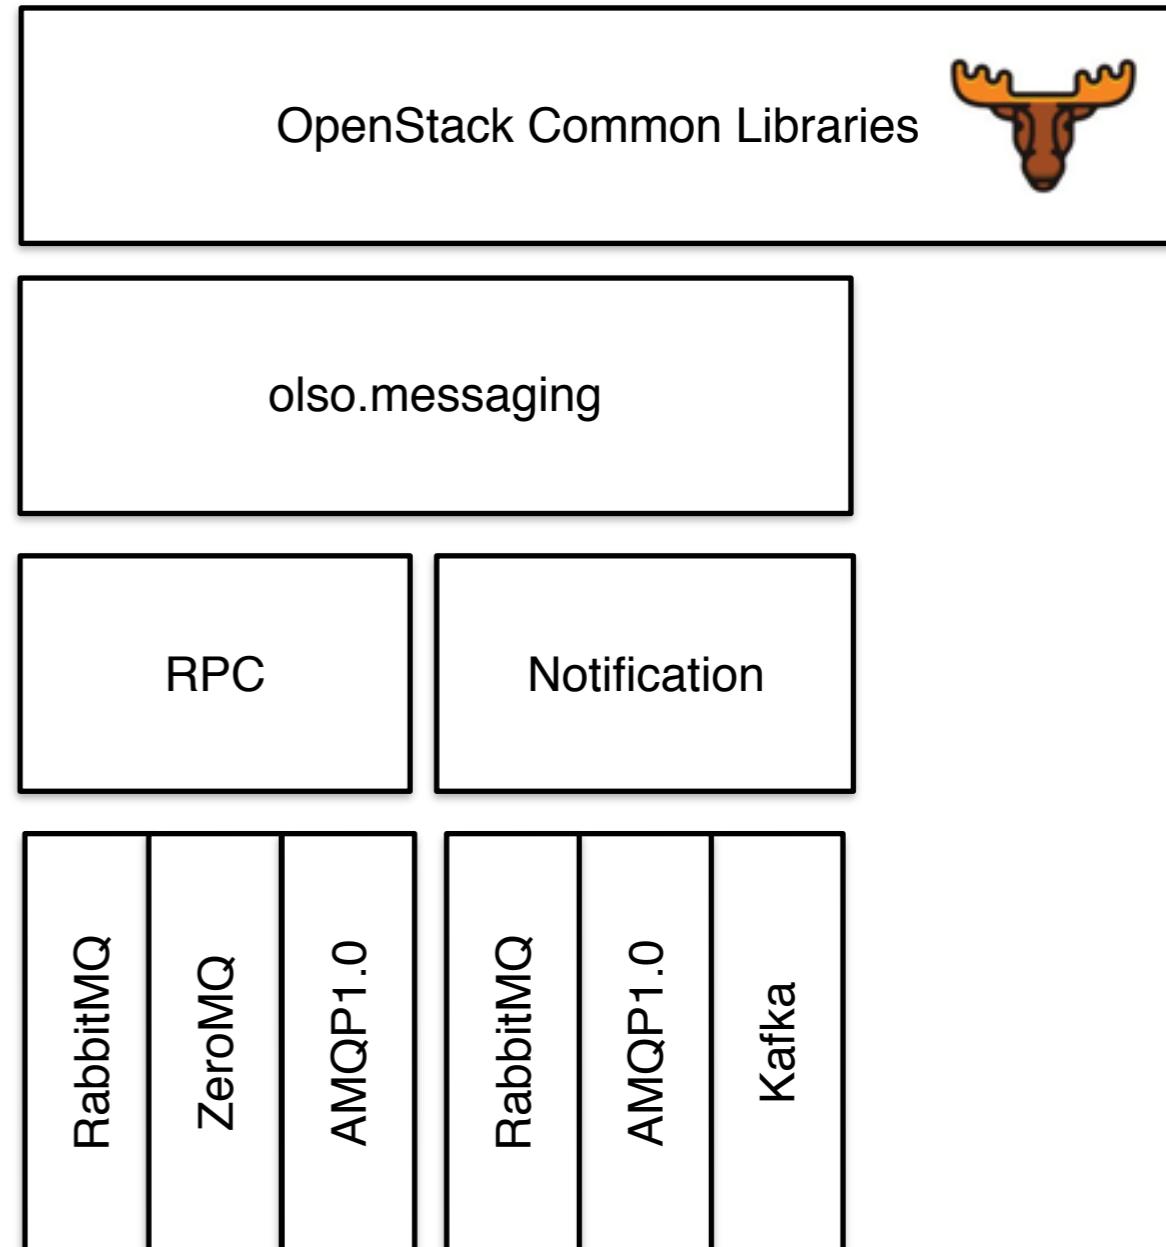
openstack.



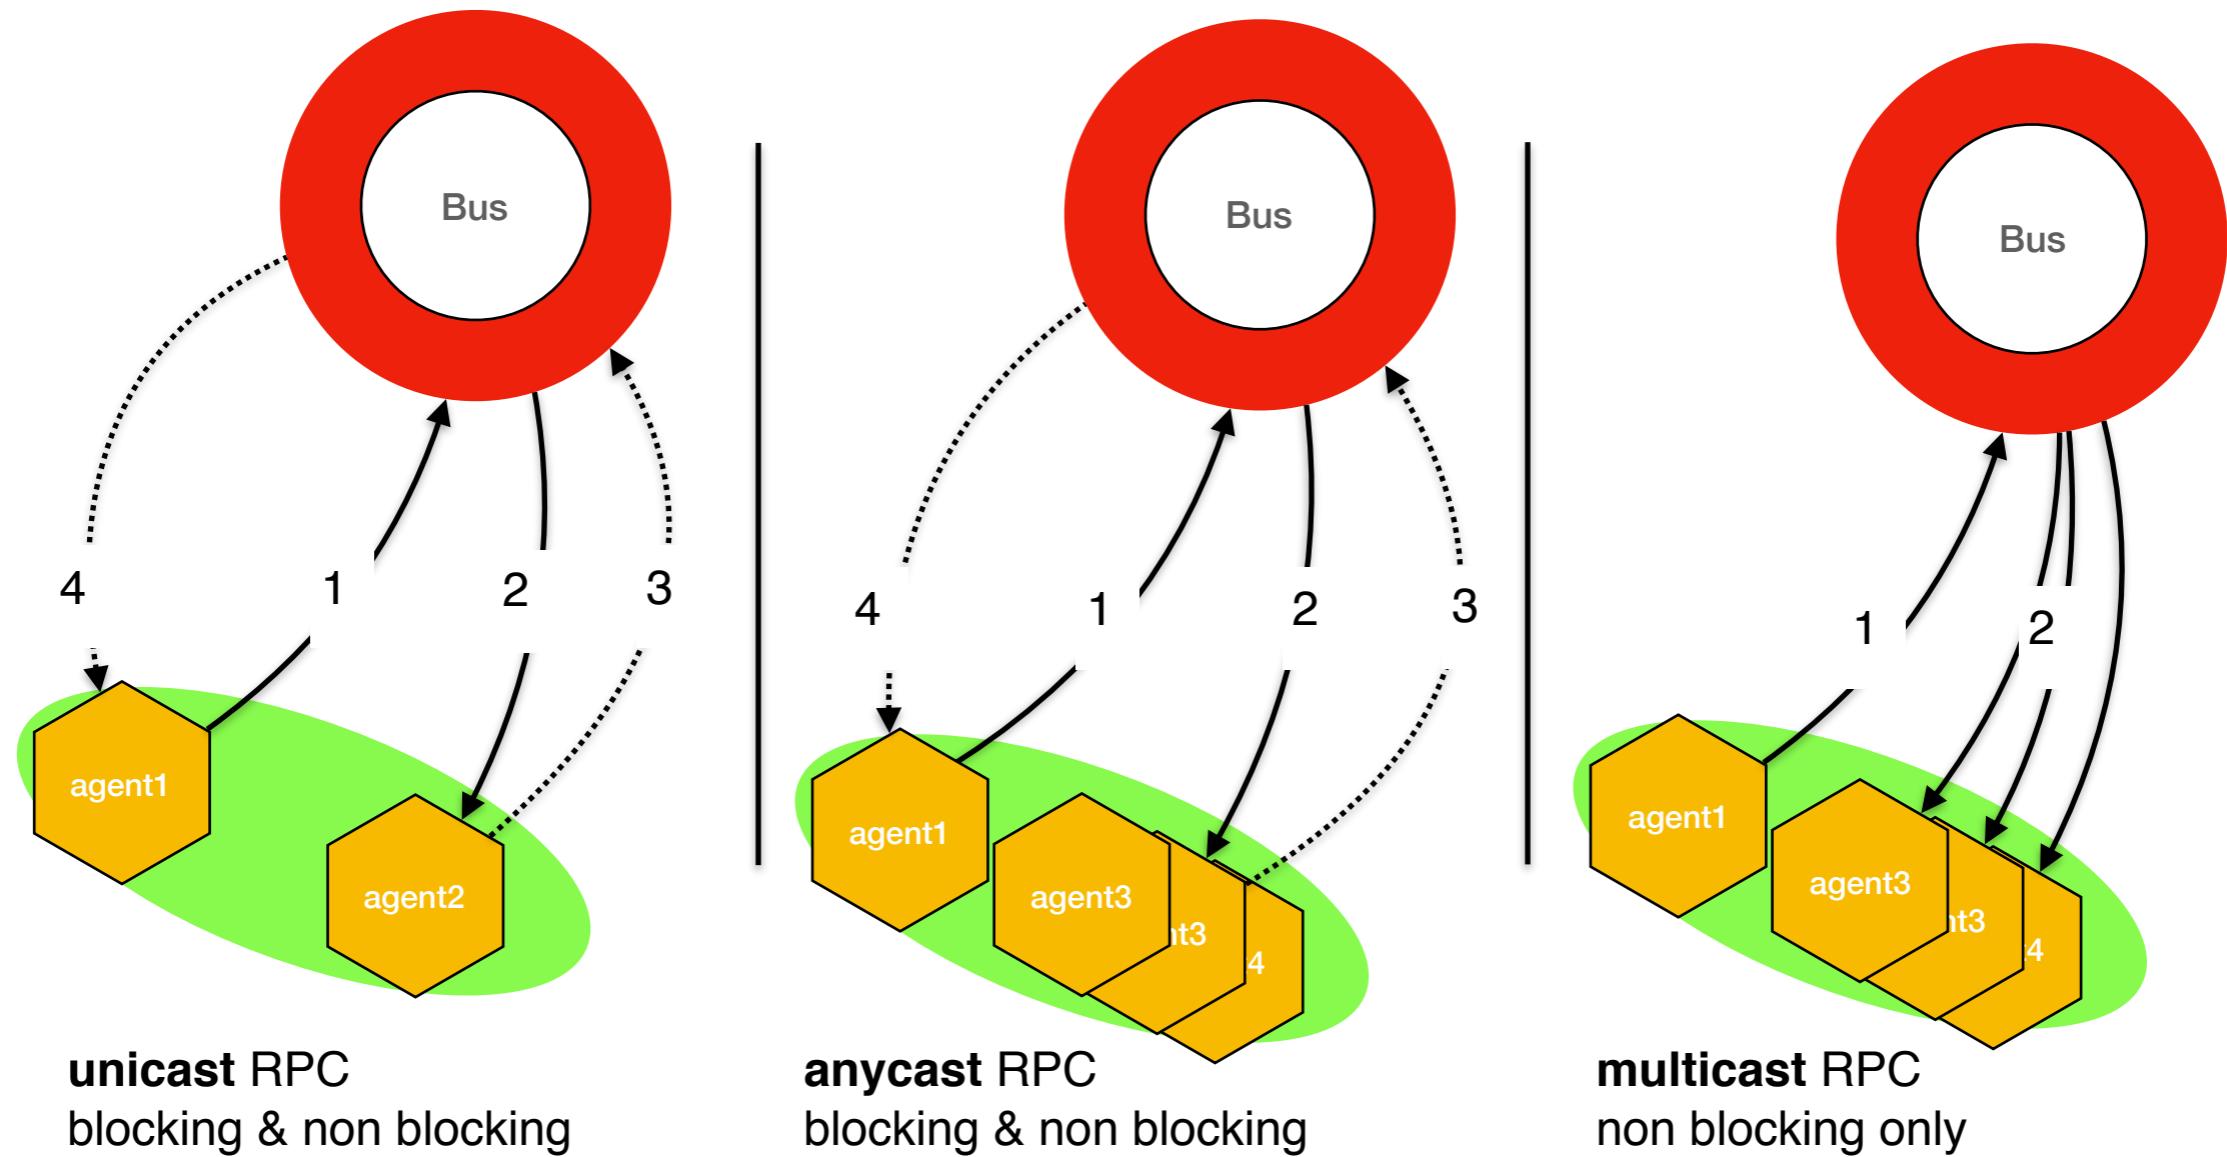
# DBs/Communication Bus



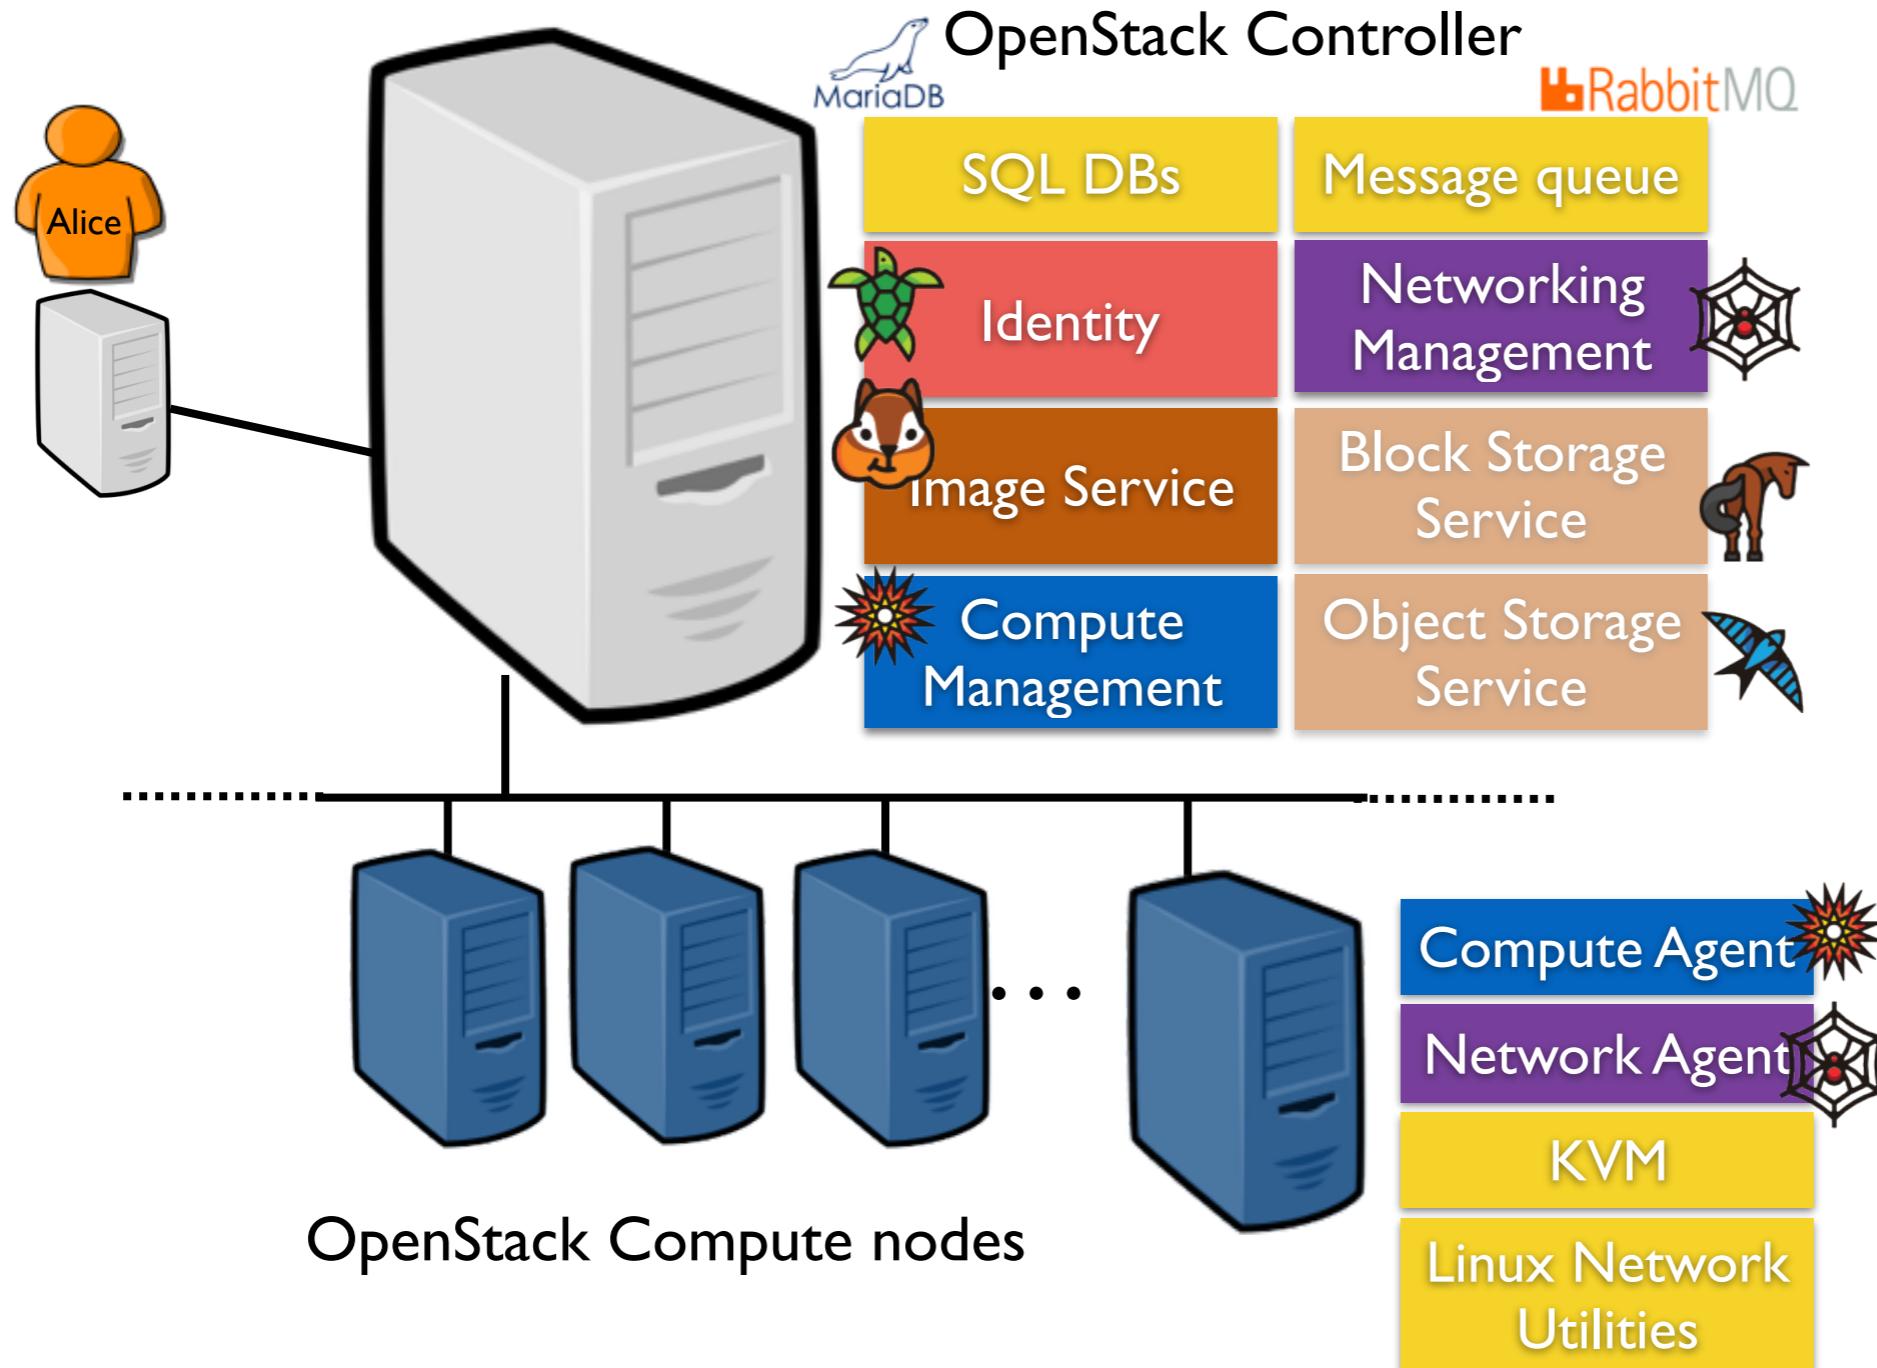
common library



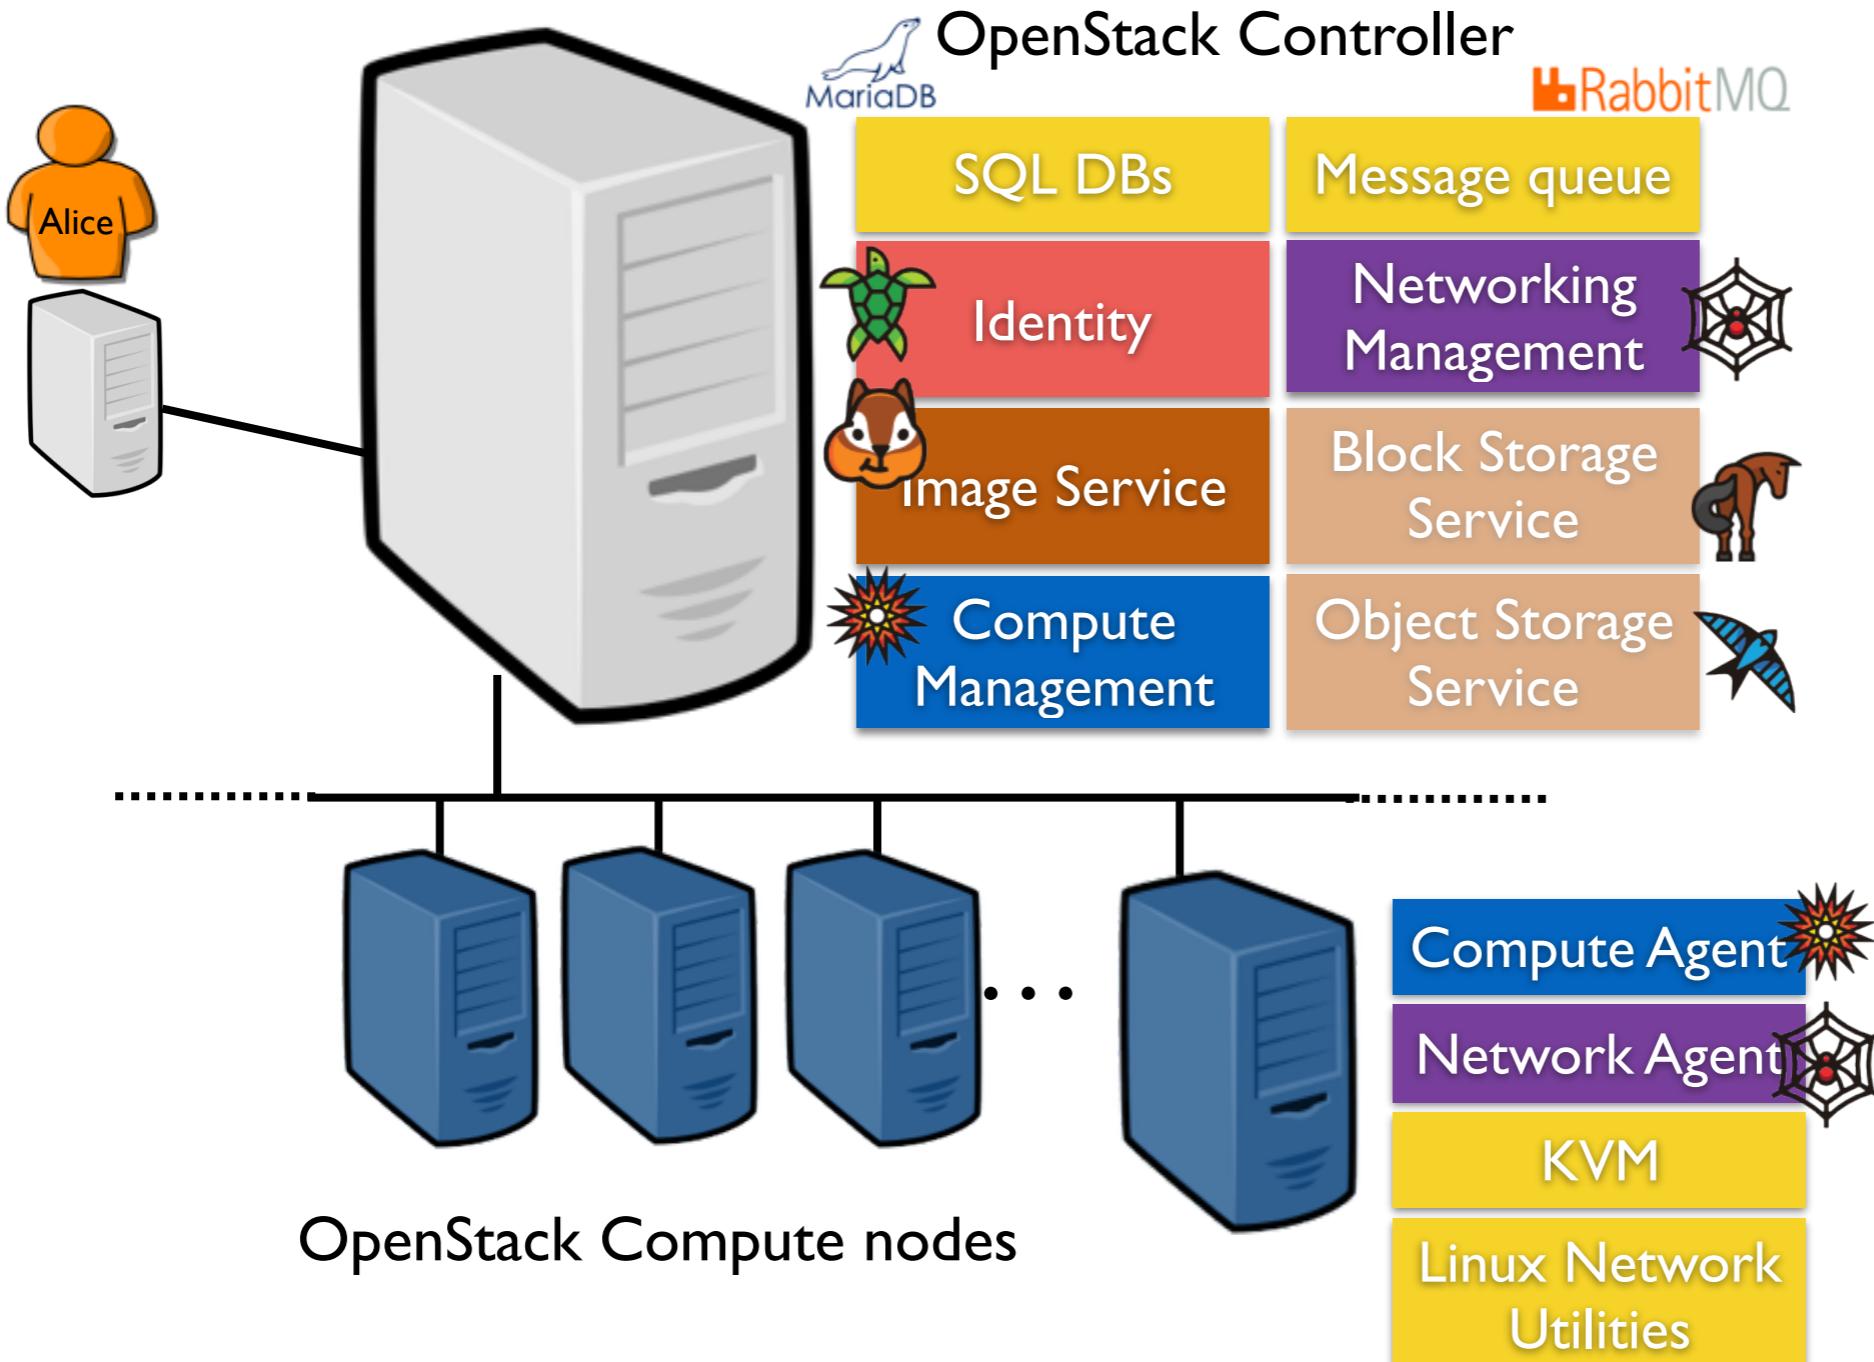
# DBs/Communication Bus



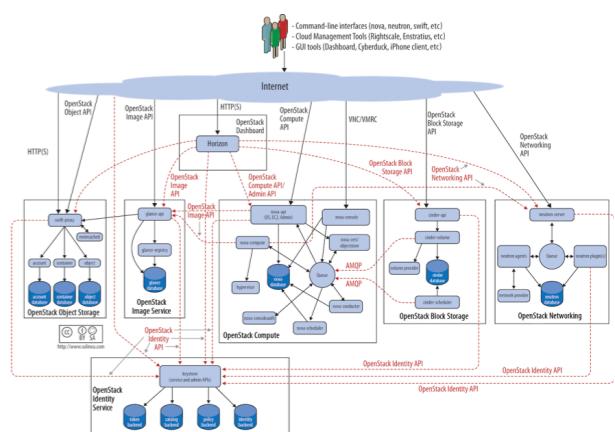
# Physically, How This Looks Like?



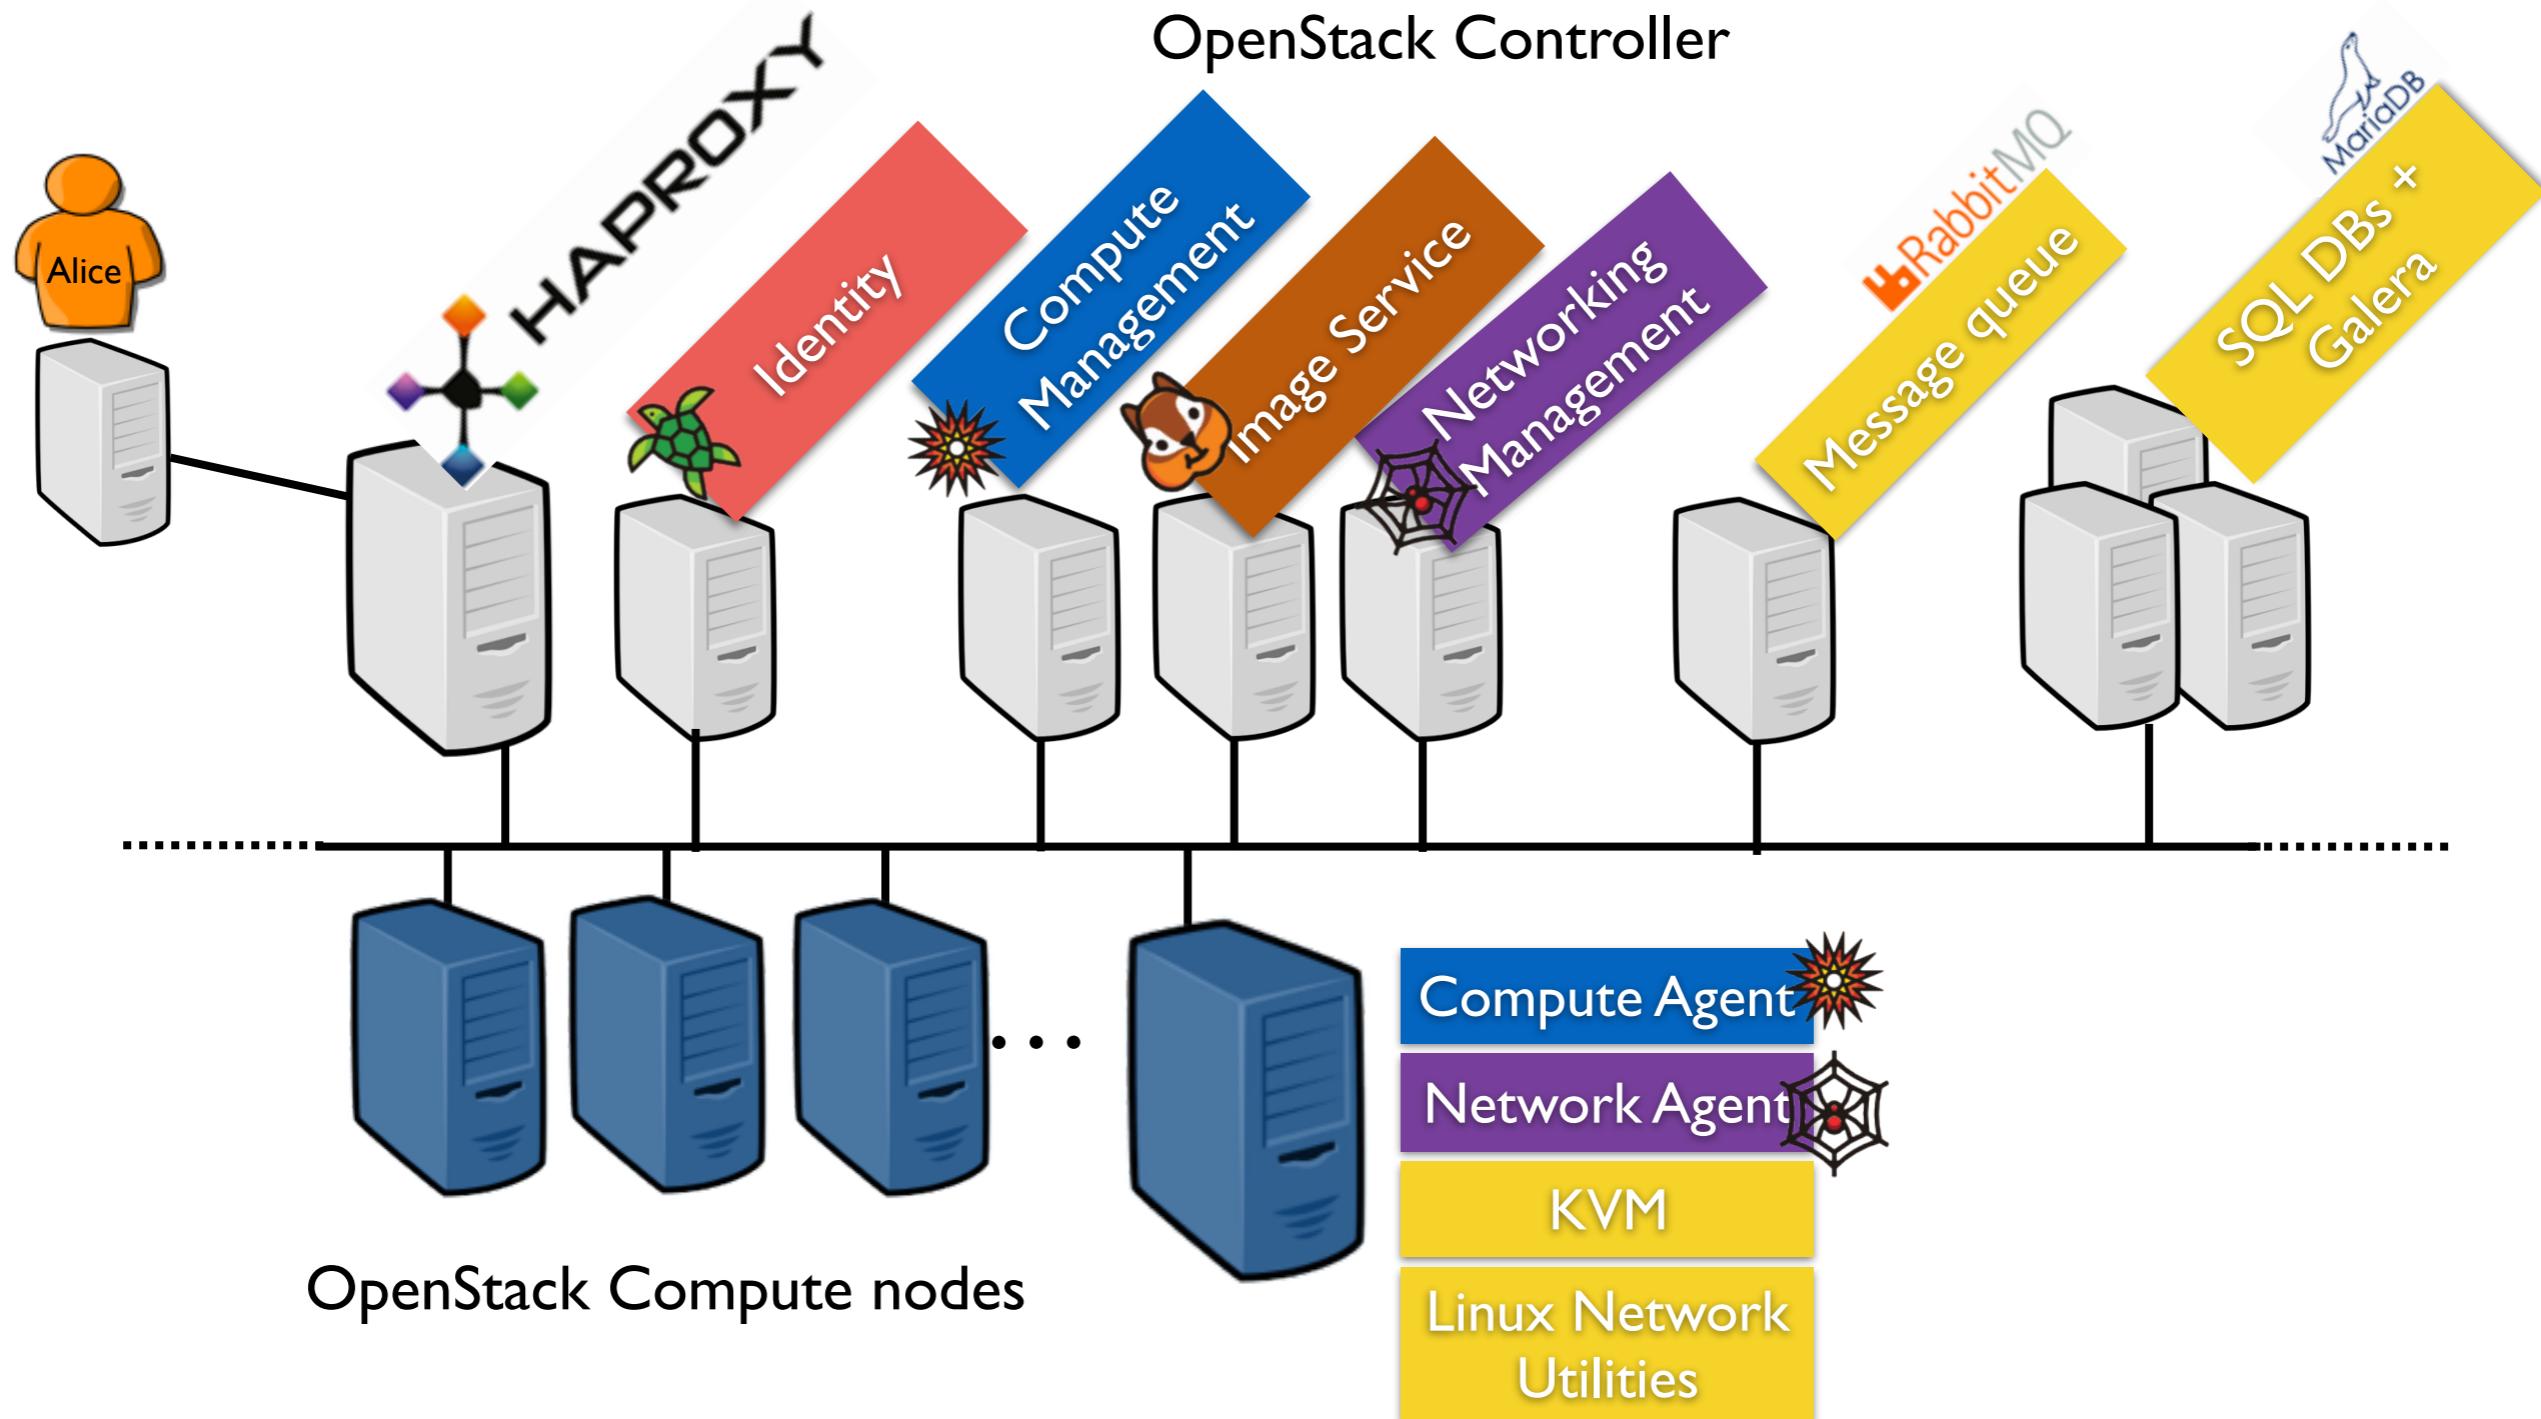
# Physically, How This Looks Like?



# A First Simplified View



# Physically, How This Looks Like?



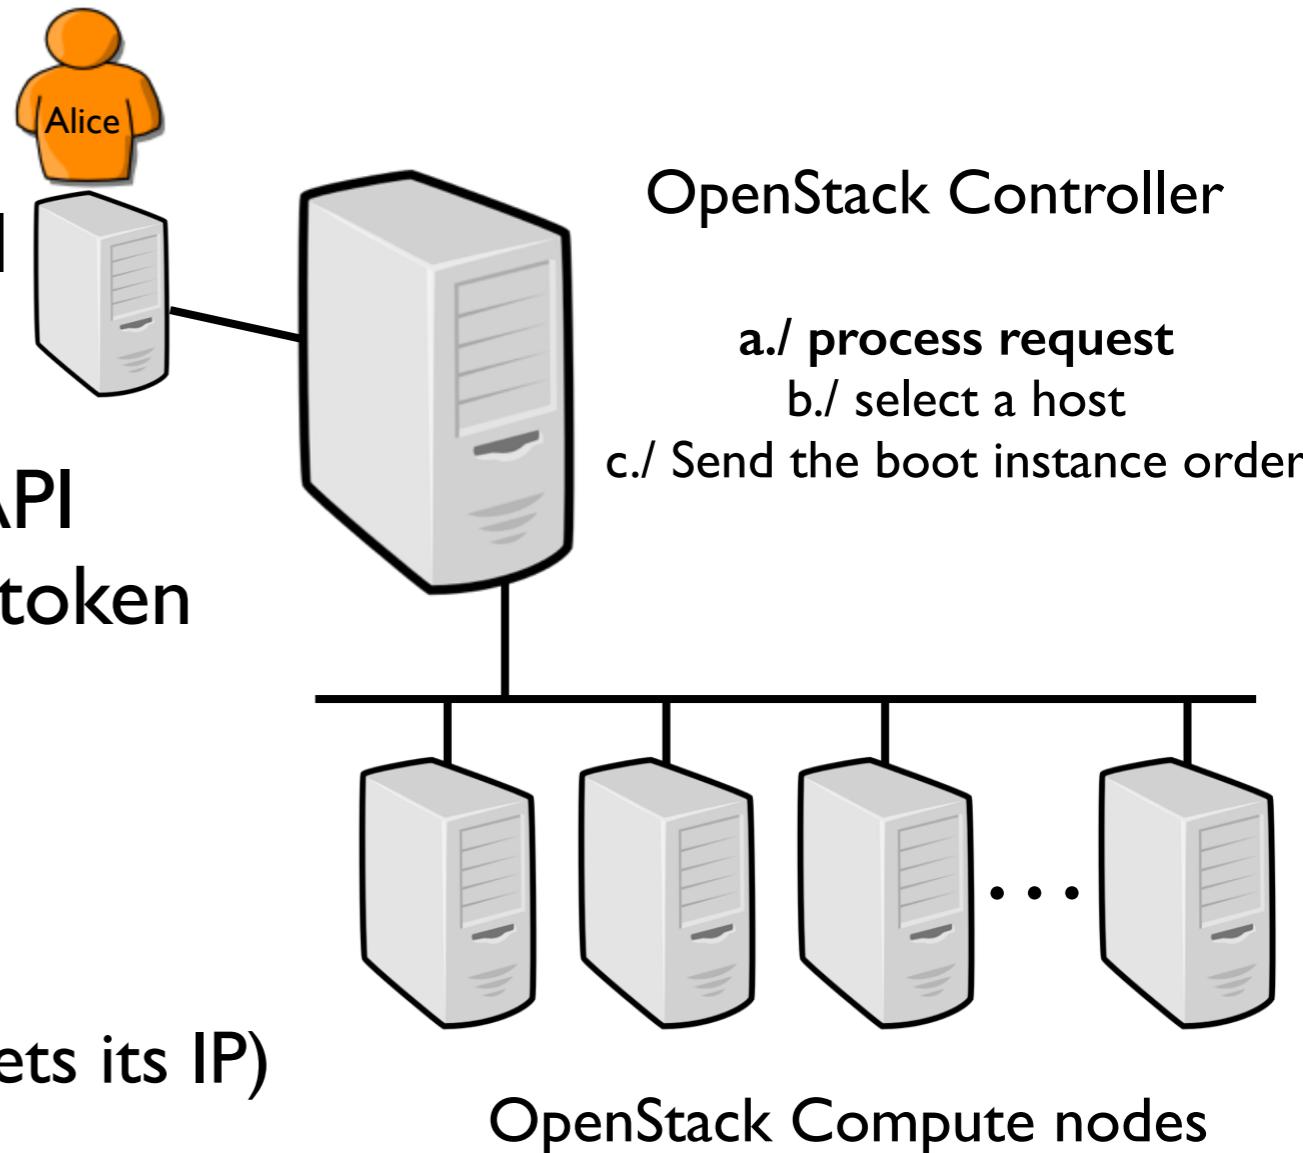
# Walkthrough of a typical Nova boot request

- Make a request to the Identity API (Keystone) to get an access token
- Make a request to the Compute API (Nova) using the aforementioned token and mentioning

The flavor (VM properties)

The image (a bootable image)

The network ID (i.e. where the VM gets its IP)



```
vagrant@enos-node:/opt/enos$ openstack server create \
--flavor m1.tiny \
--image cirros.uec \
--nic net-id=$(openstack network show private --column id --format value) \
cli-vm
```

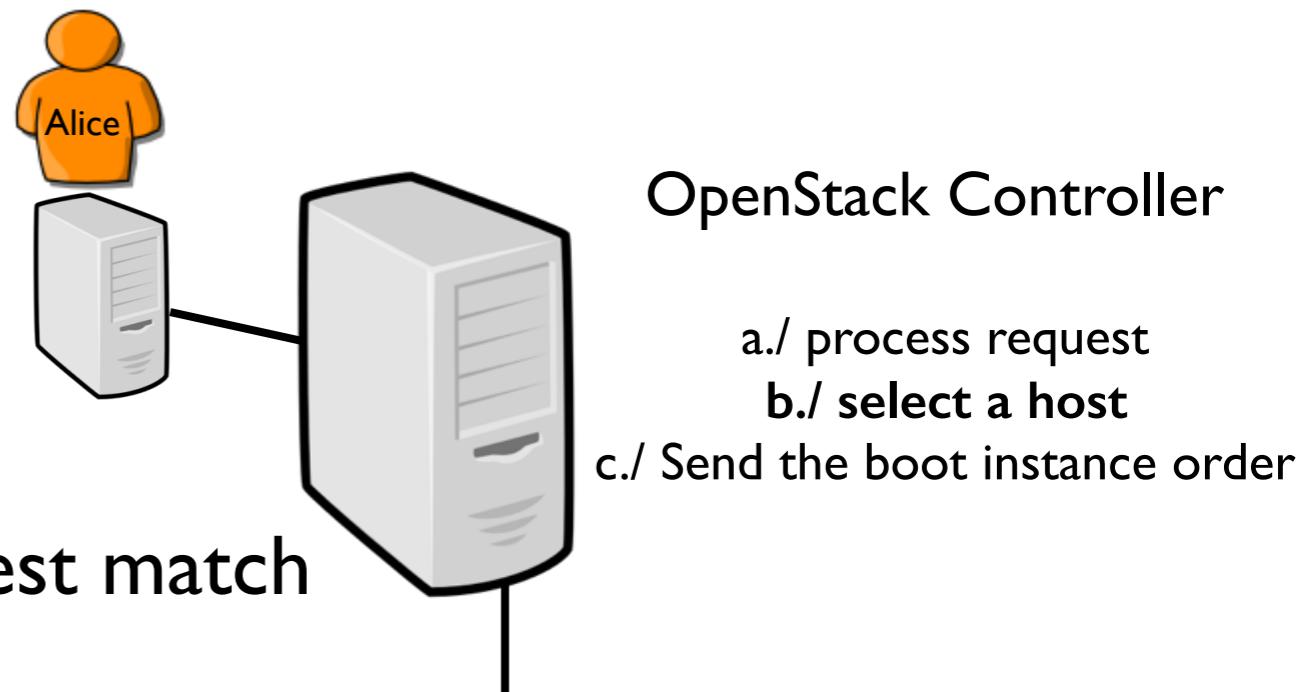
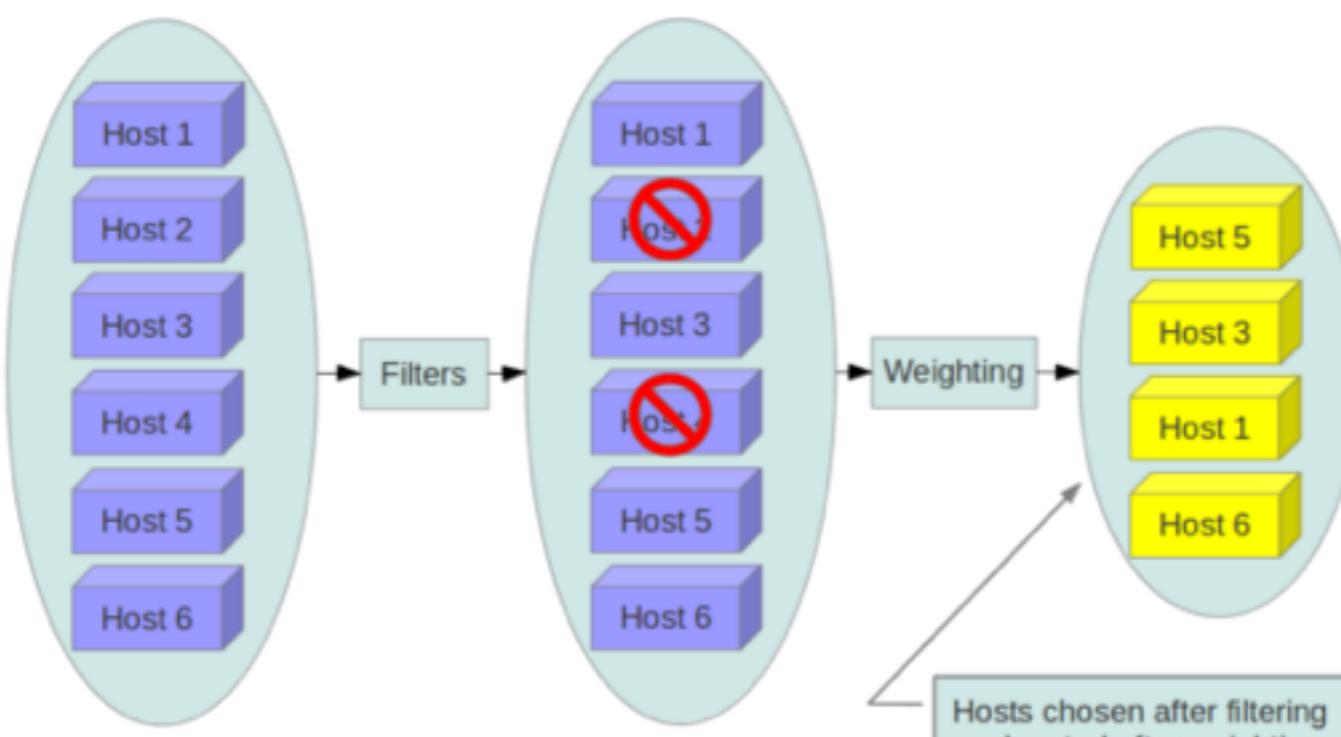
# Walkthrough of a typical Nova boot request

- Select a host (scheduler)

Ask Placement for possible hosts  
(placement API)

Apply filters

Sort hosts with weighers and take best match



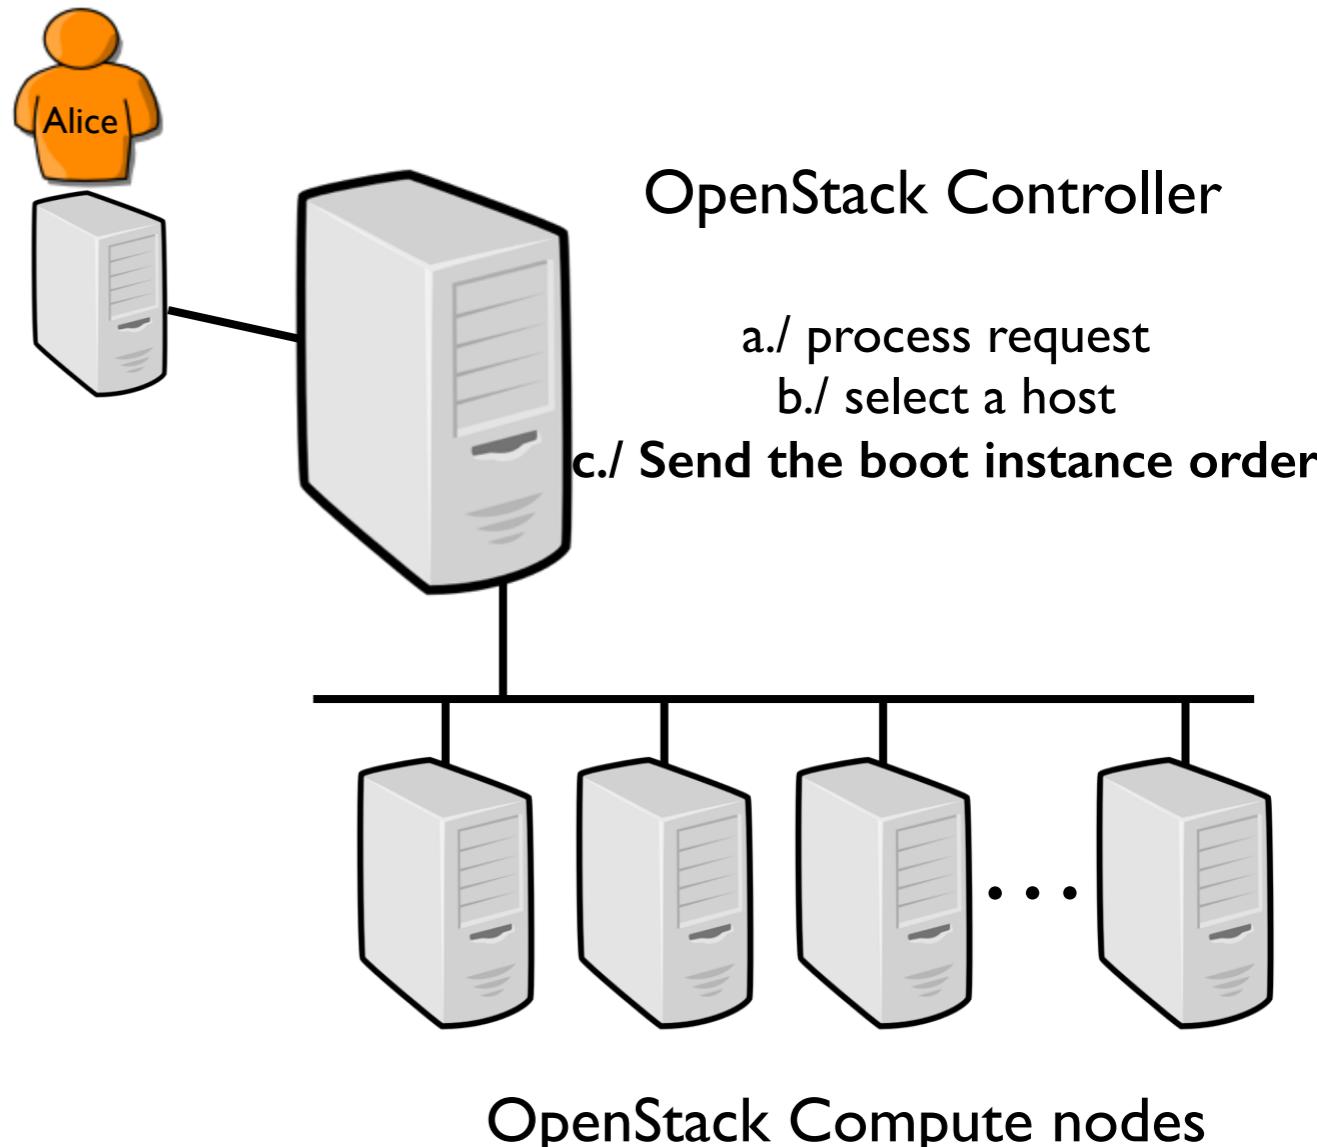
# Walkthrough of a typical Nova boot request

- When the compute node receives the request :

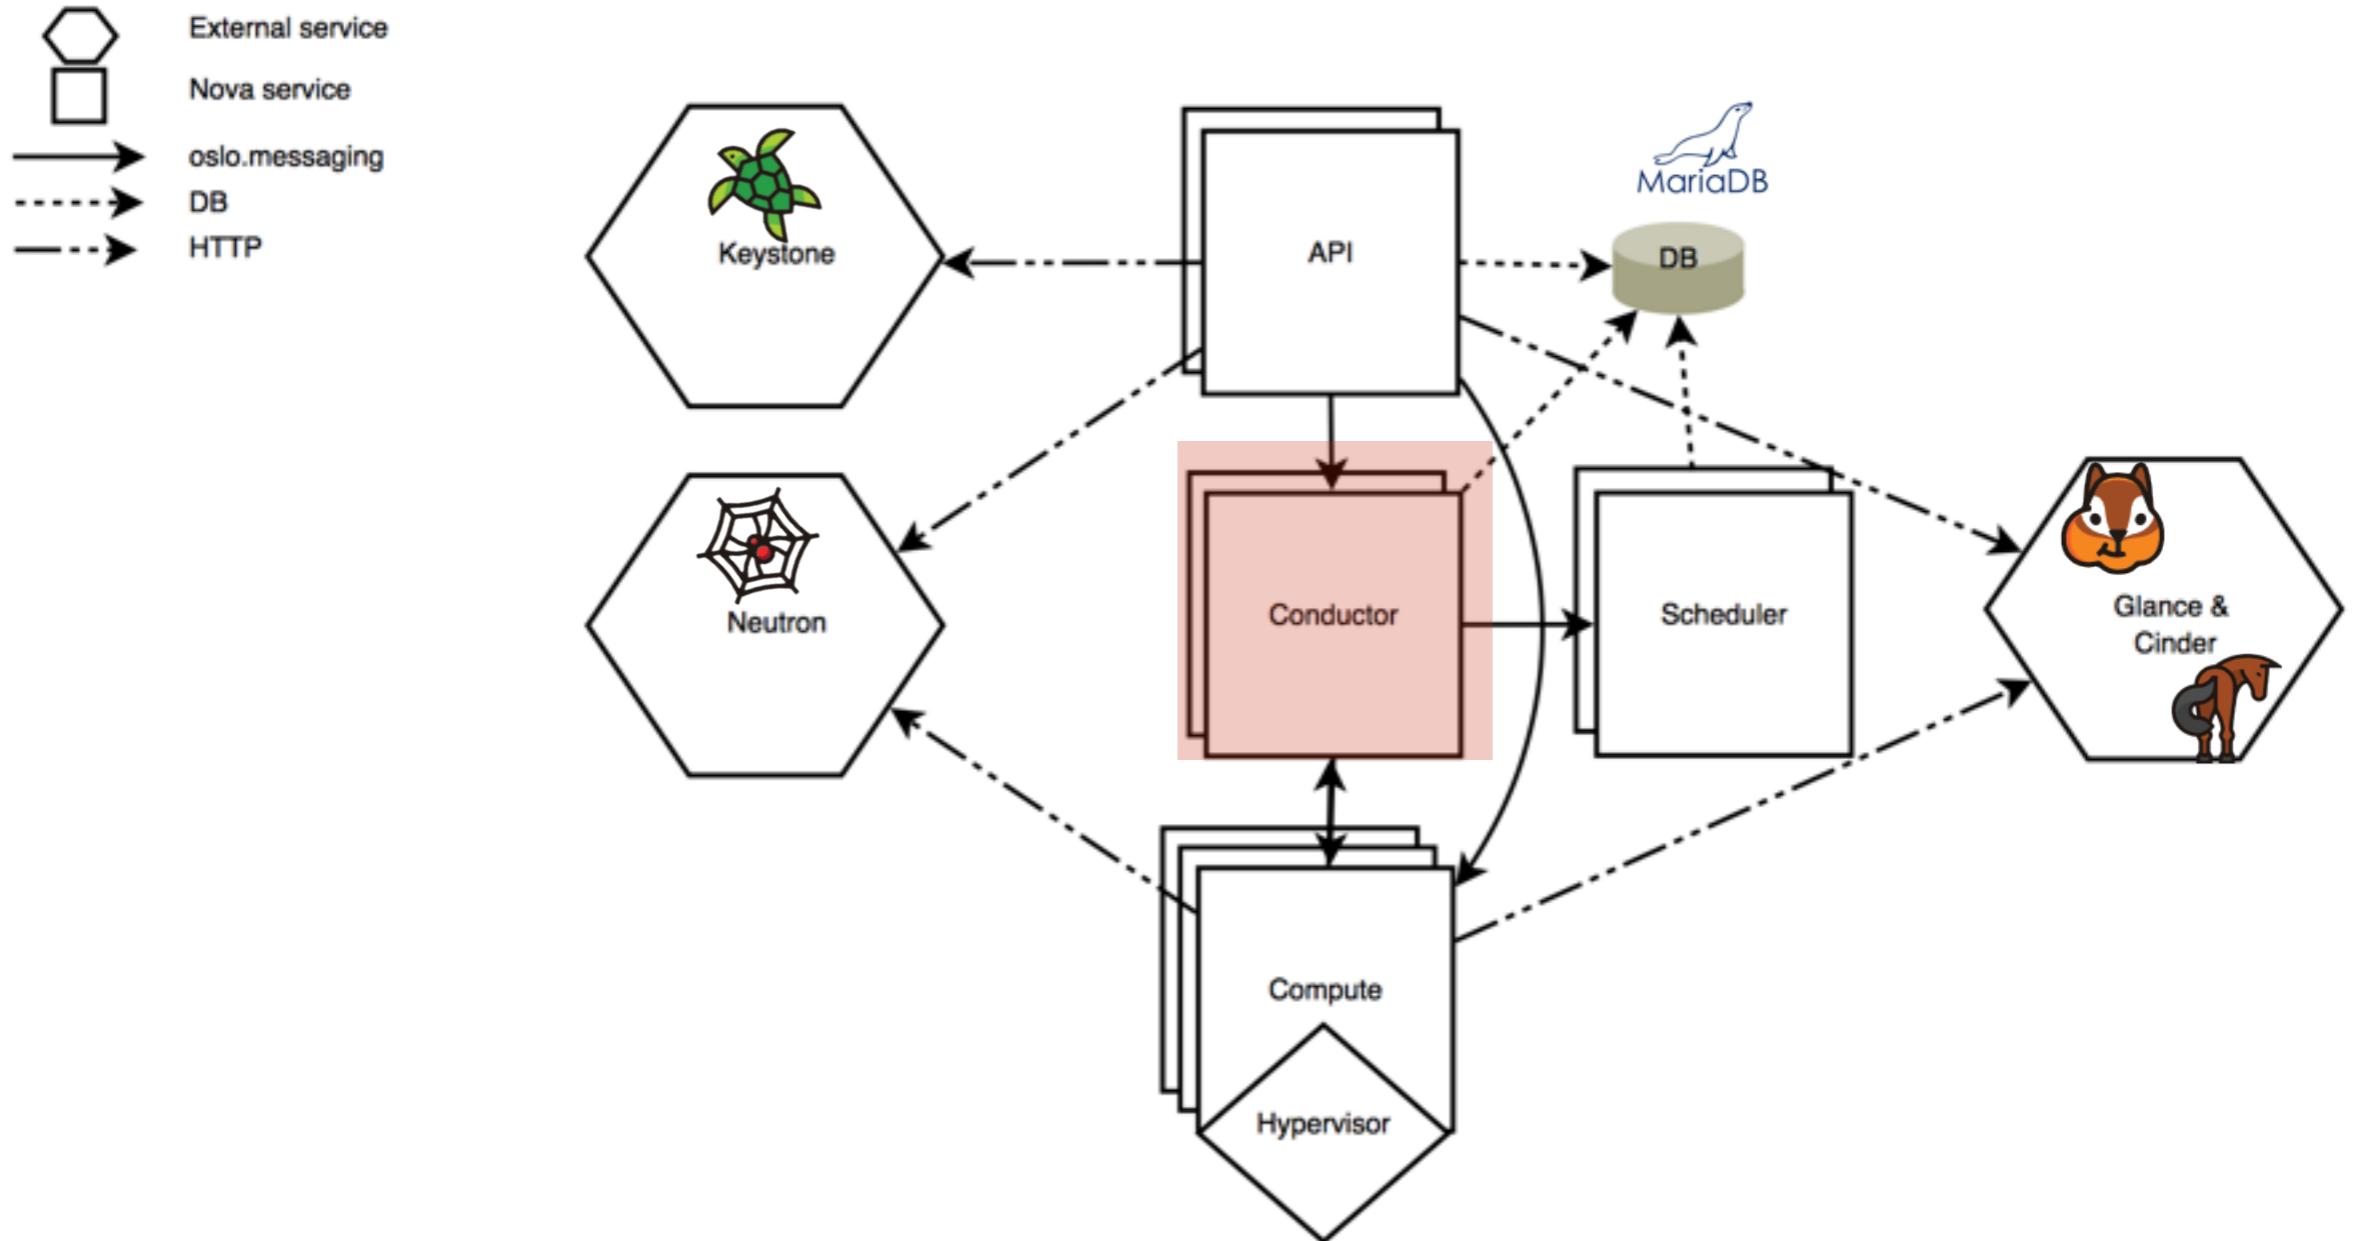
Create a port from NEUTRON

Prepare the corresponding block device

Ask the hypervisor to boot the VM

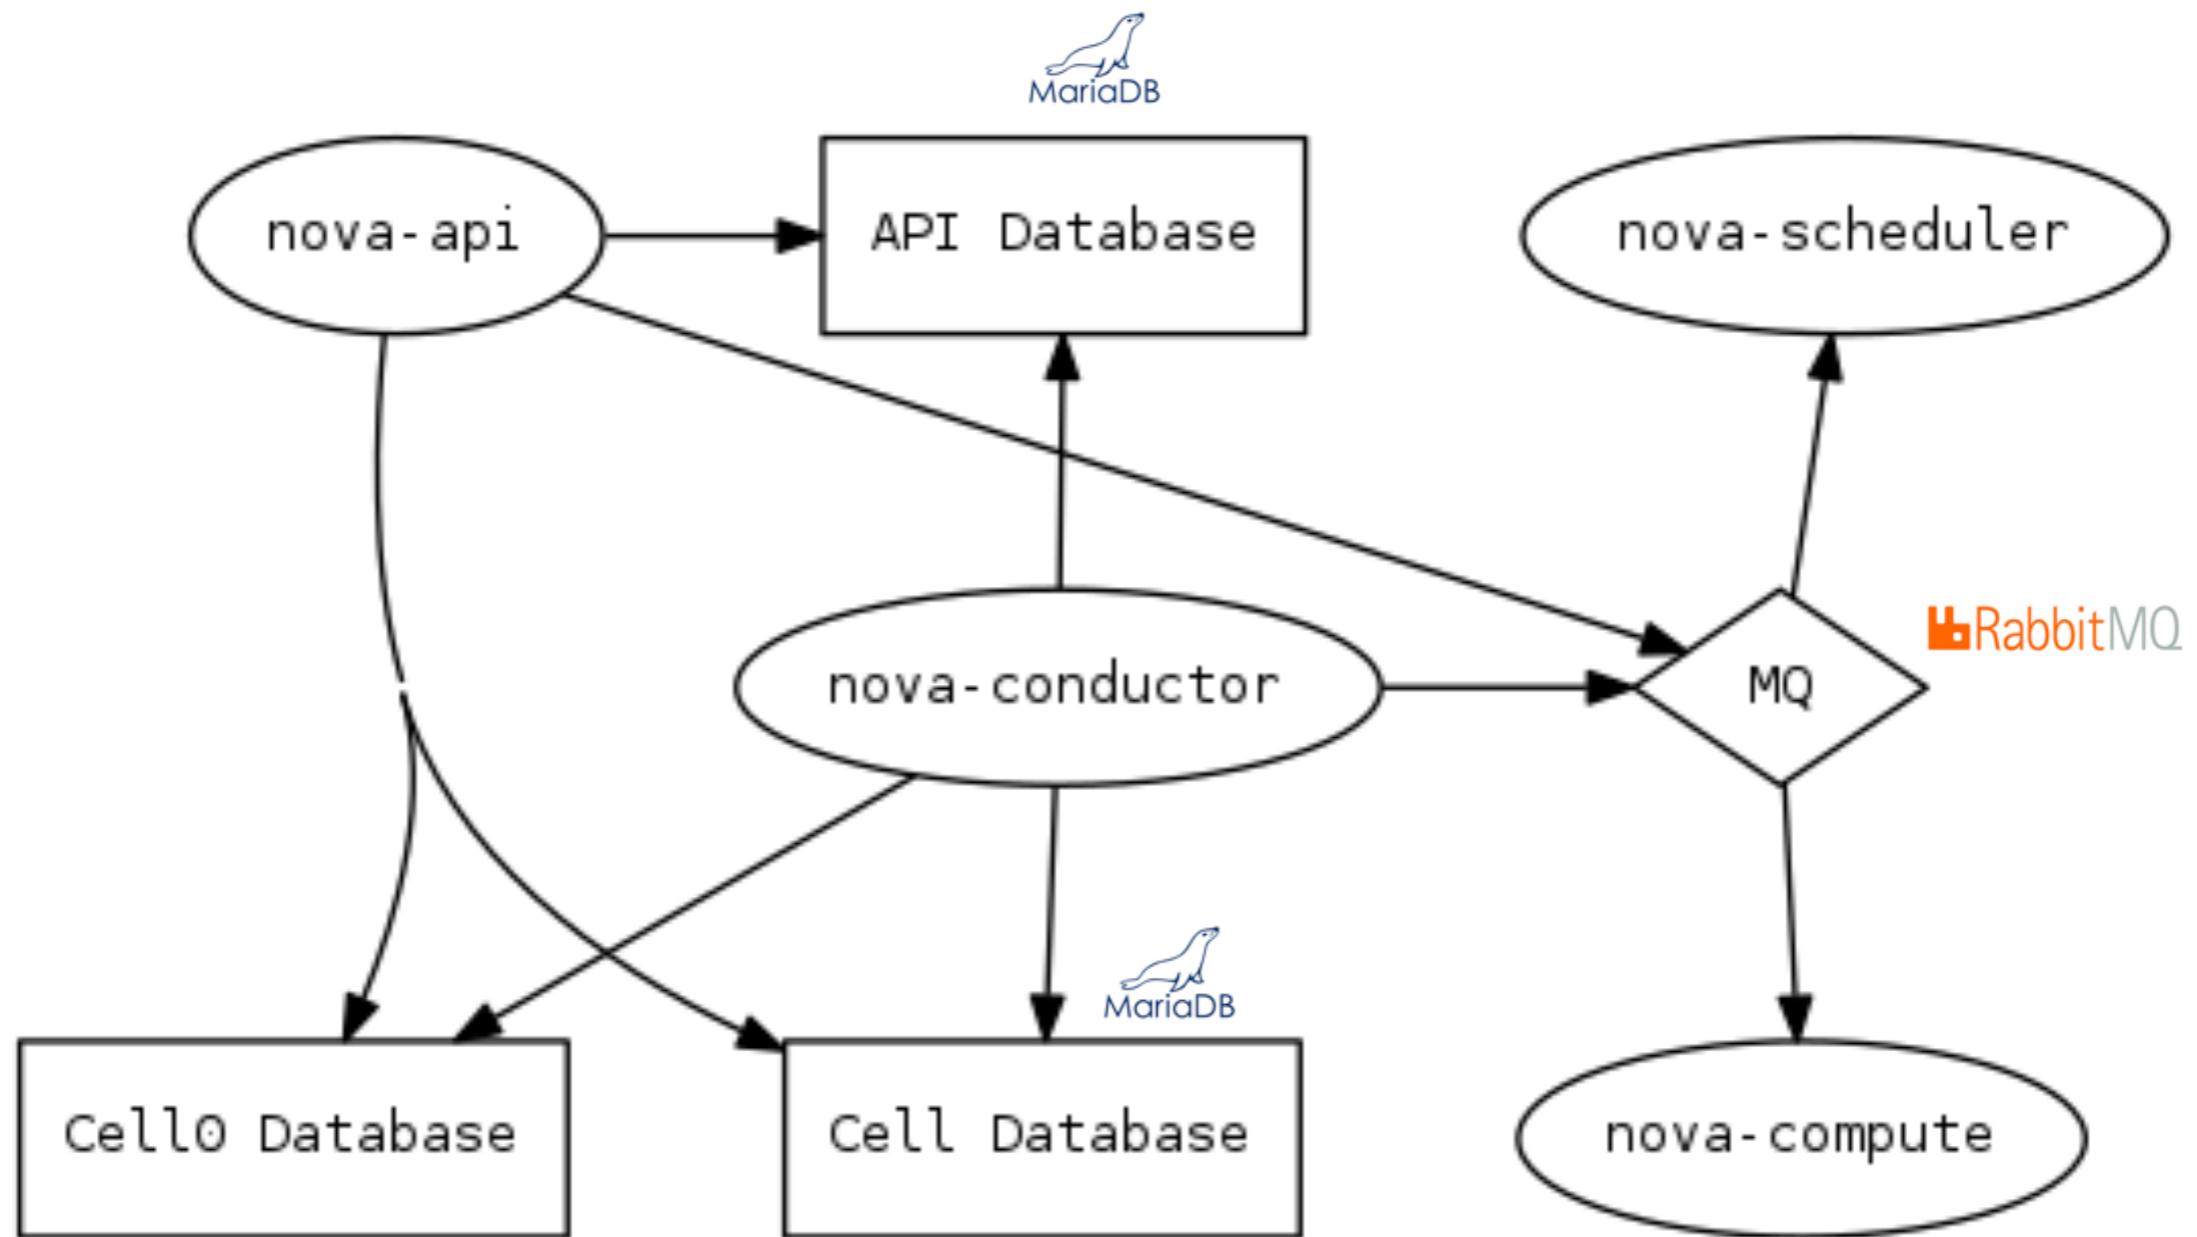


# Focus on Nova Service



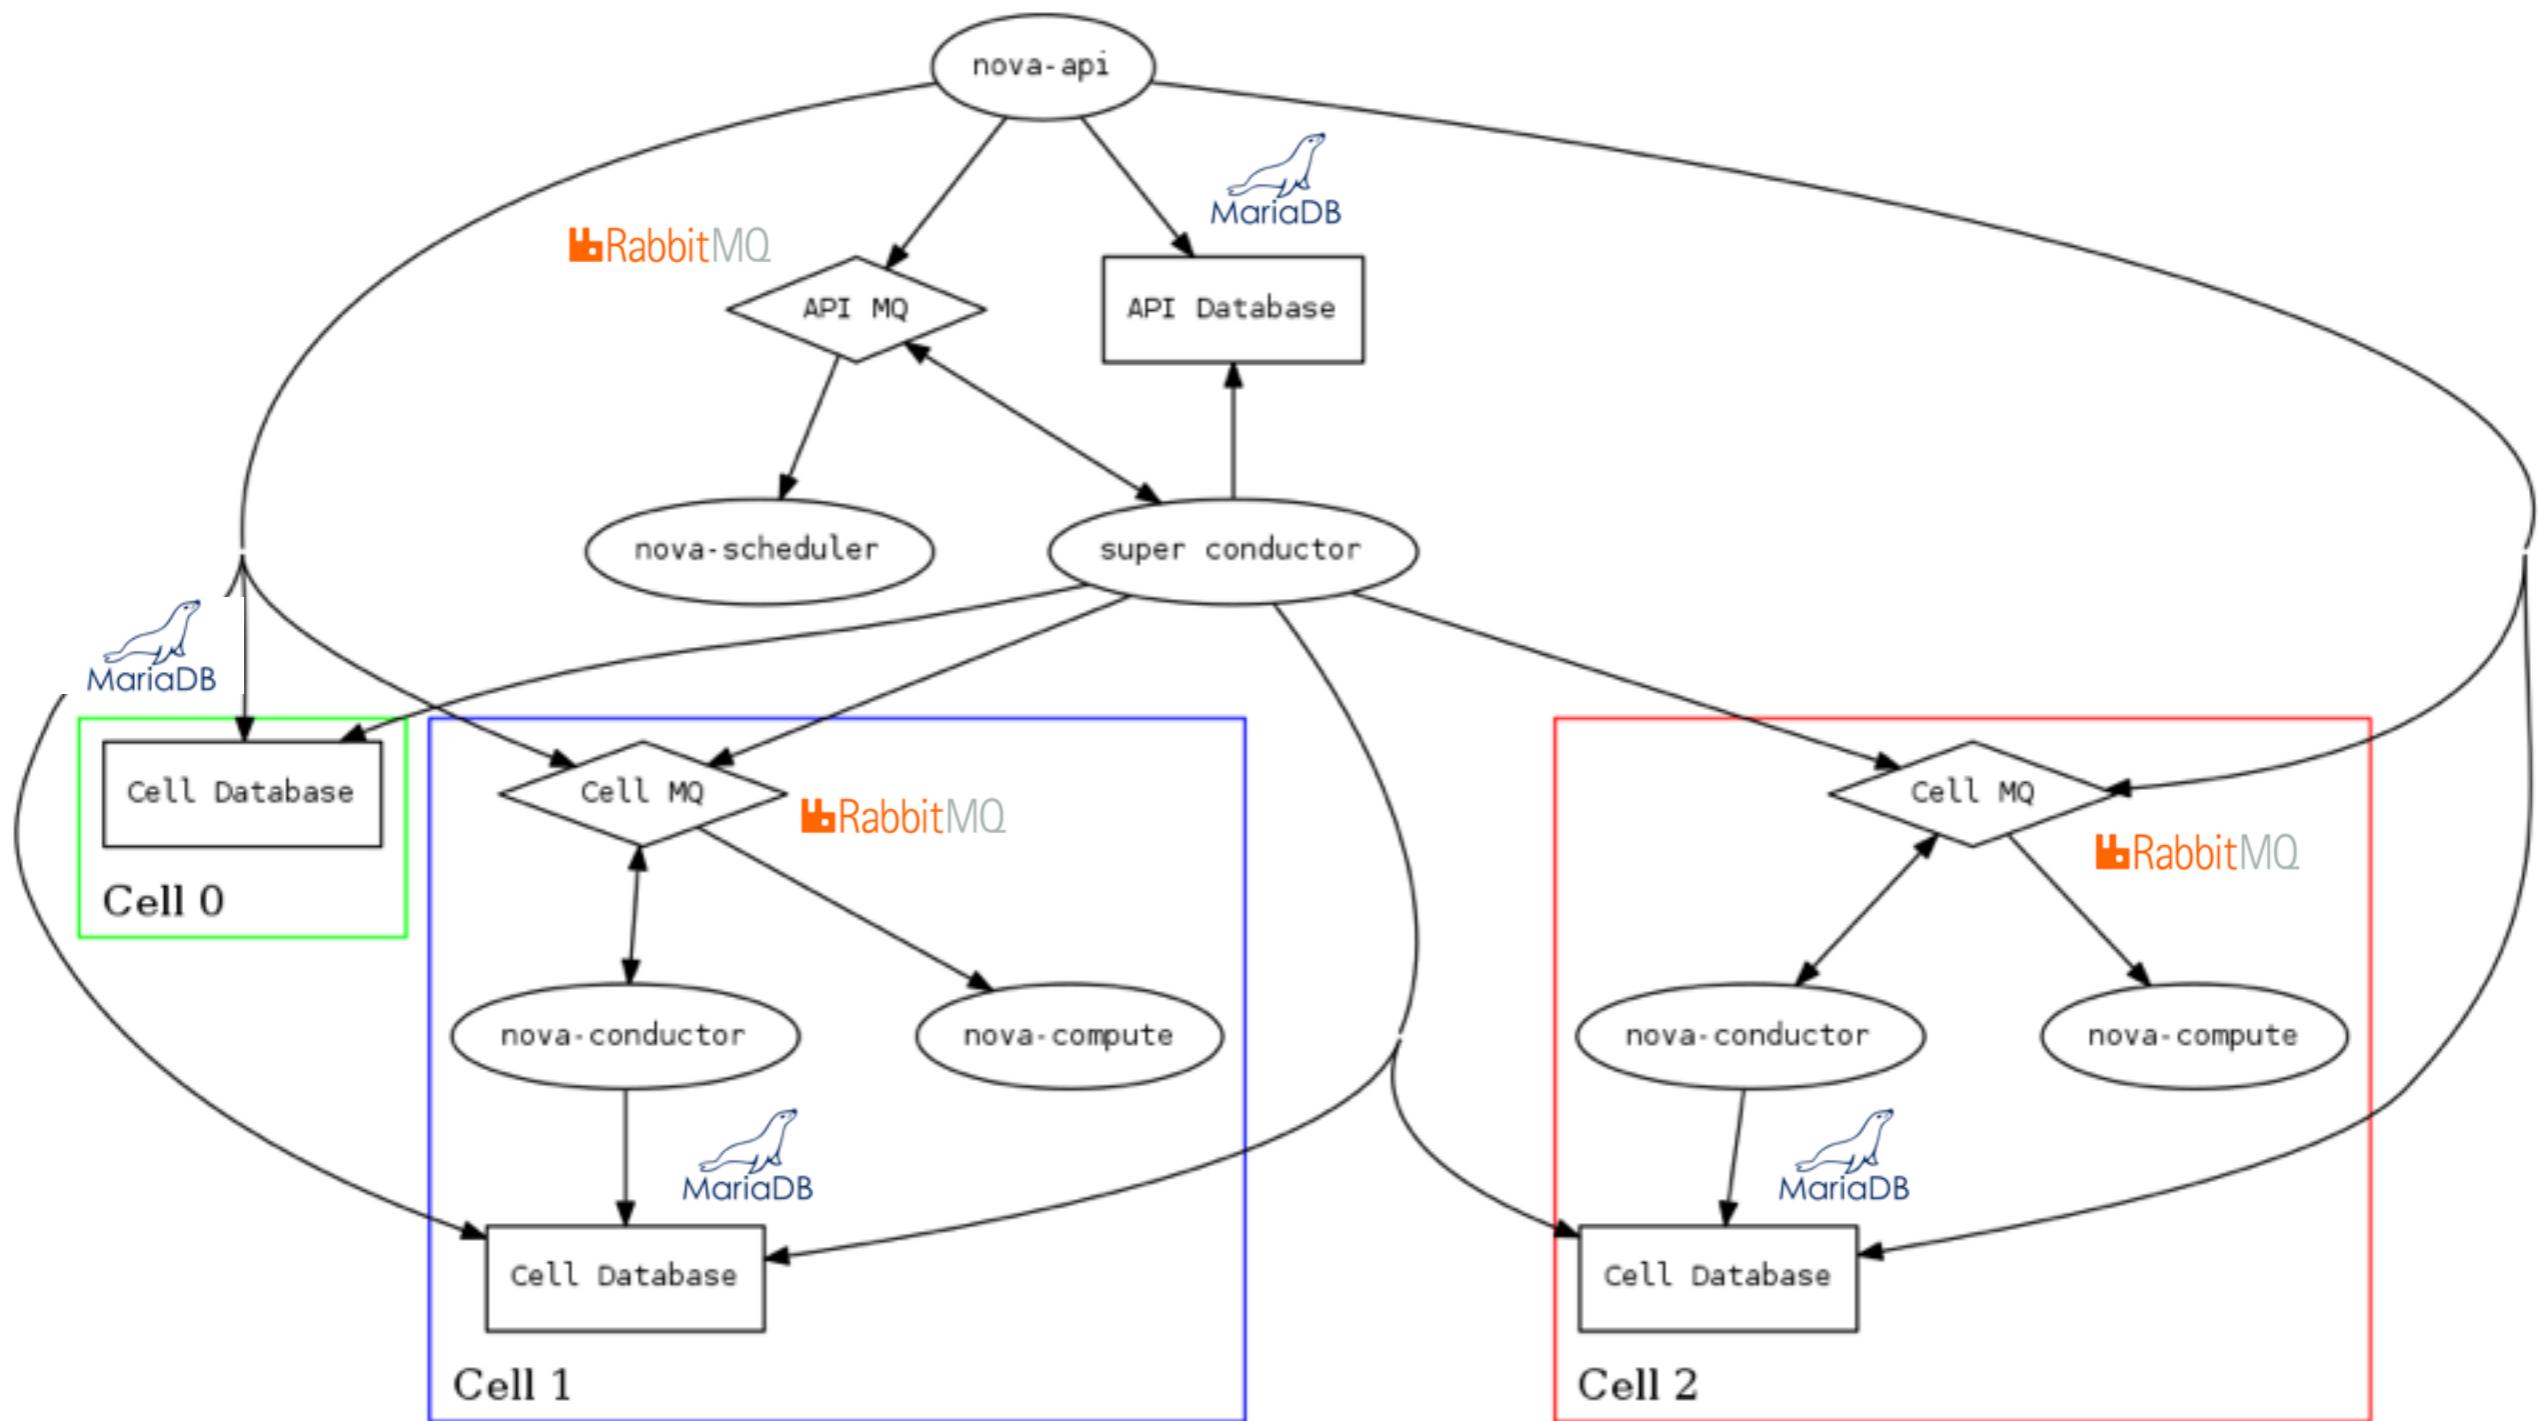
**Nova Conductor: a key element**

# Focus on Nova Service



**Cell... by default since Ocata**

# Focus on Nova Service



**Cell...a way to segregate your infrastructure**

# Focus on Neutron

- **Different kinds of networks**

Project networks provide connectivity to instances for a particular project/tenant (Private IPs)

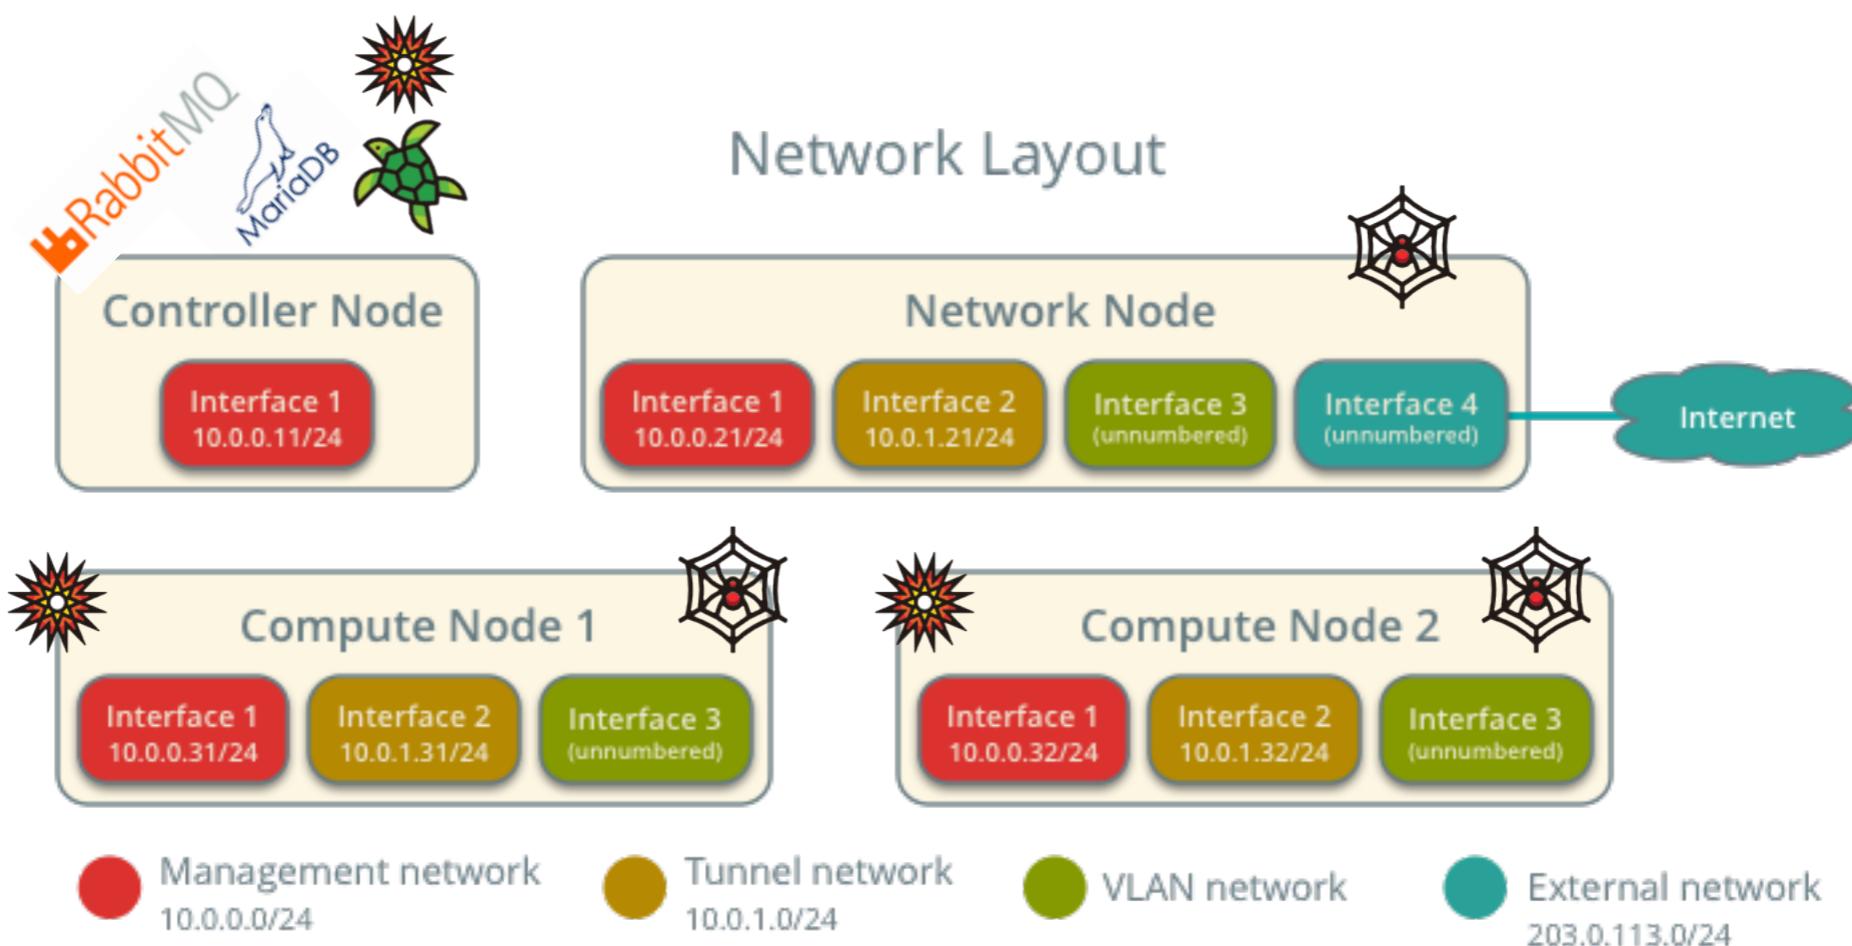
External networks provide connectivity to external networks such as the Internet (Public IPs)

Routers typically connect project and external networks.

- **Other supporting services (DHCP, ssh keys, ...)**

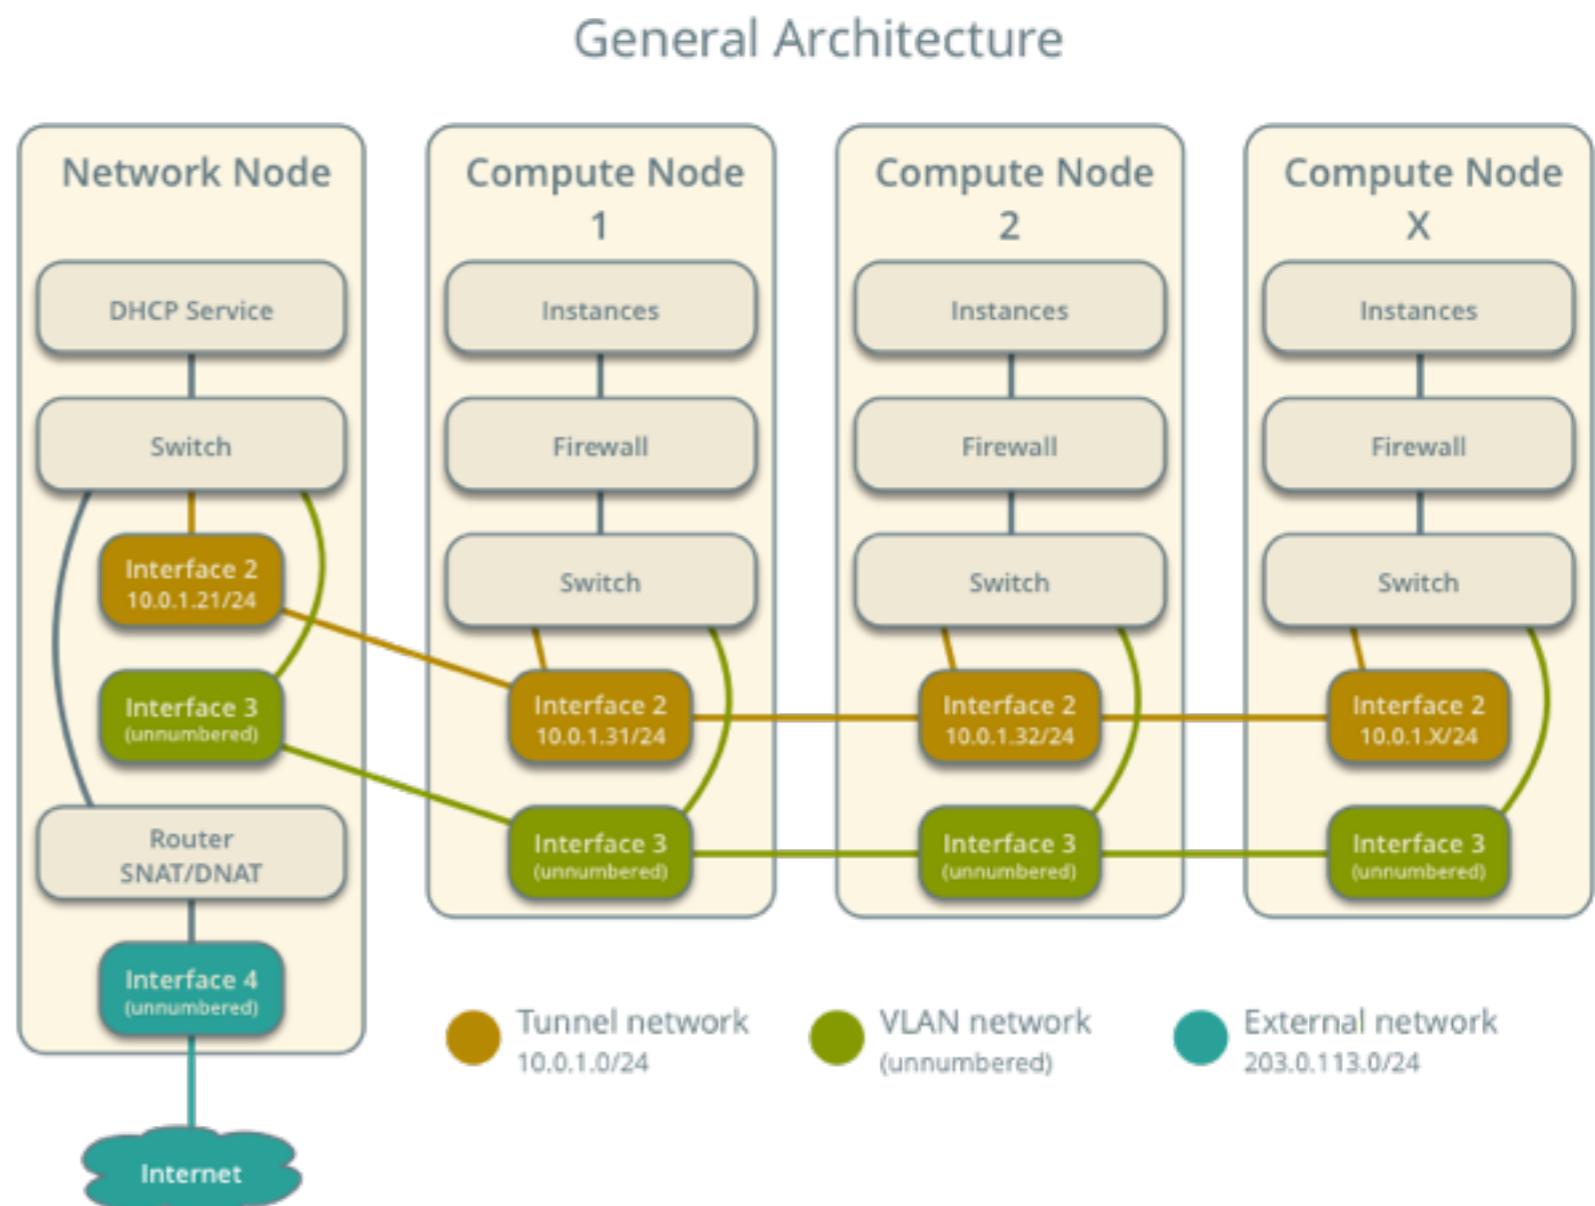
# Understanding Neutron Flows

- In the example configuration, the management network uses 10.0.0.0/24, the tunnel network uses 10.0.1.0/24, and the external network uses 203.0.113.0/24. The VLAN network does not require an IP address range because it only handles layer-2 connectivity.



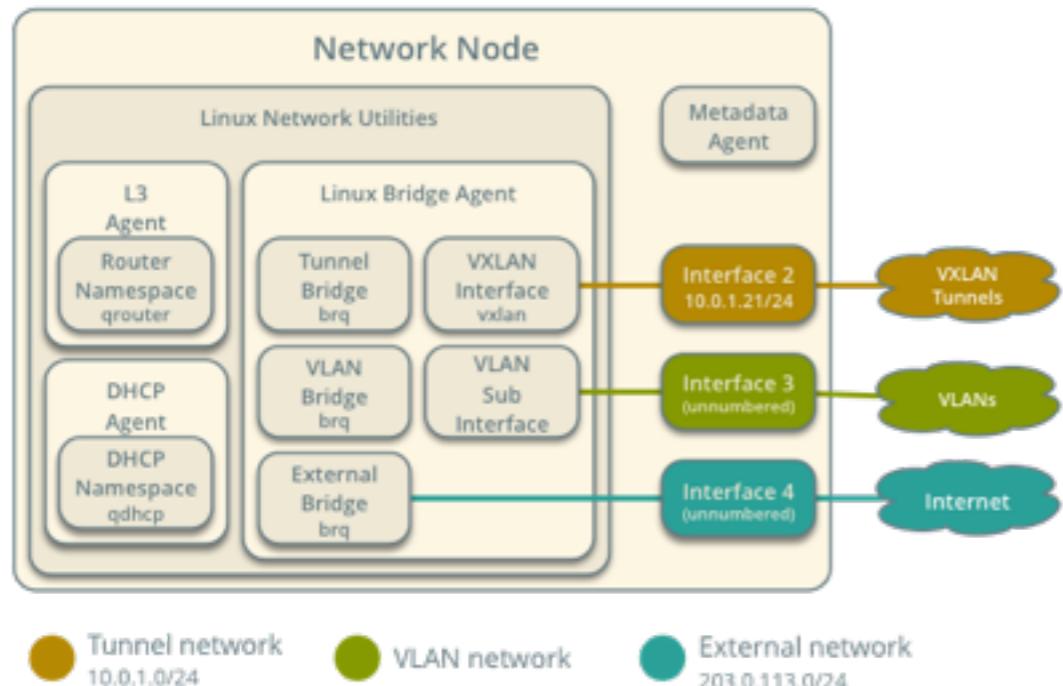
# Understanding Neutron Flows

- Routing among project and external networks resides completely on the network node.
- While this makes the management of network flows easier, it may lead to SPOF issues

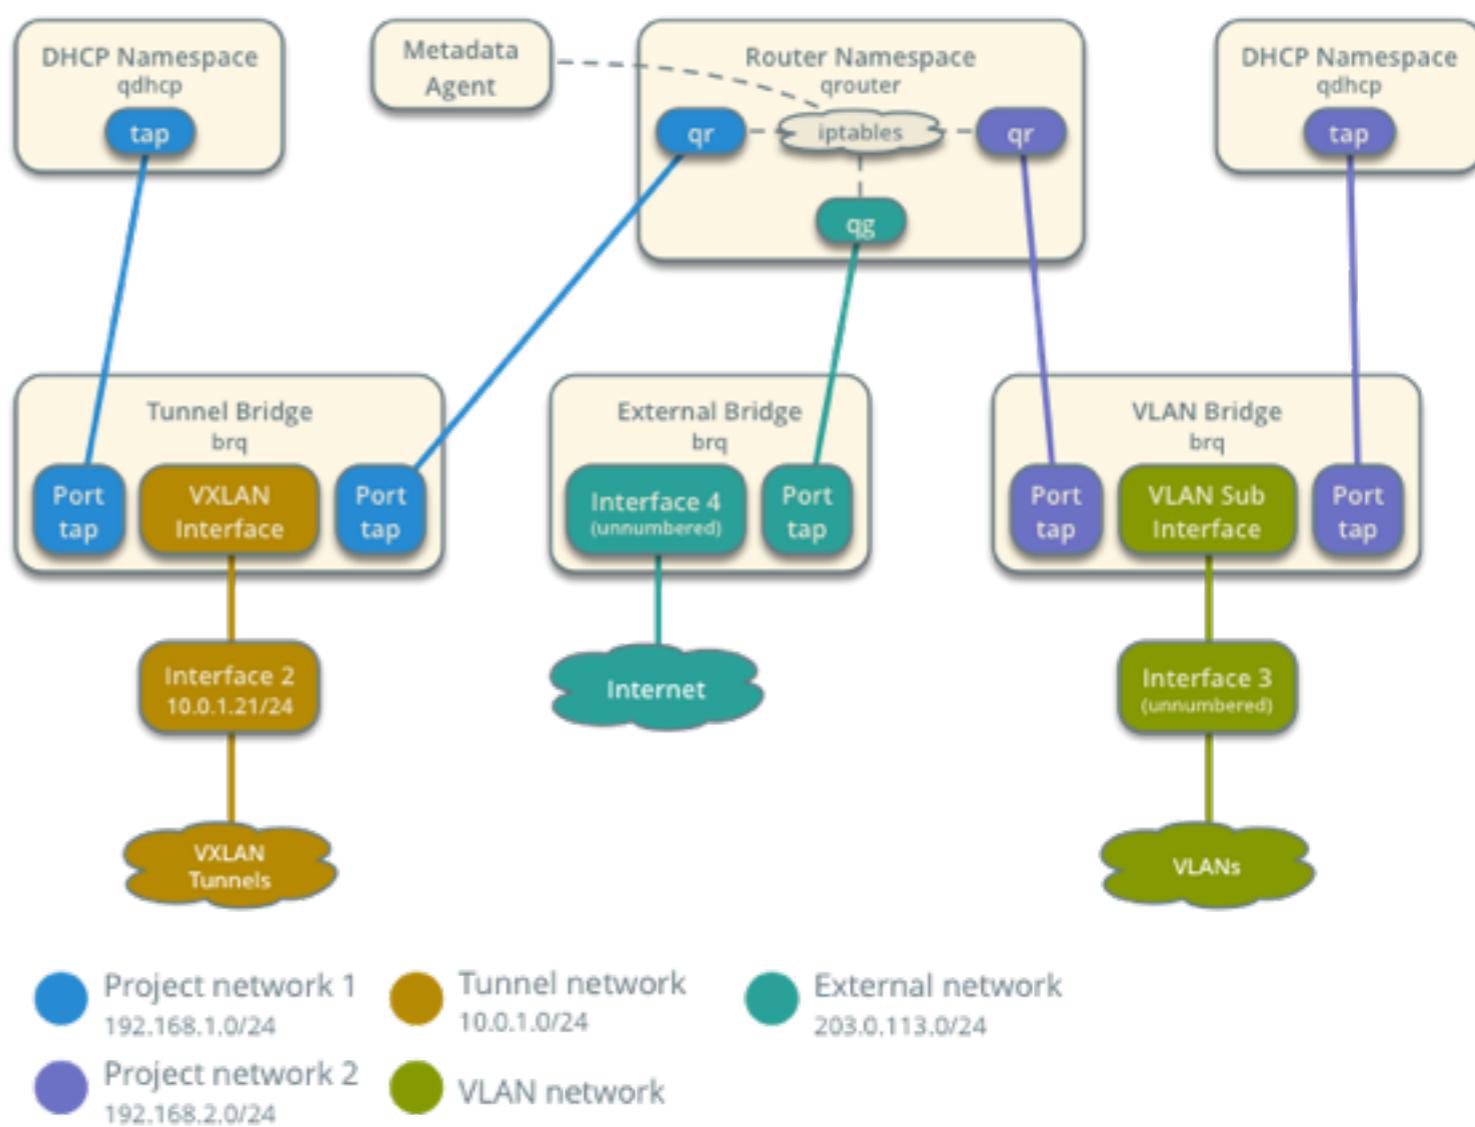


# Understanding Neutron Flows

## Network Node Overview



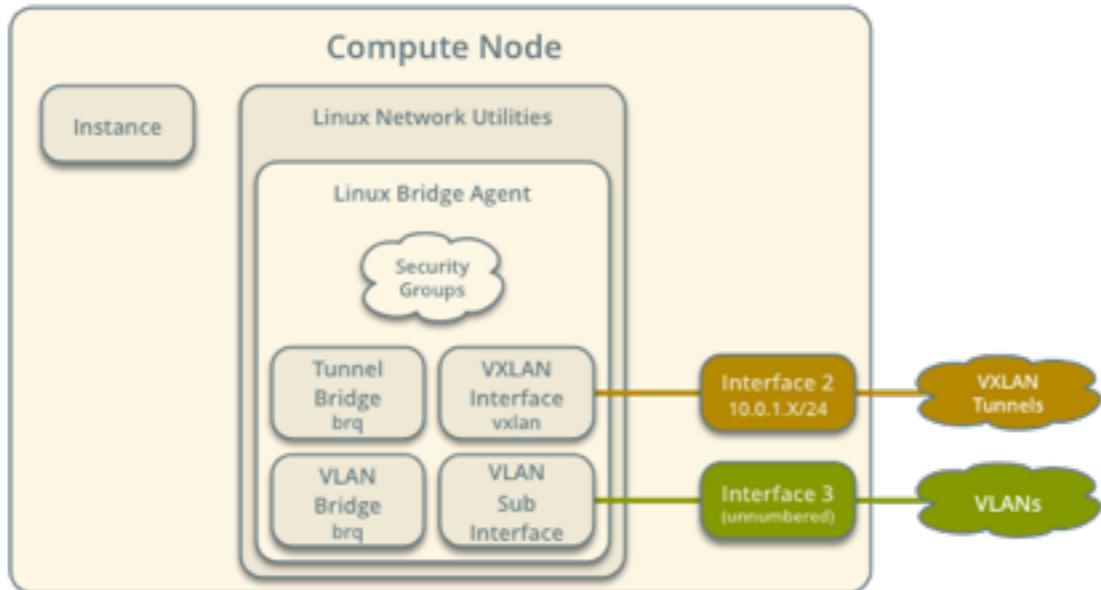
## Network Node Components



Inside the  
Network Node

# Understanding Neutron Flows

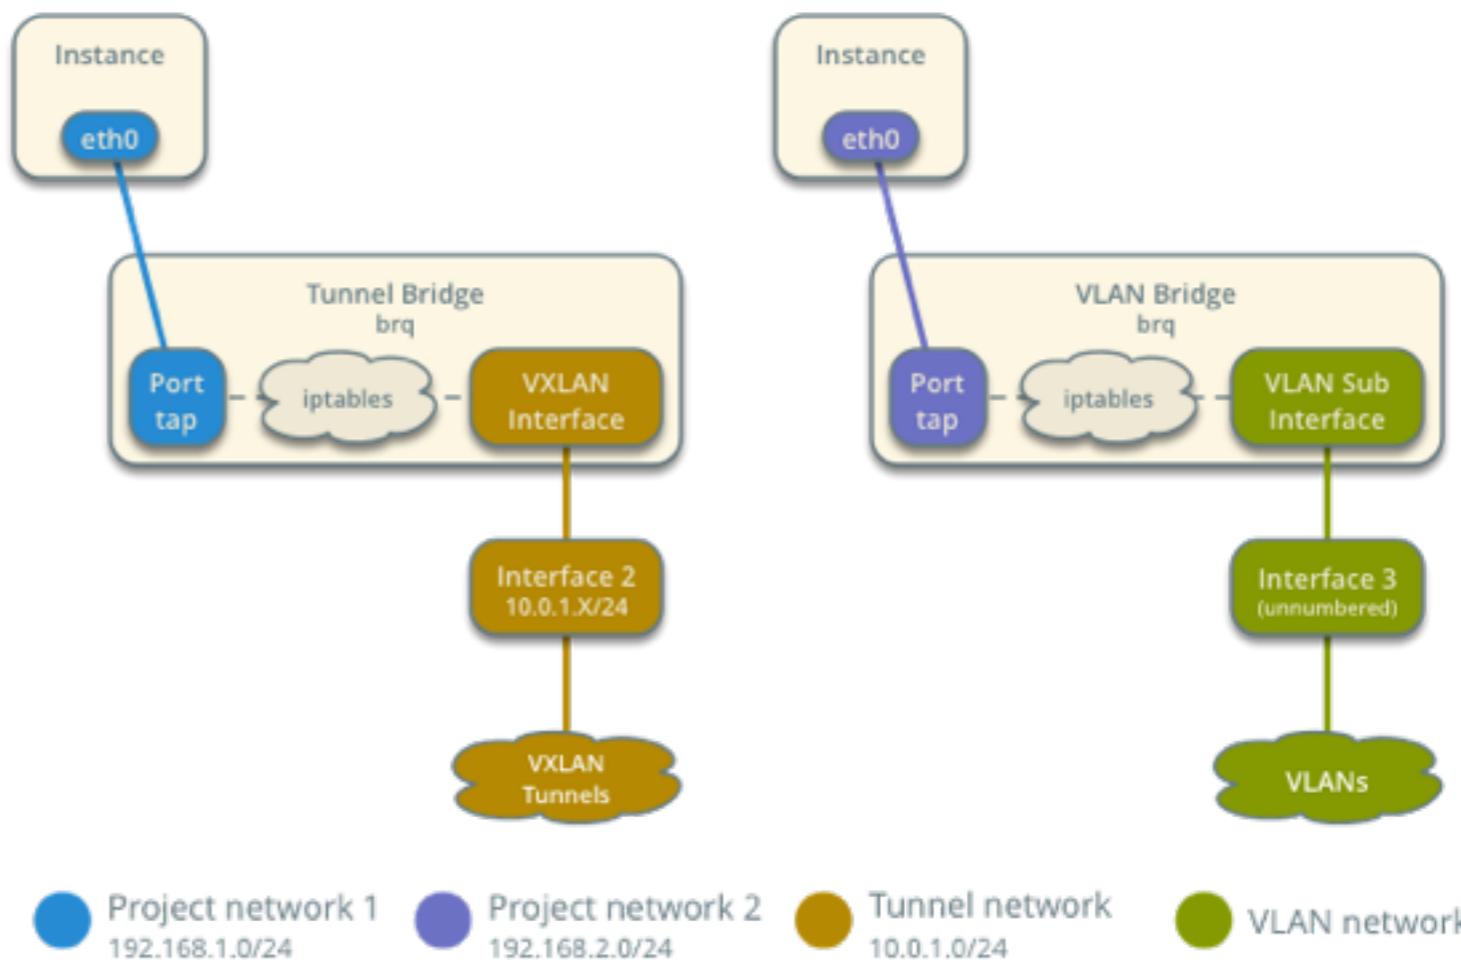
## Compute Node Overview



● Tunnel network 10.0.1.0/24   ● VLAN network

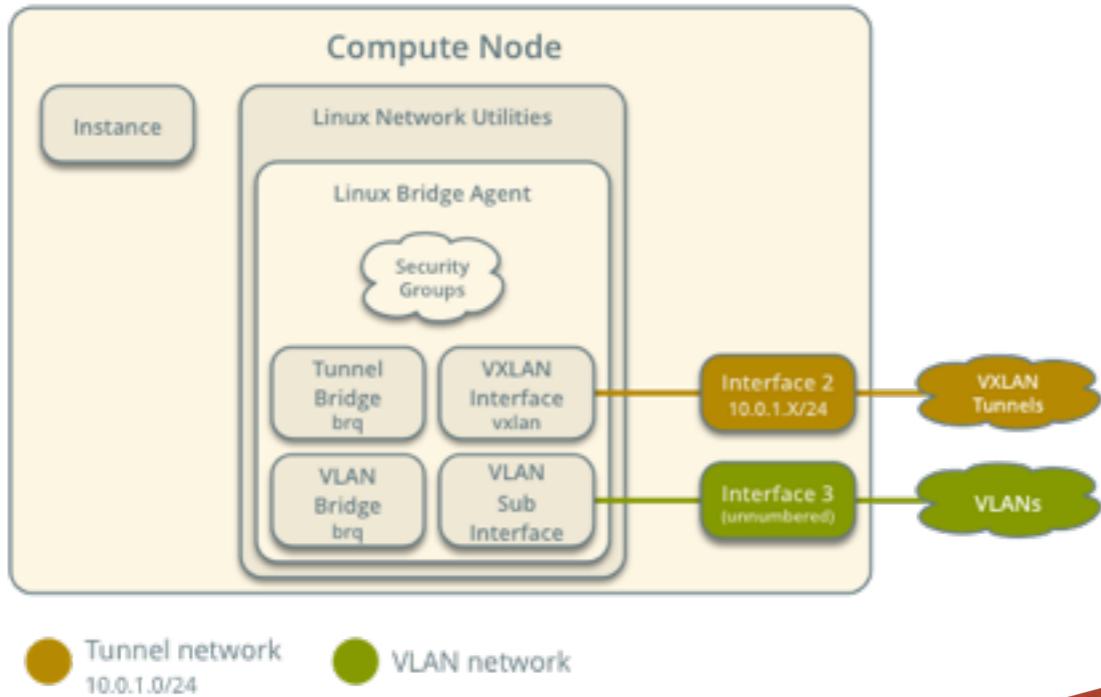


## Compute Node Components

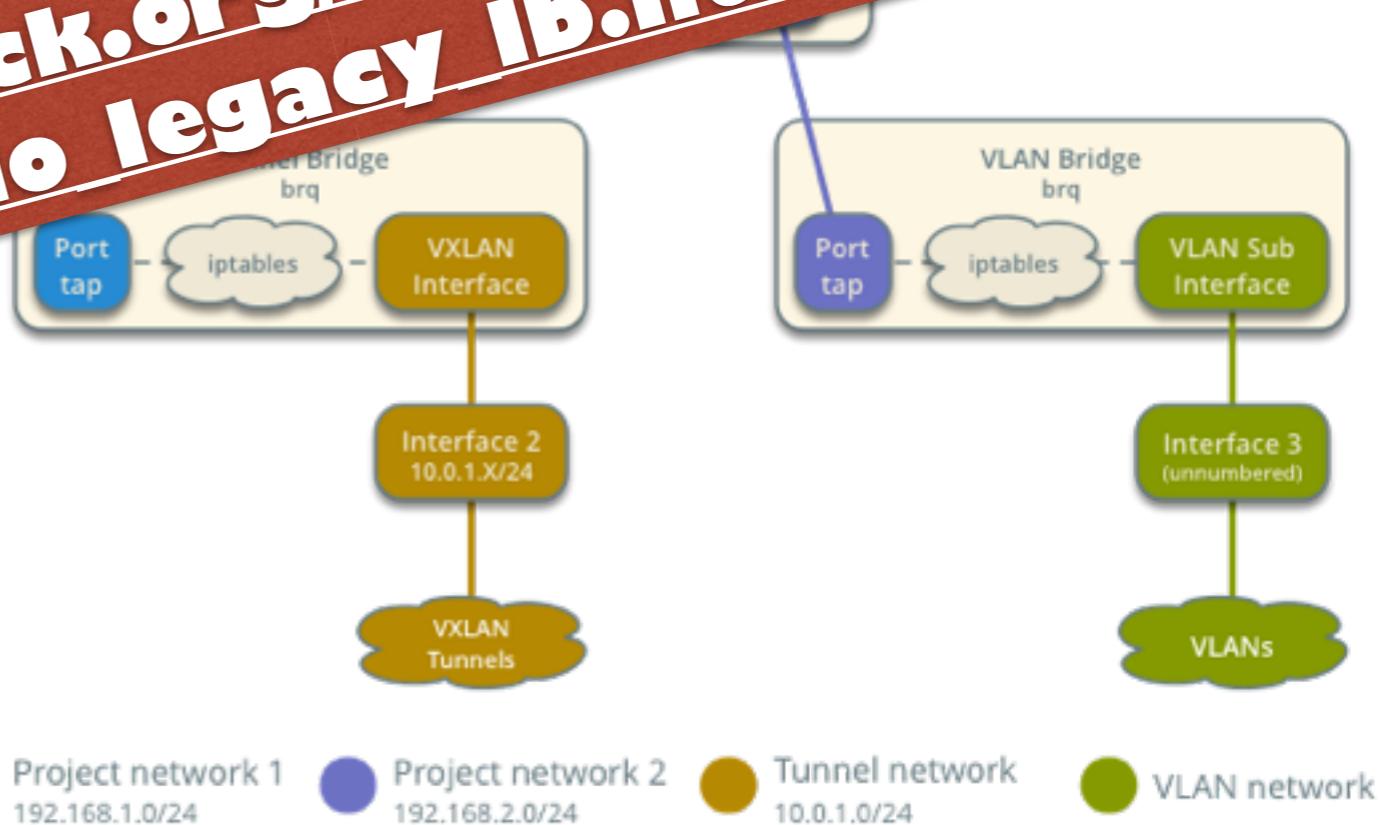


# Understanding Neutron Flows

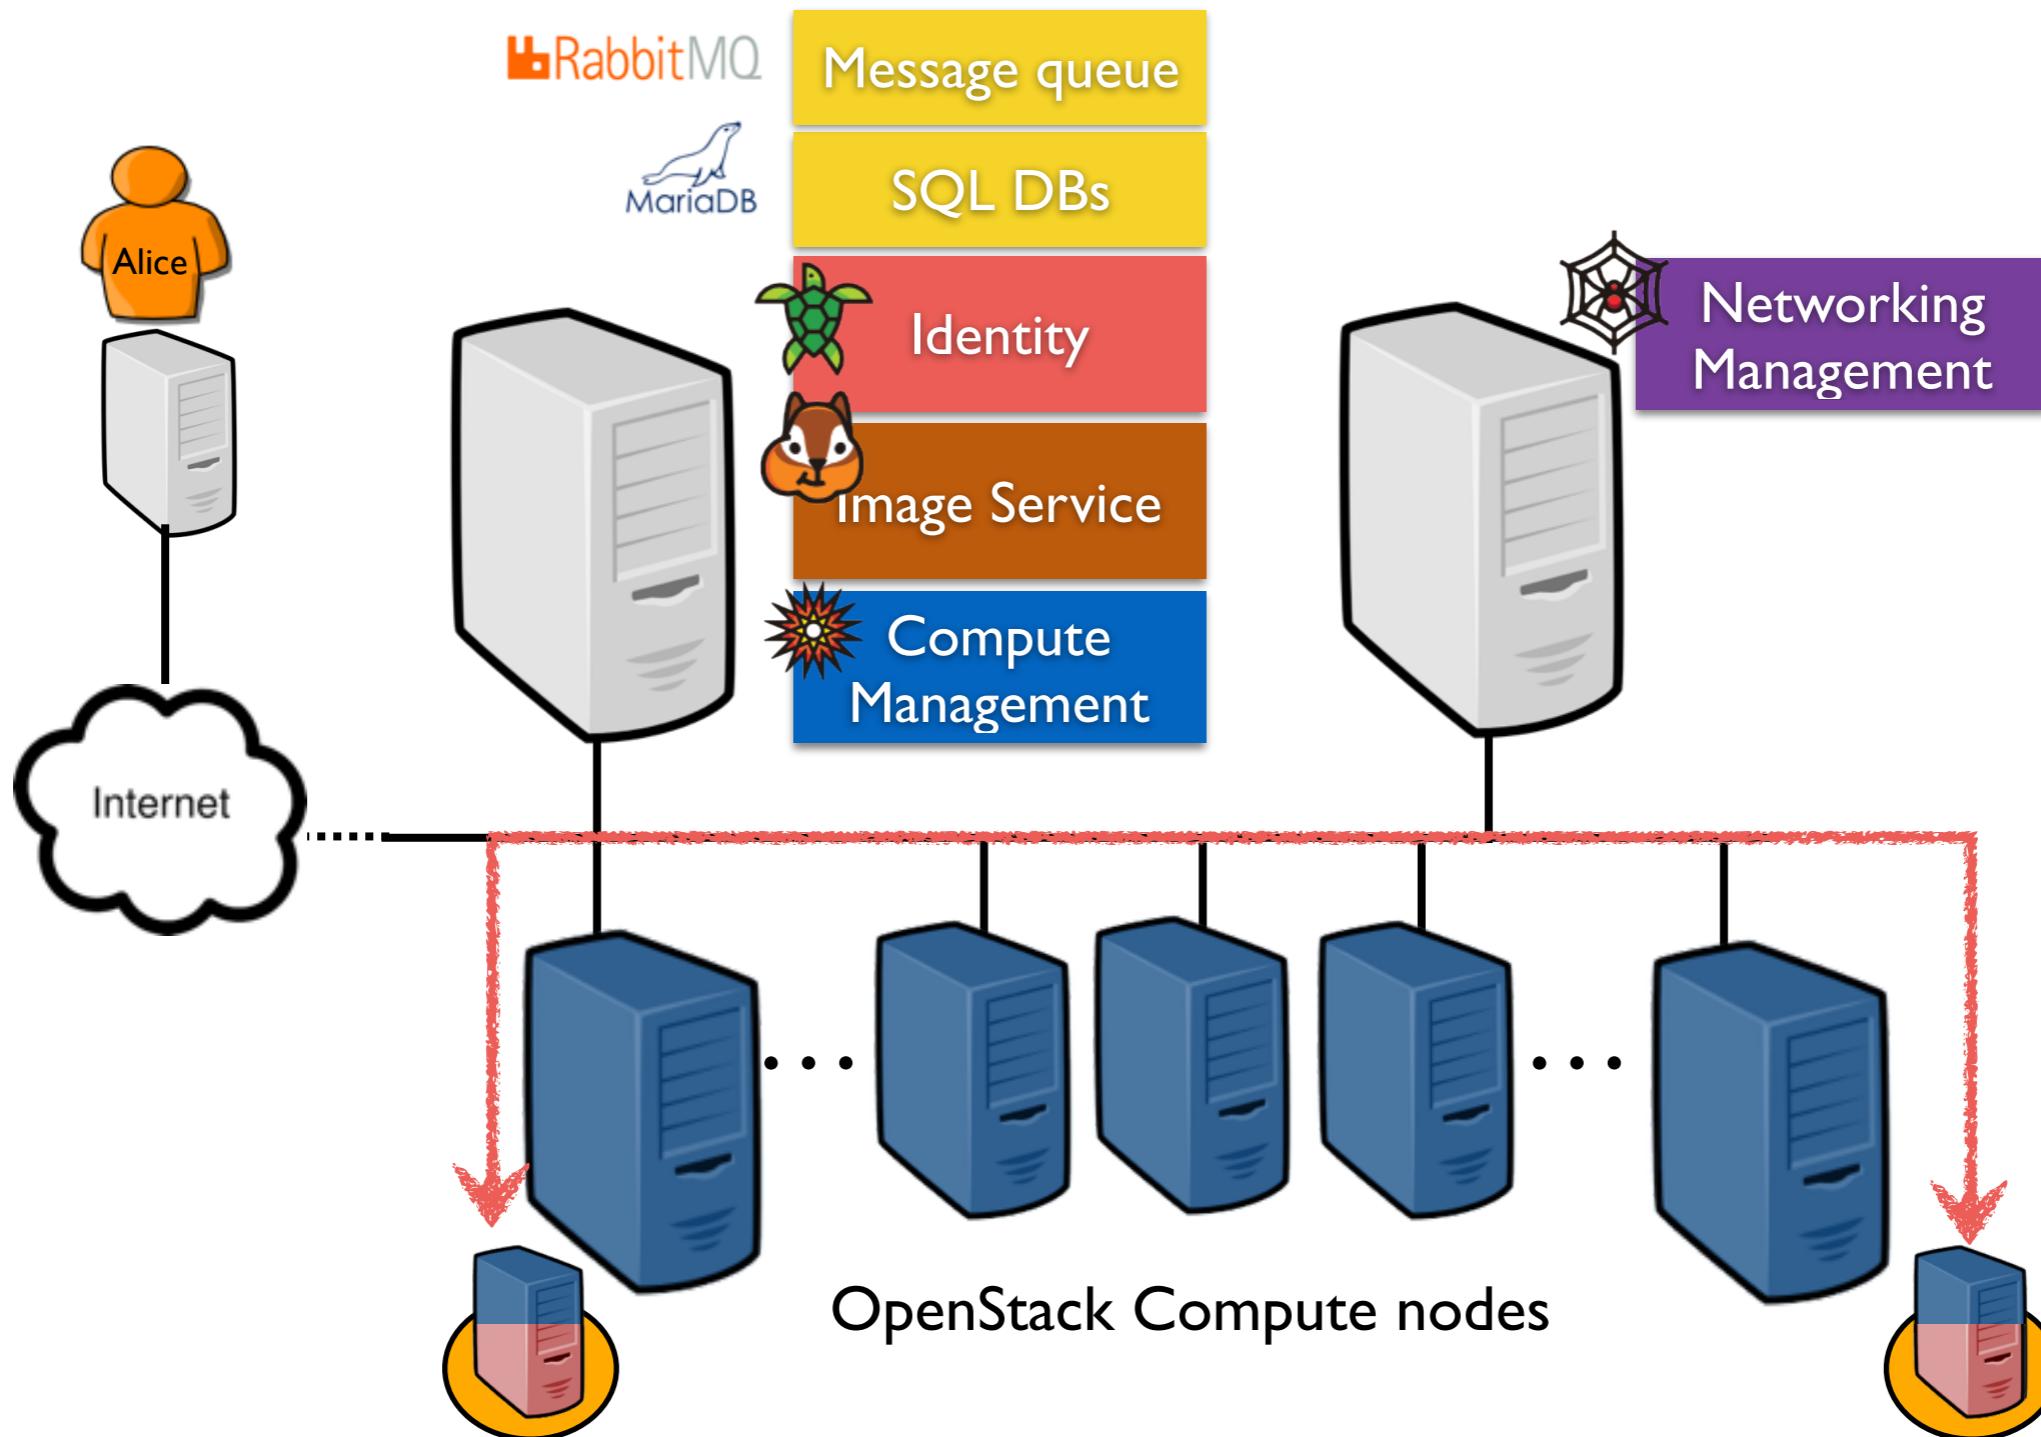
## Compute Node Overview



Interested to understand two slides, please see  
<https://docs.openstack.org/kilo/networking-lb.html>  
guide|scenario legacy\_lb.html

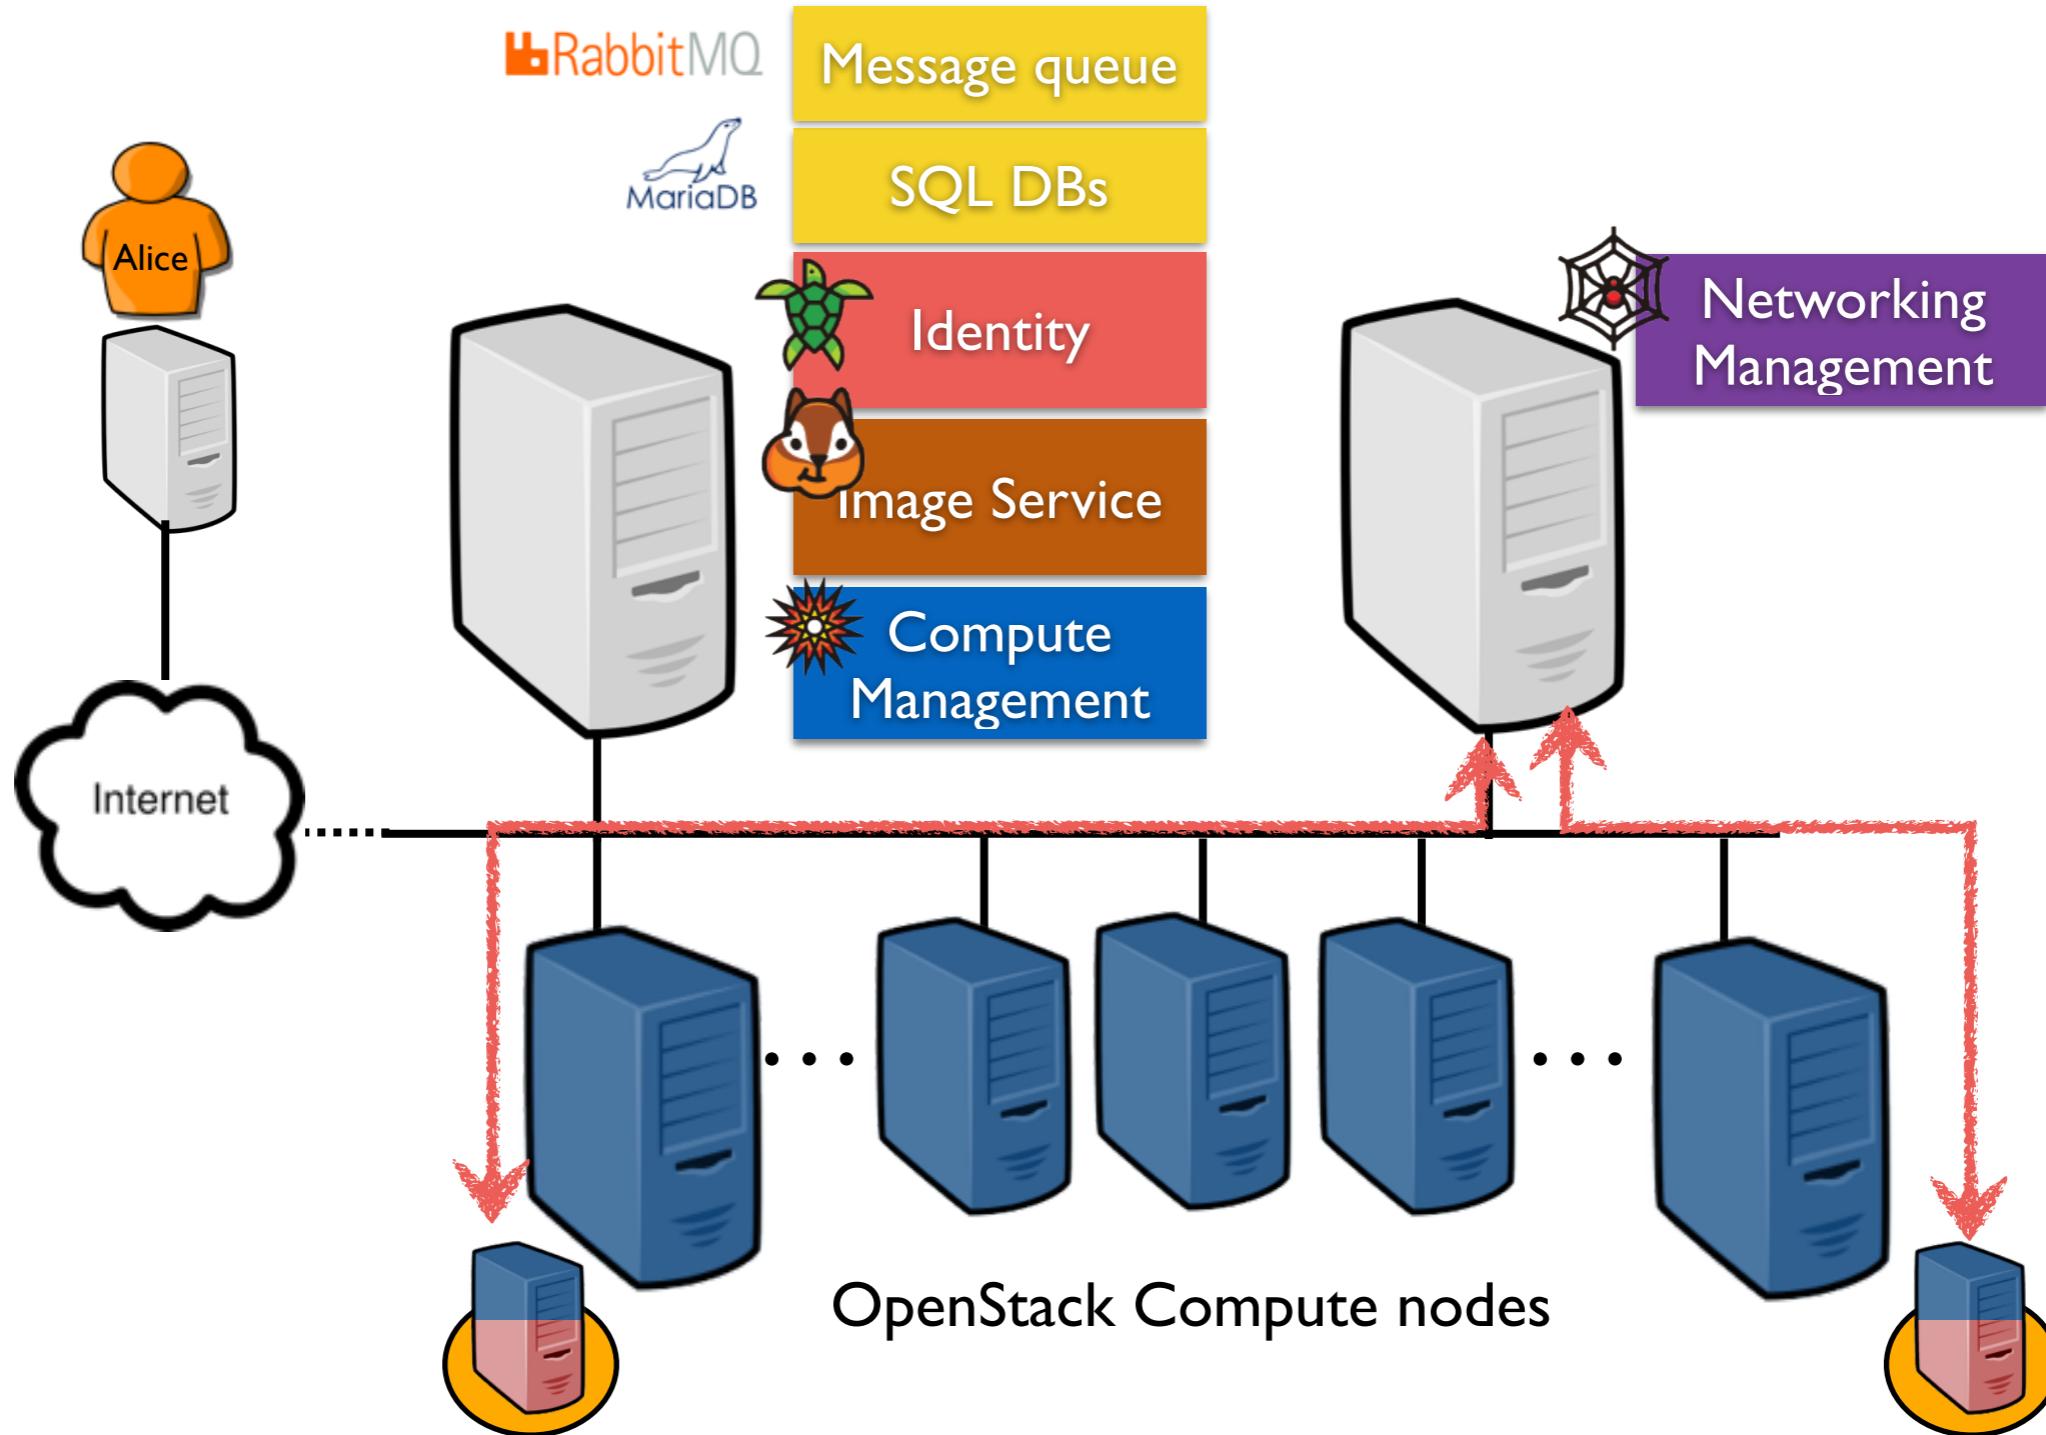


# Neutron Flows in a Nutshell



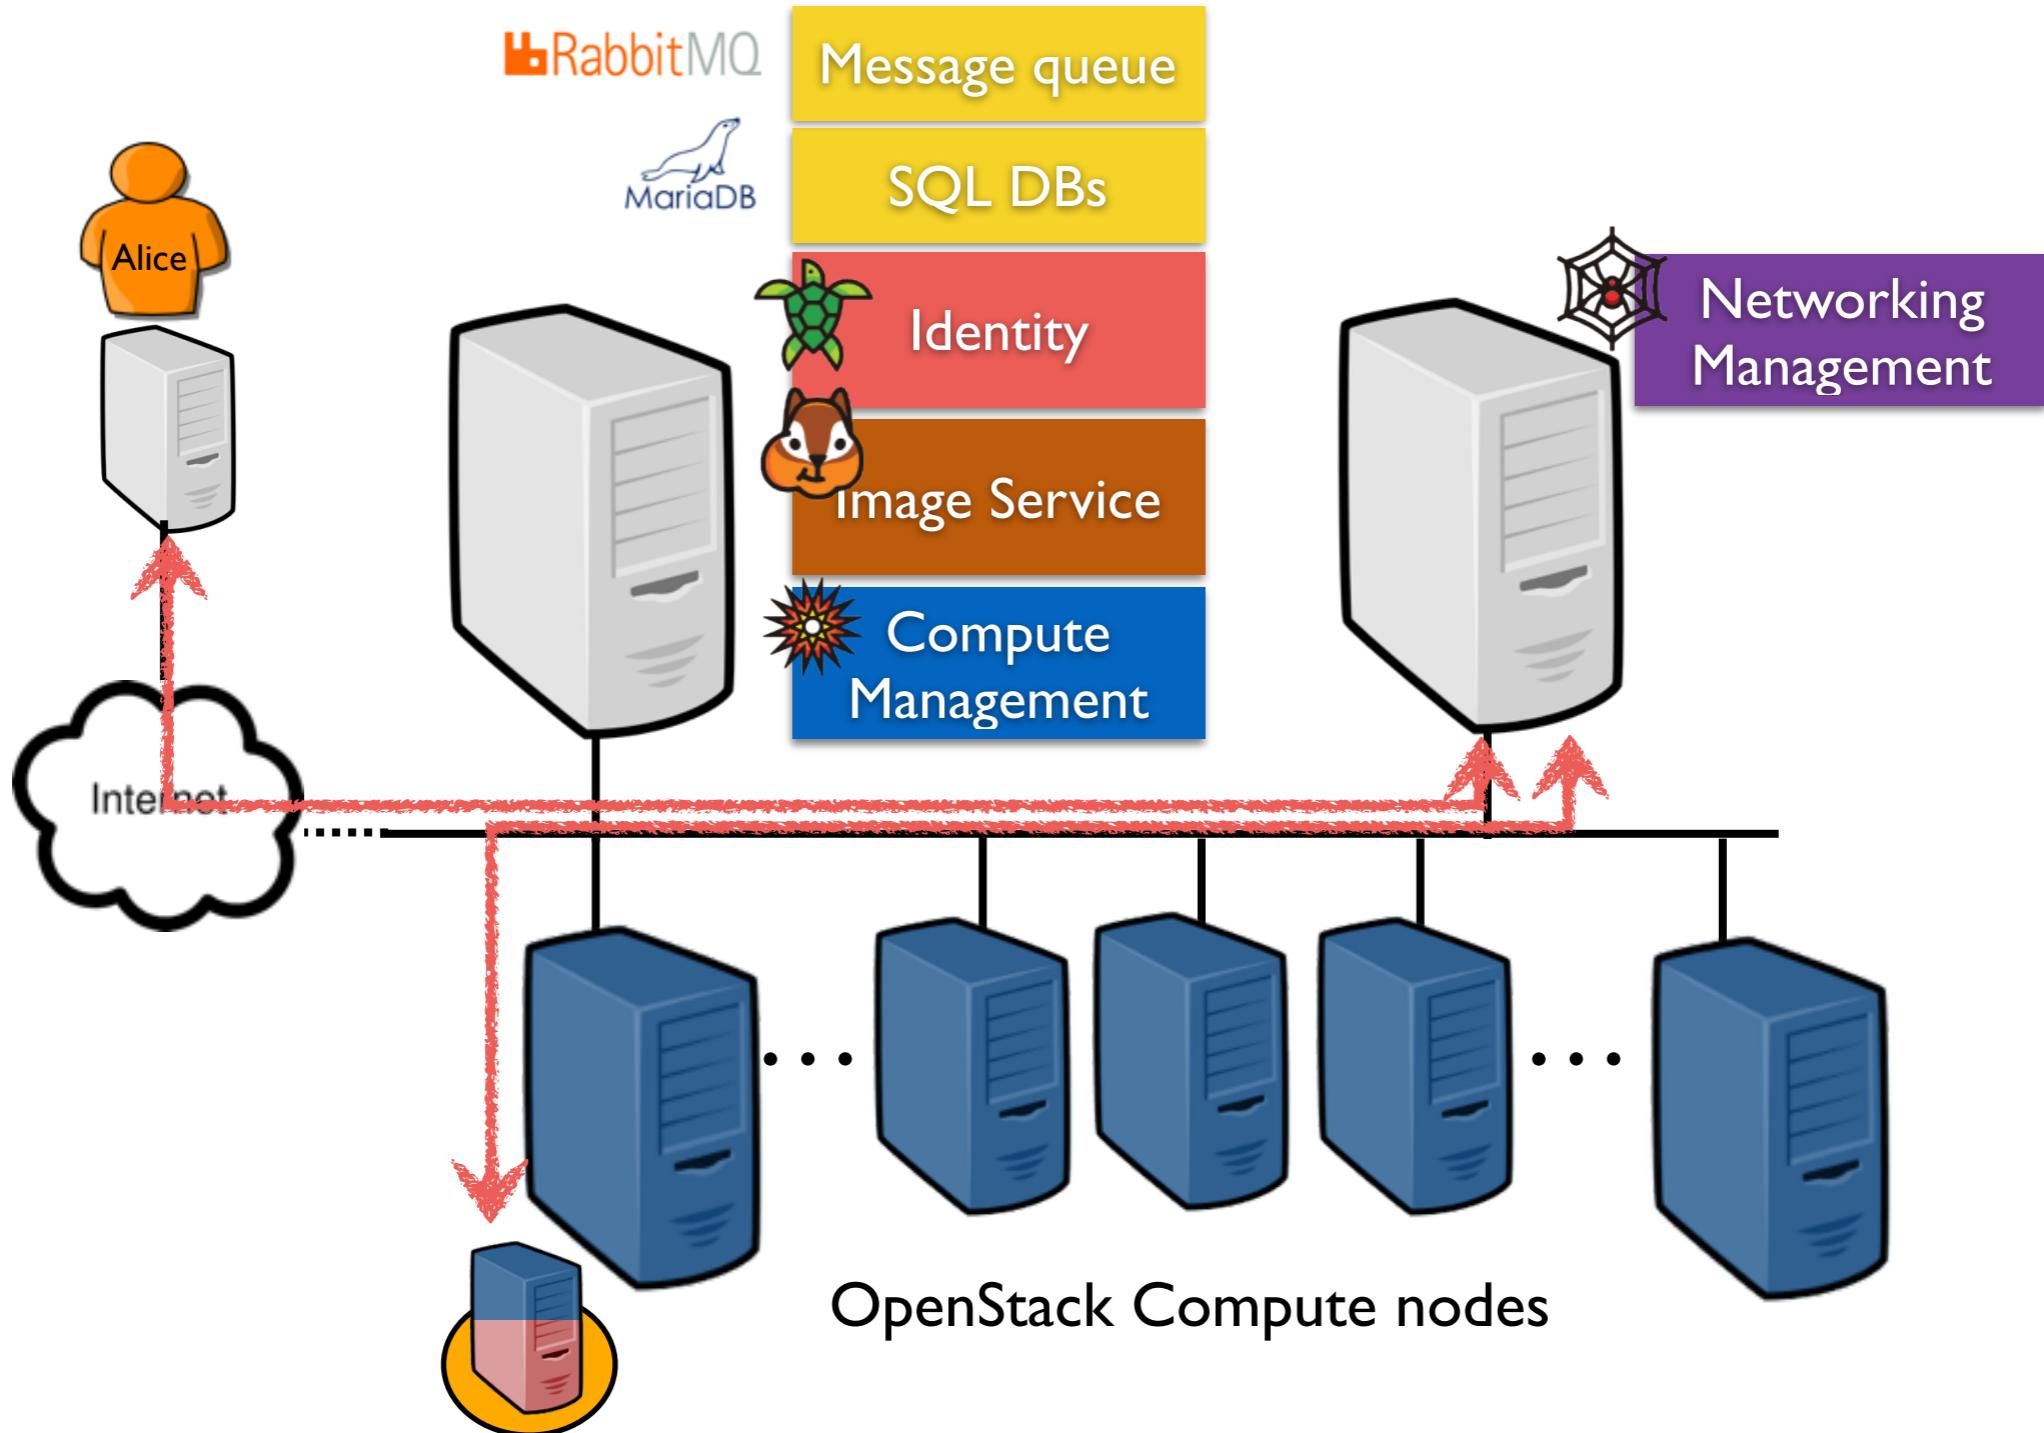
East/West - Same project network

# Neutron Flows in a Nutshell



East/West - Different project networks

# Neutron Flows in a Nutshell



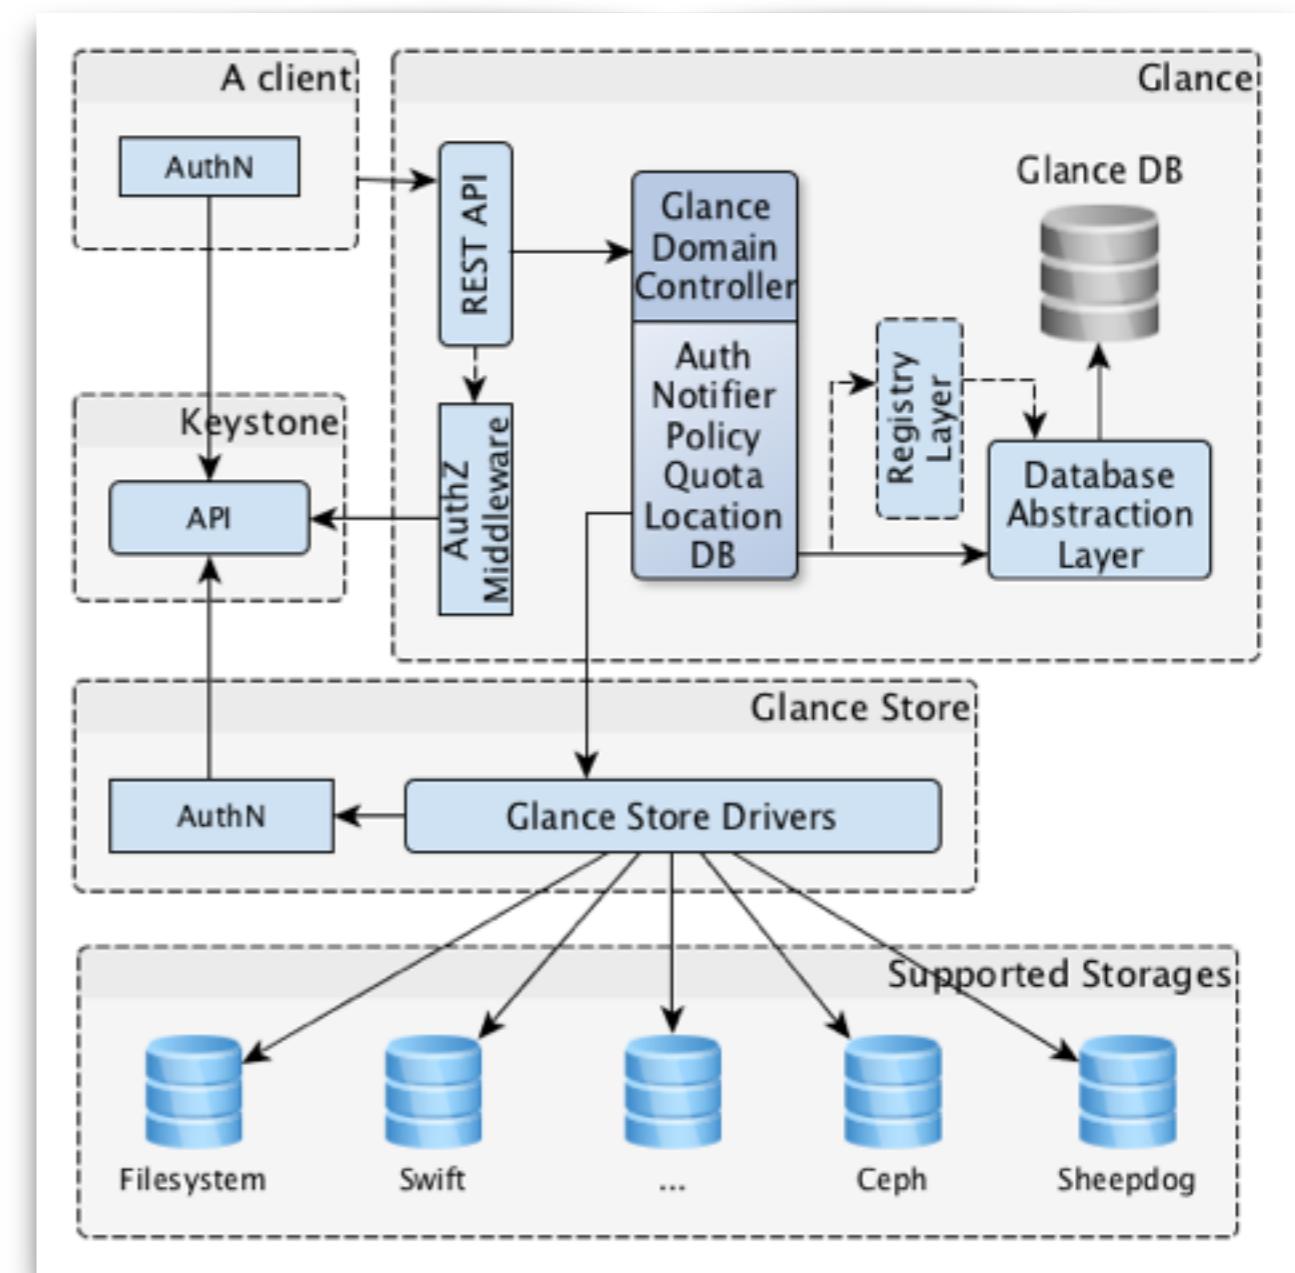
North/South - Private or Public IPs

# Focus on Glance



- For each image, one can specify properties
  - Image kind (raw, cqow2, vmdk, iso...)
  - Architecture
  - Distribution
  - Version
  - Storage space requirements
  - RAM minimal size
  - ...

- Possible backends
  - Swift/S3
  - Ceph
  - HTTP
  - Local



# Other Services



# IRONIC bare metal

- Ironic provisions bare metal (as opposed to virtual) machines.
- It may be used independently or as part of an OpenStack Cloud, and integrates with the OpenStack Identity (keystone), Compute (nova), Network (neutron), Image (glance) and Object (swift) services.
- When the Bare Metal service is appropriately configured with the Compute and Network services, it is possible to provision both virtual and physical machines through the Compute service's API.
- Although the project is mature, it is barely used...

# Other Services...

- A lot...
-  CEILOMETER: Metering & Data Collection Service
-  DESIGNATES: DNS-as-a-Service
-  OCTAVIA: load balancer
-  TROVE: a Database-as-a-service (SQL and NoSQL)
-  SAHARA: Big Data Processing Framework Provisioning
-  ZUN: Container Management Service
- ...
-  and HEAT: Orchestration

## Other Services



## HEAT orchestration

- Heat orchestrates the infrastructure resources for a cloud application based on templates (HOT) in the form of text files that can be treated like code.
- Heat provides both an OpenStack-native ReST API and a CloudFormation-compatible Query API.
- Heat also provides an autoscaling service that integrates with the OpenStack Telemetry services, so you can include a scaling group as a resource in a template.

# Focus on



# HEAT

- OpenStack-core services hide lot of complexity by automating the deployment process of a new VM/PM
  - Nova contacts Neutron for network configuration
  - Nova contacts Glance to fetch an OS image
  - Nova is in charge of booting and managing the VM
- But starting an empty OS is not enough to deliver a service...  
The hard part is to put the VM in context:
  - Install the software stack (service + its dependencies)
  - Configure the service
  - Set a floating IP address to be reachable from Internet
- A Cloud application is made of multiple services  
Multiple machines to boot, and services to deploy and configure to deliver the application

# Focus on



# HEAT



- DEVOPS Philosophy
- Infrastructure as Code
  - Scale in/out (horizontal scaling)
  - Automation
- Monitoring services/apps (instead of the infrastructure)
- Backup/Restart on demand

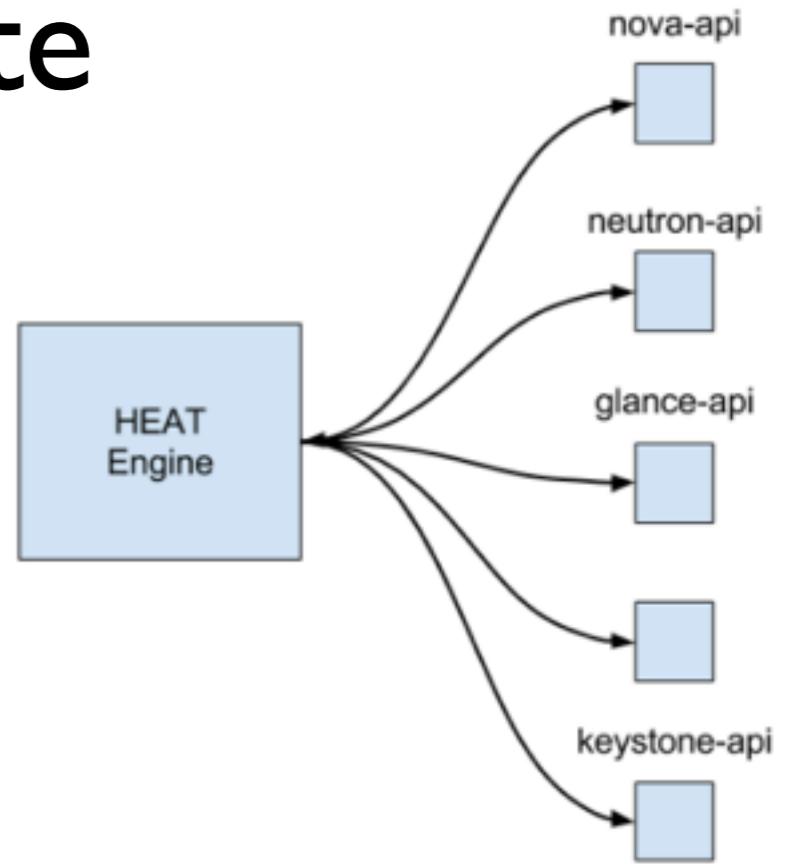
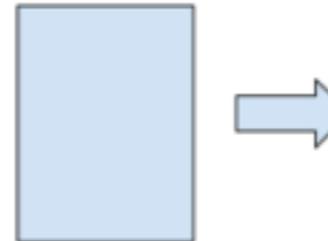
Manage Application Life-Cycle  
with code !



# HEAT Template

- Readable format:  
HEAT Orchestration Template (HOT)  
YAML file
- Describe infrastructure and applications
- Declare any OpenStack resource types  
Instances, floating IPs, volumes, images, users, ...
- Declare relationships between resources  
e.g. a VM must be booted before installing software on it
- HEAT Engine  
Take a template as input  
Parse it  
Execute tasks through OpenStack API calls

HOT  
Heat Orchestration Template



# Understanding

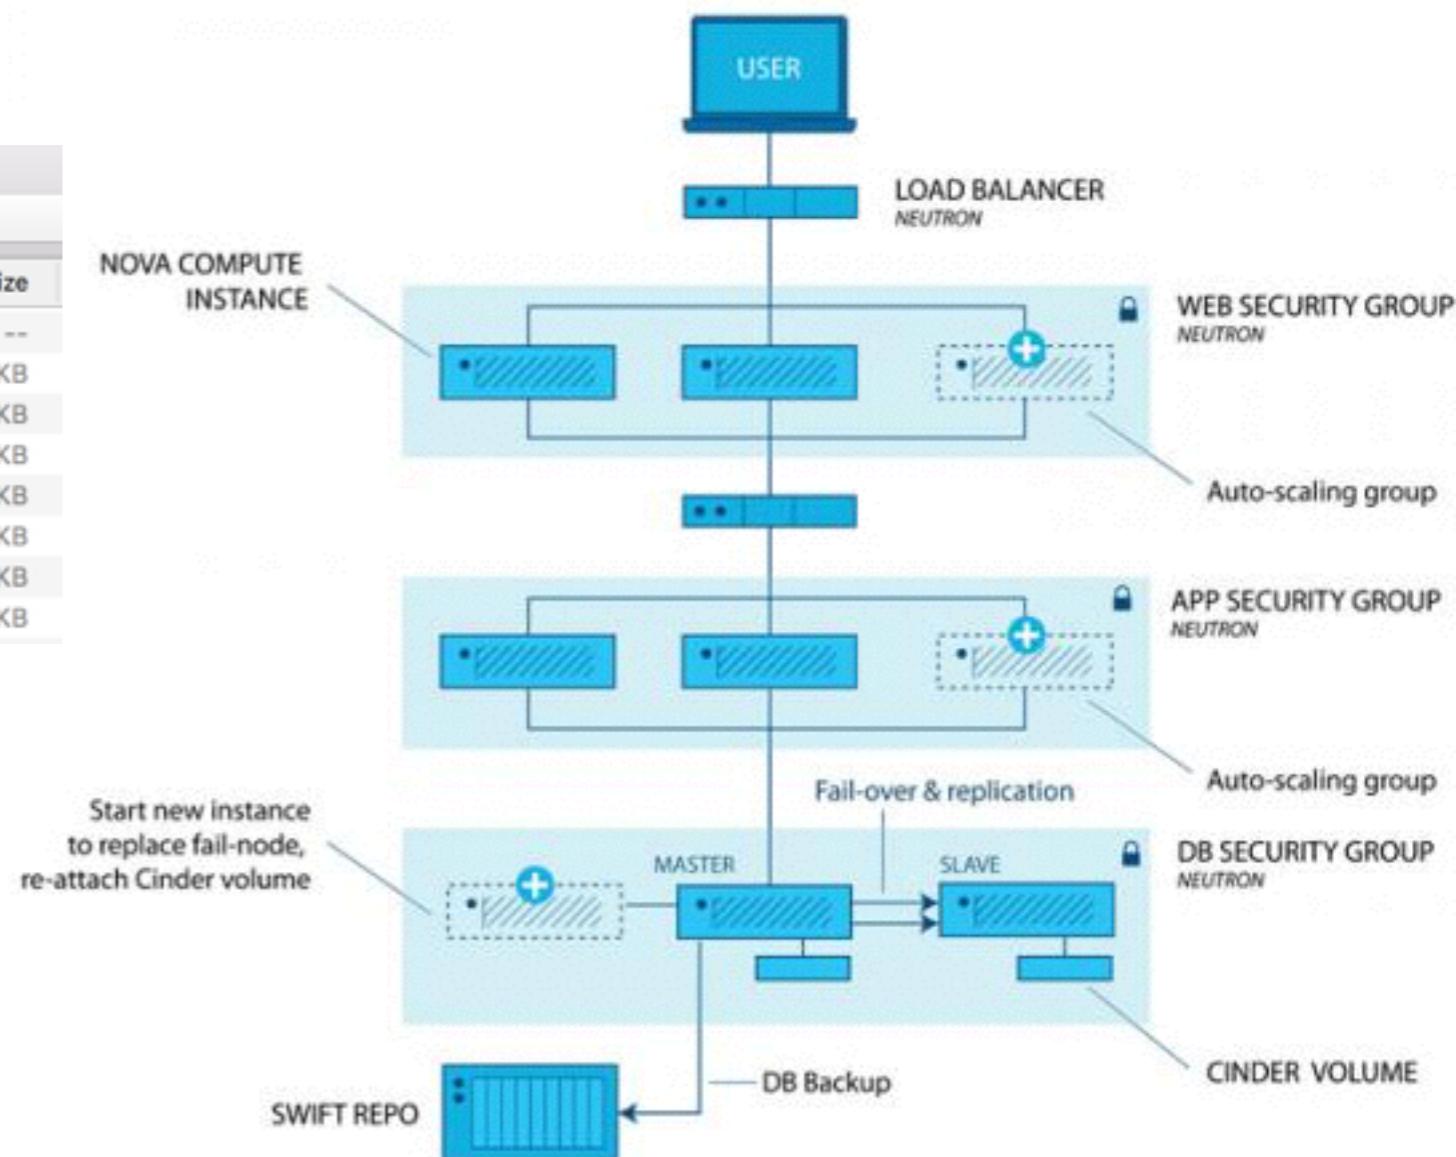


# HEAT Through an Example

- Exemple: a 3Tier-Web Application : Wordpress

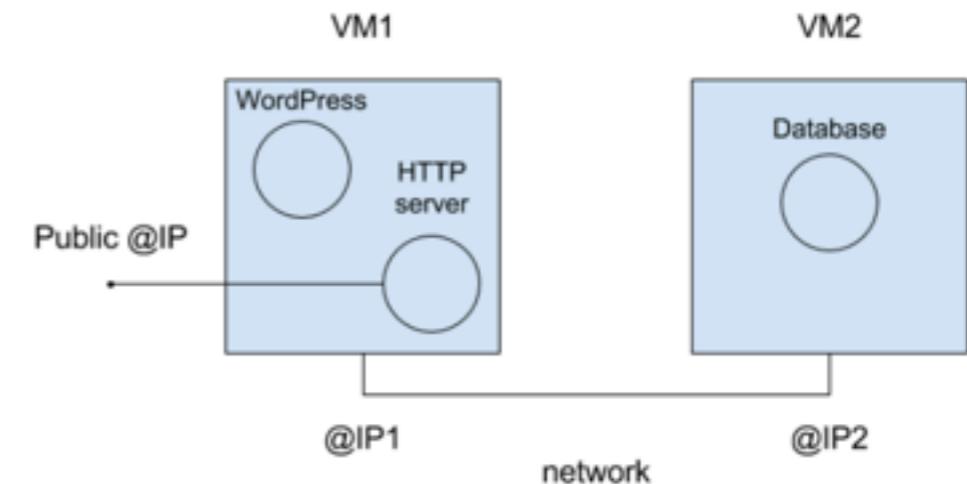
Screenshot of a file browser showing the contents of a 'web-application' directory:

web-application		
Name	Date Modified	Size
lib	nova-compute build steps	--
heat_app_tier.yaml	23 Sep 2016 15:07	5 KB
heat_sql_tier.yaml	23 Sep 2016 15:07	8 KB
heat_web_tier.yaml	23 Sep 2016 15:07	5 KB
setup_net_sg.yaml	23 Sep 2016 15:07	12 KB
README.rst	23 Sep 2016 15:15	4 KB
WebAppAutoScaling.yaml	27 Sep 2016 12:11	13 KB
WebAppStatic.yaml	23 Sep 2016 15:16	8 KB



# Understanding HEAT Through an Example

- Boot two virtual machines VM1 and VM2 Ubuntu based OS Connect them to the private network
- Install the services Install a database server (mysql) on VM2 Install a HTTP server (apache2) on VM1 Download Wordpress on VM1
- Configure services Create an appropriate database in mysql for wordpress Configure wordpress to access this database Configure apache2 to serve wordpress
- Assign a floating IP to VM1



How can I automate  
this process by leveraging HEAT?



# HEAT Template

- The Hello World example - boot a VM

```
$ cat boot.yaml
heat_template_version: 2015-04-30

description: Simple template to deploy a single compute instance

resources:                      # HEAT resources are declared here
  my_instance:                  # Name of my resource
    type: OS::Nova::Server      # Type of my resource (this resources defines a VM)
    properties:                 # Here we define the properties of this resource type
      key_name: my_key_name     # Name of an SSH key managed by Nova (or Barbican)
      image: ubuntu-trusty-x86_64 # Name of an image managed by Glance
      flavor: m1.small           # Name of a flavor managed by Nova

$ openstack stack create my_stack -f boot.yaml
```



# HEAT Template

- Let's make it a bit more complex: boot a vm with parameters

```
$ cat boot_with_parameter.yaml
heat_template_version: 2015-04-30

description: Simple template to deploy a single compute instance with a parameter

parameters:                      # Parameters definition for this template
  key_name:                      # Name of the parameter
    type: string                  # Type of the parameter
    description: Name of a KeyPair to enable SSH access to the instance

resources:
  my_instance:
    type: OS::Nova::Server
    properties:
      key_name: { get_param: key_name }    # Use an intrinsic function to get the
value of a parameter
      image: ubuntu-trusty-x86_64
      flavor: m1.small

$ openstack stack create my_stack -f boot_with_parameters.yaml --parameter
key_name=my_key
```



# HEAT Template

- Let's make it a bit more complex: boot a vm and get some outputs

```
$ cat boot_with_outputs.yaml
heat_template_version: 2015-04-30

description: Simple template to deploy a single compute instance, outputs its ip
address

resources:
  my_instance:
    type: OS::Nova::Server
    properties:
      key_name: my_key_name
      image: ubuntu-trusty-x86_64
      flavor: m1.small

outputs:                                # Definition of the outputs of this template
  instance_ip:                          # Name of the output
    description: IP address of the deployed compute instance
    value: { get_attr: [my_instance, first_address] } # Set the IP address of the
                                                       # machine as the value of the
                                                       # output instance_ip

$ openstack stack create my_stack -f boot_with_outputs.yaml
```



# HEAT Template

- Even more complex

```
$ cat boot_sql_server.yaml
heat_template_version: 2015-04-30
description: Template to deploy SQL server
parameters:
  DBRootPassword:
    type: string
resources:
  my_sql_instance:
    type: OS::Nova::Server
    properties:
      # general properties ...
      user_data:           # Definition of a boot script
      str_replace:         # Intrinsic function to replace
                           # strings in the script by parameters
      template: |          # Description of the script
        #!/bin/bash
        # do things like install mysql ...
        mysqladmin -u root password $db_rootpassword
        # do more things ...
    params:               # Description of the used parameters
    $db_rootpassword: { get_param: DBRootPassword }...
```

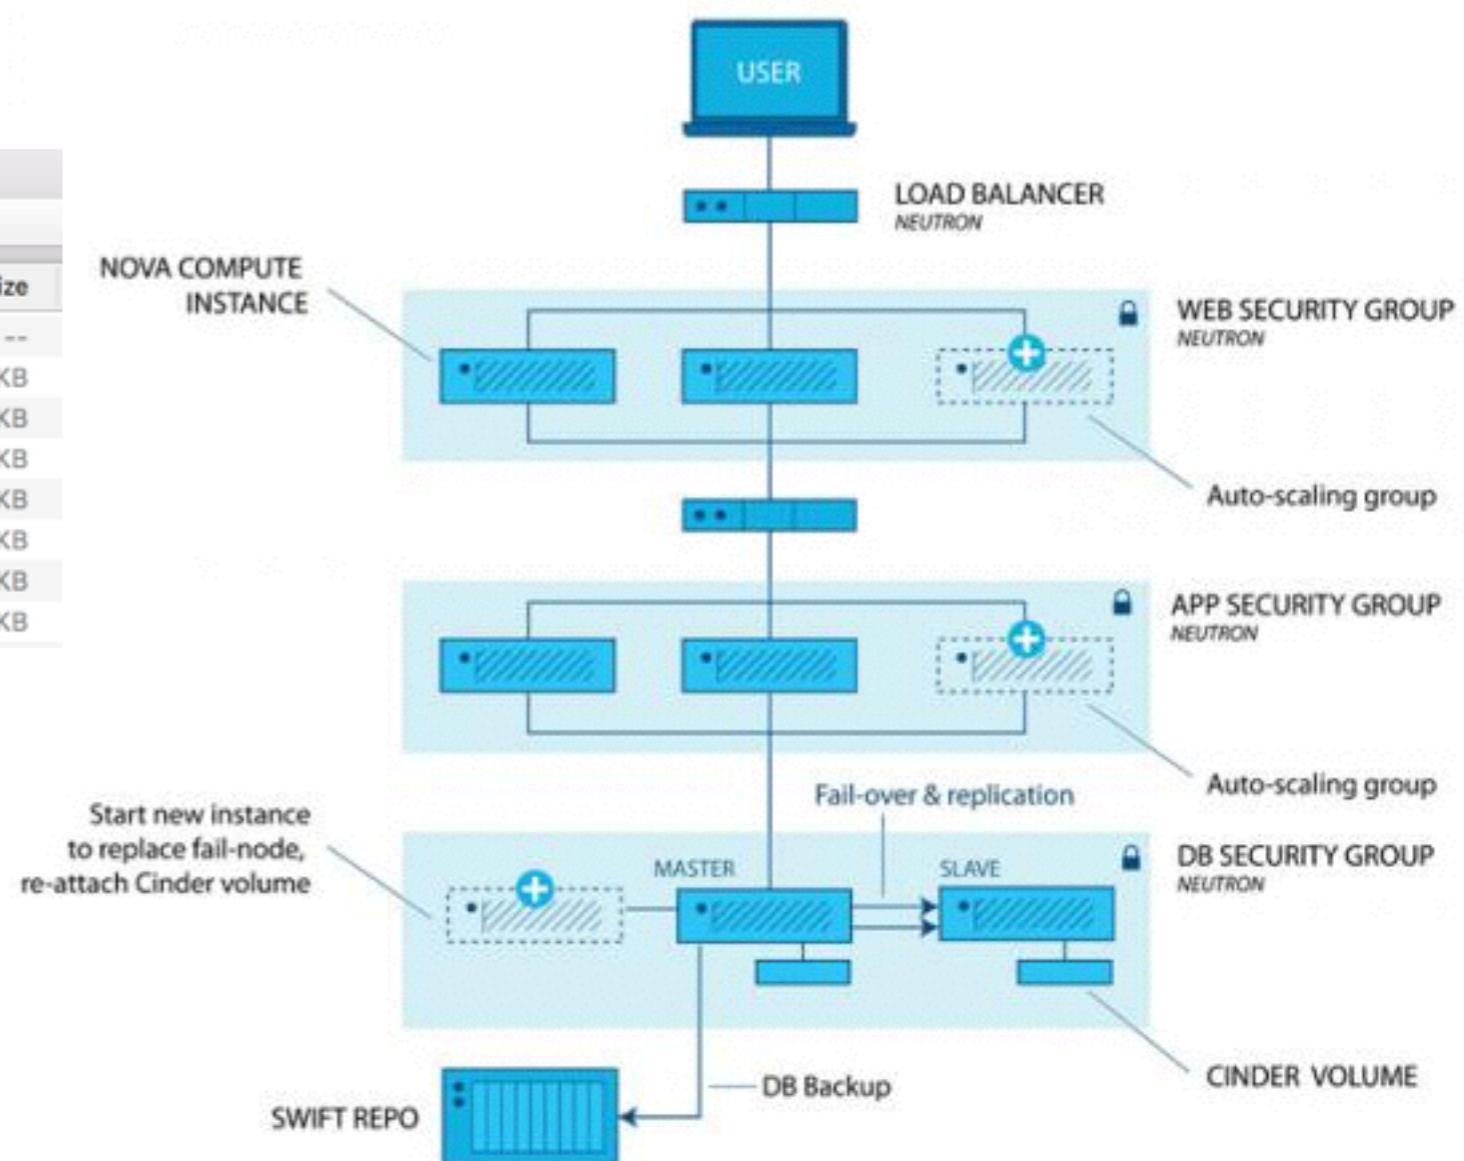
```
$ openstack stack create my_sql_server -f boot_sql_server.yaml --parameter
DBRootPassword=0p3nSt4cK
```

# Understanding HEAT Through an Example

- The complete example during the practical session !

Screenshot of a file browser showing the contents of a 'web-application' directory:

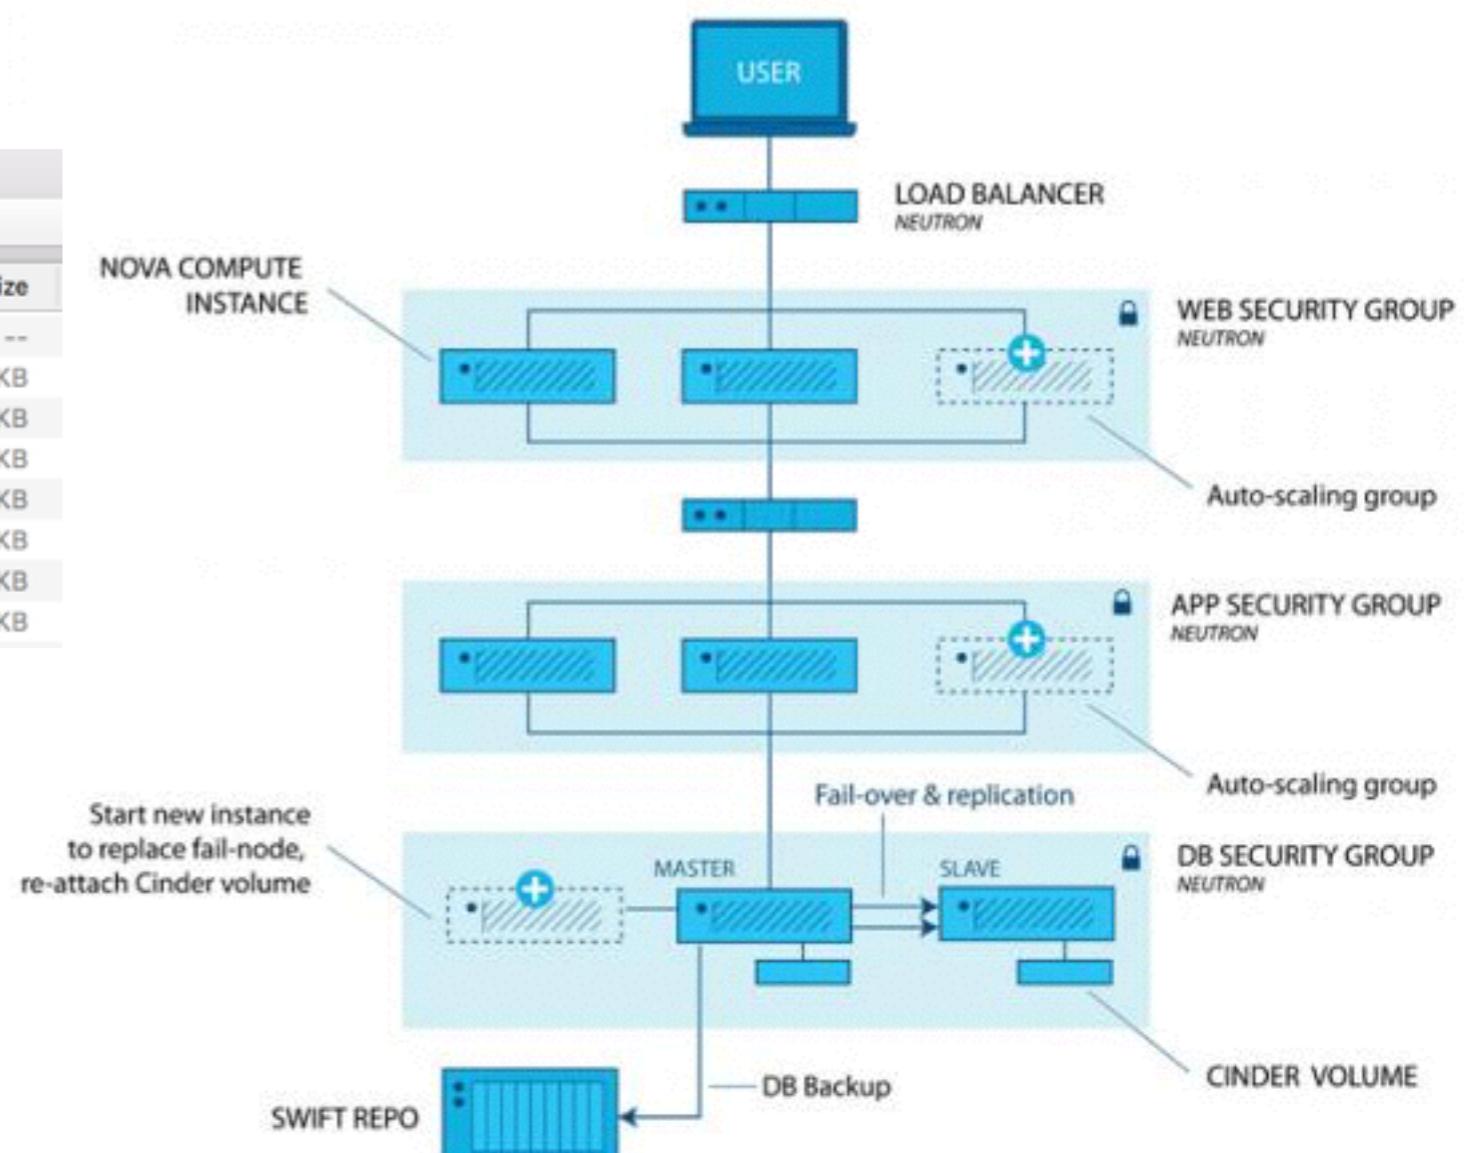
web-application		
Name	Date Modified	Size
lib	nova-compute build steps	--
heat_app_tier.yaml	23 Sep 2016 15:07	5 KB
heat_sql_tier.yaml	23 Sep 2016 15:07	8 KB
heat_web_tier.yaml	23 Sep 2016 15:07	5 KB
setup_net_sg.yaml	23 Sep 2016 15:07	12 KB
README.rst	23 Sep 2016 15:15	4 KB
WebAppAutoScaling.yaml	27 Sep 2016 12:11	13 KB
WebAppStatic.yaml	23 Sep 2016 15:16	8 KB



# Understanding HEAT Through an Example

- The complete example during the practical session !

web-application		
Name	Date Modified	Size
lib	nova-compute build steps	--
heat_app_tier.yaml	23 Sep 2016 15:07	5 KB
heat_sql_tier.yaml	23 Sep 2016 15:07	8 KB
heat_web_tier.yaml	23 Sep 2016 15:07	5 KB
setup_net_sg.yaml	23 Sep 2016 15:07	12 KB
README.rst	23 Sep 2016 15:15	4 KB
WebAppAutoScaling.yaml	27 Sep 2016 12:11	13 KB
WebAppStatic.yaml	23 Sep 2016 15:16	8 KB

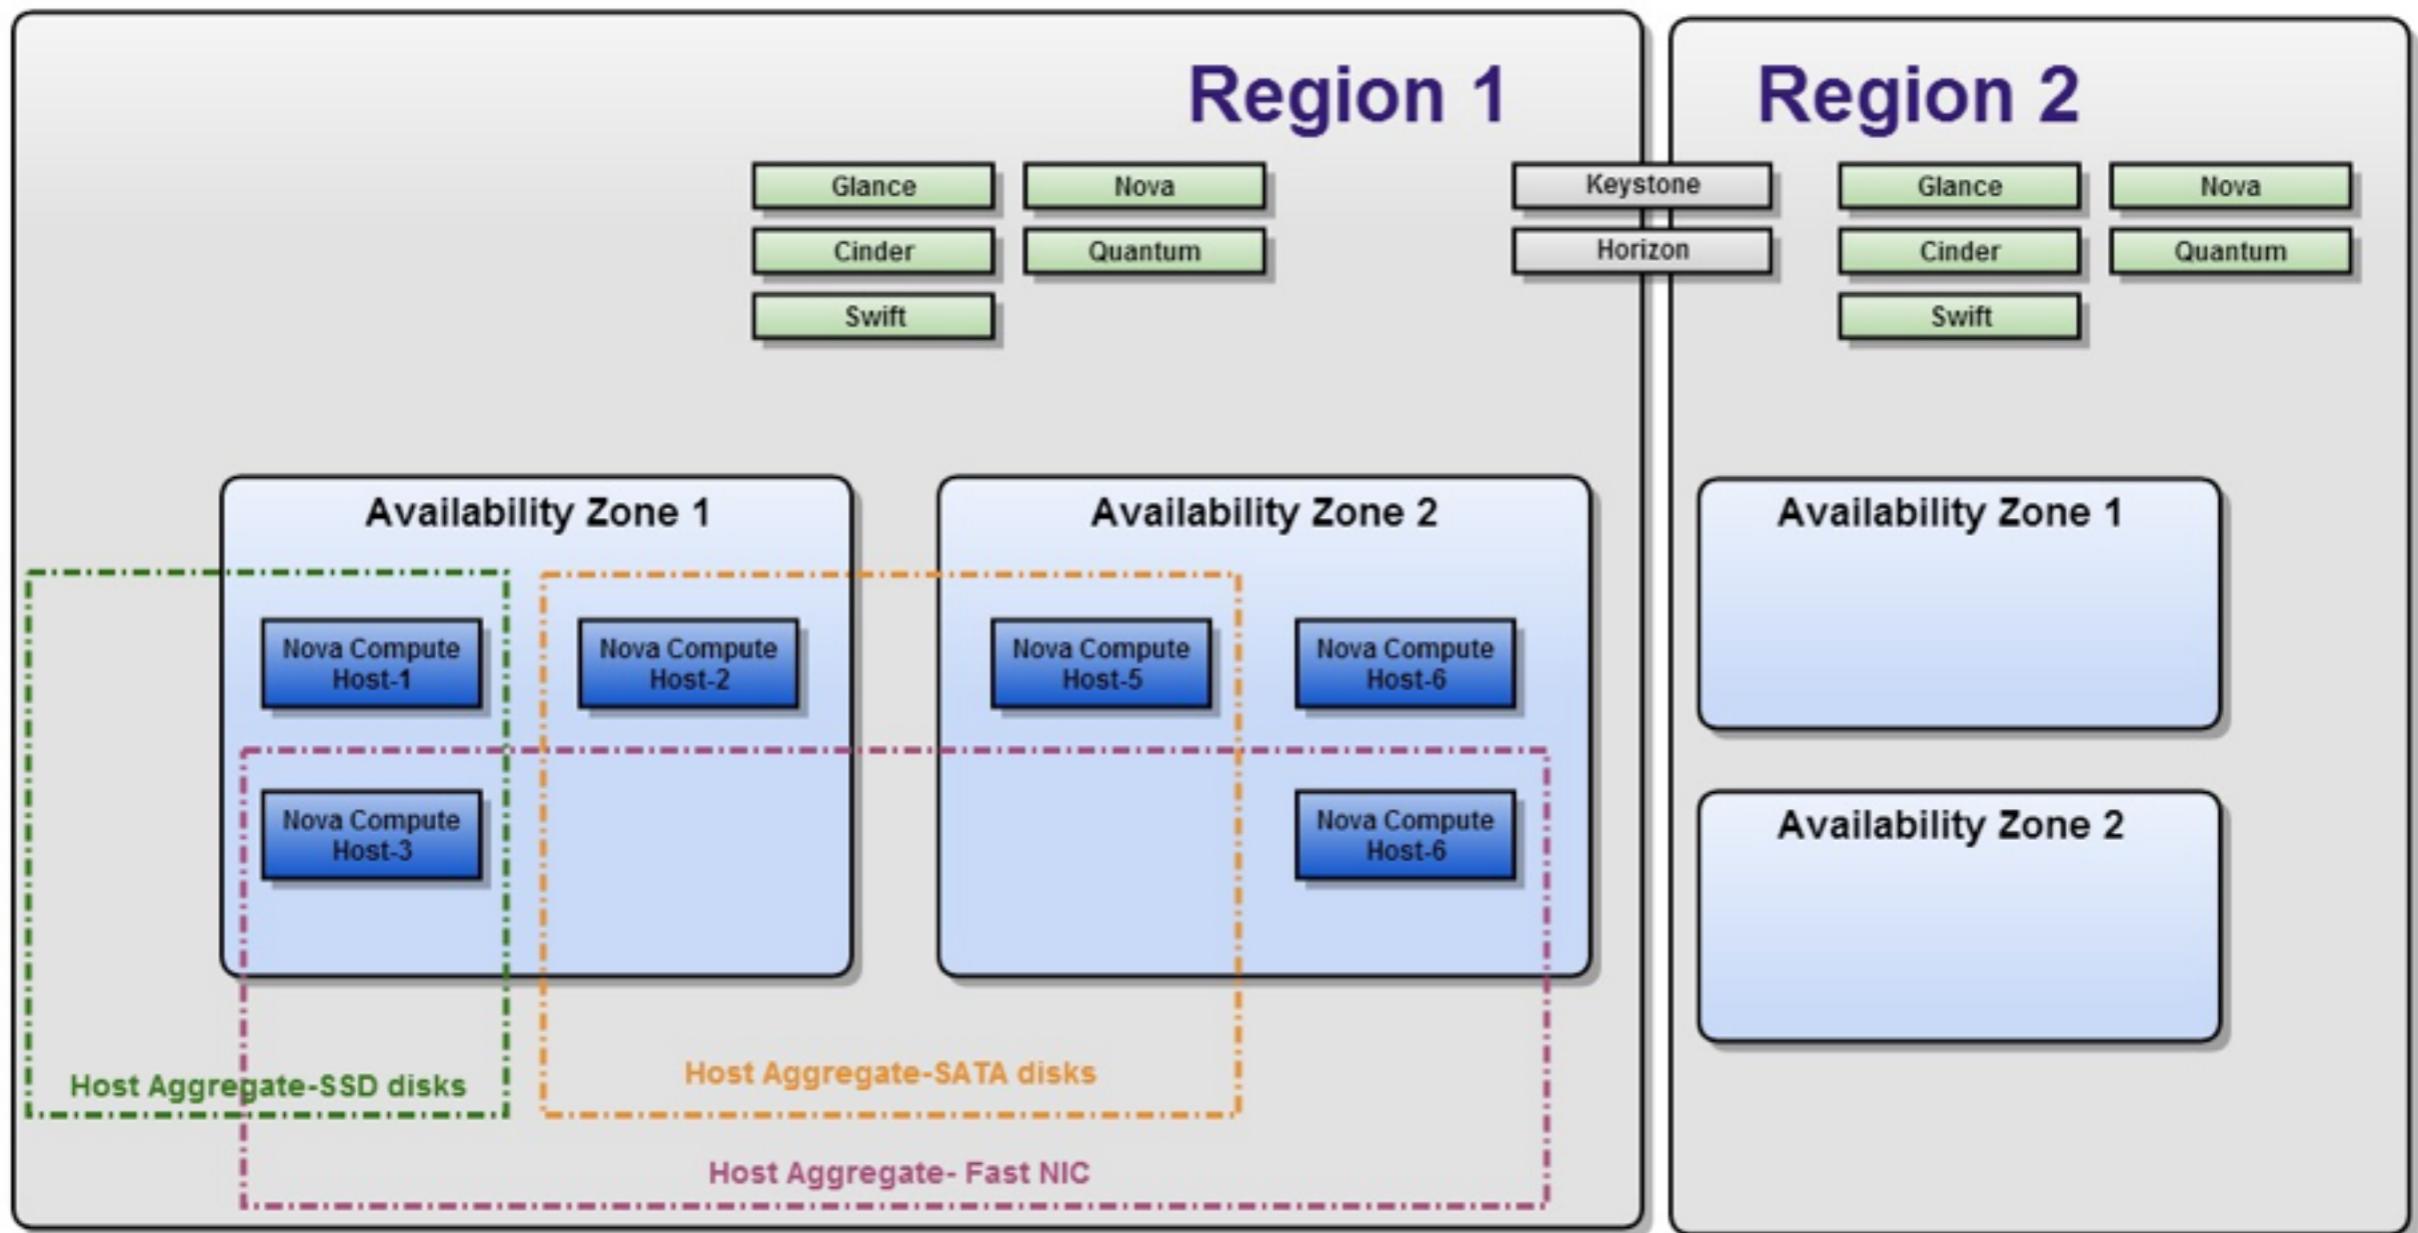


**R-A Cherrueau  
D. Pertin  
Next Week !**

# Segregation Tools

- It is sometimes a key point to segregate a large infrastructure into several subsets. OpenStack provides several ways:
  - Host aggregates - Nova: Classify compute nodes according to their specifics (Storage, GPU, ...)
  - Availability Zones - Nova, Cinder: Classify resources according to availability aspects (racks, data centers, ...)
  - Cells - Nova (Neutron soon !?): address segregation and scalability needs (one communication bus and one DB per cells)
  - Regions - a federation like approach, each region has an almost complete OpenStack (keystone/horizone are shared). Equivalent to the AWS region.

# Segregation Tools



*I want to test it !*

# DevStack



- A series of extensible scripts used to quickly bring up a complete OpenStack environment based on the latest versions of everything from git master  
`stack.sh + local.conf`
- A development environment and as the basis for much of the OpenStack project's functional testing.

# DevStack



- A series of extensible scripts used to quickly bring up a complete OpenStack environment based on the latest versions of everything from git master stack.sh + local.conf
- A development environment and as the basis for much of the OpenStack project's functional testing.

**WARNING: DevStack makes  
substantial changes to your system  
Only launch it inside a VM**

# Experimental eNvironment for OpenStack



- A dedicated framework to conduct performance analyses of OpenStack at large-scale in a reproducible manner. The framework enables engineers/researchers to conduct experiments in an automate manner on top of different testbeds such as Grid'5000, Chameleon, OpenStack...
- Developed in the context of the Discovery Initiative
- Deploy a real production system by leveraging Kolla (i.e. not DevStack)

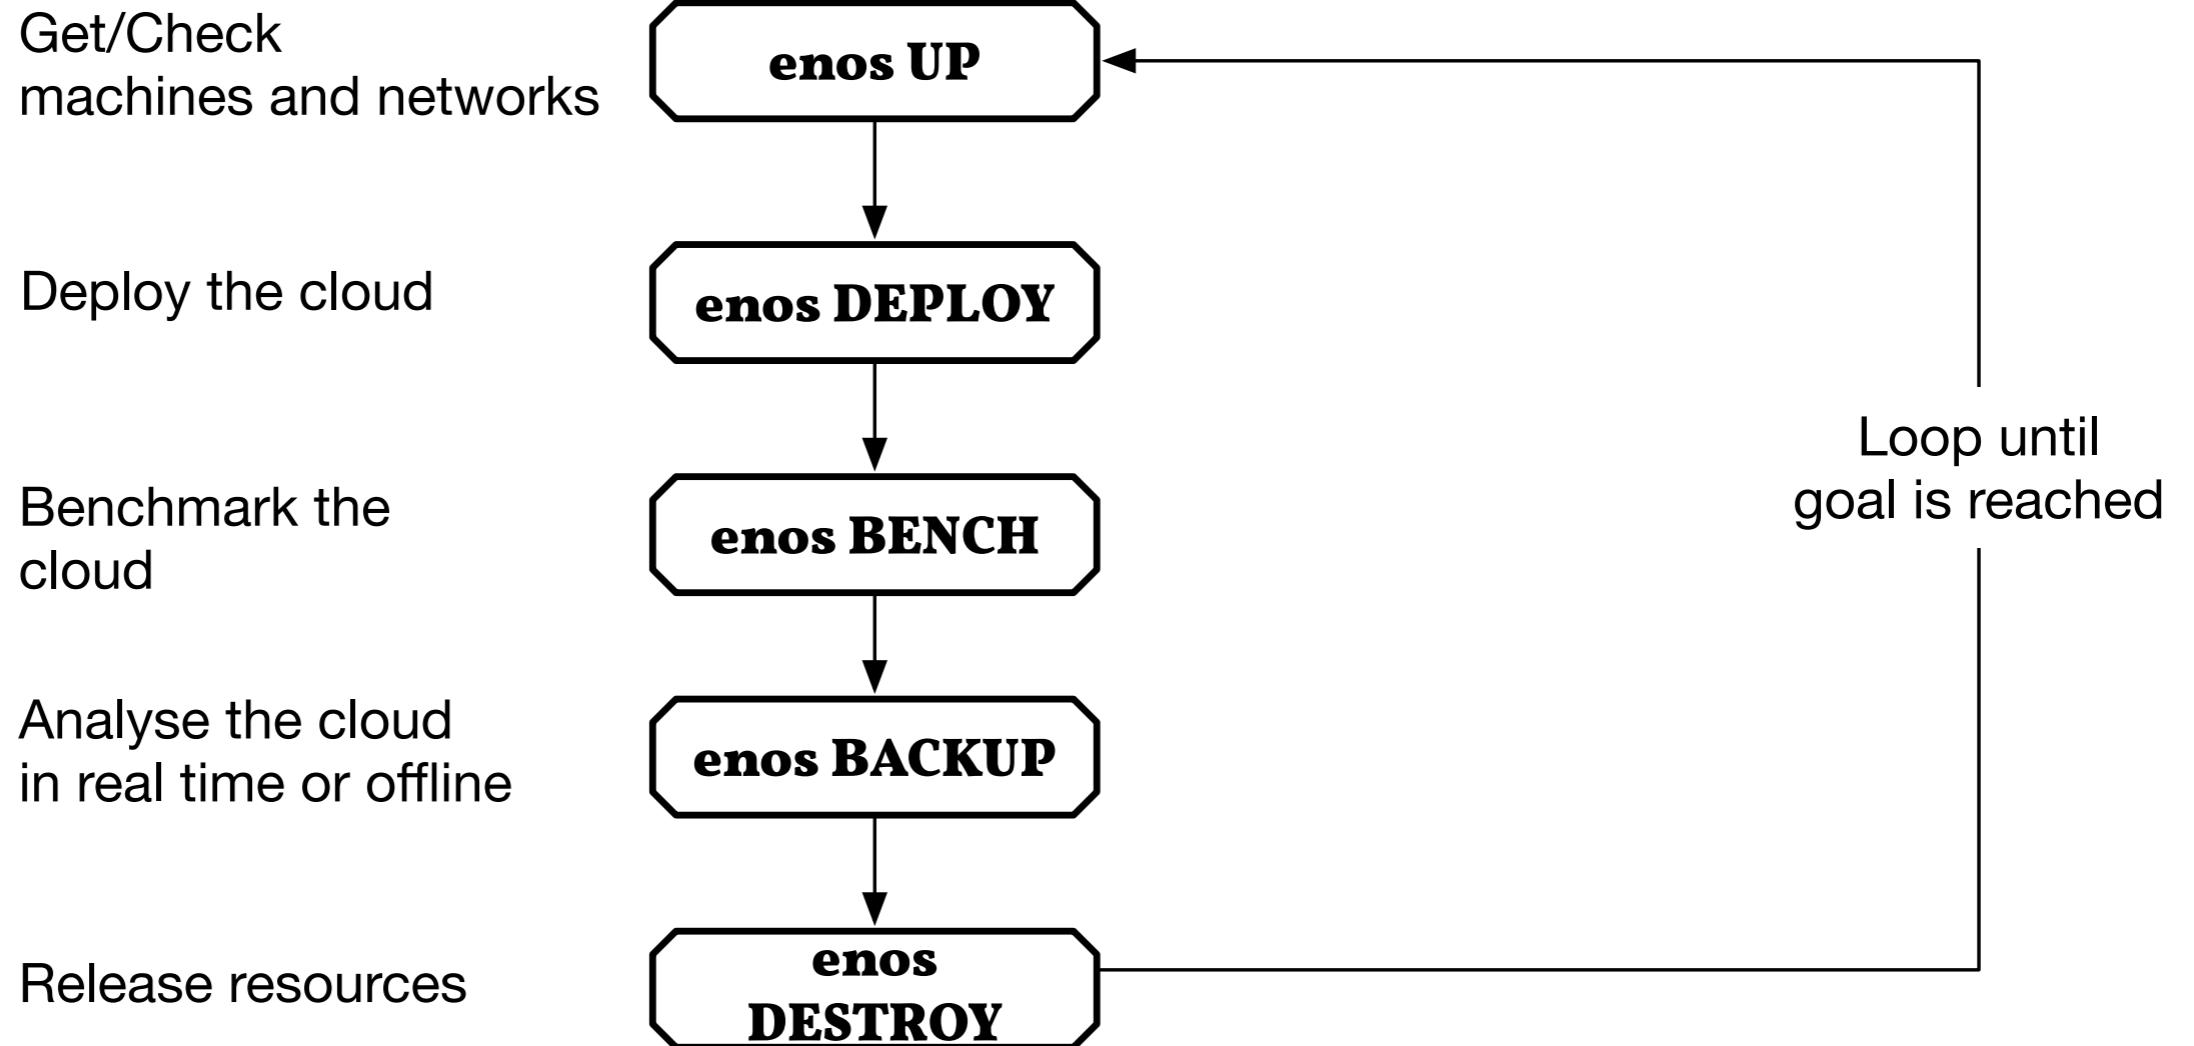
# Experimental eNvironment for OpenStack



- A dedicated framework to conduct performance analyses of OpenStack at large-scale in a reproducible manner. The framework enables engineers/researchers to conduct experiments in an automate manner on top of different testbeds such as Grid'5000, Chameleon, OpenStack...
- Developed in the context of the Discovery Initiative
- Deploy a real production system by leveraging Kolla (i.e. not DevStack)

**R-A Cherrueau / D. Pertin**  
**Next Week !**

# Experimental eNvironment for OpenStack



# Experimental eNvironment for OpenStack



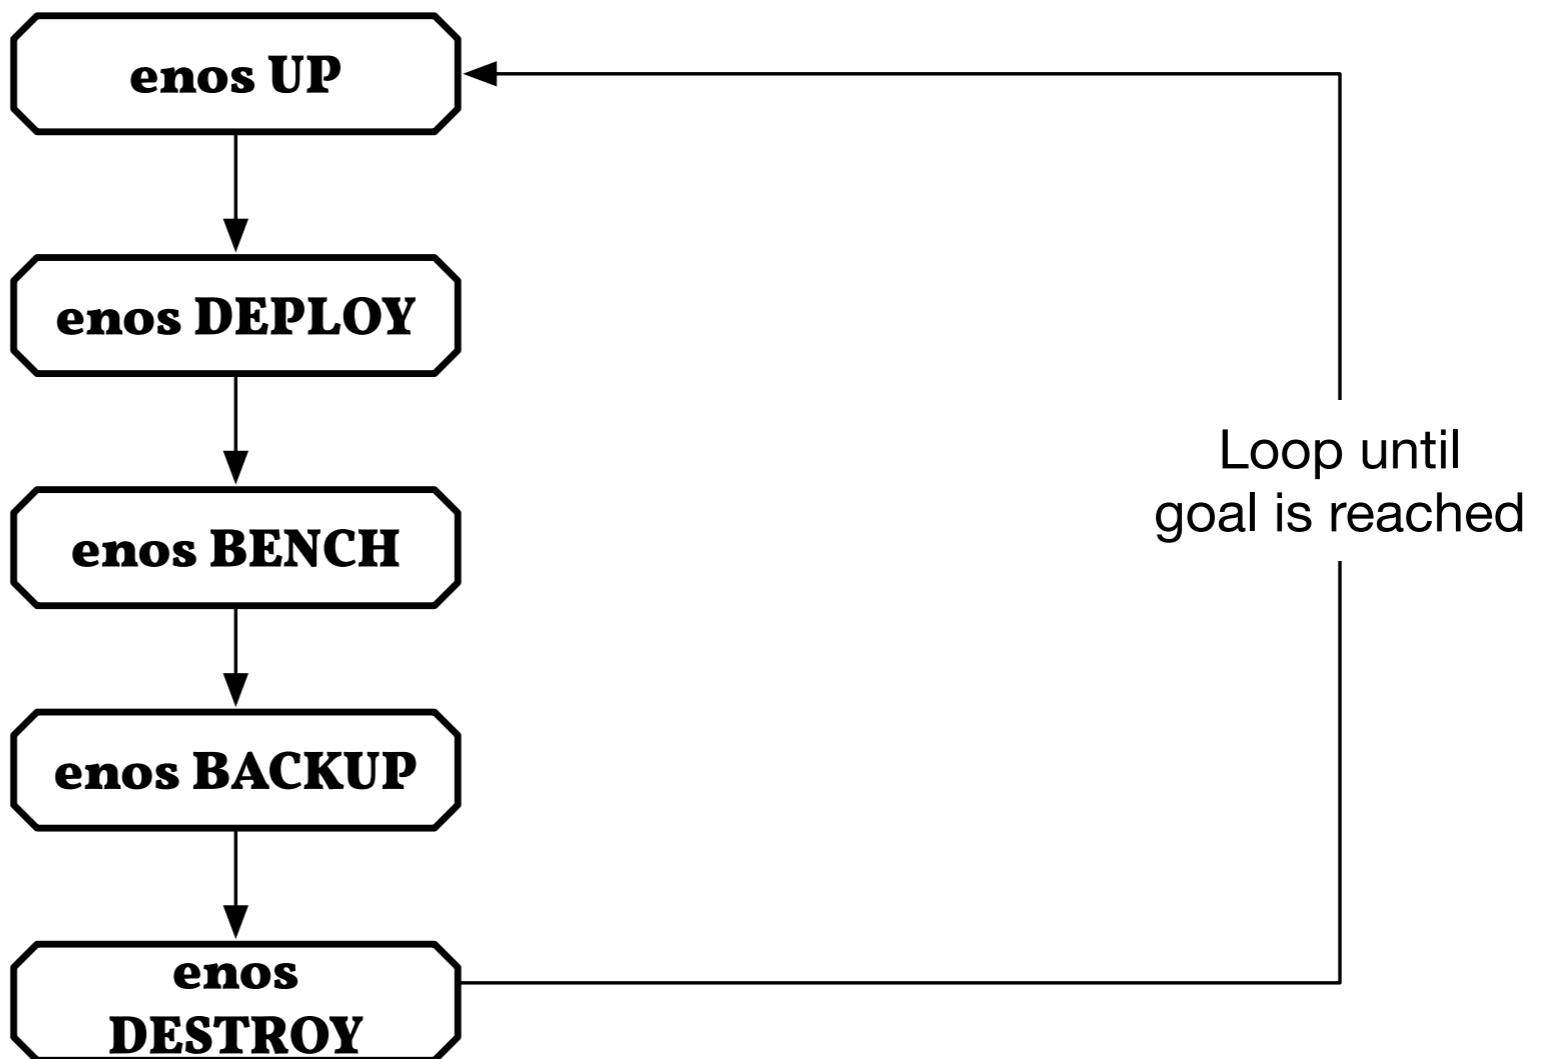
Get/Check  
machines and networks

Deploy the cloud

Benchmark the  
cloud

Analyse the cloud  
in real time or offline

Release resources



**R-A Cherruau / D. Pertin  
Next Week !**

# You cannot Wait For the Practical Session

- Install VirtualBox Vagrant EnOS
- Prepare your RAM (at least 10GB ;-))

Welcome to Enos's documentation! — Enos 3.0.1 documentation

Docs » Welcome to Enos's documentation! [Edit on GitHub](#)

## Welcome to Enos's documentation!

**Hint**

The source code is available at <https://github.com/BeyondTheClouds/enos>

Enos deploys OpenStack and targets reproducible experiments. It allows easy:

- deployment of the system
- customization of the system
- benchmarking of the system
- visualization of various metrics

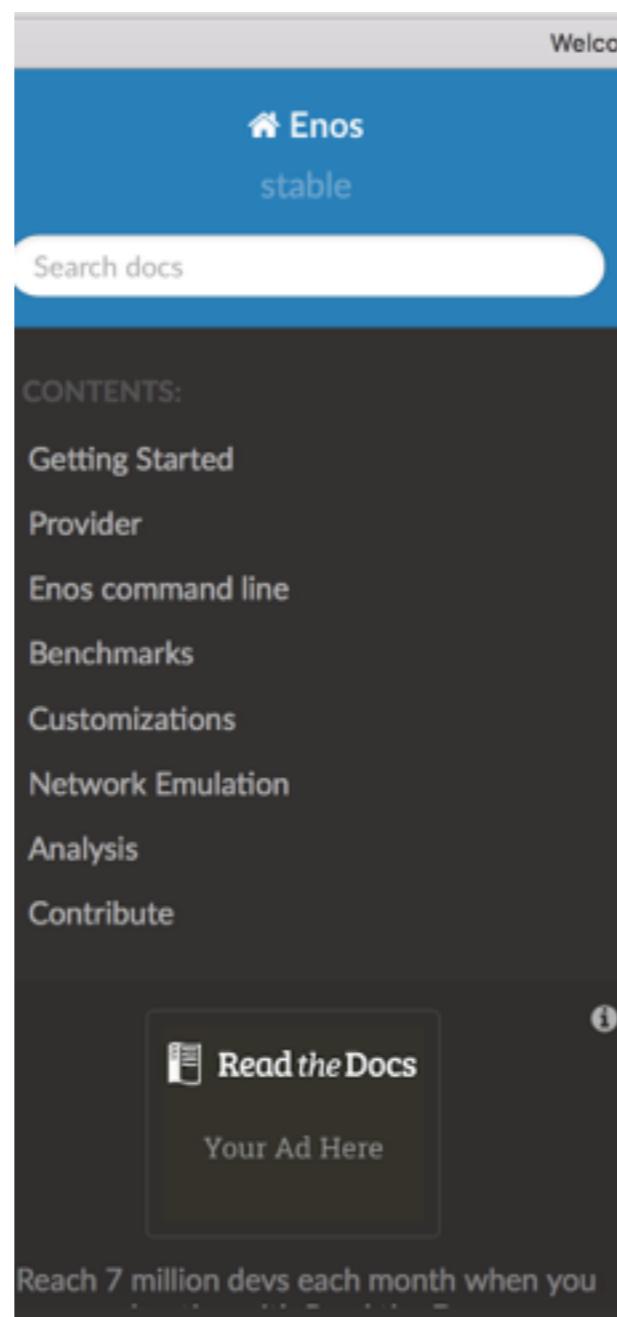
Enos is developed in the context of the [Discovery](#) initiative.

Reach 7 million devs each month when you [Read the Docs](#)

Your Ad Here

# You cannot Wait For the Practical Session

- Install VirtualBox Vagrant EnOS
- Prepare your RAM (at least 10GB ;-))



Welcome to Enos's documentation! — Enos 3.0.1 documentation

Docs » Welcome to Enos's documentation! [Edit on GitHub](#)

**Welcome to Enos's documentation!**

**Hint**

The source code is available at <https://github.com/BeyondTheClouds/enos>

Enos deploys OpenStack and targets reproducible experiments. It allows easy:

- deployment of the system
- customization of the system
- benchmarking of the system
- visualization of various metrics

Enos is developed in the context of the [Discovery](#) initiative.

**<https://enos.readthedocs.io/en/stable/>**

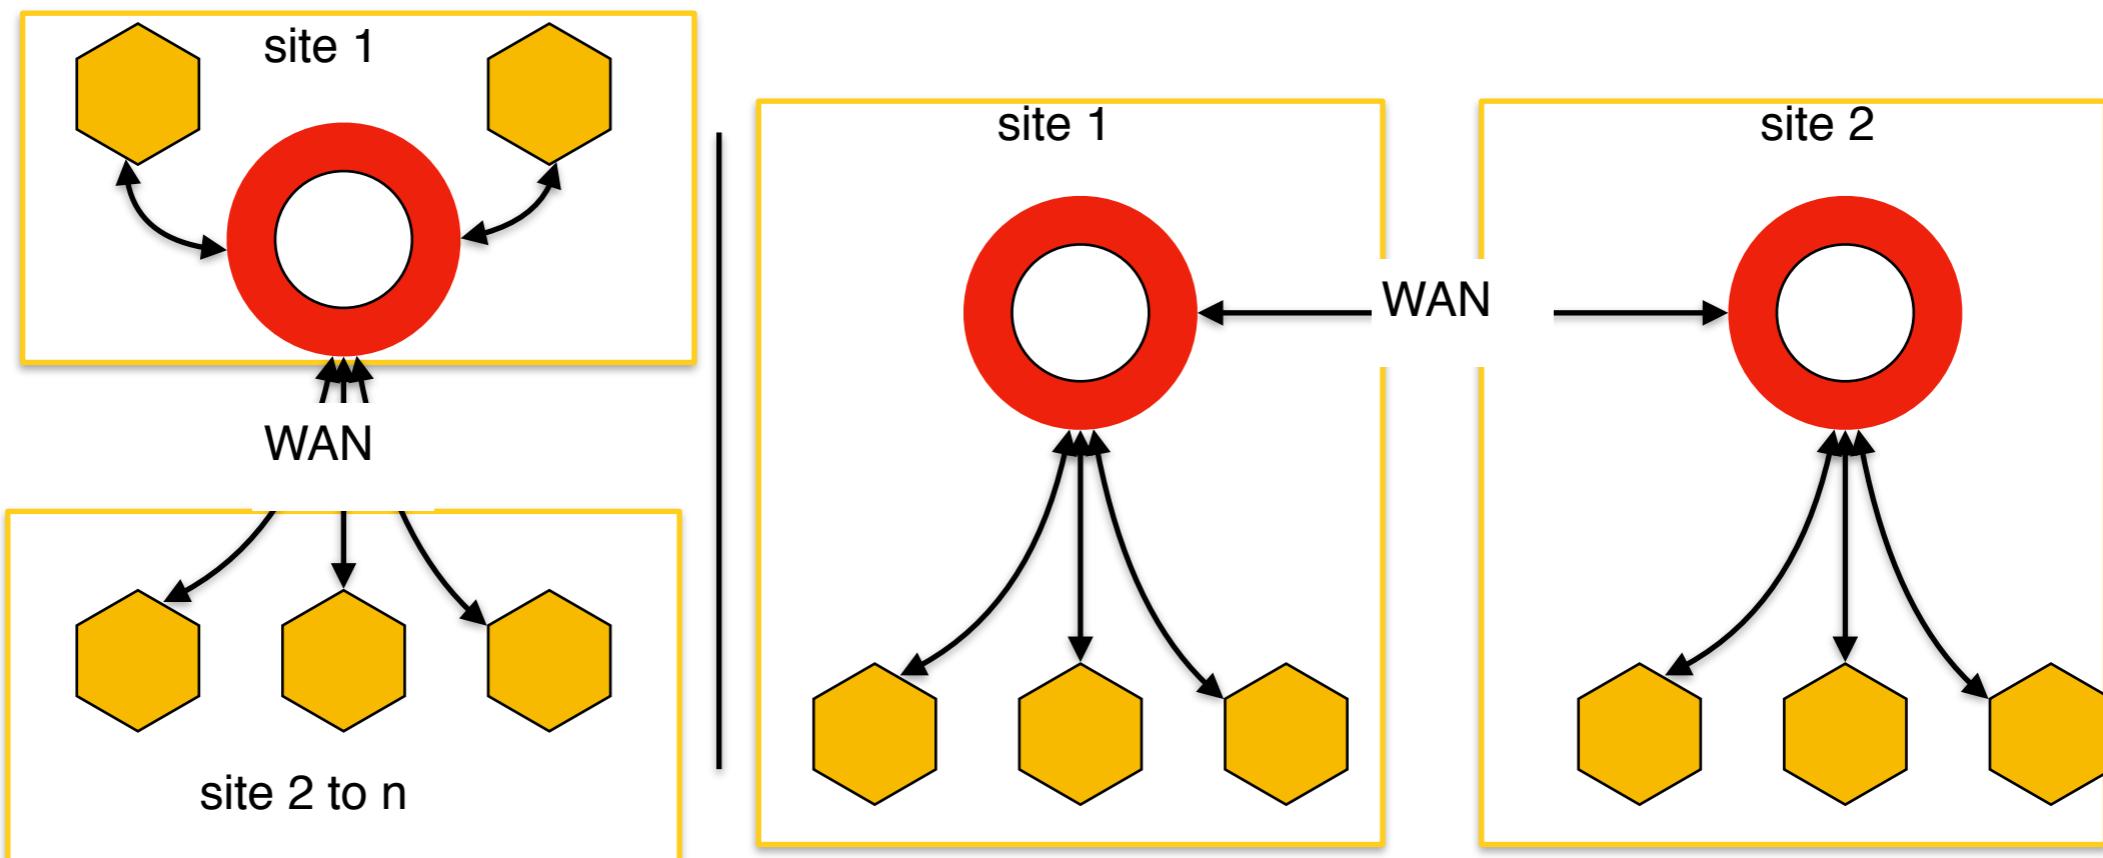
*Why diving in  
such a level of details?!*

# Fog/Edge Challenges

- **Communication Bus**

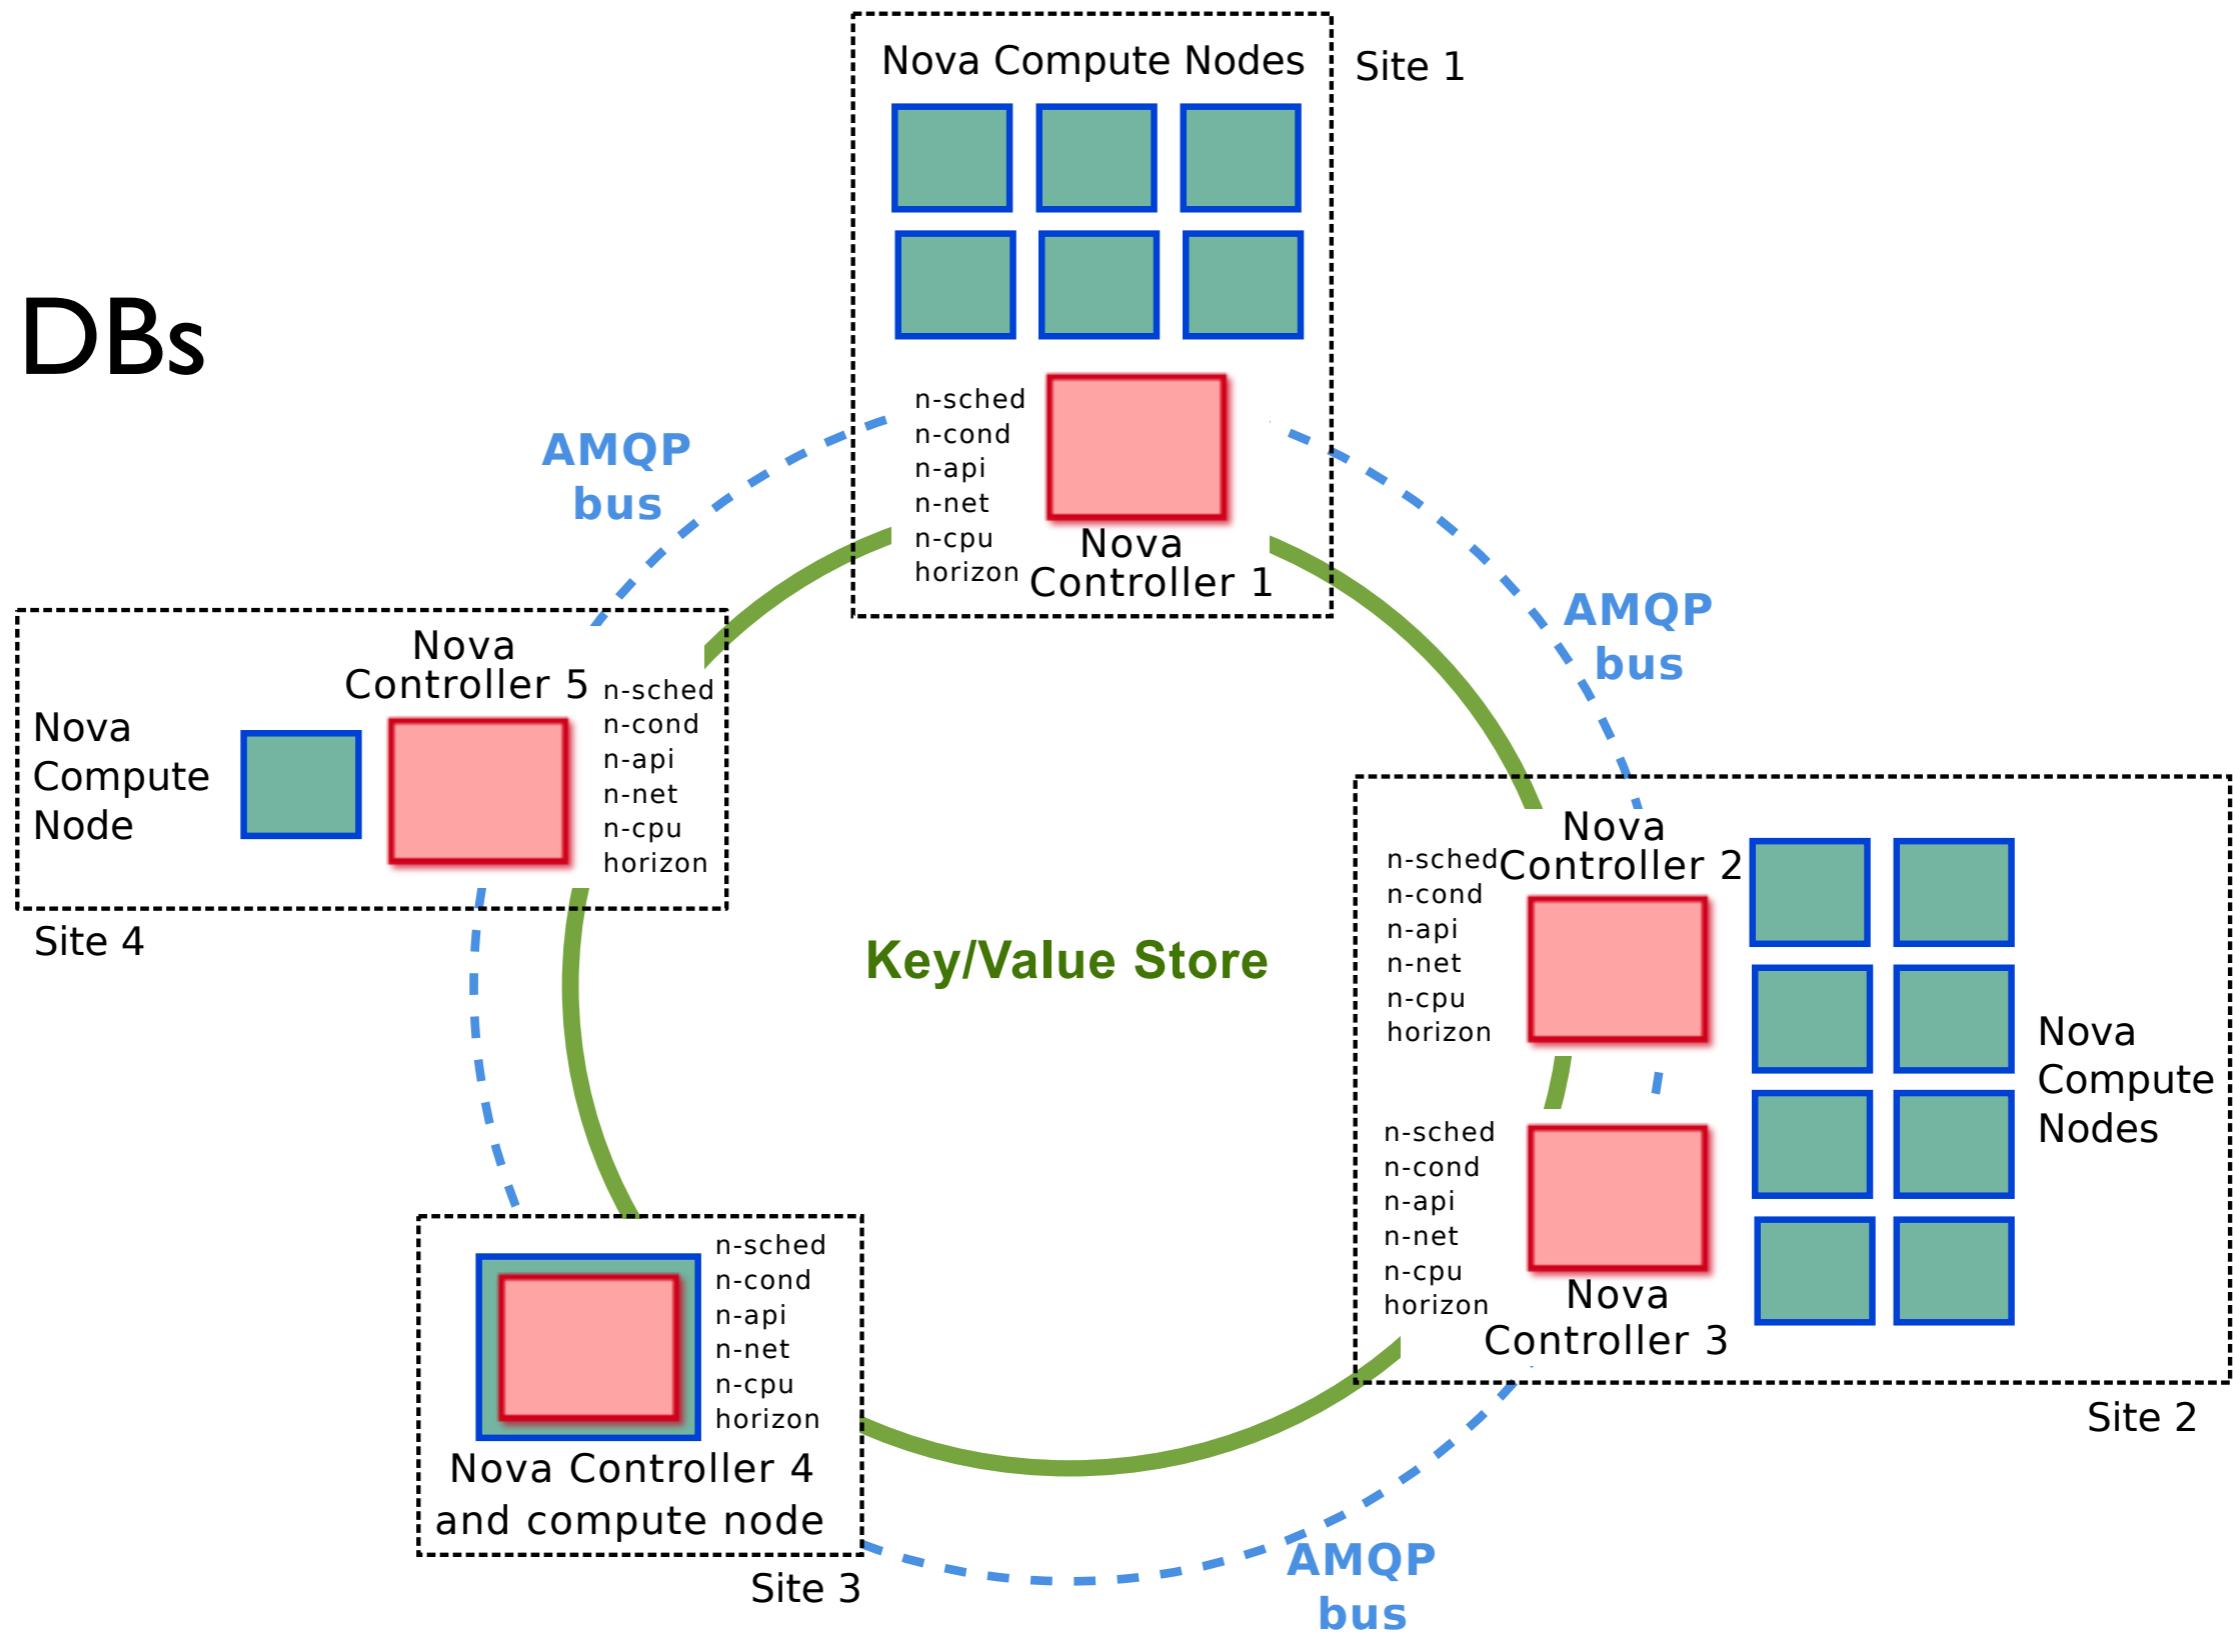
Central rabbitMQ and many edge servers

Distributed RabbitMQ through federations



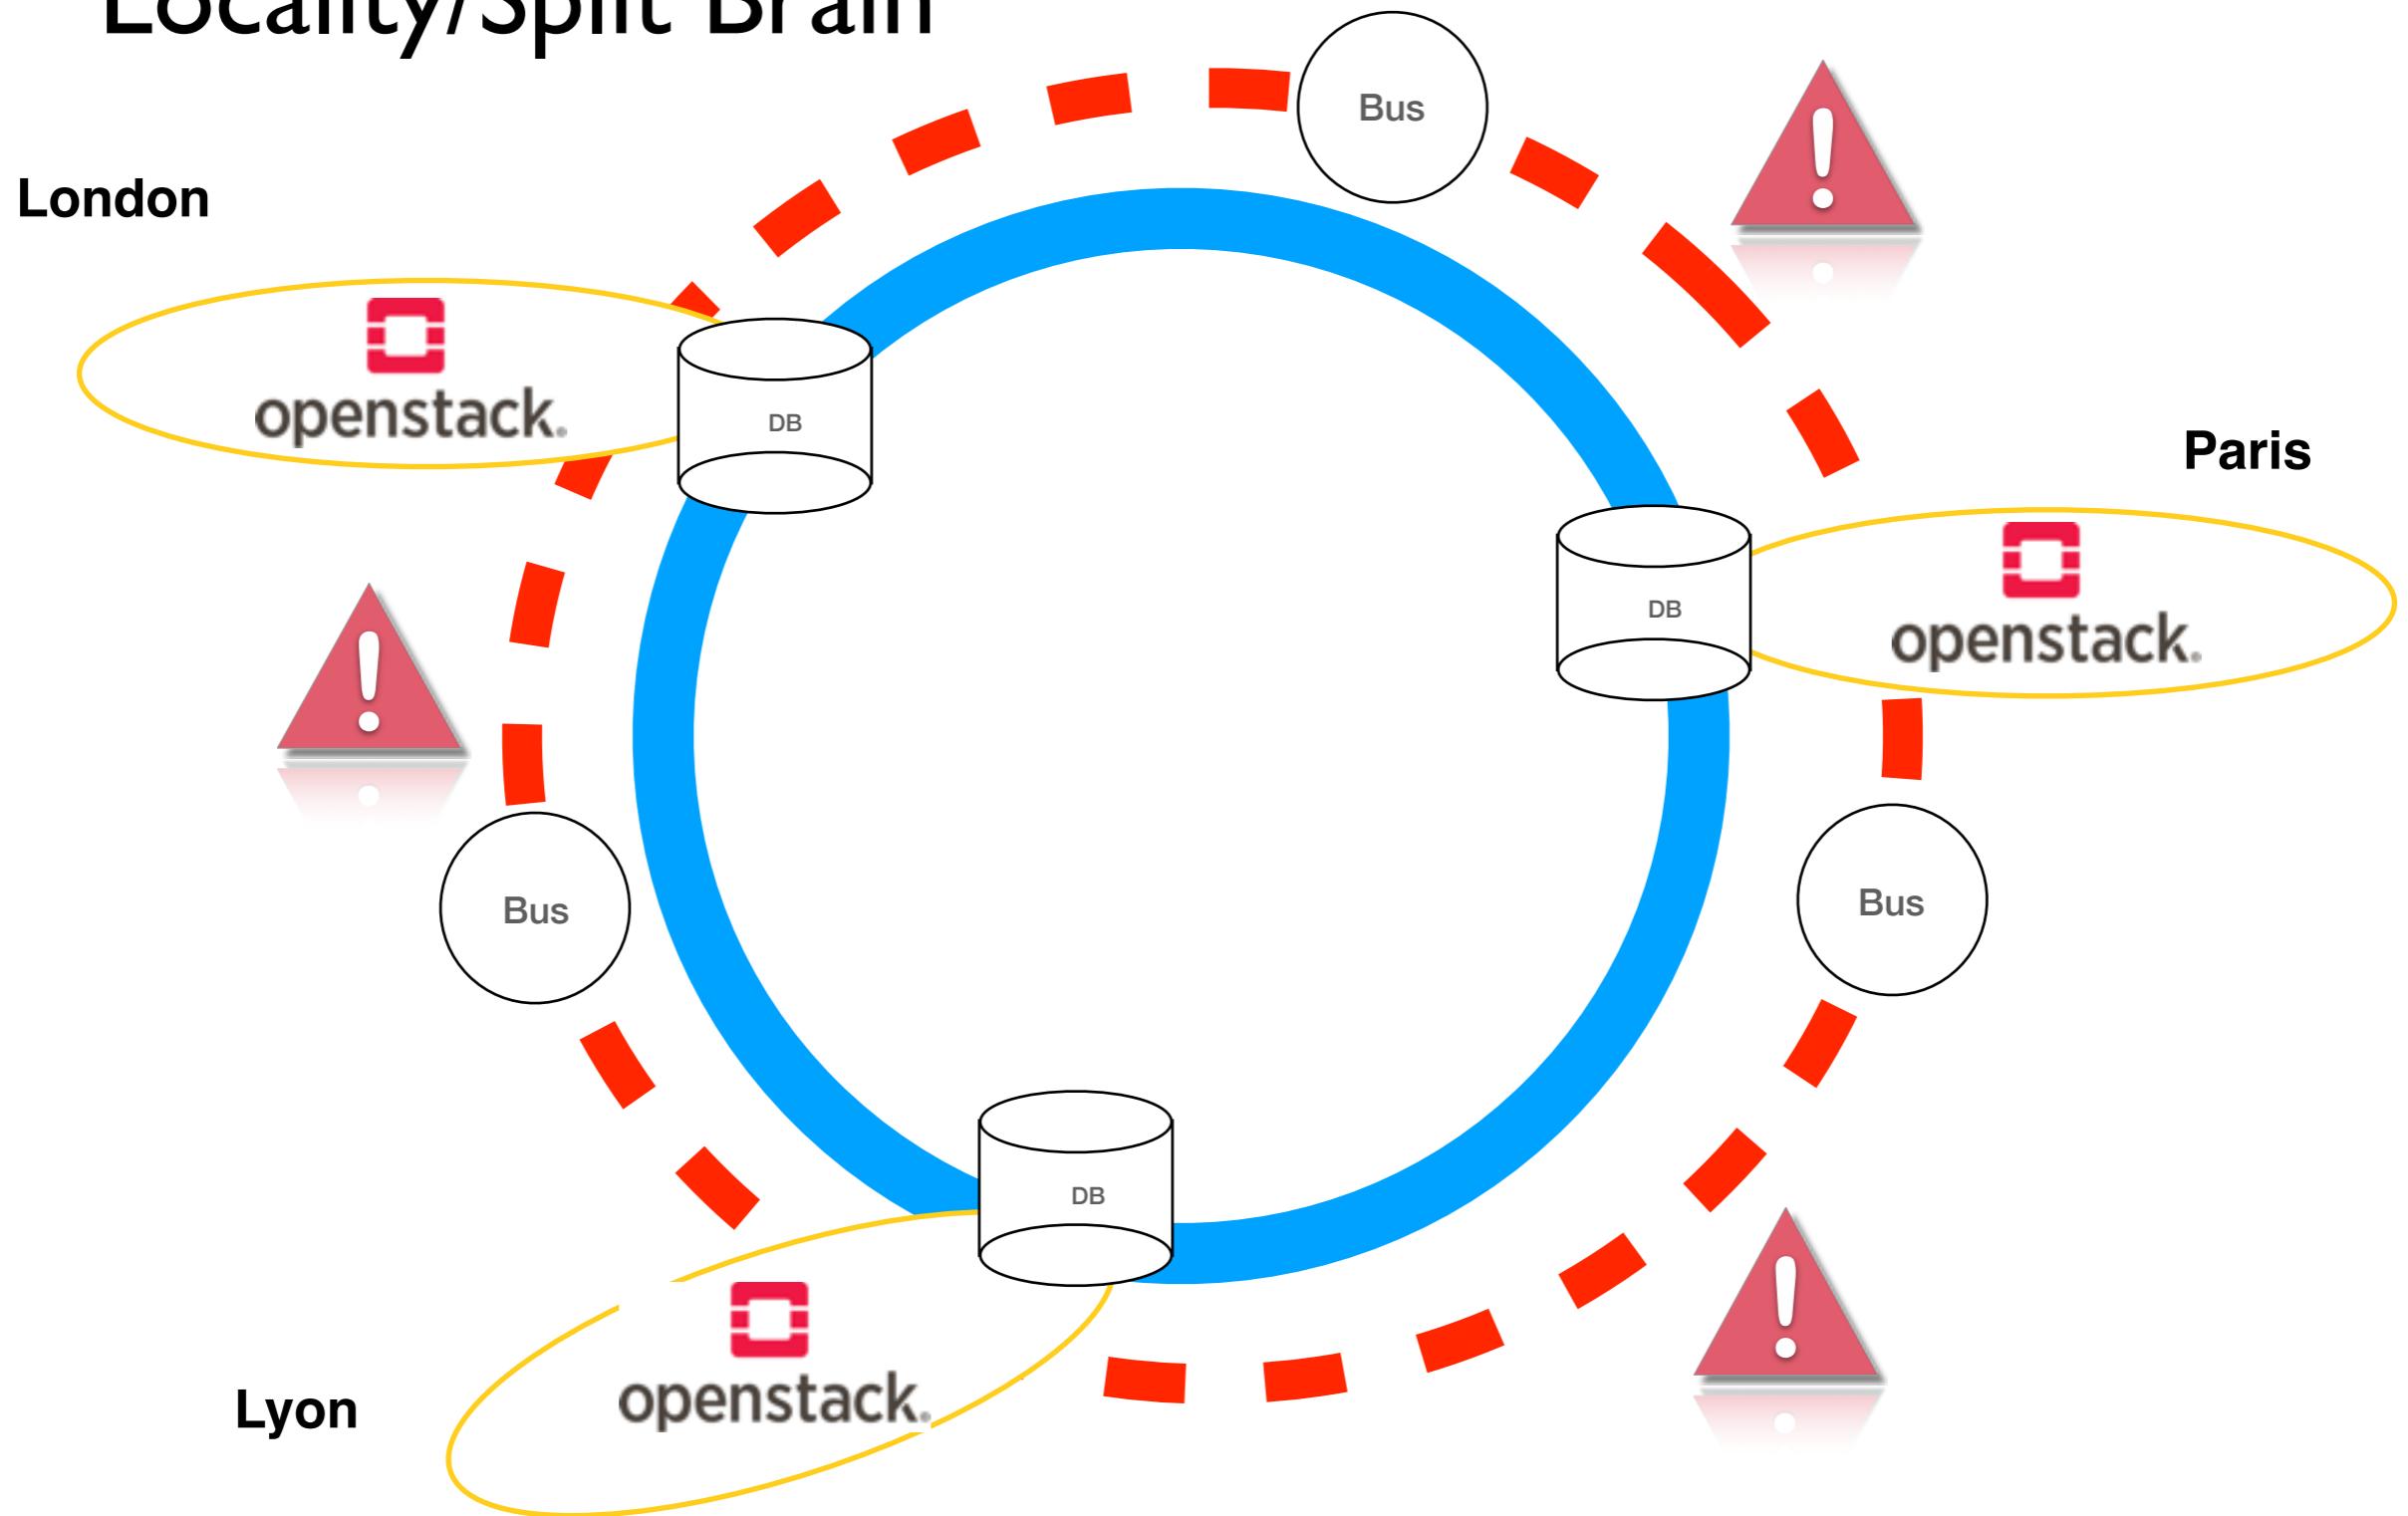
# Fog/Edge Challenges

- DBs

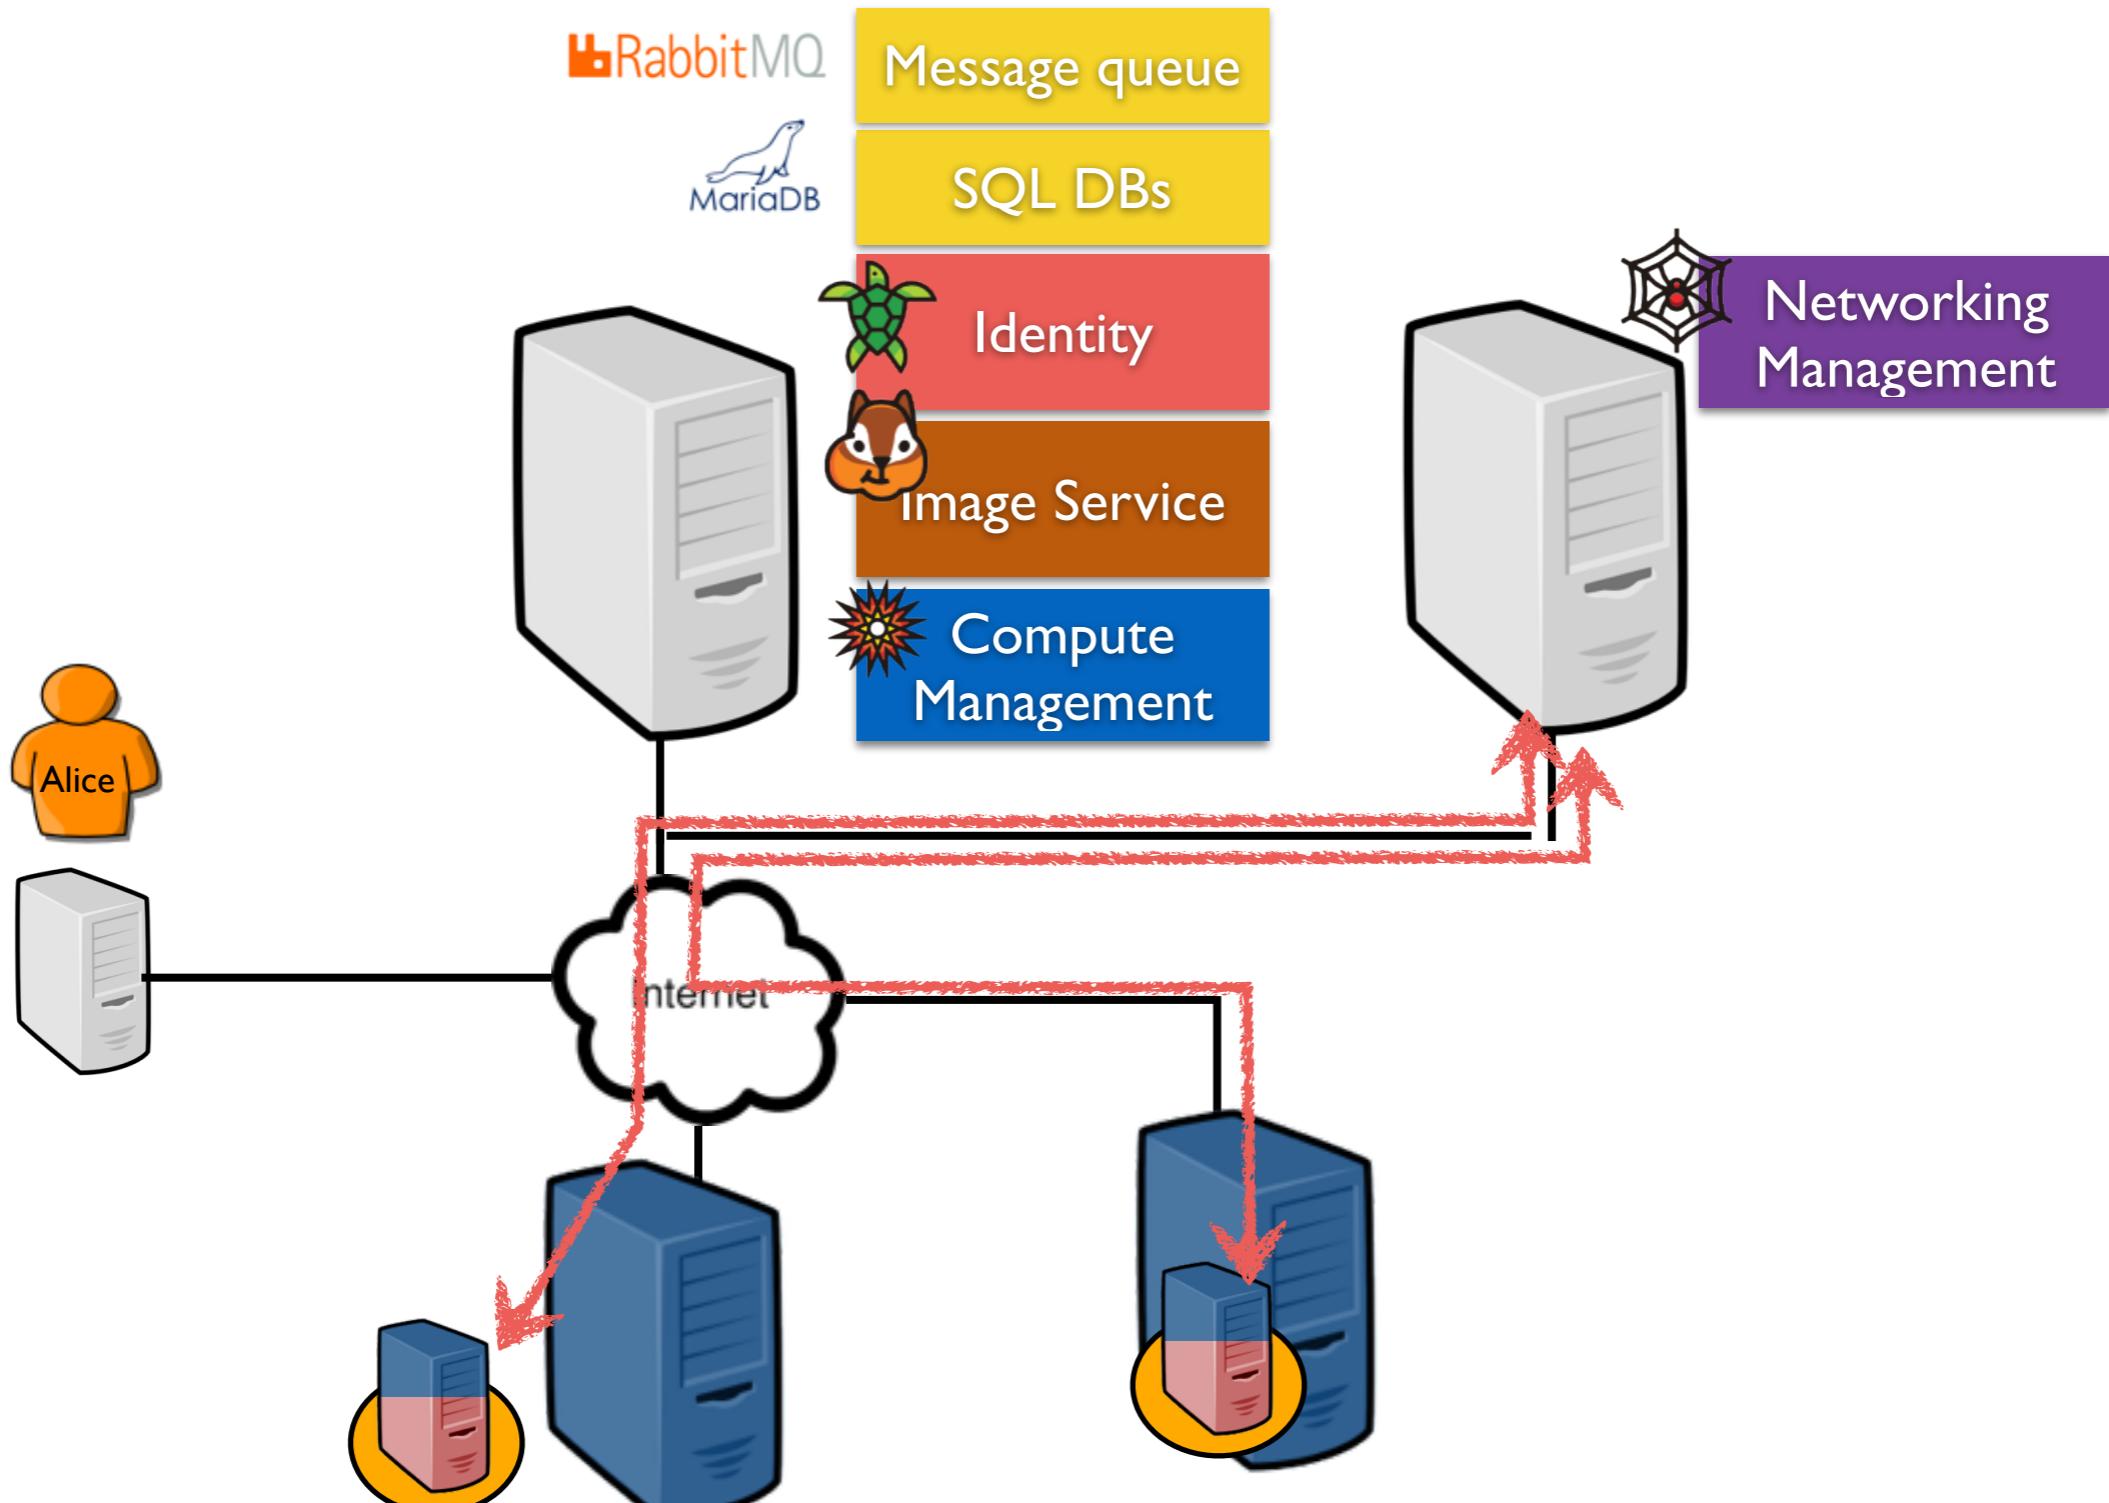


# Fog/Edge Challenges

- Locality/Split Brain

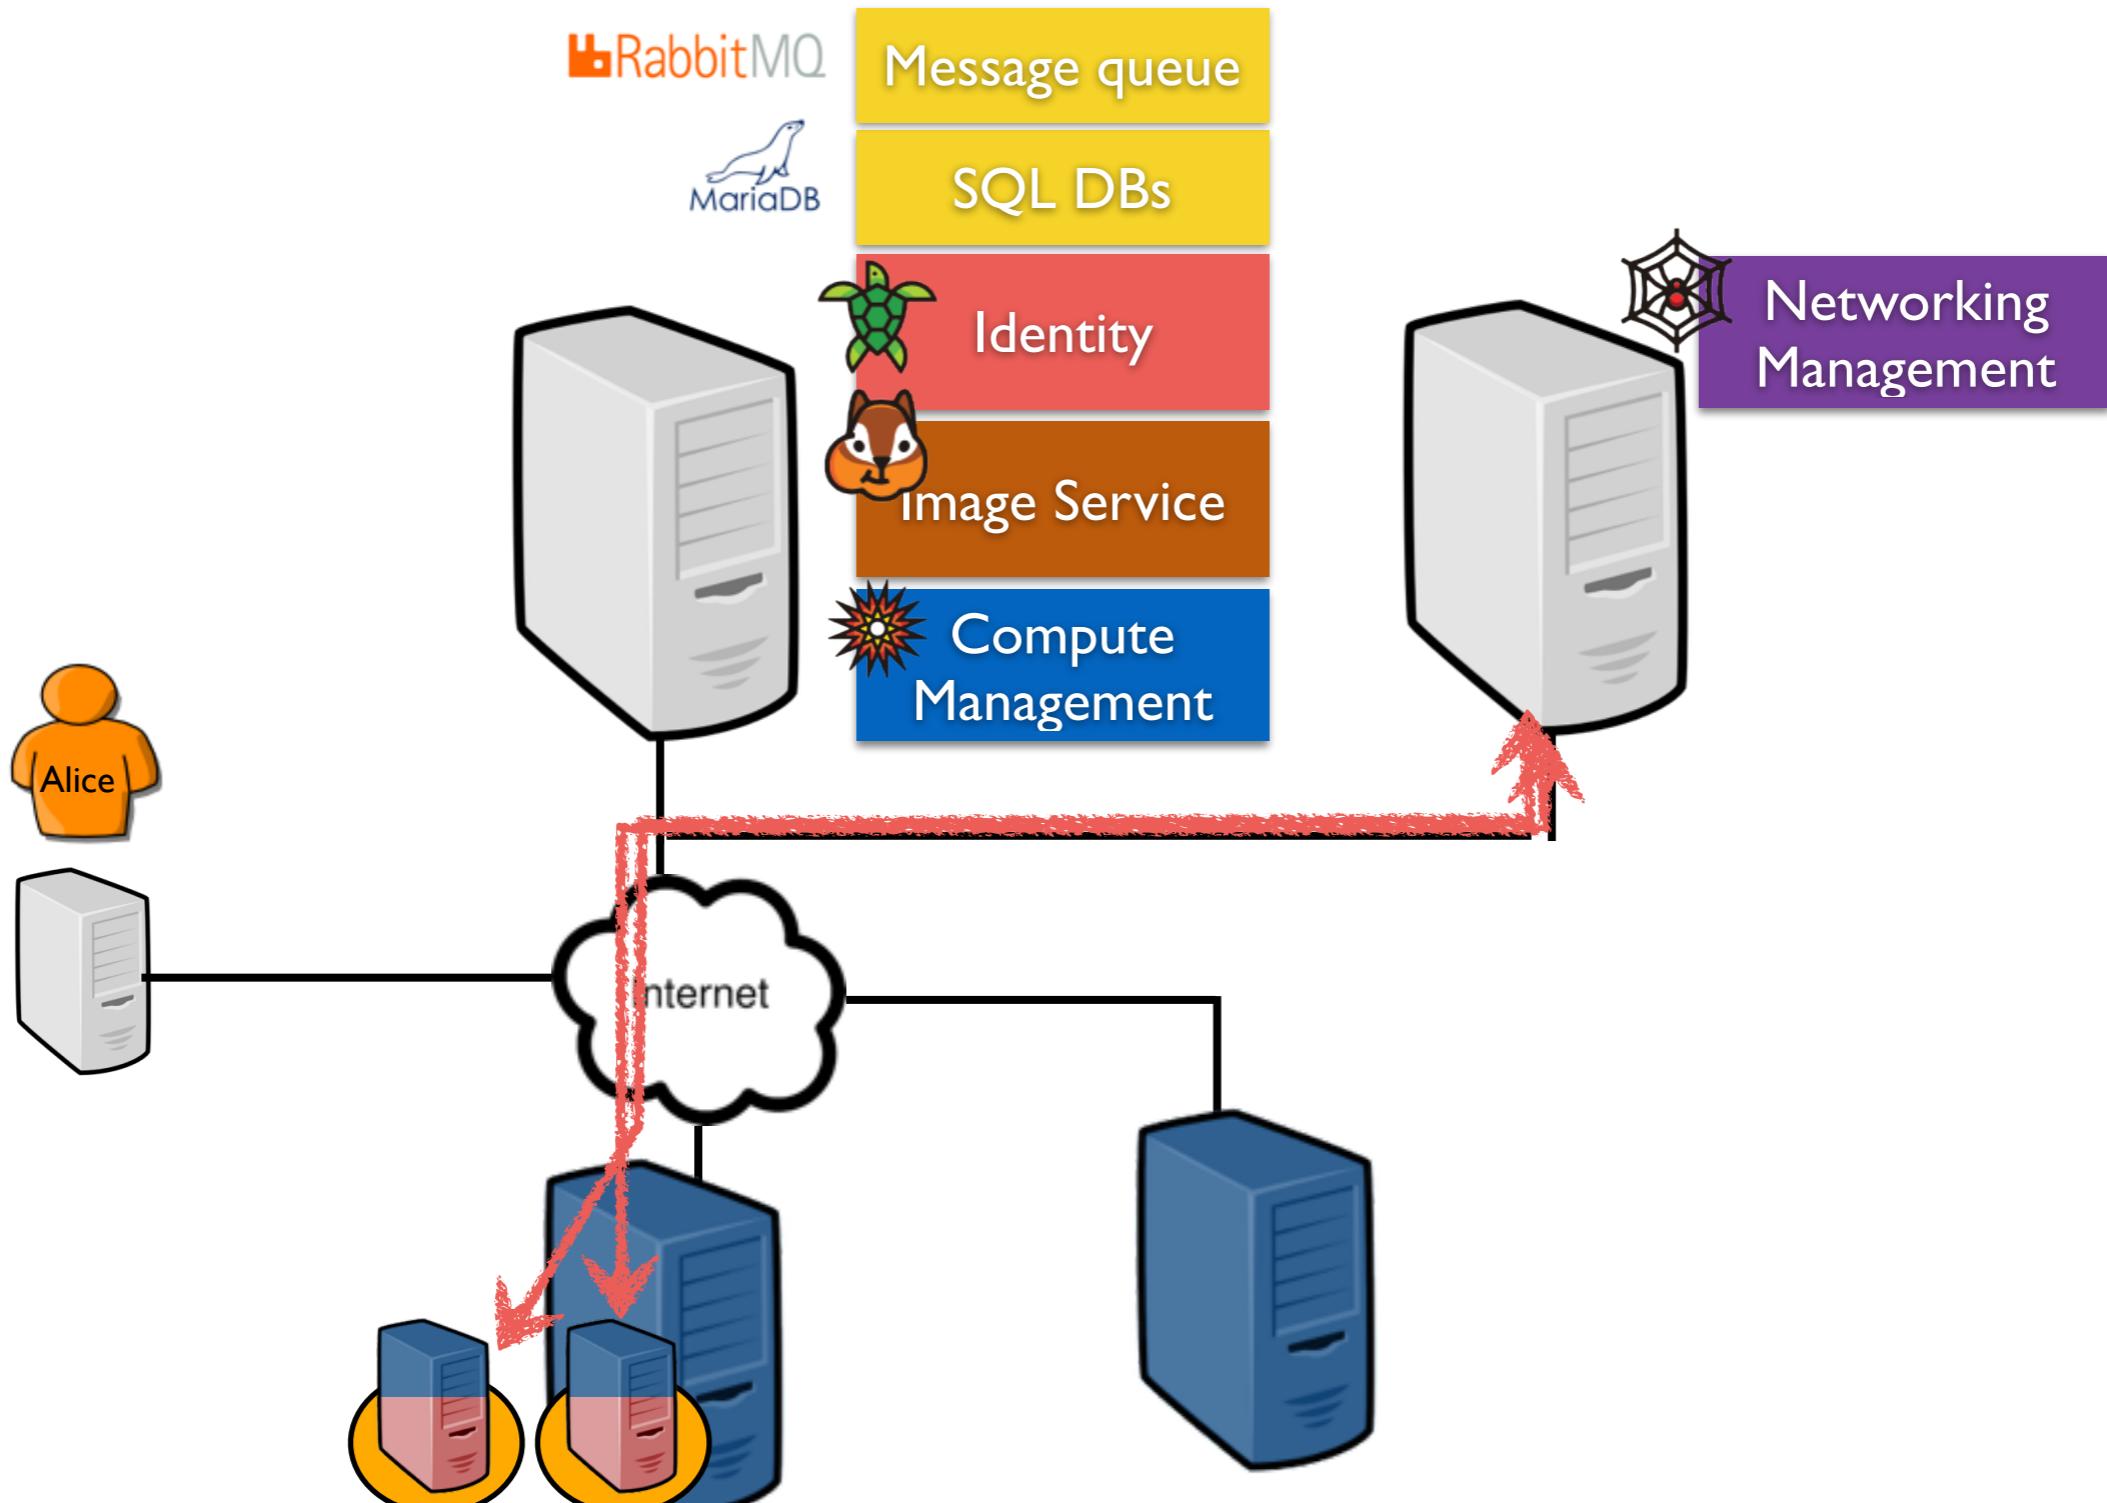


# Neutron and Fog/Edge challenges



East/West - Different project networks

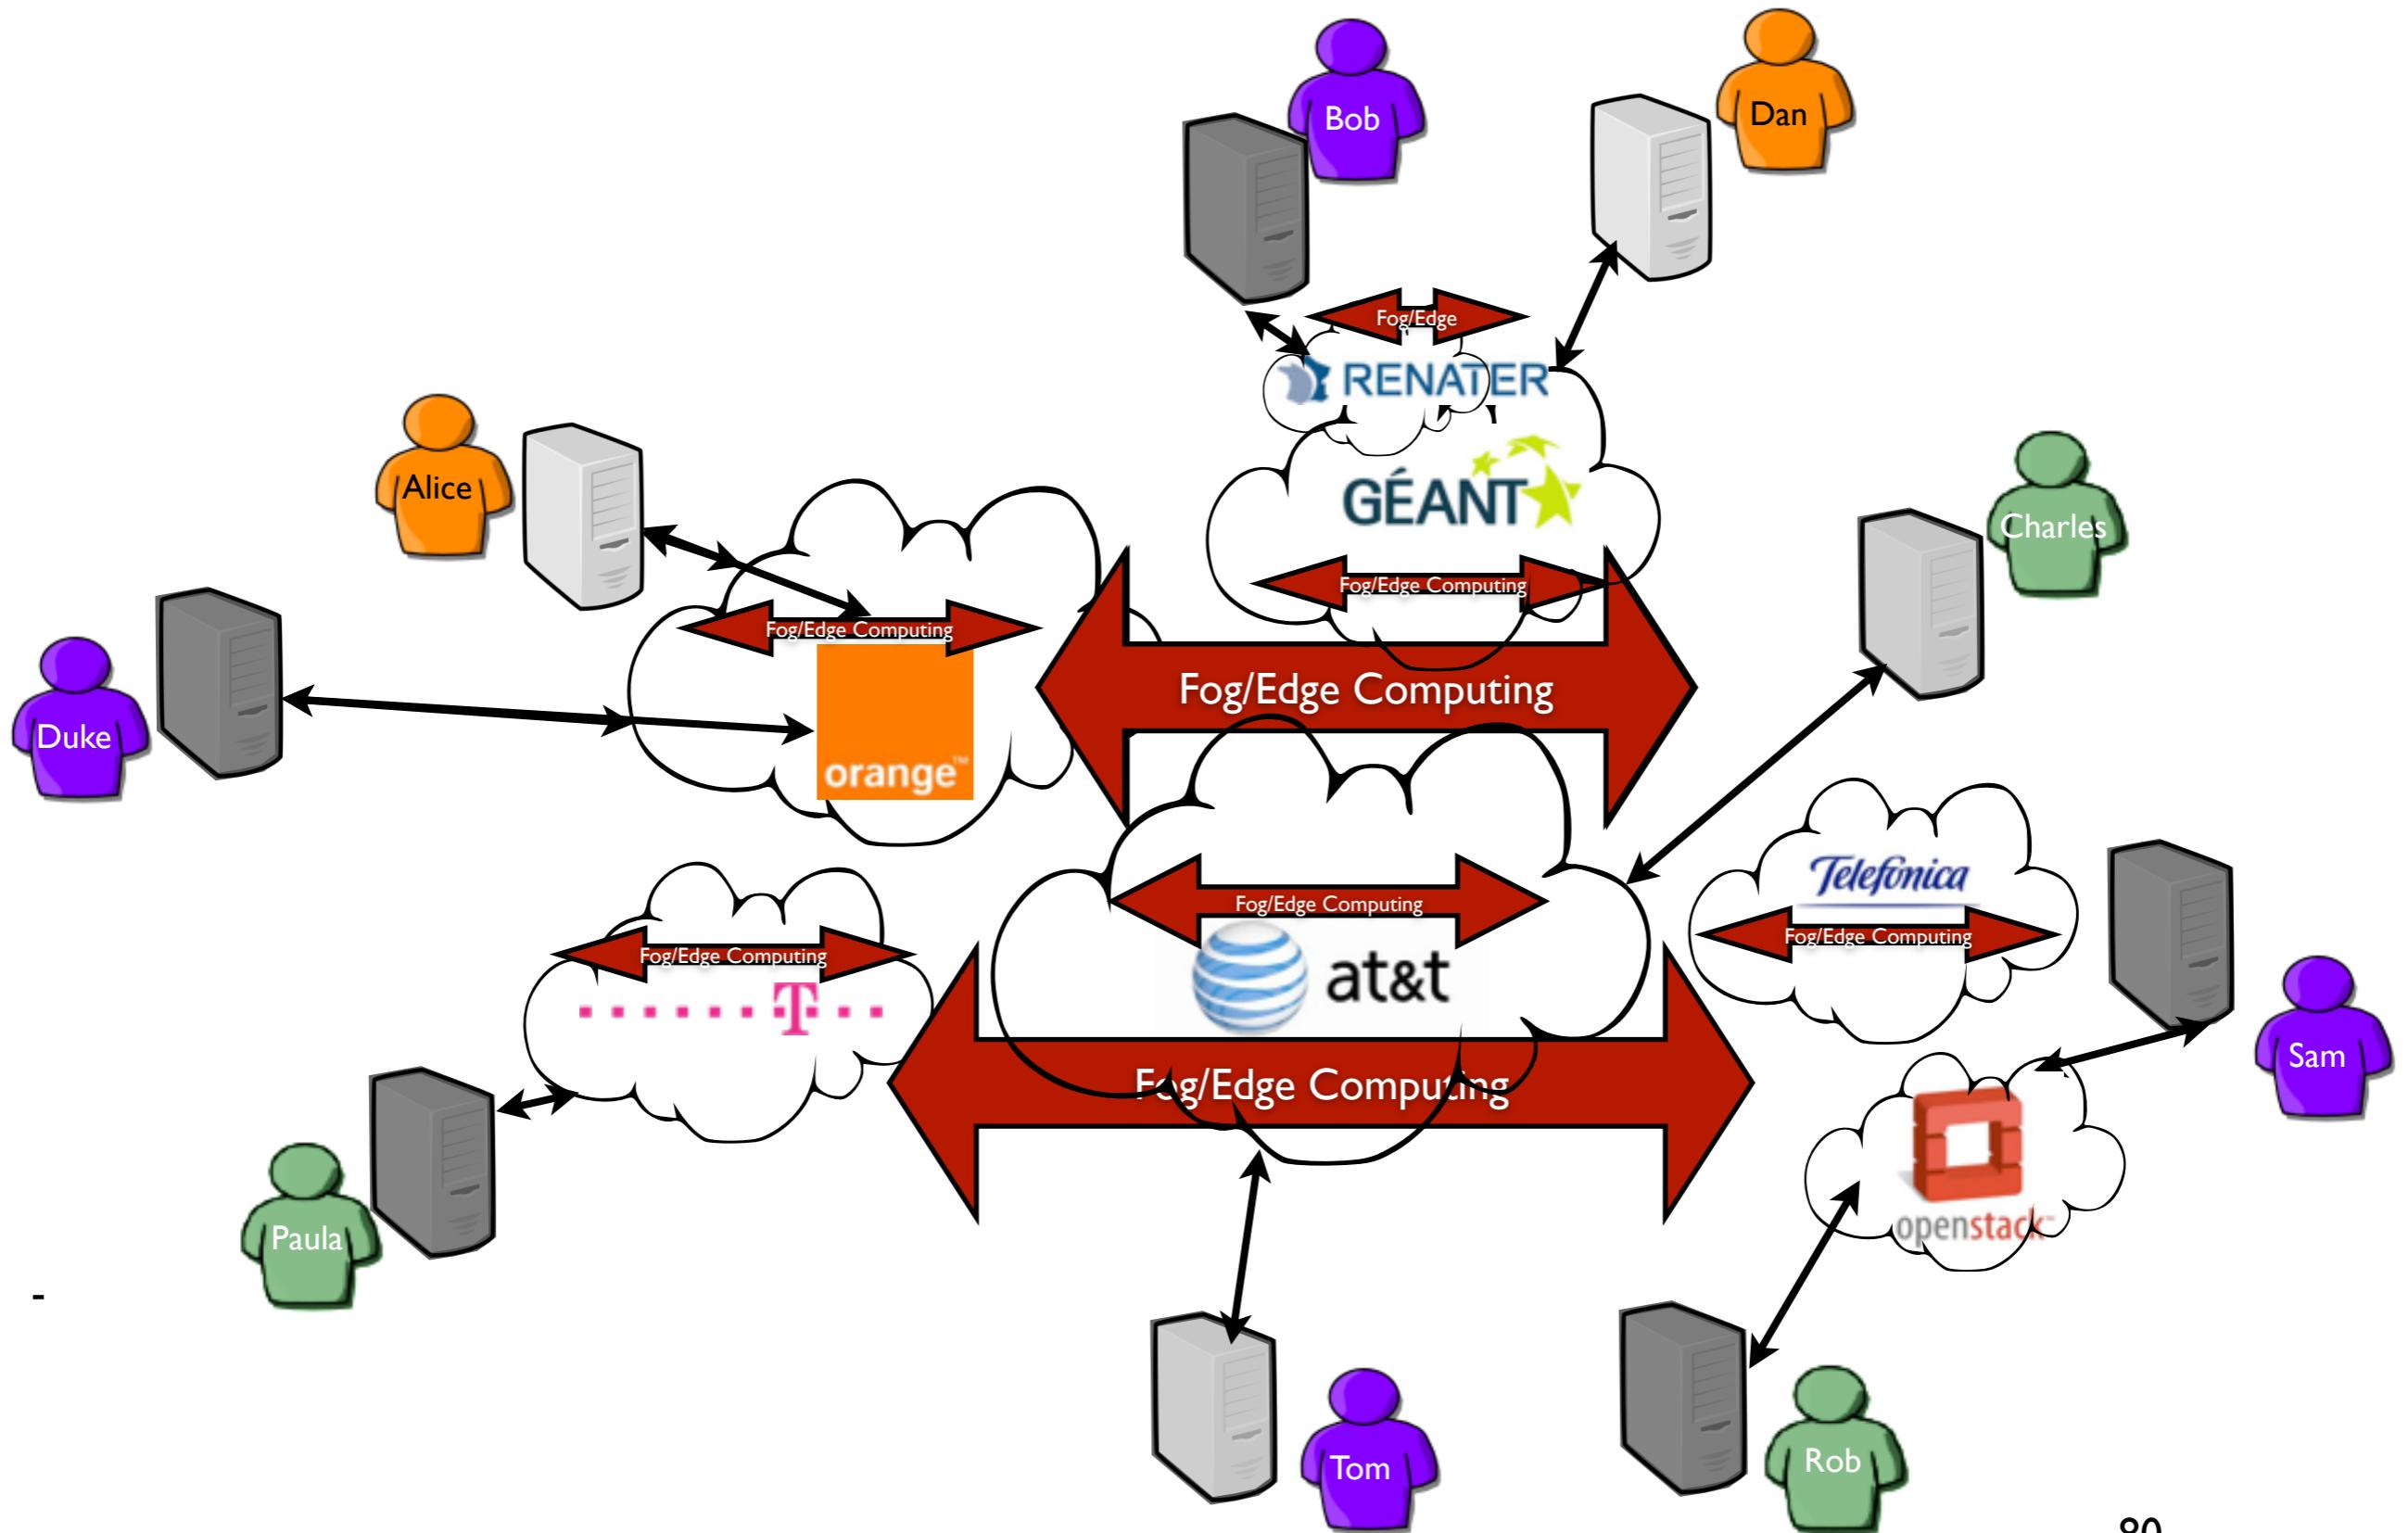
# Neutron and Fog/Edge challenges



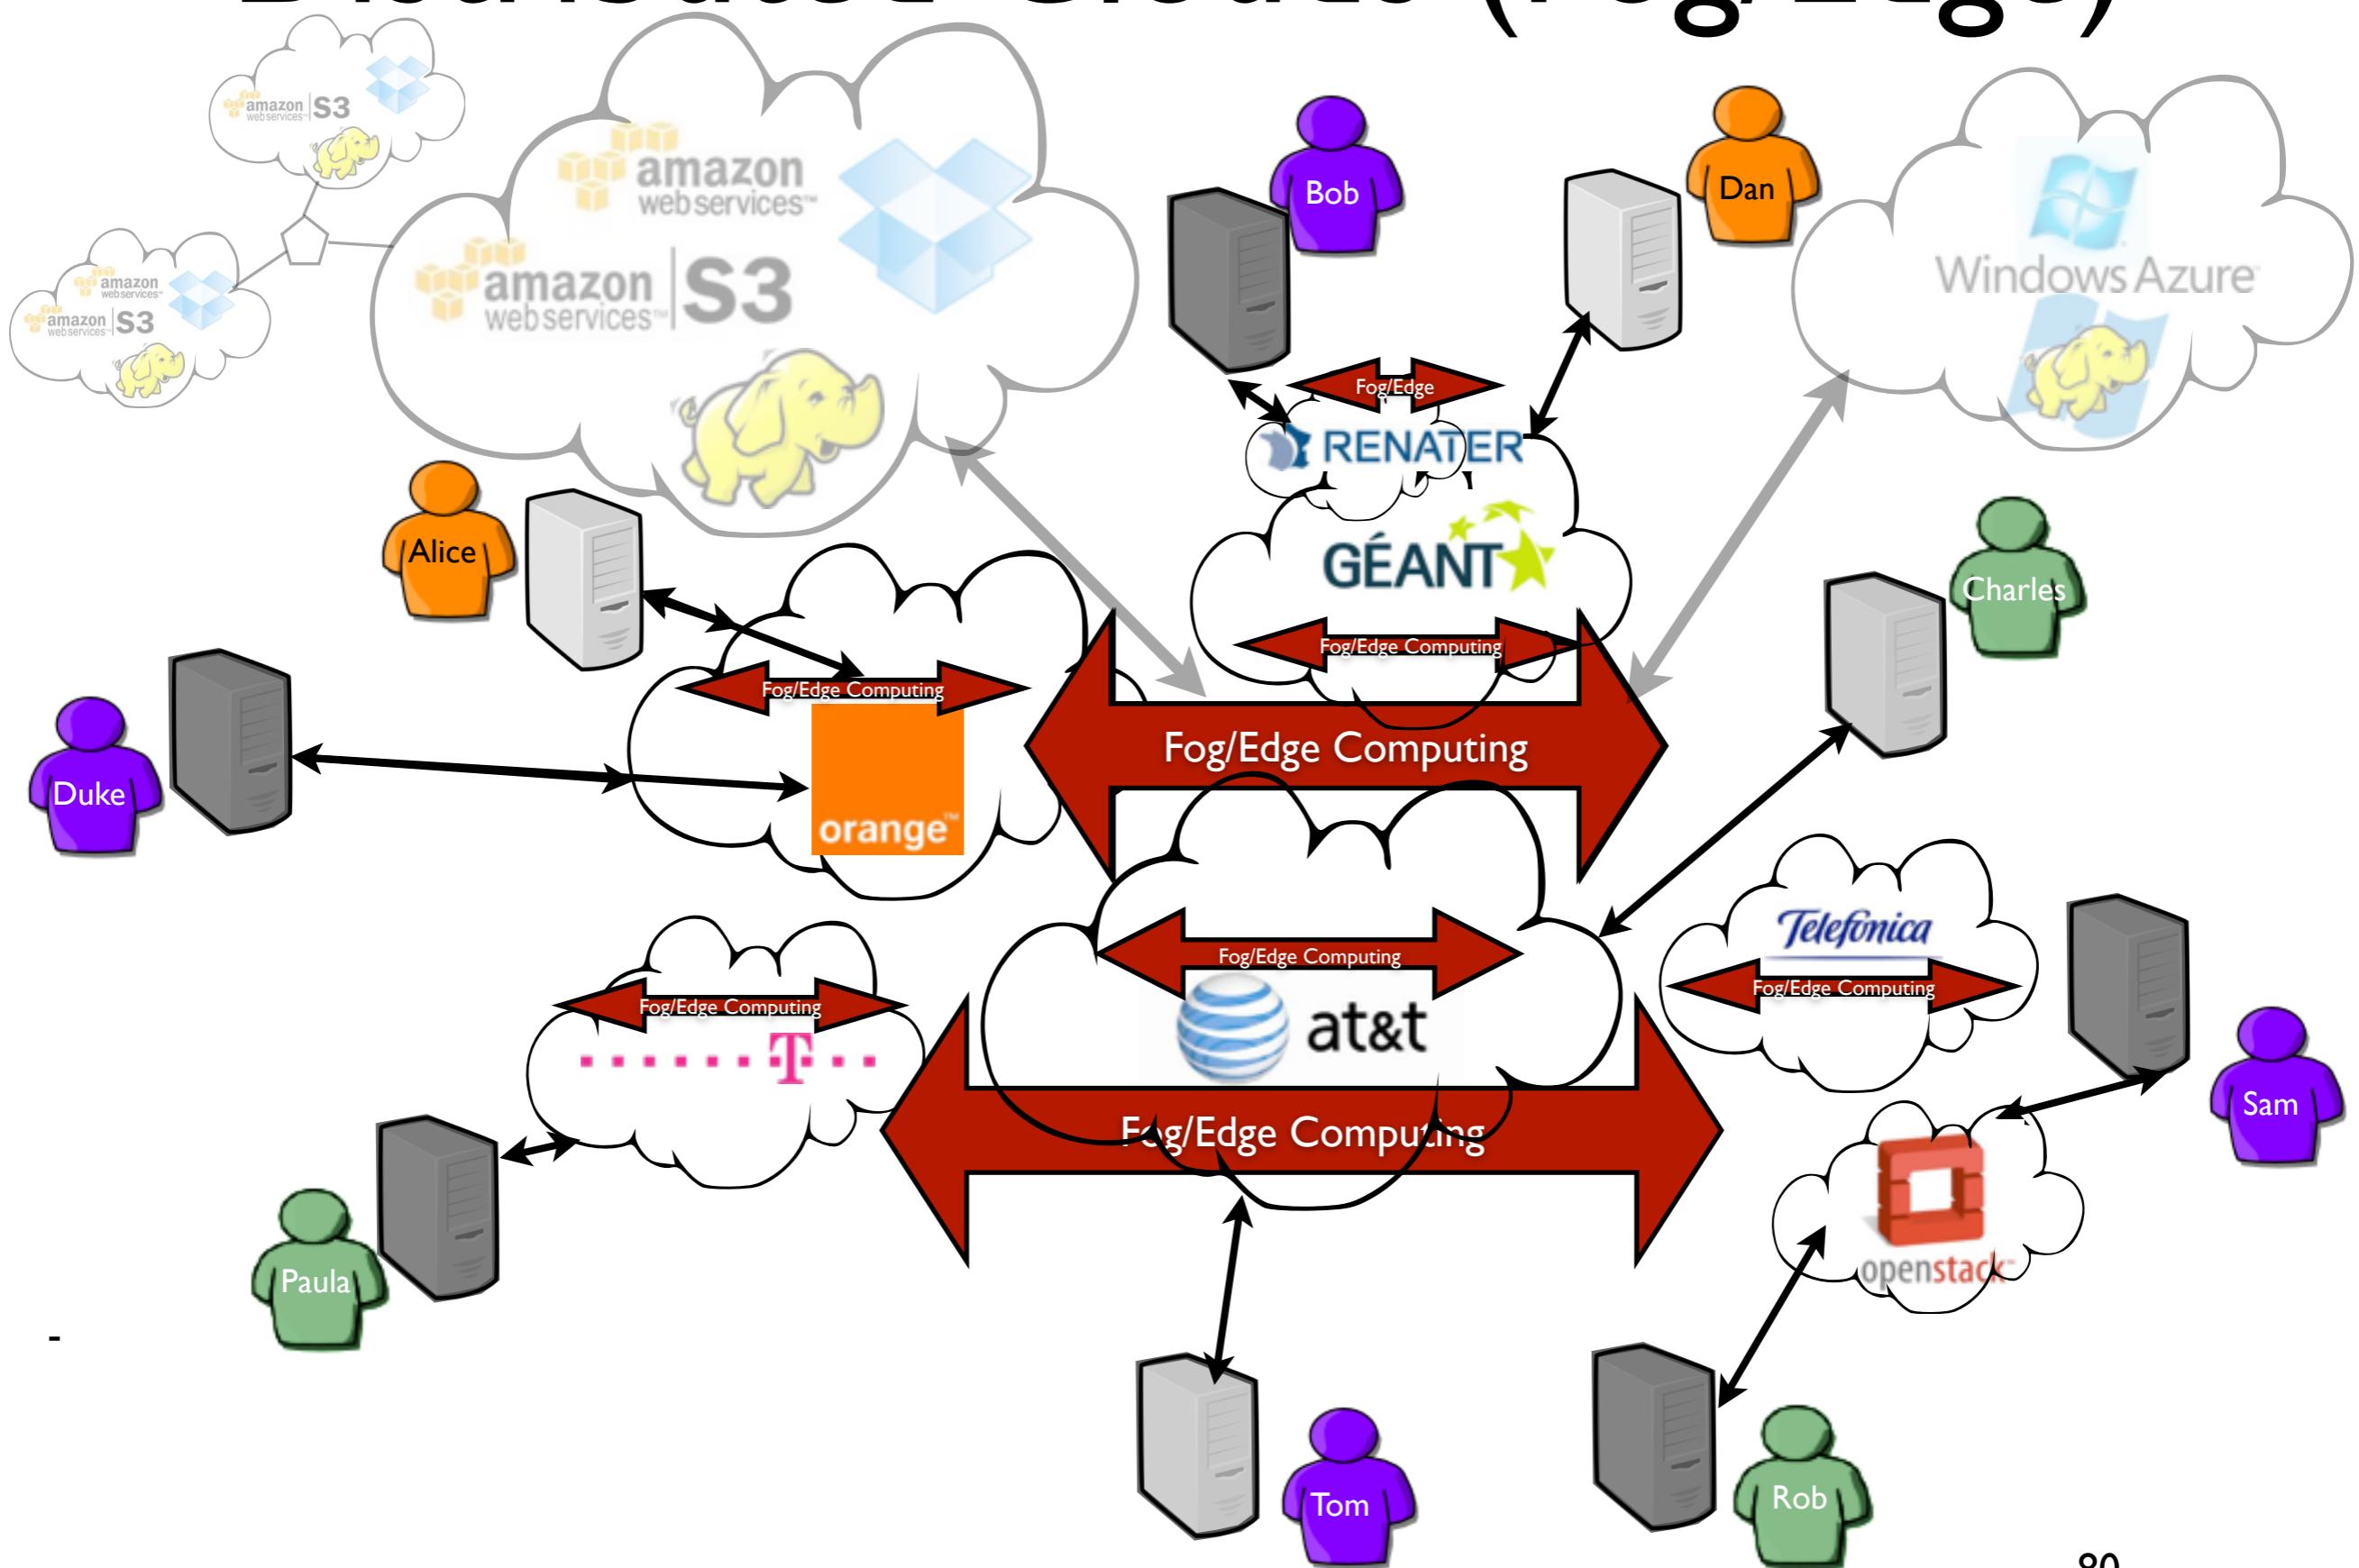
East/West - Different project networks

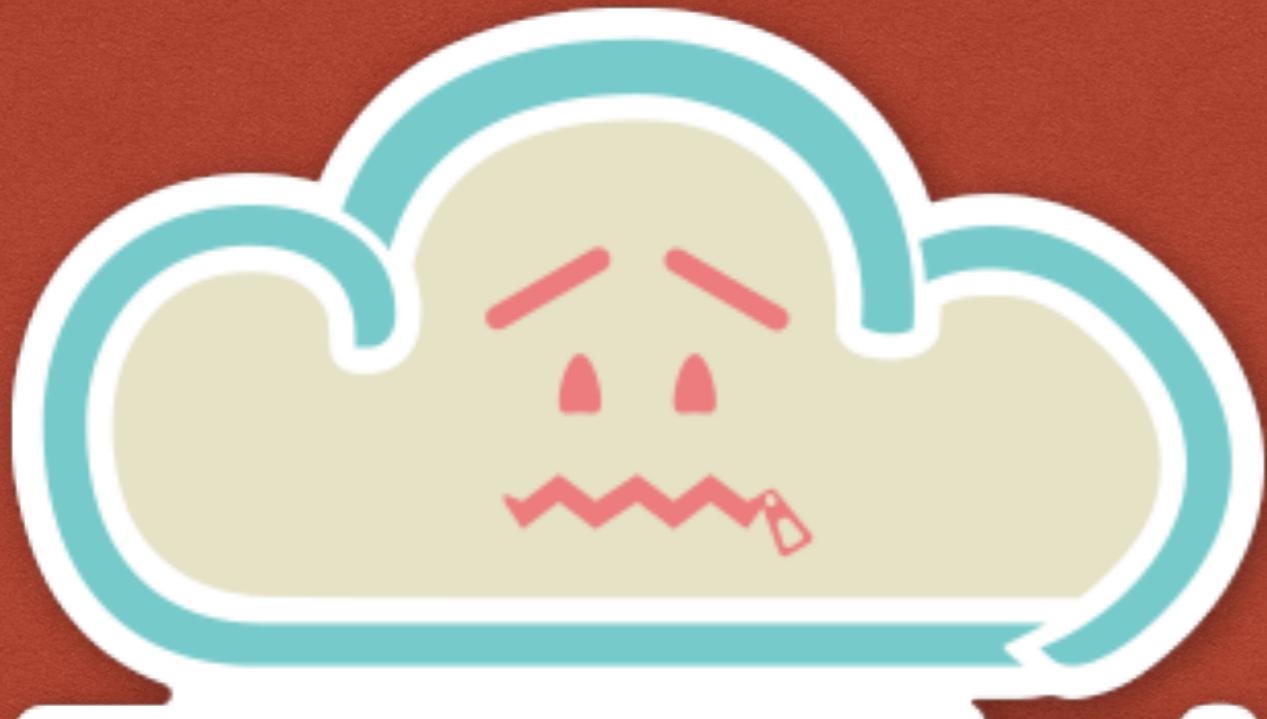
# *Takeaway Message*

# Distributed Clouds (Fog/Edge)



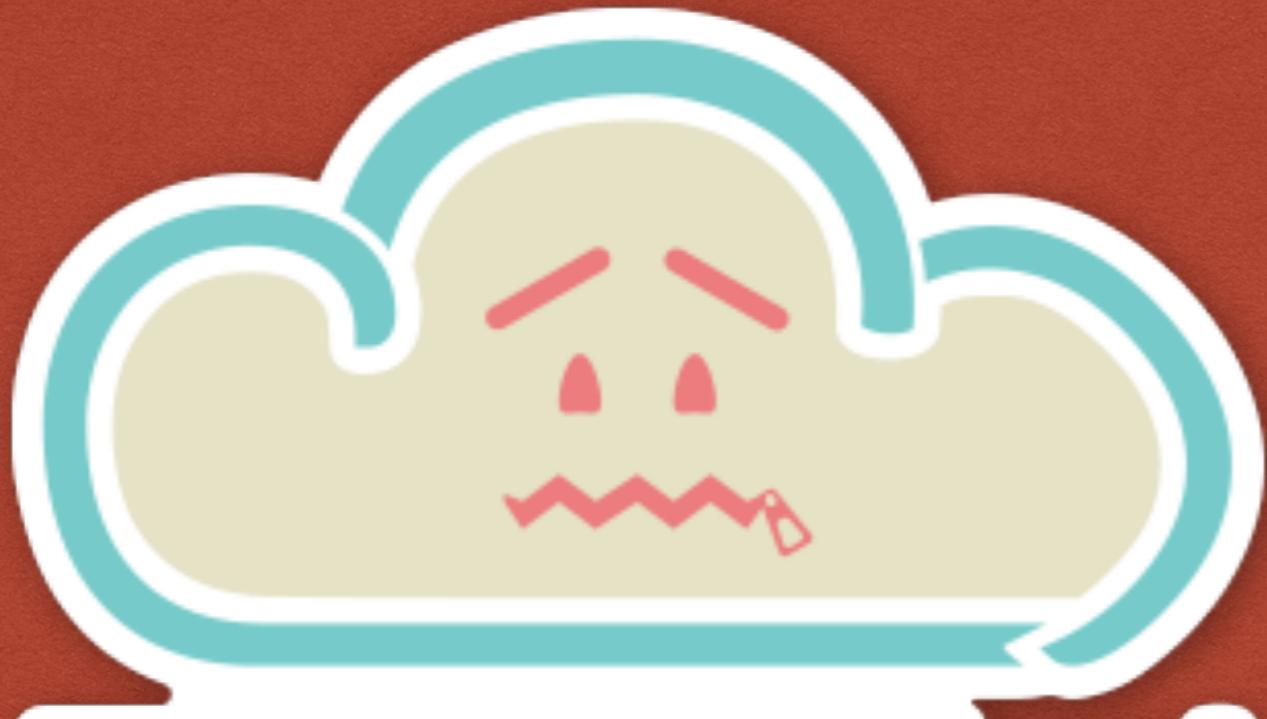
# Distributed Clouds (Fog/Edge)





**There is no cloud**  
it's just someone else's computer

Clouds hide the infrastructure...  
....by adding more layers !



**There is no cloud**  
it's just someone else's computer

and someone else's network

Clouds hide the infrastructure...  
....by adding more layers !

# Thanks

**Utility**

~~Cloud Computing technology is changing every day~~

How developers should develop new applications to benefit from geographically distributed infrastructures.

How to locate hardware/software components?

...

Do not hesitate to push the boundaries



<http://beyondtheclouds.github.io/>

We have Internship Positions

adrien.lebre@inria.fr

# Bibliography

- In a Nutshell - How OpenStack Works

A bit deprecated but good entry point to have a first idea in less than 10 minutes

<http://vmartinezdelacruz.com/in-a-nutshell-how-openstack-works/>

- OSONES SLIDES

Rich (almost up-to-date)

<https://github.com/Osones/formations>

- OpenStack official Documentation

Complete and up-to-date (at least yesterday ;-))

<https://www.openstack.org/software/>