

# Beyond the Clouds, the DISCOVERY Initiative



Credits: NASA


Localization is a key element to deliver  
**efficient** as well as **sustainable Utility Computing** Solutions



Adrien Lèbre / Ascola Project Team  
November, 2013



# Preliminary Comment



# Preliminary Comment

- Do not worry, we are not going to discuss all slides

Discovery idea, less than 3 slides


Why such an initiative, 18 slides

Interesting by additional details

Discovery in a nutshell

LRT (first POC)


What are we doing right now




*A simple Idea*

Bring Clouds back to the cloud


# The cloud from end-users



# The cloud in reality




# The DISCOVERY Initiative




*Why ?*

Let's give a look to  
the current situation


# Utility Computing - The Cloud




# Utility Computing - The Cloud




# Utility Computing - The Cloud



# Utility Computing - The Cloud



# Utility Computing - The Cloud




# The Cloud the ultimate UC model ...

- Almost mature for one site/cloud !

Open Nebula, Nimbus... vSphere... CloudStack, OpenStack  
More flexibility ! ? Infinite resources ! ?

- Current concerns

Scalability (VM Sprawl)




# The Cloud the ultimate UC model ...

- Almost mature for one site/cloud !

Open Nebula, Nimbus... vSphere... CloudStack, OpenStack  
More flexibility ! ? Infinite resources ! ?

- Current concerns

Scalability (VM Sprawl)




# The Cloud the ultimate UC model ...

- Almost mature for one site/cloud !

Open Nebula, Nimbus... vSphere... CloudStack, OpenStack  
More flexibility ! ? Infinite resources ! ?

- Current concerns

Scalability (VM Sprawl)




# The Cloud the ultimate UC model ...

- Almost mature for one site/cloud !

Open Nebula, Nimbus... vSphere... CloudStack, OpenStack  
More flexibility ! ? Infinite resources ! ?

- Current concerns

Scalability (VM Sprawl)  
Energy




# The Cloud the ultimate UC model ...

- Almost mature for one site/cloud !

Open Nebula, Nimbus... vSphere... CloudStack, OpenStack  
More flexibility ! ? Infinite resources ! ?

- Current concerns

Scalability (VM Sprawl)  
Energy  
Reliability




# The Cloud the ultimate UC model ...

- Almost mature for one site/cloud !

Open Nebula, Nimbus... vSphere... CloudStack, OpenStack  
More flexibility ! ? Infinite resources ! ?

- Current concerns

Scalability (VM Sprawl)  
Energy  
Reliability




# The Cloud the ultimate UC model ...

- Almost mature for one site/cloud !

Open Nebula, Nimbus... vSphere... CloudStack, OpenStack  
More flexibility ! ? Infinite resources ! ?

- Current concerns

Scalability (VM Sprawl)  
Energy  
Reliability



# The Cloud the ultimate UC model ...


- Almost mature for one site/cloud !

Open Nebula, Nimbus... vSphere... CloudStack, OpenStack  
More flexibility ! ? Infinite resources ! ?

- Current concerns


Scalability (VM Sprawl)  
Energy  
Reliability

Few Large  
offshore DCs




# ... Not so sure !

- Inherent limitations of the cloud computing model w.r.t public offers (or why building large offshore DCs is not appropriated).
  - I. Externalization of private applications/ data (jurisdiction concerns)



# ... Not so sure !


- Inherent limitations of the cloud computing model w.r.t public offers (or why building large offshore DCs is not appropriated).
  1. Externalization of private applications/ data (jurisdiction concerns)
  2. Overhead implied by the unavoidable use of the Internet to reach distant platforms



# ... Not so sure !

- Inherent limitations of the cloud computing model w.r.t public offers (or why building large offshore DCs is not appropriated).


- I. Externalization of private applications/ data (jurisdiction concerns)
2. Overhead implied by the unavoidable use of the Internet to reach distant platforms
3. The connectivity to the application/data cannot be insured by centralized dedicated centers (disaster recovery)



# ... Not so sure !

- Inherent limitations of the cloud computing model w.r.t public offers (or why building large offshore DCs is not appropriated).

- I. Externalization of private applications/ data (jurisdiction concerns)
2. Overhead implied by the unavoidable use of the Internet to reach distant platforms
3. The connectivity to the application/data cannot be insured by centralized dedicated centers (disaster recovery)




# ... Not so sure !

- Inherent limitations of the cloud computing model w.r.t public offers (or why building large offshore DCs is not appropriated).

- I. Externalization of private applications/ data (jurisdiction concerns)
2. Overhead implied by the unavoidable use of the Internet to reach distant platforms
3. The connectivity to the application/data cannot be insured by centralized dedicated centers (disaster recovery)

- Hybrid platforms: a promising approach  
It depends how you are going to extend the private one...



**Can we address these concerns “all in one” ? ?**


*How and where the  $\mu$ DC concept can be deployed ?*

# Locality Based Utility Computing Toward LUC Infrastructures

# Beyond the Cloud, the DISCOVERY Initiative

- Locality-based UC infrastructures

The only way to deliver highly efficient and sustainable UC services is to provide UC platforms as close as possible to the end-users.




# Beyond the Cloud, the DISCOVERY Initiative

- Locality-based UC infrastructures

The only way to deliver highly efficient and sustainable UC services is to provide UC platforms as close as possible to the end-users.

- Leveraging network backbones


Extend any point of presence of a network backbone with UC servers (from network hubs up to major DSLAMs that are operated by telecom companies and network institutions).



⇒ Operating such widely distributed resources requires the definition of a fully distributed system


# The DISCOVERY Proposal

- DIStributed and COoperative framework to manage Virtual EnviRonments autonomicallY (the LUC OS)
- Relying on a minimal (but common API)  
**libvirt** / OCCI / ...
- 3 services  
Discovery Network Tracker (DNT)  
Virtual Environments Tracker (VET)  
Local Resources Tracker (LRT)




# The DISCOVERY Proposal

- DIStributed and COoperative framework to manage Virtual EnviRonments autonomicallY (the LUC OS)
- Relying on a minimal (but common API)  
**libvirt** / OCCI / ...
- 3 services  
Discovery Network Tracker (DNT)  
Virtual Environments Tracker (VET)  
Local Resources Tracker (LRT)




# The DISCOVERY Proposal

- DIStributed and COoperative framework to manage Virtual EnviRonments autonomicallY (the LUC OS)
- Relying on a minimal (but common API)  
**libvirt** / OCCI / ...
- 3 services  
Discovery Network Tracker (DNT)  
Virtual Environments Tracker (VET)  
Local Resources Tracker (LRT)



# The DISCOVERY Proposal

- DIStributed and COoperative framework to manage Virtual EnviRonments autonomicallY (the LUC OS)
- Relying on a minimal (but common API)  
**libvirt** / OCCI / ...
- 3 services  
Discovery Network Tracker (DNT)  
Virtual Environments Tracker (VET)  
Local Resources Tracker (LRT)



# The DISCOVERY Initiative

- Focusing on the design and the implementation of a complete OS for IaaS platforms (i.e. the LUC OS)

The LUC OS

Based on VMs and VEs (group of VMs) as the fundamental granularity

**Scalability**, targeting the management of hundred thousands of VMs upon thousands of physical machines (PMs)

**Reliability**, considering “hardware failures as the norm rather the exception”

**Reactivity**, handling each reconfiguration event as swiftly as possible to maintain VEs' QoS.

- May look simple but lots of scientific/technical challenges


Cost of the DISCOVERY network !? partial view of the system !?

Impact on the others VMs !?, management of VM images !?


Which software abstractions to make the development easier and more reliable (distributed event programming) ? How to take into account locality aspects ?

- A BitTorrent like system ... but with stronger assumptions


# Beyond the Cloud, the DISCOVERY Initiative




# Beyond the Cloud, the DISCOVERY Initiative




# Beyond the Cloud, the DISCOVERY Initiative




# Beyond the Cloud, the DISCOVERY Initiative



# Beyond the Cloud, the DISCOVERY Initiative



# Beyond the Cloud, the DISCOVERY Initiative



# The Discovery Initiative Pros/Cons

- Pros

- Locality

- (jurisdiction concerns, latency-aware apps, minimize network overhead)

- Reliability/redundancy (no critical point/location/centers)

- The infrastructure is naturally distributed throughout multiple areas

- Lead time to delivery

- Leverage current PoPs and extend them according to UC demands

- Energy footprint (to be confirmed)

- Bring back part of the revenue to NRENs/Telcos*

- Cons

- Security concerns (in terms of who can access to the PoPs)

- Operate a fully IaaS in a unified but distributed manner at WAN level

# The DISCOVERY Initiative

- Leveraging former projects but still on the starting blocks!
- Strong interests of large companies  
(SAP, Orange Lab, Citrix, ...)
- RENATER
- An important actor to follow: Akamai (micro DCs)
- Preliminary works with promising results  
(especially on the LRT: a first POC)
- Long term objective: impact on the design of distributed applications in order to take advantage of the locality  
(building S3 like system)

# Conclusion

- Cloud Computing technology is changing every day
  - New features, new requirements

The main challenge of the Discovery Initiative is to ensure that such new features/mechanisms can run in a distributed manner.
- But Distributed Cloud Computing is happening !
  - Dist. CC workshop (collocated with IEEE/ACM UCC 2013)
  - FOG Computing workshop (collocated with IEEE ICC 2013)

# Amazon is on the way !

[Sign Up](#)

My Account / Console ▾

English ▾

[AWS Products & Solutions ▾](#)[AWS Product Information ▾](#)[Developers ▾](#)[Support ▾](#)

## Global Infrastructure

Amazon Web Services serves hundreds of thousands of customers in more than 190 countries.

We are steadily expanding global infrastructure to help our customers achieve lower latency and higher throughput, and to ensure that their data resides only in the Region they specify. As our customers grow their businesses, AWS will continue to provide infrastructure that meets their global requirements.

[See detailed list of offerings at all AWS locations](#)



## Europe / Middle East / Africa



### EU (Ireland) Region

EC2 Availability Zones: 3 Launched 2007

### AWS Edge Locations

Amsterdam, The Netherlands (2)	Dublin, Ireland	Frankfurt, Germany (3)	London, England (3)	Madrid, Spain
Marseille, France	Milan, Italy	Paris, France (2)	Stockholm, Sweden	Warsaw, Poland

# Amazon is on the way !

[Sign Up](#)[My Account / Console](#)

English

[AWS Products & Solutions](#)[AWS Product Information](#)[Developers](#)[Support](#)

## Amazon CloudFront

- [CloudFront Overview](#)
- [FAQs](#)
- [Pricing](#)
- [Amazon CloudFront SLA](#)
- [What's New?](#)
- [Amazon CloudFront Events](#)

## Amazon CloudFront What's New?

[Back to the CloudFront page.](#)

### What's New:

#### **Announcing New Edge Locations in Manila, Marseille and Warsaw for Amazon CloudFront and Amazon Route 53**

**Date:** Dec 15th, 2013

**Details:** We are excited to announce the launch of three new edge locations – Manila in the Philippines, Marseille in France and Warsaw in Poland. These new locations will improve performance and availability for end users of your applications being served by Amazon CloudFront and Amazon Route 53 and bring the total number of AWS edge locations to 49 worldwide. Learn more by reading our [announcement](#).

#### **Amazon CloudFront Announces Atlanta, GA PoP and Additional Pops in London and Frankfurt**

**Date:** Nov 3rd, 2013

**Details:** We're excited to announce the launch of a new Amazon CloudFront edge location in Atlanta, GA. We have also recently added third edge locations in London, UK and Frankfurt, Germany in order to increase connectivity and to provide even better service for our customers. Learn more by reading our [announcement](#).

#### **Announcing Amazon CloudFront Support for POST/PUT and other HTTP Methods**

**Date:** Oct 15th, 2013

**Details:** We are excited to announce that Amazon CloudFront has added support for five additional HTTP methods: POST, PUT, DELETE, OPTIONS and PATCH. This means you can now use CloudFront to accelerate data uploaded from end users, improving the performance of dynamic and personalized websites that have web forms, comment and login boxes, "add to cart" buttons or other features. Learn more by reading our [announcement](#) or the [Amazon CloudFront Developer Guide](#). You can also attend our webinar "[Using Amazon CloudFront to Accelerate Your Static, Dynamic, and Interactive Content](#)" on November 7, 2013 at 10AM PDT to learn more.

## Developer Resources

- [AWS Management Console](#)
- [Documentation](#)
- [Release Notes](#)
- [Sample Code & Libraries](#)
- [Developer Tools](#)
- [Community Forum](#)

## Streaming Media Awards

Amazon CloudFront receives  
Streaming Media Magazine's

# Beyond Discovery !

- From sustainable data centers to a new source of energy

The only way to deliver highly efficient and sustainable UC services is to provide UC platforms as close as possible to the end-users and to...

- Leverage “green” energy (solar, wind turbines...)

Transfer the green micro/nano DCs concept to the network PoP


Take the advantage of the geographical distribution



<http://parasol.cs.rutgers.edu>

- Leveraging the data furnaces concept

Deploy UC servers in medium and large institutions and use them as sources of heat inside public buildings such as hospitals or universities




<https://www.aoterra.de>

# Beyond Discovery !

- From sustainable data centers to a new source of energy

The only way to deliver highly efficient and sustainable UC services is to provide UC platforms as close as possible to the end-users and to...



- Leverage “green” energy (solar, wind turbines...)

Transfer the green micro/nano DCs concept to the network PoP


Take the advantage of the geographical distribution



<http://parasol.cs.rutgers.edu>

- Leveraging the data furnaces concept


Deploy UC servers in medium and large institutions and use them as sources of heat inside public buildings such as hospitals or universities



<https://www.aoterra.de>

# The DISCOVERY Initiative

- Thank you / Questions ?
- Additional materials
  - Focus on LRT (Flavien Quesnel's Phd, ended in Feb 2013)
  - Discovery internals in a nutshell
  - On going work - The discovery framework from the Software Programming viewpoint (Jonathan Pastor's Phd, 2012/2015)



<http://beyondtheclouds.github.io/>